

GRANDE FINALE V4: EXACT THRESHOLD COMPILERS, PRIMITIVE Q, AND ORDER-32 REED–SOLOMON MCA

PRZEMEK CHOJECKI

ABSTRACT. This paper is a complete companion to CAP25 v13.2 [Cho26]. It imports by citation, and does not reprove, the support-wise MCA conventions, universal cap, four deployed unsafe endpoints, identity-prefix and simple-pole constructions, aperiodic and subfield floors, prefix rigidity and exact second moment, quotient and extension reductions, balanced-core reduction, and normalized finite ledger established there.

The first half develops the full safe-side compiler. It proves exact support sparsification, syndrome–secant and fixed-union compilers, recursive affine-span and generalized-weight bounds, common-zero and directional Johnson estimates, codimension-one and arbitrary-codimension support-flag recursions, a minimum-distance sharpening using the actual direction-code support, fixed-mismatch puncturing, a quadratic mean-overlap staircase, exact quotient–remainder prefix normal forms, effective-image Fourier normalization, primitive Q after an explicit Sidon payment, saturation, a moving-root theorem for one-parameter split pencils, and the implication that a row-sharp Q theorem pays the primitive shift-pair ledger.

The second half develops a different MCA completion mechanism. A rank-two reduced intersection matrix gives an exact kernel dictionary and shows that the deployed critical order is 32. Exact dyadic descent identifies the quotient branch. Maximal-support interpolation excludes all primitive slope degrees below 18, while support-collapsed interpolation extracts every low-complexity explanation into a scalar-locator rational atom. Rational atoms satisfy collision rigidity and local-to-global coherence. A finite critical example shows that primitive spread cores can exist, so the remaining theorem is abundance rather than nonexistence. Canonical owner localization pays every small-owner atom, component-free owner pencils contribute at most seven rich points, core-compatible slopes contribute at most 58, a complete correction ray contributes at most $2(n - m + 1) + 963$, and proper or clone-tolerant correction spaces are paid in the explicit deployed ranges.

The manuscript does not assert the adjacent safe rows unconditionally. For MCA, the remaining inputs are the classification and payment of high-dimensional positive-dimensional correction components, the row-sharp image bound for one coherent large-owner atom—most sharply the Mersenne-31 factor-9.57219783037 prefix target—and complete routing of denominator-root, extension, clone, imperfect-quotient, and near-sunflower exceptions. For the list rows, the remaining inputs are the row-sharp high-nullity prefix and arbitrary-word interior bounds, extension payment, and one exact summed integer certificate.

CONTENTS

1. Scope and relation to CAP25 v13.2	3
2. Main theorem package	3
Part 1. Exact threshold refinements beyond CAP25 v13.2	4
3. Exact decomposition of the mutual layer	4
4. Syndrome–secant geometry and finite ray compilers	4
5. Coset-free codimension-one rank-flat recursion	10
6. Higher-codimension support-flag recursion	13
7. Minimum-distance sharpening and the exact primitive wedge	44
8. The quadratic mean-overlap staircase	47

Date: July 2026.

2020 Mathematics Subject Classification. 94B35, 11T71, 05B40.

Key words and phrases. Reed–Solomon codes, mutual correlated agreement, primitive Q, Sidon moments, generalized Hamming weights, order-32 rank, rational atoms, spread cores.

9.	Exact thresholds for explicit prize rows	50
10.	Asymptotic consequences of the imported prefix floor	51
Part 2.	Exact profile refinements and the profile envelope	52
11.	Profile envelope and the sequel compiler	52
12.	Quotient-remainder prefix normal form	53
13.	Effective-image Fourier normalization	59
Part 3.	Primitive Q and the residual compiler	62
14.	The primitive prefix leaf	62
15.	Prefix coordinates and the Vandermonde map	64
16.	The Sidon split	65
17.	Why the Sidon cell is necessary	67
18.	Fourier flatness as a Sidon payment	68
19.	External additive combinatorics	71
20.	Entropy inverse theory and its role	71
21.	Primitive Q after the Sidon cut	72
22.	Q , SP , and the MCA compiler	72
Part 4.	Finite deployment: Q, one-parameter BC, and the exact ledger	77
23.	Finite max-fiber calibration	77
Part 5.	Order-32 line rank and rational-atom classification	95
24.	Rank-theoretic setup	95
25.	The rank-two reduced intersection matrix	95
25.1.	Local relation spaces	95
26.	Why the deployed critical order is exactly 32	96
27.	What full relative rank would buy	97
28.	A necessary partition inequality	97
29.	Exact dyadic quotient descent	98
30.	Rational atoms and profile depth	99
31.	An exact degree-two rank defect over $\mathbb{F}_{17}$	100
32.	Maximal supports and the slope-degree barrier	101
33.	Local incidence complexity	102
34.	Support-collapsed rational interpolation	103
35.	The degree-two counterexample lies at the extraction threshold	105
36.	Collision pencils for two atom certificates	105
37.	Affine lines versus rational atoms	106
38.	Coherence or a quantified near-sunflower	107
39.	The proved part of the relative order-32 theorem	108
Part 6.	Primitive spread abundance and the exact MCA endgame	108
40.	A critical spread core exists	108
41.	The deployed spread-abundance target	110
42.	Rational atoms have a canonical owner	110
43.	A self-contained minimum-distance owner payment	112
44.	Owner pencils and a seven-point rich-incidence bound	113
45.	The exact large-owner atom target	114
46.	A conditional exact adjacent-row theorem	114
47.	Core interpolant and correction setup	115
48.	Only fifty-eight slopes can use the core interpolant	116
49.	An unconditional theorem for one complete correction ray	116
50.	Higher-dimensional correction spaces	118
50.1.	The proper-intersection compiler	118

50.2. A clone-tolerant basis-curve compiler	119
51. A rank-eleven dense-locator component route	120
52. The new spread residual	144
Part 7. Completion certificates and remaining row-sharp inputs	145
53. Integrated theorem status	145
54. The final MCA certificate	145
55. The final list certificate	146
56. Conclusion	147
References	147

1. SCOPE AND RELATION TO CAP25 v13.2

Throughout, CAP25 v13.2 [Cho26] is the foundational source for the deployed unsafe frontier and the normalized first-match ledger. Its theorems are invoked by citation rather than re-numbered here. This paper contains all of the additional threshold, compiler, rank, atom, and spread-incidence results needed for the current completion programme.

Imported from CAP25 v13.2	Proved in this paper
MCA conventions, universal cap, active unsafe endpoints	Exact sparsification, syndrome-secant, fixed-union, affine-span, generalized-weight, support-flag, and minimum-distance compilers
Identity-prefix, pole-line, aperiodic, subfield, quotient, and extension packages	Quadratic threshold staircases, exact profile normal forms, effective Fourier normalization, and primitive Q after the stated Sidon payment
Prefix rigidity, exact second moment, balanced-core reduction, normalized ledger	Saturation, moving-root BC, $Q \Rightarrow SP$, order-32 rank, rational-atom extraction and coherence, owner localization, and spread-core correction incidence

The paper uses the closed integer-ball convention of CAP25 v13.2. An adjacent certificate consists of the imported unsafe inequality at agreement a_0 and a proved upper inequality at $a_0 + 1$. Structural reductions, asymptotic estimates, and route-cut boundaries are not called adjacent certificates unless they fit the literal row budget.

2. MAIN THEOREM PACKAGE

Theorem 2.1 (Integrated theorem package). *The following statements hold.*

- (i) *The exact support, syndrome-secant, fixed-union, affine-span, generalized-weight, common-zero, codimension-one, support-flag, minimum-distance, puncturing, directional-Johnson, paving, mean-overlap, packing, and explicit prime-row results of Part I are unconditional.*
- (ii) *The quotient-remainder normal form, partial-occupancy disintegration, effective-image Fourier identities, and profile add-back of Part II are unconditional. Primitive Q follows from the Sidon-moment payment of Part III, and the shallow Fourier theorem proves that payment in its displayed range.*
- (iii) *Saturation, the one-parameter moving-root theorem, and $Q \Rightarrow SP$ are unconditional. The finite compilers pay the displayed low-nullity, low-affine-span, common-zero, support-flag, and direction-defect ranges.*
- (iv) *The rank-two kernel dictionary, critical order 32, global affine-block bound, dyadic pair descent, rational-profile converse, slope-degree barrier, low-complexity atom extraction, collision rigidity, atom coherence, owner localization, owner-pencil incidence bound, core-compatible abundance bound, complete correction-ray bound, and the stated proper and clone-tolerant correction-space bounds are unconditional.*

- (v) *The two deployed MCA adjacent rows follow from the three explicit inputs in [Theorem 46.2](#); the two list adjacent rows follow from the exact summed certificate in [Theorem 55.2](#). Those remaining inputs are not asserted in this paper.*

Part 1. Exact threshold refinements beyond CAP25 v13.2

3. EXACT DECOMPOSITION OF THE MUTUAL LAYER

Let $C \subseteq \mathbb{F}^D$ be linear, let a be an agreement threshold, and put $t = |D| - a$. We use the foundation's definitions of support-wise MCA-badness and of the challenge-restricted numerator $B_{C,\Gamma}^{\text{MCA}}(a)$ [[Cho26](#), Secs. 1–2].

A received pair is *column-far at agreement a* when no support of size at least a simultaneously explains its two coordinates. Let $B_{C,\Gamma}^{\text{CA}}(a)$ be the maximum number of challenge slopes whose affine combination is explained on some support of size at least a , maximized over column-far pairs. Define the sparse mutual numerator

$$\sigma_{C,\Gamma}(a) = \max \# \{ \gamma \in \Gamma : \gamma \text{ is MCA-bad for } (e_0, e_1) \},$$

where the maximum is over pairs with $|\text{supp}(e_0) \cup \text{supp}(e_1)| \leq t$.

Theorem 3.1 (Exact sparsification of the mutual layer). *For every linear code, every nonempty challenge set Γ , and every a ,*

$$(3.1) \quad B_{C,\Gamma}^{\text{MCA}}(a) = \max \{ B_{C,\Gamma}^{\text{CA}}(a), \sigma_{C,\Gamma}(a) \}.$$

Proof. Both terms on the right maximize over subclasses of received pairs, so the left side dominates them. Conversely fix a pair. If it is column-far, its MCA-bad and CA-bad challenge slopes coincide. Otherwise translate it by a codeword pair that explains it on one common support of size at least a . The translated pair is supported on at most t coordinates, and translation preserves, slope by slope and support by support, both affine explanation and simultaneous explanation. Its bad set is therefore bounded by $\sigma_{C,\Gamma}(a)$. $\square$

The foundation supplies the exact-agreement reduction, support uniqueness, tangent floor, and deep-regime theorem used below [[Cho26](#), Secs. 13–14]. We invoke those interfaces only by citation.

4. SYNDROME-SECANT GEOMETRY AND FINITE RAY COMPILERS

The *syndrome* of a word is the result of applying a parity-check matrix; it vanishes exactly on codewords. Thus syndrome space forgets the part of a received word already lying in the code and retains only its error class. This description separates possible error supports from distinct line parameters. Put $R = n - k$, the redundancy (and syndrome-space dimension), and choose the generalized Vandermonde parity-check matrix H with columns

$$(4.1) \quad h_x = \lambda_x(1, x, \dots, x^{R-1})^\top, \quad \lambda_x = \left(\prod_{y \in D \setminus \{x\}} (x - y) \right)^{-1}.$$

Lagrange interpolation gives $\sum_{x \in D} \lambda_x x^j = 0$ for $0 \leq j \leq n - 2$, so this matrix annihilates every degree-less-than- k evaluation. Every set of at most R columns is independent. For $E \subseteq D$, write $V_E = \text{Span}\{h_x : x \in E\}$; it is exactly the set of syndromes of words supported on E . We write $s(u) = Hu^\top$ for the syndrome of u .

Proposition 4.1 (Syndrome-line normal form). *Let $S = D \setminus E$. A word u is explained on S if and only if $s(u) \in V_E$. Consequently γ is MCA-bad on S for a pair (u_0, u_1) if and only if*

$$(4.2) \quad s(u_0) + \gamma s(u_1) \in V_E, \quad \{s(u_0), s(u_1)\} \not\subseteq V_E.$$

For fixed E , there is at most one such finite slope.

Proof. If $u = c + e$, where $c \in C$ and e is supported on E , then $s(u) = s(e) \in V_E$. Conversely, if $s(u) = \sum_{x \in E} e_x h_x$, subtract the word supported on E with values e_x ; the result has zero syndrome and belongs to C . Apply this equivalence to the line point and to both coordinates. Passing (4.2) to $\mathbb{F}^R/V_E$ shows that two distinct solutions would put both projected syndromes at zero, contradicting badness. $\square$

Since $\lambda_x \neq 0$, one has $[h_x] = [1 : x : \dots : x^{R-1}]$; these are the points of the rational normal curve in projective syndrome space, while the weights λ_x merely choose vector representatives. A *secant span* is the linear span V_E of a selected set of those points. We call the meeting of a syndrome line with V_E *transverse* when the entire line is not contained in V_E , equivalently when its two generators are not both in V_E . The next result is a *compiler* in the sense that it translates the bad-slope problem, exactly and in both directions, into this line–secant incidence problem. A *ray* means a one-dimensional direction in syndrome or coefficient space; different raw witnesses may represent the same ray and hence the same slope.

Theorem 4.2 (Exact syndrome–secant compiler). *Put $t = n - a$, and for a syndrome line $\ell_{y_0, y_1} = \{y_0 + \gamma y_1 : \gamma \in \mathbb{F}\}$ define*

$$(4.3) \quad \Theta_t(y_0, y_1) = |\{\gamma \in \mathbb{F} : \exists E \subseteq D, |E| \leq t, y_0 + \gamma y_1 \in V_E, \{y_0, y_1\} \not\subseteq V_E\}|.$$

Then

$$(4.4) \quad B_C^{\text{MCA}}(a) = \max_{y_0, y_1 \in \mathbb{F}^{n-k}} \Theta_t(y_0, y_1).$$

For each fixed E , the affine line has at most one transverse intersection parameter with V_E . For a Reed–Solomon code the spaces V_E are the spans of at most t points of the weighted rational normal curve in (4.1). Thus the higher-dimensional ray problem is exactly a transverse line–secant incidence problem, with repeated intersection parameters deduplicated.

Proof. Apply Proposition 4.1 with $S = D \setminus E$, and take the union over $|E| \leq t$. The syndrome map is surjective, so maximizing over received pairs is the same as maximizing over (y_0, y_1) . In the quotient $\mathbb{F}^{n-k}/V_E$, a transverse intersection solves $\bar{y}_0 + \gamma \bar{y}_1 = 0$ with the two projected vectors not both zero, and therefore has at most one finite solution. The last assertion is the Vandermonde description of the parity columns. $\square$

A *parity-column circuit* is a minimally linearly dependent set of columns of H . The next lemma says that many distinct transverse rays cannot all be supported inside an independent union of columns.

Lemma 4.3 (Independent-union ray bound). *Let every set of at most R parity columns be independent, put $t < R$, and fix a syndrome line. For each slope in a transverse set Z , choose an error set E_γ , $|E_\gamma| \leq t$, witnessing (4.3). If $U = \bigcup_{\gamma \in Z} E_\gamma$ has size at most R , then $|Z| \leq t + 1$. Consequently, more than $t + 1$ transverse slopes force the witness union to contain a parity-column circuit.*

Proof. If $|Z| \leq 1$, there is nothing to prove; assume that Z contains two distinct slopes. Two distinct line points lie in V_U , hence $y_0, y_1 \in V_U$. Column independence gives unique coefficient lifts $e_0, e_1 \in \mathbb{F}^U$, and every chosen lift is $e_\gamma = e_0 + \gamma e_1$. Let U' be the set of s coordinates where the pair (e_0, e_1) is nonzero. If $s \leq t$ and E_γ contained all of U' , then $y_0, y_1 \in V_{U'} \subseteq V_{E_\gamma}$, contradicting the transversality of this particular witness. Hence some $x \in U' \setminus E_\gamma$ exists, and uniqueness of the lift gives $e_0(x) + \gamma e_1(x) = 0$. Each nonzero affine coordinate function has at most one root, so $|Z| \leq s \leq t$. If $s > t$, every chosen lift of weight at most t has at least $s - t$ zero coordinates. Double counting the pairs (γ, x) with $e_0(x) + \gamma e_1(x) = 0$ gives $|Z|(s - t) \leq s$, whence $|Z| \leq \lfloor s/(s - t) \rfloor \leq t + 1$. If $|U| \leq R$, independence holds; the contrapositive proves the circuit assertion. $\square$

Theorem 4.4 (Bounded residual-kernel ray compiler). *Let H be a parity-check matrix over $\mathbb{F}$, let U be a set of columns, and put*

$$\kappa(U) = \dim_{\mathbb{F}} \ker(H_U : \mathbb{F}^U \longrightarrow \mathbb{F}^{n-k}).$$

Fix a syndrome line $y_0 + \gamma y_1$, an integer $t \geq 0$, and a set Z of finite transverse slopes. Suppose that for every $\gamma \in Z$ there is $E_\gamma \subseteq U$, $|E_\gamma| \leq t$, such that

$$y_0 + \gamma y_1 \in V_{E_\gamma}, \quad \{y_0, y_1\} \not\subseteq V_{E_\gamma}.$$

Then

$$(RC_{\ker}) \quad |Z| \leq (t+1)|\mathbb{F}|^{\kappa(U)}.$$

For a Reed–Solomon parity check of redundancy $R = n - k$, $\kappa(U) = \max\{0, |U| - R\}$. Hence a certified residual chart with $(|U| - R)_+ \log|\mathbb{F}| = o(n)$ has a direct subexponential ray payment with the displayed exact finite loss.

Proof. The assertion is immediate when $|Z| \leq 1$. Otherwise two line points belong to V_U , so $y_0, y_1 \in V_U$. Choose lifts $b_0, b_1 \in \mathbb{F}^U$ with $H_U b_i = y_i$. For every γ , choose a lift c_γ supported in E_γ and put $z_\gamma = c_\gamma - b_0 - \gamma b_1 \in \ker H_U$.

Fix $z \in \ker H_U$, and consider the slopes for which $z_\gamma = z$. Put

$$U_z = \{x \in U : (b_0(x) + z(x), b_1(x)) \neq (0, 0)\}, \quad s = |U_z|.$$

Every active coordinate is a nonzero affine function of γ , and therefore vanishes for at most one slope. Hence the total number of active zero incidences in this kernel class is at most s . If $s > t$, the weight bound on c_γ forces at least $s - t$ active zeros, so

$$|Z_z|(s - t) \leq s, \quad |Z_z| \leq \lfloor s/(s - t) \rfloor \leq t + 1.$$

If $s \leq t$, transversality forces at least one active zero: without one, $U_z = \text{supp } c_\gamma \subseteq E_\gamma$, while both $b_0 + z$ and b_1 are supported in U_z , putting y_0, y_1 in V_{E_γ} . Thus this kernel class contains at most $s \leq t$ slopes. Summing over the $|\mathbb{F}|^{\kappa(U)}$ kernel elements proves the bound. The MDS column property gives the Reed–Solomon formula for $\kappa(U)$. $\square$

Theorem 4.5 (Multiplicity-amplified fixed-union ray compiler). *Let H be a Reed–Solomon parity check of redundancy R , let $U \subseteq D$ have size $N = R + \nu$ with $\nu \geq 1$, and put $C_U = \ker H_U$. Fix $t < R$, a syndrome line $y_0 + \gamma y_1$, and a set Z of distinct finite transverse slopes. Suppose that for every $\gamma \in Z$ there is $e_\gamma \in \mathbb{F}^U$ such that*

$$H_U e_\gamma = y_0 + \gamma y_1, \quad \text{wt}(e_\gamma) \leq t, \quad \{y_0, y_1\} \not\subseteq V_{\text{supp } e_\gamma}.$$

Then

$$(4.5) \quad |Z| \leq \left\lfloor \frac{\binom{R+\nu}{\nu+1}}{R-t} \right\rfloor.$$

The estimate is independent of $|\mathbb{F}|$; it replaces the factor $|\mathbb{F}|^\nu$ in [Theorem 4.4](#) by an exact incidence multiplicity.

Proof. There is nothing to prove when $|Z| \leq 1$. Otherwise two line points lie in V_U , hence $y_0, y_1 \in V_U$. Choose lifts $b_0, b_1 \in \mathbb{F}^U$, and choose a basis of the ν -dimensional kernel C_U . Write its generator matrix as K , with row K_x at $x \in U$. Every chosen error has a unique expression

$$e_\gamma = b_0 + \gamma b_1 + K c_\gamma, \quad c_\gamma \in \mathbb{F}^\nu.$$

Put $p_\gamma = (\gamma, c_\gamma) \in \mathbb{F}^{\nu+1}$. A zero coordinate of e_γ is the affine hyperplane incidence

$$(4.6) \quad b_0(x) + \gamma b_1(x) + K_x c_\gamma = 0, \quad x \in U,$$

whose normal is $n_x = (b_1(x), K_x) \in \mathbb{F}^{\nu+1}$. Since $\text{wt}(e_\gamma) \leq t$, the point p_γ lies on at least $h = N - t$ of these hyperplanes.

We claim that the normals incident with p_γ span $\mathbb{F}^{\nu+1}$. Otherwise a nonzero pair (δ, v) is orthogonal to all of them. Then

$$q = \delta b_1 + K v$$

vanishes at every zero coordinate of e_γ , so $\text{supp } q \subseteq \text{supp } e_\gamma$ and $\text{wt}(q) \leq t$. If $\delta = 0$, then $q = K v \neq 0$ is a kernel word of weight at most $t < R + 1$, contradicting the MDS distance of $C_U = [R + \nu, \nu, R + 1]$. If $\delta \neq 0$, divide by δ . The word q/δ represents y_1 , and $e_\gamma - \gamma q/\delta$ represents y_0 ; both are supported inside $\text{supp } e_\gamma$, contradicting transversality.

Every normal n_x is nonzero. Indeed, if the coordinate functional of C_U vanished at x , then puncturing at x would leave kernel dimension ν , whereas the MDS rank of $H_{U \setminus \{x\}}$ makes that dimension $\nu - 1$. Thus the incident normals form a loopless rank- $(\nu + 1)$ matroid on at least h elements. Fix one basis. Each of the other incident elements enters a basis by a fundamental-circuit exchange, producing a distinct basis. Hence p_γ is incident with at least

$$h - (\nu + 1) + 1 = R - t$$

independent $(\nu + 1)$ -subsets of hyperplanes. An independent subset has at most one affine intersection point. Counting pairs consisting of a slope point and one of its incident bases gives

$$|Z|(R - t) \leq \binom{R + \nu}{\nu + 1},$$

which is (4.5). $\square$

Corollary 4.6 (Single MDS-circuit ray compiler, multiplicity form). *Under the hypotheses of the former single-circuit chart, namely $|U| = R + 1$,*

$$(4.7) \quad |Z| \leq \left\lfloor \frac{R(R + 1)}{2(R - t)} \right\rfloor.$$

Thus the pair-charge $\binom{R+1}{2}$ is divided by the exact number $R - t$ of independent zero-coordinate pairs forced by every transverse slope.

Proof. Apply Theorem 4.5 with $\nu = 1$. $\square$

Theorem 4.7 (Fixed-syndrome GRS–Johnson compiler). *Let H be a Reed–Solomon parity check of redundancy R , let $U \subseteq D$ have size $N = R + \nu$, and fix a syndrome y . Put*

$$\mathcal{L}_U(y, t) = \{e \in \mathbb{F}^U : H_U e = y, \text{ wt}(e) \leq t\}, \quad h = N - t,$$

where $t < R$. If $\nu = 0$, then $|\mathcal{L}_U(y, t)| \leq 1$. If $\nu \geq 1$ and

$$(4.8) \quad h^2 > N(\nu - 1),$$

then

$$(4.9) \quad |\mathcal{L}_U(y, t)| \leq \left\lfloor \frac{N(h - \nu + 1)}{h^2 - N(\nu - 1)} \right\rfloor.$$

Equivalently, the same bound holds for the number of codewords in one Reed–Solomon list whose error supports are all contained in the fixed union U .

Proof. For $\nu = 0$, the columns in U are independent, so a syndrome has at most one lift. Assume $\nu \geq 1$, choose one lift b of y , and put $C_U = \ker H_U$. This is an $[N, \nu, R + 1]$ generalized Reed–Solomon code. The set $\mathcal{L}_U(y, t)$ is the low-weight part of the affine coset $b + C_U$. Thus its members are in bijection with codewords of C_U that agree with the fixed word $-b$ on at least $h = N - t$ coordinates.

For each such codeword choose an h -subset A_i of its agreement set. Distinct codewords of the $[N, \nu]$ GRS code agree on at most $\nu - 1$ coordinates, so $|A_i \cap A_j| \leq \nu - 1$. If there are L sets, and d_x is the number containing x , then

$$\sum_x d_x = Lh, \quad \sum_x \binom{d_x}{2} \leq \binom{L}{2}(\nu - 1).$$

Cauchy–Schwarz gives

$$\sum_x \binom{d_x}{2} \geq \frac{1}{2} \left(\frac{L^2 h^2}{N} - Lh \right).$$

Combining the two inequalities and using (4.8) yields

$$L(h^2 - N(\nu - 1)) \leq N(h - \nu + 1),$$

which proves (4.9). $\square$

Theorem 4.8 (Recursive affine-span Reed–Solomon list compiler). *Let $C = \text{RS}[\mathbb{F}, D, K]$ have length n , let $u \in \mathbb{F}^D$, and put $m = K + w \leq n$. Let $\mathcal{A} = c_0 + C' \subseteq C$ be an affine subspace of codewords whose direction space has dimension $s \geq 0$. Then*

$$(4.10) \quad |\mathcal{L}_{\mathcal{A}}(u, m)| \leq \left\lfloor \frac{(n - K + s)^s}{(w + 1)^s} \right\rfloor = \left\lfloor \frac{\binom{n-K+s}{s}}{\binom{w+s}{s}} \right\rfloor.$$

The empty products at $s = 0$ are one. In particular, a one-dimensional codeword pencil contains at most $\lfloor (n - K + 1)/(w + 1) \rfloor$ members of the list.

Proof. We induct on s . The case $s = 0$ is immediate. Assume $s \geq 1$, and let

$$G = \{x \in D : c(x) = 0 \text{ for every } c \in C'\}, \quad z = |G|.$$

The polynomial space representing C' has dimension s . Division by the locator of G embeds it into the polynomials of degree less than $K - z$, so

$$(4.11) \quad z \leq K - s.$$

Let $g \leq z$ be the number of coordinates in G at which the common value of the affine flat agrees with u . Every member of $\mathcal{L}_{\mathcal{A}}(u, m)$ therefore has at least $m - g$ agreements outside G .

For $x \in D \setminus G$, the evaluation functional on C' is nonzero. Hence

$$\mathcal{A}_x = \{c \in \mathcal{A} : c(x) = u(x)\}$$

is either empty or an affine subspace of dimension $s - 1$. By the induction hypothesis, the number of listed codewords in $\mathcal{A}_x$ is at most

$$B_{s-1} := \frac{\binom{n-K+s-1}{s-1}}{\binom{w+s-1}{s-1}}.$$

Double-counting pairs (c, x) with $c \in \mathcal{L}_{\mathcal{A}}(u, m)$, $x \notin G$, and $c(x) = u(x)$ gives

$$(4.12) \quad |\mathcal{L}_{\mathcal{A}}(u, m)|(m - g) \leq (n - z)B_{s-1}.$$

Put $b = z - g$ and $q = K - s - z$. Both are nonnegative by (4.11), and the exact identity

$$(4.13) \quad \begin{aligned} (m - g)(n - K + s) - (n - z)(w + s) \\ = b(n - K + s) + q(n - m) \geq 0 \end{aligned}$$

shows that

$$\frac{n - z}{m - g} \leq \frac{n - K + s}{w + s}.$$

Combining this with (4.12) yields

$$|\mathcal{L}_{\mathcal{A}}(u, m)| \leq B_{s-1} \frac{n - K + s}{w + s} = \frac{\binom{n-K+s}{s}}{\binom{w+s}{s}}.$$

Taking the integer floor proves (4.10). $\square$

Theorem 4.9 (Generalized-weight rank-flat list compiler). *Retain the hypotheses of Theorem 4.8, assume $s \geq 1$, and let*

$$d_j = d_j(C') := \min_{\substack{V \subseteq C' \\ \dim V = j}} |\text{supp}(V)|, \quad 1 \leq j \leq s,$$

be the generalized Hamming weights of the direction code. Put

$$R = n - K, \quad t = n - m,$$

let G be the common-zero set of C' , let $z = |G| = n - d_s$, let g be the number of common agreements with u on G , and put $b = z - g$. Then

$$(4.14) \quad |\mathcal{L}_{\mathcal{A}}(u, m)| \leq \left\lfloor \frac{d_s^s}{\prod_{j=1}^s (d_j - t + b)} \right\rfloor.$$

Every factor in the denominator is positive. More precisely,

$$(4.15) \quad d_j \geq R + j, \quad d_j - t + b \geq w + j + b.$$

Proof. Choose a basis of C' and identify the affine flat with $\mathbb{F}^s$. For each active coordinate $x \in D \setminus G$, agreement with u is an affine hyperplane whose nonzero normal is the evaluation functional $\ell_x \in (C')^*$. Every listed point is incident with at least

$$h = m - g = d_s - t + b$$

of these hyperplanes.

Suppose that r independent incident normals have already been chosen and span $W \leq (C')^*$. The coordinates whose normals lie in W are common zeros of the $(s - r)$ -dimensional subcode $W^\perp$. This subcode has support at least d_{s-r} ; after the common-zero set G is removed, at most

$$(n - d_{s-r}) - z = d_s - d_{s-r}$$

active normals lie in W . Hence there are at least

$$h - (d_s - d_{s-r}) = d_{s-r} - t + b$$

choices for the next independent incident normal. Each listed point is therefore incident with at least

$$\prod_{j=1}^s (d_j - t + b)$$

ordered bases of coordinate hyperplanes. There are at most d_s^s ordered coordinate bases globally, and an independent affine-hyperplane basis has at most one common point. Double counting proves (4.14).

Finally, if $V \leq C'$ has dimension j and common-zero set of size z_V , division by its locator embeds the representing polynomial space into the polynomials of degree less than $K - z_V$. Thus $j \leq K - z_V$, so $|\text{supp}(V)| = n - z_V \geq n - K + j = R + j$. This proves (4.15) and positivity. $\square$

Corollary 4.10 (Over-budget rank-flat concentration). *In Theorem 4.9, let $s \geq 2$ and let B be a nonnegative integer. Define*

$$(4.16) \quad \Theta_{s,B}(d_s, b) = \max \left\{ q \in \mathbb{Z}_{\geq 0} : q^{s-1} \leq \left\lfloor \frac{d_s^s}{(B+1)(d_s - t + b)} \right\rfloor \right\}.$$

If $|\mathcal{L}_{\mathcal{A}}(u, m)| > B$, then some proper direction subcode has dimension $1 \leq j < s$ and

$$(4.17) \quad d_j \leq t - b + \Theta_{s,B}(d_s, b).$$

A row-independent coarse version is obtained by replacing d_s^s by n^s and $d_s - t + b$ by $w + s$.

Proof. If the list contains at least $B + 1$ points, (4.14) implies

$$\prod_{j=1}^{s-1} (d_j - t + b) \leq \left\lfloor \frac{d_s^s}{(B+1)(d_s - t + b)} \right\rfloor.$$

At least one of the $s - 1$ integer factors is at most the integer root in (4.16), proving (4.17). The coarse form follows from $d_s \leq n$ and (4.15). $\square$

Theorem 4.11 (Common-zero Johnson list compiler). *In the notation of Theorem 4.9, put*

$$N = d_s, \quad h = m - g = d_s - t + b, \quad \alpha = d_s - d_1.$$

If

$$(4.18) \quad h^2 > N\alpha,$$

then

$$(4.19) \quad |\mathcal{L}_{\mathcal{A}}(u, m)| \leq \left\lfloor \frac{N(h - \alpha)}{h^2 - N\alpha} \right\rfloor.$$

Equivalently, write $z = K - \nu$. Then $N = R + \nu$, $h = w + \nu + b$, and the generalized Singleton bound gives the universal specialization

$$(4.20) \quad |\mathcal{L}_{\mathcal{A}}(u, m)| \leq \left\lfloor \frac{(R + \nu)(w + b + 1)}{(w + \nu + b)^2 - (R + \nu)(\nu - 1)} \right\rfloor$$

whenever the displayed denominator is positive.

Proof. All codewords of the affine flat have the same values on G . Delete those coordinates and, for each listed codeword, choose an h -subset of its active agreement set. If two distinct codewords are chosen, their difference is a nonzero word of C' , has weight at least d_1 , and therefore vanishes on at most $N - d_1 = \alpha$ active coordinates. The standard Johnson incidence count now gives

$$L(h^2 - N\alpha) \leq N(h - \alpha),$$

which proves (4.19).

If $z = K - \nu$, then $N = n - z = R + \nu$ and $h = m - (z - b) = w + \nu + b$. By (4.15), $d_1 \geq R + 1$, so $\alpha \leq \nu - 1$. The Johnson fraction is nondecreasing in α while its denominator is positive. Replacing α by $\nu - 1$ gives (4.20). $\square$

5. COSET-FREE CODIMENSION-ONE RANK-FLAT RECURSION

The rank-flat concentration theorem produces a low-support direction subcode, but a naive partition into its affine cosets introduces the field size. The next theorem removes that loss when the support closure has codimension one. It counts the whole affine chart at once by two disjoint resources: bases using an agreement in the new support layer, and bases whose final coordinate lies in the old support.

Theorem 5.1 (Two-resource codimension-one recursion). *Let*

$$C = \text{RS}[\mathbb{F}, D, K], \quad n = |D|, \quad m = K + w, \quad t = n - m,$$

and let $\mathcal{A} = c_0 + C' \subseteq C$ be an affine codeword flat with $\dim C' = j + 1$, where $j \geq 1$. Let $V < C'$ be a hyperplane, put

$$U = \text{supp}(V), \quad S = \text{supp}(C'), \quad T = S \setminus U, \quad H = D \setminus S,$$

and assume $e := |T| \geq 1$. Write $d = |U|$. Then

$$(5.1) \quad V = \{v \in C' : \text{supp}(v) \subseteq U\};$$

in particular every word of $C' \setminus V$ is nonzero at every coordinate of T .

Fix a received word $u \in \mathbb{F}^D$. Let b_0 be the number of coordinates in H on which the common value of $\mathcal{A}$ disagrees with u . For $0 \leq b \leq e$, let L_b count the members $c \in \mathcal{A}$ satisfying

$$\text{agr}(u, c) \geq m, \quad |\{x \in T : c(x) \neq u(x)\}| = b.$$

Let $f_i = d_i(V)$, $1 \leq i \leq j$, be the generalized Hamming weights of V , and put

$$(5.2) \quad P_b = \prod_{i=1}^j (f_i - t + b_0 + b), \quad Q = w + 1 + b_0.$$

Then every factor in P_b is positive and

$$(5.3) \quad \sum_{b=0}^e L_b(e - b) P_b \leq e d^j,$$

$$(5.4) \quad \sum_{b=0}^e L_b(Q - e + b)_+ P_b \leq d^{j+1}.$$

Moreover, with $L = \sum_b L_b$,

$$(5.5) \quad L \leq \left\lfloor \frac{(n - j)d^j}{QP_0} \right\rfloor.$$

More generally, if $x, y \geq 0$ satisfy

$$(5.6) \quad x(e - b)P_b + y(Q - e + b)_+ P_b \geq 1 \quad (0 \leq b \leq e),$$

then

$$(5.7) \quad L \leq \left\lfloor xe d^j + y d^{j+1} \right\rfloor.$$

Thus the affine cosets of V are never summed separately.

Proof. For $x \in T$, every word of V vanishes at x , whereas some word of C' is nonzero there. Evaluation at x therefore induces a nonzero functional on the one-dimensional quotient C'/V . It is nonzero on every nonzero quotient class, proving (5.1).

Parameterize $\mathcal{A}$ by the direction space C' . Agreement at a coordinate x is an affine hyperplane with normal $\ell_x \in (C')^*$. A listed point counted by L_b has at least

$$m - (|H| - b_0) - (e - b) = d - t + b_0 + b$$

agreements in U . Restrict the corresponding normals to V . After r independent restricted normals have been selected, their span annihilates a $(j - r)$ -dimensional subcode of V , whose support has size at least f_{j-r} . Hence at most $d - f_{j-r}$ normals of U lie in the old span, and there are at least

$$f_{j-r} - t + b_0 + b$$

choices for the next independent normal. Every listed point counted by L_b is therefore incident with at least P_b ordered V -normal bases. The generalized Singleton inequalities $f_i \geq n - K + i$ show at the same time that all factors in P_b are at least $w + i + b_0 + b$.

The span in $(C')^*$ of a lifted V -normal basis meets the one-dimensional annihilator $V^\perp$ trivially. Every agreeing coordinate of T has a nonzero normal in $V^\perp$, and therefore extends the basis. A point in class b supplies at least $(e - b)P_b$ ordered bases of this type. Globally there are at most $e d^j$ ordered choices, and an independent affine hyperplane basis cuts out at most one point. This proves (5.3).

Fix now one lifted V -normal basis and let $h \in C' \setminus V$ generate its one-dimensional annihilator. A coordinate normal extends the basis exactly when h is nonzero at that coordinate. The nonzero Reed–Solomon word h has at most $K - 1$ zeros. All coordinates of H are zeros of h , and exactly $|H| - b_0$ of them are agreements. Thus among the at least m agreement coordinates of the listed point, at most $K - 1 - b_0$ fail to extend the basis. There are consequently at least

$$m - (K - 1 - b_0) = Q$$

extending agreement coordinates. All $e - b$ agreeing coordinates of T already extend, so at least $(Q - e + b)_+$ extending agreements lie in U . Double-counting the resulting ordered $(j + 1)$ -tuples inside U proves (5.4).

The same argument without separating U from T gives at least $QP_b \geq QP_0$ full bases through every listed point. There are at most $(n - j)d^j$ ordered choices of a V -basis and one additional coordinate, proving (5.5). Finally multiply (5.3) by x , multiply (5.4) by y , and use (5.6); this gives (5.7). $\square$

Lemma 5.2 (Profile interpolation). *Let*

$$P(z) = \prod_{i=1}^j (a_i + z), \quad a_i > 0,$$

and let $e, Q > 0$. For every integer $0 \leq b \leq e$,

$$(5.8) \quad \frac{(e - b)P(b)}{eP(0)} + \frac{(Q - e + b)_+ P(b)}{QP(e)} \geq 1$$

under either of the following conditions:

- (a) $e \leq Q$;
- (b) $e > Q$ and

$$(5.9) \quad \left(\frac{P(e)}{P(0)} \right)^{1-Q/e} \geq \frac{e}{Q}.$$

If $j \geq 2$, $A = \max_i a_i$, and $(j - 1)Q \geq A$, then (5.9) holds for every $e \geq Q$.

Proof. Put $r = P(e)/P(0) \geq 1$. Since $\log P$ is concave, with $\theta = b/e$ one has

$$\frac{P(b)}{P(0)} \geq r^\theta, \quad \frac{P(b)}{P(e)} \geq r^{-(1-\theta)}.$$

If $e \leq Q$, then $(Q - e + b)/Q \geq \theta$, and weighted AM–GM gives

$$(1 - \theta)r^\theta + \theta r^{-(1-\theta)} \geq 1.$$

This proves (5.8) in case (a).

Assume $e > Q$, put $c = Q/e$, and first take $0 \leq \theta \leq 1 - c$. The second summand in (5.8) vanishes, while concavity gives the lower bound $(1 - \theta)r^\theta$. The logarithm of this expression is concave in θ ; it is zero at 0, and at $1 - c$ it is nonnegative exactly when (5.9) holds.

For $1 - c \leq \theta \leq 1$, write $\alpha = 1 - \theta \in [0, c]$. After multiplication by r^α , the lower bound supplied by log-concavity is

$$\alpha r + 1 - \frac{\alpha}{c}.$$

Convexity of r^α on $[0, c]$ gives

$$r^\alpha \leq 1 + \frac{\alpha}{c}(r^c - 1).$$

The former expression dominates the latter precisely when $cr \geq r^c$, which is again (5.9). This proves case (b).

Finally, if $A = \max_i a_i$, then

$$r \geq (1 + e/A)^j.$$

For $e \geq Q$, set

$$\phi(e) = j(1 - Q/e) \log(1 + e/A) - \log(e/Q).$$

One has $\phi(Q) = 0$, and the elementary inequality $\log(1 + x) \geq x/(1 + x)$ gives

$$\phi'(e) \geq \frac{j}{A + e} - \frac{1}{e} = \frac{(j - 1)e - A}{e(A + e)} \geq 0.$$

Thus $\phi(e) \geq 0$, which is (5.9). □

Corollary 5.3 (MDS-soft codimension-one compiler). *Retain the hypotheses of Theorem 5.1, and put*

$$(5.10) \quad \Pi_b = (d - t + b_0 + b) \prod_{i=1}^{j-1} (w + i + b_0 + b).$$

Suppose the profile interpolation inequality (5.8) holds for $P = \Pi$, with $Q = w + 1 + b_0$. Then

$$(5.11) \quad L \leq \left\lfloor \frac{d^j}{\Pi_0} + \frac{d^{j+1}}{Q\Pi_e} \right\rfloor.$$

Independently of profile interpolation,

$$(5.12) \quad L \leq \left\lfloor \frac{(n - j)d^j}{Q\Pi_0} \right\rfloor.$$

In particular, (5.11) is a bound for the whole affine chart, not a per-coset estimate.

Proof. The generalized Singleton inequalities give $f_i \geq n - K + i$ for $i < j$, while $f_j = d$. Hence $P_b \geq \Pi_b$. The two resource inequalities of Theorem 5.1 remain true after replacing P_b by Π_b . Under profile interpolation, the choices

$$x = \frac{1}{e\Pi_0}, \quad y = \frac{1}{Q\Pi_e}$$

are feasible in (5.6), and (5.7) gives (5.11). The all-extender estimate follows directly from (5.5). □

6. HIGHER-CODIMENSION SUPPORT-FLAG RECURSION

The codimension-one theorem removes the field-size loss for a single support-saturated hyperplane. The same principle extends to arbitrary codimension, but one must keep all possible choices of “new-layer” and “old-support” extenders. The resulting object is a finite exact linear program with one resource inequality for each binary flag pattern. In particular, no affine coset is ever counted separately.

Definition 6.1 (Support saturation). *If $V \leq W \leq C$, the support closure of V in W is*

$$\overline{V}^W = \{c \in W : \text{supp}(c) \subseteq \text{supp}(V)\}.$$

We call V support-saturated in W when $\overline{V}^W = V$.

Lemma 6.2 (Generalized-weight minimizers are support-saturated). *Let W be a finite-length linear code, and let $V \leq W$ have dimension j and support size $d_j(W)$. Then V is support-saturated in W . More generally, the generalized Hamming weights satisfy*

$$d_1(W) < d_2(W) < \cdots < d_{\dim W}(W).$$

Proof. Let $Y \leq W$ have dimension $i + 1$ and support size $d_{i+1}(W)$. Choose a coordinate on which some word of Y is nonzero. The kernel of evaluation at that coordinate is an i -dimensional subcode whose support omits that coordinate, so

$$d_i(W) \leq d_{i+1}(W) - 1.$$

This proves strict increase. If the support closure of a j -dimensional minimizer V had dimension $k > j$, then

$$d_k(W) \leq |\text{supp}(V)| = d_j(W),$$

contradicting strict increase. $\square$

Lemma 6.3 (Existence of a support-saturated flag). *Let V be support-saturated in W and put $r = \dim W - \dim V$. There is a flag*

$$V = V_0 < V_1 < \cdots < V_r = W, \quad \dim V_\ell = \dim V + \ell,$$

such that $V_{\ell-1}$ is support-saturated in V_ℓ for every $1 \leq \ell \leq r$. Consequently, with

$$U_\ell = \text{supp}(V_\ell), \quad T_\ell = U_\ell \setminus U_{\ell-1},$$

every layer T_ℓ is nonempty.

Proof. Construct the flag downward. Suppose $V < X \leq W$. Since V is support-saturated in W , some word of $X \setminus V$ is nonzero at a coordinate $x \notin \text{supp}(V)$. Put

$$X' = \{c \in X : c(x) = 0\}.$$

Then $V \leq X'$, X' is a hyperplane of X , and $x \notin \text{supp}(X')$. If $c \in X$ has support contained in $\text{supp}(X')$, then $c(x) = 0$, hence $c \in X'$. Thus X' is support-saturated in X . Repeating until the dimension reaches $\dim V$ gives the claimed flag; the terminal subspace contains V and has the same dimension, so it is V . A proper support-saturated inclusion cannot have equal supports, proving $T_\ell \neq \emptyset$. $\square$

Theorem 6.4 (Exact support-flag resource recursion). *Let*

$$C = \text{RS}[\mathbb{F}, D, K], \quad n = |D|, \quad m = K + w, \quad t = n - m,$$

and let $\mathcal{A} = c_0 + C' \subseteq C$ be an affine codeword flat. Assume that

$$V_0 < V_1 < \cdots < V_r = C'$$

is a support-saturated flag with

$$\dim V_0 = j \geq 1, \quad \dim V_\ell = j + \ell.$$

Put

$$U_\ell = \text{supp}(V_\ell), \quad d_\ell = |U_\ell|, \quad T_\ell = U_\ell \setminus U_{\ell-1}, \quad e_\ell = |T_\ell|,$$

and $H = D \setminus U_r$. Fix a received word $u \in \mathbb{F}^D$, and let b_0 be the number of coordinates of H on which the common value of $\mathcal{A}$ disagrees with u .

For a profile $\mathbf{b} = (b_1, \dots, b_r)$ with $0 \leq b_\ell \leq e_\ell$, let $L_{\mathbf{b}}$ count the words $c \in \mathcal{A}$ satisfying

$$\text{agr}(u, c) \geq m, \quad |\{x \in T_\ell : c(x) \neq u(x)\}| = b_\ell \quad (1 \leq \ell \leq r).$$

Let $f_i = d_i(V_0)$, $1 \leq i \leq j$, and define

$$(6.1) \quad P_{\mathbf{b}} = \prod_{i=1}^j \left(f_i - t + b_0 + \sum_{q=1}^r b_q \right),$$

$$(6.2) \quad Q_\ell(\mathbf{b}) = w + 1 + b_0 + \sum_{q>\ell} b_q,$$

$$(6.3) \quad A_\ell(\mathbf{b}) = e_\ell - b_\ell, \quad B_\ell(\mathbf{b}) = (Q_\ell(\mathbf{b}) - e_\ell + b_\ell)_+.$$

For a binary pattern $\sigma \in \{T, U\}^r$, put

$$(6.4) \quad C_\sigma(\mathbf{b}) = P_{\mathbf{b}} \prod_{\ell: \sigma_\ell = T} A_\ell(\mathbf{b}) \prod_{\ell: \sigma_\ell = U} B_\ell(\mathbf{b}),$$

$$(6.5) \quad G_\sigma = d_0^j \prod_{\ell: \sigma_\ell = T} e_\ell \prod_{\ell: \sigma_\ell = U} (d_{\ell-1} - (j + \ell - 1)).$$

Then every factor in $P_{\mathbf{b}}$ is positive and, for every $\sigma \in \{T, U\}^r$,

$$(6.6) \quad \boxed{\sum_{\mathbf{b}} L_{\mathbf{b}} C_\sigma(\mathbf{b}) \leq G_\sigma.}$$

Consequently, if nonnegative weights x_σ satisfy

$$(6.7) \quad \sum_{\sigma \in \{T, U\}^r} x_\sigma C_\sigma(\mathbf{b}) \geq 1 \quad \text{for every profile } \mathbf{b},$$

then, with $L = \sum_{\mathbf{b}} L_{\mathbf{b}}$,

$$(6.8) \quad \boxed{L \leq \left\lfloor \sum_{\sigma \in \{T, U\}^r} x_\sigma G_\sigma \right\rfloor.}$$

The bound is for the whole affine chart and contains no factor depending on $|\mathbb{F}|$ or on the number of affine cosets of V_0 .

Proof. Parameterize $\mathcal{A}$ by C' . Agreement at a coordinate x is an affine hyperplane with normal $\ell_x \in (C')^*$. A point of profile $\mathbf{b}$ has at least

$$d_0 - t + b_0 + \sum_{q=1}^r b_q$$

agreements in U_0 . Applying the ordered-basis argument of [Theorem 4.9](#) inside V_0 shows that it is incident with at least $P_{\mathbf{b}}$ ordered coordinate bases whose restrictions form a basis of V_0^* .

Suppose an adapted basis has already been chosen through level $V_{\ell-1}$. Its annihilator in V_ℓ is one-dimensional; let $h_\ell \in V_\ell \setminus V_{\ell-1}$ generate it. At a coordinate of T_ℓ , evaluation vanishes on $V_{\ell-1}$ but not on V_ℓ ; on the one-dimensional quotient $V_\ell/V_{\ell-1}$ it is therefore a nonzero functional. Hence h_ℓ is nonzero at every coordinate of T_ℓ . Moreover, h_ℓ vanishes outside U_ℓ . The number of mismatches outside U_ℓ is

$$b_0 + \sum_{q>\ell} b_q.$$

Since a nonzero degree- $< K$ word has at most $K - 1$ zeros, at most $K - 1 - b_0 - \sum_{q>\ell} b_q$ agreement coordinates fail to extend the current basis. Hence there are at least $Q_\ell(\mathbf{b})$ extending agreement coordinates in total. The $A_\ell(\mathbf{b})$ agreeing coordinates of T_ℓ all extend, so at least $B_\ell(\mathbf{b})$ extending agreement coordinates lie in $U_{\ell-1}$.

Fix a pattern σ . Multiplying the preceding lower bounds over the flag levels gives at least $C_\sigma(\mathbf{b})$ adapted ordered bases of pattern σ through every point of profile $\mathbf{b}$. Globally, the initial ordered basis has at most d_0^j choices. At a T -level there are at most e_ℓ coordinate choices. At a U -level, all $j + \ell - 1$ previously selected coordinates lie in $U_{\ell-1}$, and an extending normal cannot repeat one of them; there are therefore at most $d_{\ell-1} - (j + \ell - 1)$ choices. An independent affine-hyperplane basis cuts out at most one parameter point. Double counting proves (6.6). Multiplying those inequalities by x_σ , summing over σ , and using (6.7) proves (6.8). Positivity follows from $f_i \geq n - K + i$. $\square$

Corollary 6.5 (All-extender flag bound). *In the notation of Theorem 6.4, put*

$$P_0 = \prod_{i=1}^j (f_i - t + b_0), \quad Q = w + 1 + b_0.$$

Then

$$(6.9) \quad L \leq \left\lfloor \frac{d_0^j \prod_{\ell=1}^r (d_\ell - (j + \ell - 1))}{P_0 Q^r} \right\rfloor.$$

In particular, the numerator may be replaced by $d_0^j (n - j)^r$.

Proof. Every point has at least P_0 initial bases. At level ℓ it has at least $Q_\ell(\mathbf{b}) \geq Q$ extending agreement coordinates, without separating the two resources. The global number of adapted extensions at that level is at most $d_\ell - (j + \ell - 1)$. Multiplication and double counting give the first bound. Since $d_\ell \leq n$, each factor is at most $n - j - \ell + 1$, giving the second statement. $\square$

Lemma 6.6 (One-level interpolation defect). *For $e, Q > 0$ and $\rho \geq 1$, define*

$$(6.10) \quad \eta(e, Q, \rho) = \begin{cases} 1, & e \leq Q, \\ \min \left\{ 1, \frac{Q}{e} \rho^{1-Q/e} \right\}, & e > Q. \end{cases}$$

Then, for every real $0 \leq b \leq e$,

$$(6.11) \quad \rho^{b/e} \left(\frac{e-b}{e} + \frac{(Q-e+b)_+}{Q\rho} \right) \geq \eta(e, Q, \rho).$$

Thus the right side is 1 precisely in the profile-interpolation range used in Lemma 5.2.

Proof. If $e \leq Q$, the proof is the weighted AM–GM argument of Lemma 5.2. Suppose $e > Q$, put $c = Q/e$ and $\theta = b/e$. On $0 \leq \theta \leq 1 - c$, the logarithm of the left side is $\log(1 - \theta) + \theta \log \rho$, a concave function, so its minimum is at an endpoint and equals at least $\min\{1, c\rho^{1-c}\}$. On $1 - c \leq \theta \leq 1$, put $\alpha = 1 - \theta \in [0, c]$. The left side becomes

$$\rho^{-\alpha} \left(\alpha\rho + 1 - \frac{\alpha}{c} \right).$$

Its logarithm is a linear function plus the logarithm of a positive affine function, hence is concave. Its endpoint values are again 1 and $c\rho^{1-c}$. This proves the claim. $\square$

Corollary 6.7 (Tensorized support-flag certificate). *Retain the notation of Theorem 6.4. Put*

$$E = \sum_{\ell=1}^r e_\ell, \quad P(z) = \prod_{i=1}^j (f_i - t + b_0 + z), \quad \rho = \frac{P(E)}{P(0)},$$

and define

$$\rho_\ell = \rho^{e_\ell/E}, \quad \eta_\ell = \eta(e_\ell, Q, \rho_\ell), \quad Q = w + 1 + b_0.$$

Then

$$(6.12) \quad L \leq \left\lfloor \frac{d_0^j}{P(0) \prod_{\ell=1}^r \eta_\ell} \prod_{\ell=1}^r \left(1 + \frac{d_{\ell-1} - (j + \ell - 1)}{Q\rho_\ell} \right) \right\rfloor.$$

If every $\eta_\ell = 1$, this is a closed higher-codimension analogue of (5.11). For $r = 1$ it specializes to the codimension-one two-resource certificate, with the explicit defect factor when profile interpolation fails.

Proof. Let $B = \sum b_\ell$. Concavity of $\log P$ on $[0, E]$ gives

$$\frac{P(B)}{P(0)} \geq \rho^{B/E} = \prod_{\ell=1}^r \rho_\ell^{b_\ell/e_\ell}.$$

Also $Q_\ell(\mathbf{b}) \geq Q$, so $B_\ell(\mathbf{b}) \geq (Q - e_\ell + b_\ell)_+$. By Lemma 6.6,

$$\frac{P_{\mathbf{b}}}{P(0)} \prod_{\ell=1}^r \left(\frac{A_\ell(\mathbf{b})}{e_\ell} + \frac{B_\ell(\mathbf{b})}{Q\rho_\ell} \right) \geq \prod_{\ell=1}^r \eta_\ell.$$

Use the tensor-product dual weights

$$x_\sigma = \frac{1}{P(0) \prod_\ell \eta_\ell} \prod_{\ell: \sigma_\ell = T} \frac{1}{e_\ell} \prod_{\ell: \sigma_\ell = U} \frac{1}{Q\rho_\ell}.$$

They satisfy (6.7); substituting them into (6.8) and factoring the sum over binary patterns gives (6.12). $\square$

Remark 6.8 (What the flag theorem closes). The uncontrolled field-size coset sum is now removed in every codimension. For a concrete finite chart, (6.6) is an exact rational linear-program certificate, while (6.12) is a closed analytic certificate. What can still fail is numerical payment: a low-mismatch flag with several short layers may have a certified bound above the row budget. Such a flag is a specific finite residual atom, not an unproved multiplicity principle.

Proposition 6.9 (Counterexample to the affine-span transverse MCA compiler). *The affine-span MCA bound formerly printed here is false even under the direction-separation hypothesis. There is a received line for*

$$C = \text{RS}[\mathbb{F}_{1009}, \{0, \dots, 99\}, 1], \quad (n, K, m, w, s) = (100, 1, 21, 20, 1),$$

with 31 distinct selected slopes and explanations in the one-dimensional constant code, while

$$\max_{c \in C} \text{agr}(r_1, c) = 20 < 21 = m$$

and the claimed bound was

$$\left\lfloor \max \left\{ \frac{100^2}{21 \cdot 20}, \frac{100^2}{20^2} \right\} \right\rfloor = 23 < 31.$$

Every selected maximal agreement support has size exactly 21 and is same-support pair-noncontained.

Proof. Partition the domain into sets C_0, E, T, W of sizes 20, 30, 21, 29. Put $(r_0, r_1) = (0, 0)$ on C_0 . Index E by $1 \leq i \leq 30$ and put $(r_0, r_1) = (-i^2, i)$ at its i -th point. For slope i , the zero codeword then agrees exactly on C_0 and that point of E .

On T , put $r_0 = 1$ and choose 21 distinct fresh nonzero direction values avoiding $1 + ir_1 = 0$ for every $1 \leq i \leq 30$. At slope zero, the constant-one codeword agrees exactly on T . On W , choose fresh direction values and choose each base value outside $\{0, 1\} \cup \{-ir_1 : 1 \leq i \leq 30\}$. All choices exist in $\mathbb{F}_{1009}$, and they remove every unintended selected agreement.

The direction is zero on C_0 and takes distinct nonzero values elsewhere, so its maximum agreement with a constant is exactly 20. The selected explanations are zero and one, hence have affine span one. On a zero-explanation support, the core forces a simultaneous constant pair to be $(0, 0)$, but the remaining direction value is nonzero. On T , the direction values are distinct. Thus every support is pair-noncontained.

The proof failure is the asserted proper-subspace occupancy bound for the incident normals. On each zero-explanation support, 20 normals lie on one line and one is transverse. There are only 40 ordered bases, not the claimed lower count $mw = 420$. The complete deterministic reconstruction is in `experimental/verify_mca_affine_span_incidence_counterexample_v1.py`. $\square$

Remark 6.10 (Scope of the route cut). The counterexample retracts only the MCA incidence compiler and statements that use its ordered-basis denominator. It does not refute the ordinary affine-span list theorem, common-core cancellation, the directional Johnson compiler, codeword-direction gauge equivalence, or the selector-free all-LineRay error-affine-core set-pair theorem, whose proof uses zero-mask restriction and the Bollobas set-pair inequality instead.

Theorem 6.11 (Pair-noncontained proper-subspace MCA compiler). *Let $C = \text{RS}[\mathbb{F}, D, K]$ have length n , let $r_\gamma = r_0 + \gamma r_1$, put $w = m - K > 0$, and let $\mathcal{A} = c_0 + C' \subseteq C$ have $\dim C' = s$, where $1 \leq s \leq K$. Let $Z \subseteq \mathbb{F}$ be a set of distinct slopes such that for every $\gamma \in Z$, some $c_\gamma \in \mathcal{A}$ has agreement at least m , and its maximal agreement support is same-support pair-noncontained. Define*

$$e = \min_{b \in C} \text{wt}(r_1 - b), \quad L = \max\{1, e - (n - m)\}.$$

Then

$$(6.13) \quad |Z| \leq \left\lfloor \frac{1}{L} \max \left\{ \frac{n^{s+1}}{m(w+1)^{\overline{s-1}}}, \frac{(n-K+s)^{s+1}}{(w+1)^{\overline{s}}} \right\} \right\rfloor.$$

For $s = 1$, the empty rising product in the first denominator is one.

Proof. Choose a basis $c_1, \dots, c_s$ of C' . As in the rejected proof, the agreement hyperplane at coordinate x has normal

$$v_x = (r_1(x), -c_1(x), \dots, -c_s(x)) \in \mathbb{F}^{s+1}.$$

The normals incident with every selected parameter point span $\mathbb{F}^{s+1}$. Indeed, an annihilator with zero slope coordinate would give a nonzero codeword with at least $m > K - 1$ roots. An annihilator with nonzero slope coordinate would give simultaneous codeword explanations for r_0 and r_1 on the same maximal support, contrary to pair noncontainment.

Let $W < \mathbb{F}^{s+1}$ have dimension $j < s$, and put $X_W = \{x : v_x \in W\}$. If the slope coordinate vanishes on $W^\perp$, then $s + 1 - j$ independent codewords vanish on X_W . Otherwise the kernel of that coordinate supplies $s - j$ independent codewords. The MDS generalized-weight formula gives in either case

$$|X_W| \leq K - s + j.$$

Consequently, after $j = 1, \dots, s - 1$ independent incident normals have been selected, there are at least $w + s - j$ choices outside their span. These factors multiply to $(w + 1)^{\overline{s-1}}$.

After s independent normals have been selected, local full rank leaves at least one final choice. If their hyperplane has an annihilator with nonzero slope coordinate, it contains at most $n - e$ coordinate normals and leaves at least $e - (n - m)$ incident choices. If that slope coordinate is zero, the ordinary root bound leaves at least $w + 1$. Interpolation on K coordinates gives $e \leq n - K$, so every case leaves at least L final choices.

Let z be the number of zero normals and let $g \leq z$ be those whose agreement hyperplane contains every parameter point. Common-zero MDS gives $z \leq K - s$. Every selected point is therefore incident with at least

$$(m - g)(w + 1)^{\overline{s-1}}L$$

ordered normal bases, while there are at most $(n - z)^{\overline{s+1}}$ ordered independent coordinate tuples. For fixed z , the resulting ratio is largest at $g = z$. Its successive ratio in z has the sign of $n - (s + 1)m + sz$, so its maximum on $0 \leq z \leq K - s$ is at an endpoint. The endpoints $z = 0$ and $z = K - s$ are precisely the two terms in (6.13). $\square$

Theorem 6.12 (Support-transverse affine-span MCA compiler). *Retain the code, received line, affine explanation space, and parameters of Theorem 6.11, with $n \geq m = K + w$, $w > 0$, and $1 \leq s \leq K$. Let $\emptyset \neq Z \subseteq \mathbb{F}$ be a set of distinct slopes. Suppose that for every $\gamma \in Z$ there are a codeword $c_\gamma \in \mathcal{A}$ and an exact m -set $S_\gamma \subseteq D$ such that $r_\gamma = c_\gamma$ on S_γ , while no pair in C^2 explains (r_0, r_1) simultaneously on that same set. Define*

$$(6.14) \quad \theta_0 = \min_{\substack{\gamma \in Z \\ b \in C'}} |\{x \in S_\gamma : r_1(x) \neq b(x)\}|, \quad \theta = \min\{\theta_0, w + 1\}.$$

Then $\theta \geq 1$, and

$$(6.15) \quad |Z| \leq \left\lceil \max \left\{ \frac{n^{s+1}}{m\theta(w+1)^{s-1}}, \frac{(n-K+s)^{s+1}}{\theta(w+1)^s} \right\} \right\rceil.$$

The first rising product is empty when $s = 1$. If e, L are the global direction-coset quantities of [Theorem 6.11](#), then $\theta \geq L$. Thus (6.15) refines that theorem whenever the selected supports have more local transversality than is forced by the global direction distance.

Proof. Choose a basis $c_1, \dots, c_s$ of C' and use the same incident normal count as in [Theorem 6.11](#). The lower-dimensional proper-normal-subspace argument is unchanged and supplies the factors $(w+1)^{s-1}$. At the final step, let (δ, μ) span the annihilator of the selected s -space. If $\delta = 0$, the ordinary root bound leaves at least $w+1 \geq \theta$ incident choices outside it. If $\delta \neq 0$, put $b = \delta^{-1} \sum_i \mu_i c_i \in C'$. Every point of the actual support S_γ on which $r_1 \neq b$ has normal outside the selected space, so (6.14) supplies at least θ final choices. Pair noncontainment gives $\theta_0 \geq 1$: if $r_1 = b$ on S_γ , then $(c_\gamma - \gamma b, b)$ explains (r_0, r_1) there.

Keeping the zero-normal parameters z, g from the preceding proof, every selected point therefore owns at least $(m-g)\theta(w+1)^{s-1}$ ordered bases, while there are at most $(n-z)^{s+1}$ ordered independent coordinate tuples. For fixed z , the ratio is largest at $g = z$; its successive ratio in z again has the sign of $n - (s+1)m + sz$. The two endpoints $z = 0, K-s$ give (6.15). Finally, for every $b \in C' \subseteq C$,

$$|\{x \in S_\gamma : r_1(x) \neq b(x)\}| \geq \text{wt}(r_1 - b) - (n - m) \geq e - (n - m).$$

Moreover $e \leq n - K$, so $e - (n - m) \leq w$. If this difference is at most one, then $L = 1 \leq \theta$; otherwise the preceding inequality and the cap $w+1$ give $\theta \geq e - (n - m) = L$. Hence $\theta \geq L$. $\square$

Theorem 6.13 (Support-wise MCA error-rank gauge). *Let $C \leq \mathbb{F}^D$ be a linear code, let $r_\gamma = r_0 + \gamma r_1$, and let Z contain at least two distinct slopes. For every $\gamma \in Z$, choose $h_\gamma \in C$ and a support S_γ on which $r_\gamma = h_\gamma$, with no simultaneous C^2 -explanation of (r_0, r_1) on that same support. Put*

$$e_\gamma = r_\gamma - h_\gamma, \quad a = \dim \text{Span}\{e_\gamma - e_{\gamma_0} : \gamma \in Z\}.$$

Then $a \geq 1$, and there is a codeword $b \in C$ such that the two-way gauge

$$(6.16) \quad r'_0 = r_0, \quad r'_1 = r_1 - b, \quad h'_\gamma = h_\gamma - \gamma b$$

preserves every finite slope, error word, agreement support, and same-support containment predicate, while

$$\dim \text{Aff}\{h'_\gamma : \gamma \in Z\} = a - 1.$$

This is a reversible codeword translation, not identity of the received-line representative.

Proof. First $r_1 \notin C$. Otherwise, on every selected support the pair $(h_\gamma - \gamma r_1, r_1)$ would explain (r_0, r_1) , contrary to the hypothesis. Fix γ_0 , and consider

$$E = \text{Span}\{(\gamma - \gamma_0, h_\gamma - h_{\gamma_0}) : \gamma \in Z\} \leq \mathbb{F} \oplus C.$$

The map $(\delta, c) \mapsto \delta r_1 - c$ is injective on E : a kernel element with $\delta \neq 0$ would put r_1 in C , while $\delta = 0$ forces $c = 0$. Its image is the affine error-difference space, so $\dim E = a$. The slope projection on E is nonzero; choose $(1, b) \in E$. Under (6.16), the explanation differences are the image of E under $(\delta, c) \mapsto c - \delta b$, namely the kernel of the slope projection. Their dimension is $a - 1$.

The identity

$$r_0 + \gamma(r_1 - b) - (h_\gamma - \gamma b) = r_0 + \gamma r_1 - h_\gamma$$

proves error and support preservation. Subtracting b from the second codeword of a simultaneous explanation, and adding it back for the inverse, proves preservation of same-support containment. $\square$

Remark 6.14 (Relation to the full-explanation lifted-rank split). When the selected explanations have full affine rank (K), the lifted pair space (E) above has dimension (K) or ($K+1$). The gauge theorem gives explanation rank ($K-1$) in the first branch and (K) in the second. Thus it is the arbitrary-rank existence adapter underlying the two numerical cases of [Theorem 6.72](#) below. That theorem is sharper at full rank: it classifies every gauge, rather than selecting one.

Remark 6.15 (Exact first shortened-row calibration). For a shortened row $(n, K, m) = (R + K, K, d + K)$, the last factor is $L = \max\{1, e - (R - d)\} \leq d$. At the first residual dimensions from the common-core staircase, exact integer evaluation gives

row	K	paid for every e	additional paid walls
KoalaBear	14	$s \leq 9$	$s = 10, 11, 12, 13 : e \geq 981108, 981153, 981861, 992852$
Mersenne-31	6	$s \leq 1$	$s = 2, 3, 4, 5 : e \geq 981144, 981363, 984779, 1037876$

At top rank, the required factors are $182530 > d = 67472$ and $882143 > d = 67448$, respectively, so this compiler does not pay those cells. The exact replay is in `experimental/verify_mca_proper_subspace_occupancy_compiler_v1.py`.

Theorem 6.16 (Sparse-direction punctured Johnson MCA profile). *Let $C = \text{RS}[\mathbb{F}, D, K]$ have*

$$N = R + K, \quad m = d + K,$$

and let $r_\gamma = r_0 + \gamma r_1$. Suppose

$$r_1 = b + q, \quad b \in C, \quad E = \text{supp}(q), \quad |E| = e < d.$$

For every selected slope, assume an exact size- m , same-support pair-noncontained agreement witness with explanation c_γ , and put $a_\gamma = c_\gamma - \gamma b$. For each distinct selected transformed explanation define

$$h_a = m - |\{x \in D \setminus E : a(x) = r_0(x)\}|.$$

Then $1 \leq h_a \leq e$, and one deficit- h explanation owns at most $\lfloor e/h \rfloor$ selected slopes.

Assume

$$(6.17) \quad \Delta_e = (m - e)^2 - (N - e)(K - 1) > 0.$$

Put $J_0 = 0$ and, for $1 \leq h \leq e$,

$$(6.18) \quad J_h = \left\lfloor \frac{(N - e)(m - h - K + 1)}{(m - h)^2 - (N - e)(K - 1)} \right\rfloor.$$

The selected slope set Z satisfies

$$(6.19) \quad |Z| \leq \sum_{h=1}^e (J_h - J_{h-1}) \left\lfloor \frac{e}{h} \right\rfloor \leq (e - 1)J_{\lfloor e/2 \rfloor} + J_e.$$

For $e = 1$, the term indexed by zero in the last expression is zero.

Proof. Outside E , the direction equals b . If a agreed with r_0 on at least m outside coordinates, an m -subset there would be simultaneously explained by (a, b) , contrary to same-support pair noncontainment. Hence $h_a \geq 1$. A size- m witness can use at most e coordinates of E , so $h_a \leq e$.

For $x \in E$, agreement at slope γ is equivalent to

$$\frac{a(x) - r_0(x)}{q(x)} = \gamma.$$

These ratio fibers are disjoint. A deficit- h explanation needs at least h inside agreements, and therefore owns at most $\lfloor e/h \rfloor$ slopes.

Puncture E and fix h . Each explanation of deficit at most h has an agreement set of size at least $A = m - h$ in length $n' = N - e$. Agreement sets of two distinct degree- $< K$ polynomials intersect in at most $K - 1$ coordinates. If there are L explanations and s_x is the incidence multiplicity at coordinate x , then

$$\sum_x s_x \geq LA, \quad \sum_x \binom{s_x}{2} \leq (K - 1) \binom{L}{2}.$$

Cauchy–Schwarz gives

$$\sum_x \binom{s_x}{2} \geq \frac{1}{2} \left(\frac{L^2 A^2}{n'} - LA \right).$$

Since $A^2 > n'(K-1)$, comparison yields $L \leq J_h$. Condition (6.17) is the weakest, $h = e$, denominator and hence makes every J_h legal. The Johnson fractions are nondecreasing in h . Applying their cumulative caps to the nonincreasing owner weights proves the first inequality in (6.19).

Finally put $u = \lfloor e/2 \rfloor$. Bound all weights with $h \leq u$ by e ; every weight with $h > u$ equals one. The two telescoping blocks give

$$eJ_u + (J_e - J_u) = (e-1)J_u + J_e.$$

□

Remark 6.17 (Exact sparse-direction support walls). At the first shortened KoalaBear and Mersenne-31 rows, exact integer evaluation gives

row	K	last paid e	Δ_e	$(e-1)J_{\lfloor e/2 \rfloor} + J_e$
KoalaBear	14	63908	1218	4607583
Mersenne-31	6	65236	2794	2605443

against budgets 274980728111395087 and 16777215. At the adjacent supports the denominators are -5924 and -1636 , so this theorem makes no claim there. The exact scan is in `experimental/verify_mca_sparse_direction_punctured_johnson_profile_v1.py`.

Theorem 6.18 (Near-Johnson centered-Gram list compiler). *Let $S_1, \dots, S_L$ be distinct A -subsets of an n -set such that $|S_i \cap S_j| \leq c$ for $i \neq j$, where $A > c$. Put*

$$g = nc - A^2 \geq 0, \quad G = (A - c)^2 - cg.$$

If $G > 0$, then

$$(6.20) \quad L \leq \left\lfloor \frac{nA(A-c)}{G} \right\rfloor.$$

Proof. Let B be the L -by- n incidence matrix and put

$$H = BB^\top - c\mathbf{1}\mathbf{1}^\top.$$

Every row of B has sum A , so $\mathbf{1} = A^{-1}B\mathbf{1}_n$. The images of both summands defining H therefore lie in $\text{col}(B)$, and

$$(6.21) \quad \text{rank}(H) \leq \text{rank}(B) \leq n.$$

Write $\delta_{ij} = c - |S_i \cap S_j|$. Then $0 \leq \delta_{ij} \leq c$ is integral and

$$\text{tr}(H) = L(A - c), \quad \text{tr}(H^2) = L(A - c)^2 + 2 \sum_{i < j} \delta_{ij}^2.$$

If s_x is the incidence multiplicity at coordinate x , then Cauchy–Schwarz gives

$$2 \sum_{i < j} |S_i \cap S_j| = \sum_x s_x^2 - LA \geq \frac{L^2 A^2}{n} - LA.$$

Thus, for $\Delta = \sum_{i < j} \delta_{ij}$,

$$2\Delta \leq L \left(\frac{Lg}{n} + A - c \right).$$

Since $\delta_{ij}^2 \leq c\delta_{ij}$,

$$\text{tr}(H^2) \leq L \left((A - c)^2 + c(A - c) + \frac{cLg}{n} \right).$$

The symmetric trace-rank inequality $\text{rank}(H) \geq \text{tr}(H)^2 / \text{tr}(H^2)$, together with (6.21), yields

$$L((A - c)^2 - cg) \leq nA(A - c),$$

which is (6.20). □

Corollary 6.19 (Sparse-direction near-Johnson MCA payment). *Retain the hypotheses of Theorem 6.16, put*

$$u = \lfloor e/2 \rfloor, \quad n = N - e, \quad c = K - 1, \quad A = m - e,$$

and assume

$$(m - u)^2 - nc > 0, qquad g = nc - A^2 \geq 0, qquad G = (A - c)^2 - cg > 0.$$

Define

$$J_u = \left\lfloor \frac{n(m - u - K + 1)}{(m - u)^2 - nc} \right\rfloor, \quad Q_e = \left\lfloor \frac{nA(A - c)}{G} \right\rfloor.$$

Then

$$(6.22) \quad |Z| \leq (e - 1)J_u + Q_e.$$

Proof. Let N_u be the number of transformed explanations with outside deficit at most u , and let N_e be their total number. The deficit-owner bound gives

$$|Z| \leq eN_u + (N_e - N_u) = (e - 1)N_u + N_e.$$

The ordinary Johnson argument on the punctured row gives $N_u \leq J_u$. For each explanation counted by N_e , choose exactly $A = m - e$ outside agreement coordinates. Distinct degree- $< K$ explanations give distinct blocks with intersections at most $c = K - 1$, so Theorem 6.18 gives $N_e \leq Q_e$. $\square$

Remark 6.20 (Exact centered-Gram support walls). Together with the preceding Johnson prefix, exact integer evaluation of (6.22) expands the low-support walls to

row	K	last paid e	endpoint bound
KoalaBear	14	64037	198047217
Mersenne-31	6	65418	16759641

against budgets 274980728111395087 and 16777215. At KoalaBear $e = 64038$, the centered-Gram denominator is -36911 . At Mersenne $e = 65419$, it remains positive, but the valid bound 18212004 exceeds budget. The exact replay is in `experimental/verify_mca_sparse_direction_near_johnson_gram_rank_v1.py`.

Theorem 6.21 (Mean-centered Gram list compiler). *Let $S_1, \dots, S_L$ be distinct A -subsets of an n -set with $|S_i \cap S_j| \leq c$ for $i \neq j$, where $A > c$. Put*

$$g = nc - A^2, \quad T = (n - A)^2 - (n - 1)g.$$

If

$$g \geq 0, \quad 2A^2 \geq nc, \quad T > 0,$$

then

$$(6.23) \quad L \leq \left\lfloor \frac{(n - 1)n^2(A - c)}{AT} \right\rfloor.$$

Proof. Let B be the incidence matrix and put

$$P = I_n - \frac{1}{n}J_n, \quad H = BPB^\top.$$

Then H is positive semidefinite and has rank at most $n - 1$. Set $p = A^2/n$. The diagonal entries are $A - p$, while an off-diagonal entry is

$$x_{ij} = |S_i \cap S_j| - p \in [-p, c - p].$$

Convexity gives the endpoint-chord inequality

$$x_{ij}^2 \leq (c - 2p)x_{ij} + p(c - p).$$

The assumption $2A^2 \geq nc$ makes the chord slope nonpositive. Since H is positive semidefinite,

$$2 \sum_{i < j} x_{ij} \geq -L(A - p).$$

With $C = p(c - p) = A^2g/n^2$, the preceding two displays imply

$$\text{tr}(H^2) \leq CL^2 + A(A - c)L.$$

Now $\text{tr}(H) = L(A - p)$. The PSD trace-rank inequality gives

$$L((A - p)^2 - (n - 1)C) \leq (n - 1)A(A - c).$$

Finally,

$$(A - p)^2 - (n - 1)C = \frac{A^2T}{n^2},$$

which proves (6.23). $\square$

Corollary 6.22 (Exact mean-centered sparse-direction profile). *Retain the sparse-direction notation of Theorem 6.16. For each deficit threshold $1 \leq h \leq e$, put $A_h = m - h$ and let C_h be the ordinary Johnson cap when $A_h^2 > (N - e)(K - 1)$, and otherwise the cap in Theorem 6.21 whenever its three hypotheses hold. Assume every C_h is defined, and put*

$$B_0 = 0, \quad B_h = \min_{h \leq v \leq e} C_v.$$

Then

$$(6.24) \quad |Z| \leq \sum_{h=1}^e (B_h - B_{h-1}) \left\lfloor \frac{e}{h} \right\rfloor.$$

Proof. Let N_h count distinct transformed explanations with outside deficit at most h . Puncturing the direction support and applying the appropriate ordinary-list theorem gives $N_h \leq C_h$. For every $v \geq h$, also $N_h \leq N_v \leq C_v$, hence $N_h \leq B_h$. The suffix minima are nondecreasing. Apply the deficit owner weight $\lfloor e/h \rfloor$ and sum by parts. $\square$

Remark 6.23 (Exact mean-centered support walls). The complete exact profile in (6.24) expands the walls to

row	K	last paid e	endpoint profile
KoalaBear	14	64047	181731868
Mersenne-31	6	65454	16101127

At KoalaBear $e = 64048$, the endpoint denominator is $T = -1499457466$. At Mersenne $e = 65455$, every cap remains defined, but the exact profile is $17120123 > 16777215$. The replay is in `experimental/verify_mca_sparse_direction_mean_centered_gram_profile_v1.py`.

Theorem 6.24 (Terminal-deficit affine-line cap). *Retain the sparse-direction notation of Theorem 6.16, assume $e \geq K$, and put*

$$n = N - e, \quad A = m - e, \quad c = K - 1.$$

If $A > c$, then the number L_e of transformed explanations that own a selected slope and have exact outside deficit e satisfies

$$(6.25) \quad L_e \leq \left\lfloor \frac{n - c}{A - c} \right\rfloor.$$

Proof. An exact-deficit- e explanation has exactly $A = m - e$ outside agreements. To own a selected slope it must therefore agree at all e coordinates of the direction support E . If two such explanations have slopes $\gamma \neq \delta$, their normalized difference is a nonzero degree- $< K$ codeword p whose restriction to E is the gauged direction $q|_E$. Every further terminal explanation gives the same normalized difference: the two codewords agree on E , and evaluation on E is injective because $e \geq K$. Thus all terminal explanations lie on one affine codeword line.

Outside E , let G be the common agreement core on the zeros of p , and write $g = |G| \leq K - 1 = c$. Away from G , a coordinate can be an agreement for at most one affine-line parameter. Since each terminal explanation has $A - g$ agreements away from the core,

$$L_e(A - g) \leq n - g.$$

The ratio $(n - g)/(A - g)$ increases with g , so $g \leq c$ gives (6.25). The case $L_e \leq 1$ is immediate. $\square$

Corollary 6.25 (Prefix-plus-terminal sparse-direction profile). *For every $1 \leq h < e$, let C_h be a proved cumulative cap on the number N_h of transformed explanations with outside deficit at most h , and put*

$$B_0 = 0, \quad B_h = \min_{h \leq v < e} C_v.$$

Then

$$(6.26) \quad |Z| \leq \sum_{h=1}^{e-1} (B_h - B_{h-1}) \left\lfloor \frac{e}{h} \right\rfloor + \left\lfloor \frac{N - e - (K - 1)}{m - e - (K - 1)} \right\rfloor.$$

Proof. For $h < e$, monotonicity gives $N_h \leq \min_{h \leq v < e} C_v = B_h$. Sum the nonincreasing owner weights $\lfloor e/h \rfloor$ by parts over the prefix layers, then apply [Theorem 6.24](#) to the exact terminal layer, whose owner weight is one. $\square$

Remark 6.26 (Exact terminal-line support walls). Using the Johnson/mean-centered cumulative caps for $h < e$, (6.26) pays one further support in each deployed row:

row	K	last paid e	terminal cap	total profile
KoalaBear	14	64048	287	181326343
Mersenne-31	6	65455	493	16100647

At KoalaBear $e = 64049$, the cumulative cap at $h = e - 1$ is unavailable. At Mersenne $e = 65456$, the valid profile is $17119507 > 16777215$. The remaining full-lift intervals are therefore $64049 \leq e \leq 1044238$ and $65456 \leq e \leq 1044241$. This is a support-wall refinement, not a proof of either deployed row. The exact replay is in `experimental/verify_mca_sparse_direction_terminal_deficit_line_payment_v1.py`.

Theorem 6.27 (Top-third exact-layer affine-line cap). *Retain the sparse-direction notation of [Theorem 6.16](#). Assume $e \geq K$, put*

$$s = \left\lfloor \frac{e - K}{3} \right\rfloor, \quad H = e - s - 1,$$

and assume $N - m > s$. For $0 \leq r \leq s$, let L_r be the number of selected explanations with exact outside deficit $h = e - r$. Then

$$(6.27) \quad L_r \leq \left\lfloor \frac{N - e - (K - 1)}{m - e + r - (K - 1)} \right\rfloor.$$

Every explanation counted by L_r owns at most one selected slope.

Proof. For an explanation a_i in this layer, let $S_i \subseteq E$ be its inside agreement set with its assigned slope. Then

$$|E \setminus S_i| \leq r.$$

For explanations at distinct slopes γ_i, γ_j , the normalized difference

$$p_{ij} = \frac{a_i - a_j}{\gamma_i - \gamma_j}$$

is a degree- $< K$ codeword and equals the gauged direction q on $S_i \cap S_j$. Any three inside agreement sets have intersection at least $e - 3r \geq K$. Hence p_{ij} and p_{ik} agree on at least K evaluation coordinates and are equal. Fixing two anchors puts the entire exact layer on one affine codeword line.

Each member has $A_r = m - e + r$ outside agreements. The line has a common agreement core of size $g \leq K - 1$, and outside that core each coordinate agrees for at most one line parameter. Thus

$$L_r(A_r - g) \leq N - e - g.$$

The hypothesis $N - m > s$ gives $N - e > A_r$, so the ratio increases with g ; substituting $g \leq K - 1$ proves (6.27). Finally $r < e/2$, so the deficit-owner bound is $\lfloor e/(e - r) \rfloor = 1$. $\square$

Corollary 6.28 (Top-third sparse-direction profile). *For $1 \leq h \leq H$, let C_h be a proved cumulative explanation cap and put*

$$B_0 = 0, \quad B_h = \min_{h \leq v \leq H} C_v.$$

Then

$$(6.28) \quad |Z| \leq \sum_{h=1}^H (B_h - B_{h-1}) \left\lfloor \frac{e}{h} \right\rfloor + \sum_{r=0}^s \left\lfloor \frac{N - e - (K - 1)}{m - e + r - (K - 1)} \right\rfloor.$$

If the positive punctured-Johnson caps J_h are defined through H and $u = \lfloor e/2 \rfloor$, the first sum is at most

$$(e - 1)J_u + J_H.$$

Proof. Apply cumulative summation by parts through H . The remaining exact layers are $h = e - r$, $0 \leq r \leq s$; apply [Theorem 6.27](#). For the coarse prefix estimate, give layers through u weight at most e and all later prefix layers weight one. $\square$

Remark 6.29 (Complete sparse-direction branch payment). For the Johnson rational cap

$$F(n, A) = \frac{n(A - (K - 1))}{A^2 - n(K - 1)},$$

the function increases with n and decreases with A . Uniformly over $K \leq e < d$, conservative endpoint substitution gives

$$J_{\lfloor e/2 \rfloor} \leq 31, \quad J_H \leq 47$$

on both deployed rows. Each affine-line summand is nondecreasing in e , so its sum is maximal at $e = d - 1$. The exact endpoint calculations are

row	line sum	uniform total	budget
KoalaBear	9405342	11496959	274980728111395087
Mersenne-31	9405365	11496238	16777215

Consequently every sparse-direction support $1 \leq e < d$ is paid. The remaining full-lift intervals are

$$67472 \leq e \leq 1044238 \quad \text{and} \quad 67448 \leq e \leq 1044241.$$

This closes the low-support branch, not either deployed row. The exact replay is in `experimental/verify_mca_sparse_direction_top_third_affine_line_payment_v1.py`.

Theorem 6.30 (Pair-noncontained total-core line cap). *Retain the notation of [Theorem 6.27](#), but allow $e \geq d$. Put*

$$r_{\min} = \max\{0, e - m\}, \quad t = N - m.$$

For $r_{\min} \leq r \leq s$, set $A_r = m - e + r$. Then the exact-layer count satisfies

$$(6.29) \quad L_r \leq Q_r, \quad Q_r = \begin{cases} t + 1, & A_r \leq K - 1, \\ \left\lfloor \frac{N - e - (K - 1)}{A_r - (K - 1)} \right\rfloor, & A_r > K - 1. \end{cases}$$

Proof. The triple-overlap proof of [Theorem 6.27](#) still puts the exact layer on an affine explanation line $a_\gamma = A + \gamma p$. A coordinate agrees for every line parameter exactly when $A = r_0$ and $p = q$ there. If this total common agreement core had at least m coordinates, the codeword pair $(A, b + p)$ would explain the received pair on an m -set, contradicting pair noncontainment. Thus its size g is at most $m - 1$. Away from the core, each coordinate agrees for at most one parameter, and therefore

$$L_r(m - g) \leq N - g, \quad L_r \leq N - m + 1 = t + 1.$$

When $A_r > K - 1$, outside-coordinate packing gives the second branch as in [Theorem 6.27](#). That branch is no larger than $t + 1$: after clearing its positive denominator, the difference is $t(A_r - K) + r \geq 0$. $\square$

Corollary 6.31 (Full-lift common-core profile). *If the positive punctured-Johnson caps exist through $H = e - s - 1$, then*

$$(6.30) \quad |Z| \leq (e - 1)J_{\lfloor e/2 \rfloor} + J_H + \sum_{r=r_{\min}}^s Q_r.$$

Proof. Deficits through H use the two-threshold Johnson prefix. Deficits above H and at most m are exactly the layers $r_{\min} \leq r \leq s$; apply [Theorem 6.30](#). $\square$

Remark 6.32 (Exact full-lift common-core walls). Exact integer evaluation of (6.30) extends the complete walls to

row	last paid e	endpoint bound	next bound
KoalaBear	95943	27414298	undefined
Mersenne-31	67452	16266965	17248067

At KoalaBear $e = 95944$, the prefix Johnson denominator at H is -1037 . The adjacent Mersenne bound exceeds budget by 470852. The remaining full-lift intervals are

$$95944 \leq e \leq 1044238 \quad \text{and} \quad 67453 \leq e \leq 1044241.$$

Neither adjacent failure is an unsafe certificate. The replay is in `experimental/verify_mca_full_lift_top_third_common_core_payment_v1.py`.

Theorem 6.33 (Cross-layer top-third global-line cap). *Retain the full-lift notation above and put*

$$s = \left\lfloor \frac{e - K}{3} \right\rfloor, \quad r_{\min} = \max\{0, e - m\}.$$

All selected explanations in the union of exact layers

$$r_{\min} \leq r \leq s, \quad h = e - r,$$

lie on one common affine codeword line. The number L_{high} of explanations in this entire union satisfies

$$(6.31) \quad L_{\text{high}} \leq N - m + 1.$$

Each explanation in the union owns at most one selected slope.

Proof. For an explanation a_i , let $r_i \leq s$ be its missed-coordinate allowance on the direction support E , and let $S_i \subseteq E$ be its inside agreement set. For any three explanations,

$$|S_i \cap S_j \cap S_k| \geq e - (r_i + r_j + r_k) \geq e - 3s \geq K.$$

For distinct slopes, the normalized pair difference

$$p_{ij} = \frac{a_i - a_j}{\gamma_i - \gamma_j}$$

is a degree- $< K$ codeword equal to the gauged direction q on $S_i \cap S_j$. Hence $p_{ij} = p_{ik}$ by restriction injectivity on the displayed triple intersection. Fixing two anchors synchronizes all pair directions and puts the entire cross-layer union on one affine codeword line. Families of size at most two are immediate.

Let g be the line's total common agreement core. As in [Theorem 6.30](#), pair noncontainment gives $g \leq m - 1$. Away from this core, each coordinate agrees for at most one line parameter. Every selected explanation has at least $m - g$ agreements off the common core when all N coordinates are counted, so

$$L_{\text{high}}(m - g) \leq N - g.$$

The right-hand ratio increases with g , and $g \leq m - 1$ yields $L_{\text{high}} \leq N - m + 1$. Finally $r_i \leq s < e/2$, so the deficit-owner bound is one. $\square$

Corollary 6.34 (Full-lift global-line profile). *Put $H = e - s - 1$ and $u = \lfloor e/2 \rfloor$. If the positive punctured-Johnson caps J_u, J_H exist, then*

$$(6.32) \quad |Z| \leq (e - 1)J_u + J_H + (N - m + 1).$$

Proof. Deficits at most H form the Johnson prefix and contribute at most $(e-1)J_u + J_H$. Every remaining possible deficit belongs to the cross-layer union in [Theorem 6.33](#); charge its line cap once. $\square$

Remark 6.35 (Exact cross-layer global-line walls). Exact integer evaluation of (6.32) gives

row	last paid e	endpoint bound	largest bound
KoalaBear	95943	6336049	6336049 at $e = 95943$
Mersenne-31	97908	6682339	6683188 at $e = 97907$

At the adjacent supports $e = 95944$ and $e = 97909$, the respective H -prefix Johnson denominators are -1037 and -965 . The remaining full-lift intervals are therefore

$$95944 \leq e \leq 1044238 \quad \text{and} \quad 97909 \leq e \leq 1044241.$$

This replaces the per-layer top-third charge by one global charge. It does not repair the failed prefix cap, certify either adjacent support unsafe, or close a deployed row. The exact replay is in `experimental/verify_mca_full_lift_top_third_global_line_payment_v1.py`.

Corollary 6.36 (Full-lift mean-centered global-line profile). *Retain the notation of [Theorem 6.33](#), put*

$$H = e - s - 1, \quad c = K - 1,$$

and assume $m - H > c$. For every $1 \leq h \leq H$, put $A_h = m - h$, and let C_h be the punctured ordinary-Johnson cap when $A_h^2 > (N - e)c$, and otherwise the cap in [Theorem 6.21](#) whenever its three hypotheses hold. Assume every C_h is defined, and put

$$B_0 = 0, \quad B_h = \min_{h \leq v \leq H} C_v.$$

Then

$$(6.33) \quad |Z| \leq \sum_{h=1}^H (B_h - B_{h-1}) \left\lfloor \frac{e}{h} \right\rfloor + (N - m + 1).$$

Proof. Let N_h count transformed explanations with outside deficit at most h . The earlier sparse-direction corollary was calibrated for $e < d$, but its ordinary set-system proof has no such hypothesis. After puncturing the direction support, choose exactly $A_h = m - h$ outside agreement coordinates for every explanation. The guard $m - H > c$ gives $A_h > c$ throughout the prefix, and distinct degree- $< K$ explanations intersect in at most c coordinates. Thus $N_h \leq C_h$. For $v \geq h$, also $N_h \leq N_v \leq C_v$, so $N_h \leq B_h$. Summation by parts against the owner weights $\lfloor e/h \rfloor$ gives the displayed prefix charge. Every remaining deficit belongs to the cross-layer high union; apply [Theorem 6.33](#) once. $\square$

Remark 6.37 (Exact full-lift mean-centered walls). Exact evaluation of (6.33) gives

row	last paid e	endpoint bound	largest bound
KoalaBear	96150	479693401	479693401 at $e = 96150$
Mersenne-31	98229	16488216	16489118 at $e = 98228$

At KoalaBear $e = 96151$, the first undefined cap is at $h = H = 64105$, where the mean-centered denominator is -4625043784 . At Mersenne-31 $e = 98230$, every cap remains defined but the profile is 17415873, exceeding budget by 638658. The remaining full-lift intervals are

$$96151 \leq e \leq 1044238 \quad \text{and} \quad 98230 \leq e \leq 1044241.$$

Neither adjacent failure is an unsafe certificate, and no deployed v4 atom moves. The exact replay is in `experimental/verify_mca_full_lift_mean_centered_global_line_profile_v1.py`.

Corollary 6.38 (One-layer boundary-anchor continuation). *Retain the hypotheses and notation of [Corollary 6.36](#), and put*

$$q = e - K - 3s.$$

For $0 \leq J \leq H$, define

$$B_0^{(J)} = 0, \quad B_h^{(J)} = \min_{h \leq v \leq J} C_v, \quad P_J = \sum_{h=1}^J (B_h^{(J)} - B_{h-1}^{(J)}) \left\lfloor \frac{e}{h} \right\rfloor, \quad P_0 = 0.$$

If $H \geq 2$, $q \geq 1$, and $2(s+1) < e$, then

$$(6.34) \quad |Z| \leq \max\{P_H + 1, P_{H-1} + N - m + 1\}.$$

Proof. Let $\mathcal{A}$ be the selected explanations with missed-coordinate allowance at most s , and let $\mathcal{D}$ be the exact boundary layer with allowance $s+1$. If $|\mathcal{A}| \leq 1$, the complete prefix through H and the one-slope ownership guard give $|Z| \leq P_H + 1$.

Suppose $|\mathcal{A}| \geq 2$, and choose distinct anchors $a, b \in \mathcal{A}$. For each $c \in \mathcal{D}$,

$$|S_a \cap S_b \cap S_c| \geq e - s - s - (s+1) = K + q - 1 \geq K.$$

The normalized pair differences for (c, a) and (c, b) agree with the gauged direction on this intersection. Restriction injectivity for degree- $< K$ codewords makes them identical, so c lies on the affine codeword line already containing $\mathcal{A}$. Thus $\mathcal{A} \cup \mathcal{D}$ is one pair-noncontained line. The guard $2(s+1) < e$ gives one-slope ownership throughout this union, and [Theorem 6.30](#) charges it by $N - m + 1$. The remaining prefix is bounded independently by P_{H-1} . $\square$

Remark 6.39 (First Mersenne boundary support). For Mersenne-31 at $e = 98230$,

$$(s, q, H) = (32741, 1, 65488), \quad P_H = 16434744, \quad P_{H-1} = 15506184,$$

and $N - m + 1 = 981129$. Hence [\(6.34\)](#) gives

$$|Z| \leq \max\{16434745, 16487313\} = 16487313 < 16777215,$$

with slack 289902. At $e = 98231$ the same legal bound is 17492173, which exceeds budget by 714958; the residue-two refinement below repairs that proof-method failure. The exact replay is in `experimental/verify_mca_full_lift_boundary_anchor_continuation_v1.py`.

Corollary 6.40 (Residue-two boundary-layer continuation). *Retain the hypotheses and notation of [Corollary 6.38](#). Suppose in addition that*

$$q = 2, \quad H \geq 3, \quad 2(s+2) < e, \quad m - H > c,$$

and put

$$Q_H = \left\lfloor \frac{N - e - c}{m - H - c} \right\rfloor, \quad \Delta_H = \left\lfloor \frac{e}{s+1} \right\rfloor.$$

Then

$$(6.35) \quad |Z| \leq \max \left\{ \begin{array}{l} P_{H-2} + N - m + 1, \\ P_{H-1} + Q_H + 1, \\ P_{H-1} + 2, \\ P_{H-1} + Q_H, \\ P_{H-1} + \Delta_H \end{array} \right\}.$$

Proof. Let $\mathcal{A}$ be the top-third union, and let $\mathcal{D}_1, \mathcal{D}_2$ be the exact layers of deficits $H, H-1$. Choose inside agreement sets at their layer sizes, so their missed sets have sizes $s+1, s+2$. The guard $2(s+2) < e$ gives one-slope ownership.

If $|\mathcal{A}| \geq 2$, two top anchors meet every member of $\mathcal{D}_1 \cup \mathcal{D}_2$ on at least

$$e - s - s - (s+2) = K + q - 2 = K$$

common coordinates. Restriction injectivity puts all three layers on one affine codeword line. Total-core packing charges it by $N - m + 1$, leaving P_{H-2} .

Suppose $\mathcal{A}$ has one member. If $|\mathcal{D}_1| \geq 2$, that member and any two boundary members have at least $e - s - 2(s+1) = K$ common coordinates, so $\mathcal{D}_1$ is one affine line. Charge the top member separately. Every boundary member has exactly $m - H$ outside agreements, and outside-core line packing bounds the layer by Q_H . This gives $P_{H-1} + Q_H + 1$. If $|\mathcal{D}_1| \leq 1$, direct charging gives $P_{H-1} + 2$.

Finally suppose $\mathcal{A}$ is empty, and write R_i for the size- $(s+1)$ missed set of a member of $\mathcal{D}_1$. If some R_a, R_b intersect, then every third member has

$$|S_a \cap S_b \cap S_j| \geq e - (3(s+1) - 1) = K + q - 2 = K.$$

The fixed pair synchronizes the whole layer, giving $P_{H-1} + Q_H$. Otherwise the R_i are pairwise disjoint, whence $|\mathcal{D}_1|(s+1) \leq e$ and the charge is $P_{H-1} + \Delta_H$. These cases are exhaustive. $\square$

Remark 6.41 (Second Mersenne boundary support). At Mersenne-31 $e = 98231$,

$$(s, q, H) = (32741, 2, 65489), \quad P_{H-2} = 15505282, \quad P_{H-1} = 16433719,$$

while $N - m + 1 = 981129$, $Q_H = 484$, and $\Delta_H = 3$. The five cases in (6.35) are

$$16486411, \quad 16434204, \quad 16433721, \quad 16434203, \quad 16433722.$$

Thus $|Z| \leq 16486411 < 16777215$, with slack 290804. At $e = 98232$ the residue resets to $q = 0$, so this corollary is inapplicable; that is not an unsafe certificate. The Mersenne full-lift residual interval is now $98232 \leq e \leq 1044241$. The exact replay is in `experimental/verify_mca_full_lift_two_boundary_layer_continuation_v1.py`.

Corollary 6.42 (Boundary direction-class router). *Retain the pair-noncontained full-lift hypotheses and notation of Theorem 6.33, put*

$$H = e - s - 1, \quad c = K - 1,$$

For $1 \leq h \leq H - 1$, define C_h by the punctured Johnson or mean-centered rule in Corollary 6.36, and suppose every such cap is defined. Define the independently truncated prefix

$$B_0 = 0, \quad B_h = \min_{h \leq v \leq H-1} C_v, \quad P_{H-1} = \sum_{h=1}^{H-1} (B_h - B_{h-1}) \left\lfloor \frac{e}{h} \right\rfloor.$$

Let $\mathcal{D}$ be the exact deficit- H layer, and write q_{dir} for the gauged direction on its support E . Put

$$A = 2H - e, \quad J = \left\lfloor \frac{e(A - c)}{A^2 - ec} \right\rfloor, \quad Q = \left\lfloor \frac{N - e - c}{m - H - c} \right\rfloor.$$

If

$$A > 0, \quad A^2 > ec, \quad N - e > m - H > c,$$

then

$$(6.36) \quad |\mathcal{D}| \leq 1 + J(Q - 1).$$

If $\mathcal{T}$ is the synchronized union of deficits greater than H , then consequently

$$(6.37) \quad |Z| \leq P_{H-1} + 1 + J(Q - 1) + |\mathcal{T}|.$$

Proof. The claim is immediate when $\mathcal{D}$ is empty. Otherwise fix an anchor $(\gamma_0, c_0) \in \mathcal{D}$. For every other member put

$$p_\gamma = \frac{c_\gamma - c_0}{\gamma - \gamma_0}.$$

The chosen size- H inside agreement sets meet in at least $A = 2H - e$ coordinates, and on this intersection $p_\gamma = q_{\text{dir}}$. Thus the intrinsic set

$$U_p = \{x \in E : q_{\text{dir}}(x) = p(x)\}$$

has size at least A for every represented direction p . Such a p is nonzero: otherwise the two explanations would be the same codeword, while on their nonempty inside intersection that codeword would equal received words whose difference is $(\gamma - \gamma_0)q_{\text{dir}} \neq 0$.

For distinct represented directions p, p' , the intersection $U_p \cap U_{p'}$ lies in the zero set of the nonzero degree- $< K$ codeword $p - p'$, and hence has size at most c . Choosing A points from every U_p and applying the constant-block Johnson double count gives at most J represented directions.

For a fixed p , its members and the anchor lie on the affine codeword line $c_0 + (\gamma - \gamma_0)p$. Outside E the received direction vanishes, so the common agreement core of this line lies in the zero set of the nonzero codeword p and has size at most c . Every member has $m - H$ outside

agreements. Since $N - e > m - H > c$, outside-core packing bounds the line by Q . The anchor is repeated in each line closure, so each of the at most J classes contributes at most $Q - 1$ nonanchors. This proves (6.36). The boundary and top layers have one-slope ownership; the independently truncated prefix charges all lower deficits, proving (6.37). $\square$

Remark 6.43 (First residue-zero Mersenne router). At Mersenne-31 $e = 98232$,

$$(s, H, A) = (32742, 65489, 32746), \quad J = 3, \quad Q = 484,$$

so $|\mathcal{D}| \leq 1450$. The exact prefix is $P_{H-1} = 16432695$, and therefore

$$|Z| \leq 16434145 + |\mathcal{T}|.$$

An unsafe family would need $|\mathcal{T}| \geq 343071$. If g is the total common core of this affine line, line packing gives

$$g \geq \left\lceil \frac{|\mathcal{T}|m - N}{|\mathcal{T}| - 1} \right\rceil \geq 67452 = m - 2.$$

Thus the first residue-zero obstruction is reduced to a near-maximal-core line. The router alone is not a safety or unsafety certificate at $e = 98232$; the following corollary absorbs the remaining layers into this line. The exact router replay is in `experimental/verify_mca_full_lift_residue_zero_direction_router_v1.py`, with an independent endpoint audit in `experimental/audit_mca_full_lift_residue_zero_direction_router_v1.py`.

Corollary 6.44 (First residue-zero Mersenne payment). *Under the pair-noncontained full-lift hypotheses, the Mersenne-31 support $e = 98232$ is safe:*

$$|Z| \leq 16777215.$$

Consequently its full-lift residual interval starts at $e = 98233$.

Proof. Suppose instead that $|Z| > 16777215$. By the preceding router, the synchronized top line $\mathcal{T}$ has at least 343071 members and total common agreement core G of size at least $67452 = m - 2$. Write the line as $c_\gamma = A_0 + \gamma p$. Its direction p is nonzero. Outside the gauged direction support E , a coordinate common to two line members must be a zero of p , so

$$(1) \quad |G \cap E| \geq 67452 - (K - 1) = 67447 =: u.$$

Take any selected slope δ whose assigned explanation has outside deficit $h \geq e - u + K = 30791$. Ownership supplies an inside agreement set S_δ of size h . Choose two distinct top anchors. By (1),

$$|G \cap E \cap S_\delta| \geq u + h - e \geq K.$$

On these coordinates, both normalized differences from c_δ to the two anchors equal the gauged direction. Restriction injectivity makes the two degree- $< K$ differences equal. Subtracting their defining equations then identifies this common difference with p , and hence $c_\delta = A_0 + \delta p$. Thus every assigned pair with $h \geq 30791$ joins the same affine line. The nonzero direction makes its slope parameter injective, and total-core line packing charges all these slopes once by

$$(2) \quad N - m + 1 = 981129.$$

Every remaining assigned explanation has $h \leq 30790$, so after puncturing E it agrees with the outside received word on at least $m - 30790 = 36664$ coordinates. The ordinary Johnson count gives at most

$$(3) \quad \left\lfloor \frac{950350(36664 - 5)}{36664^2 - 950350 \cdot 5} \right\rfloor = 26$$

distinct explanations. Each owns at most $e = 98232$ slopes. Combining (2) and (3) yields

$$|Z| \leq 98232 \cdot 26 + 981129 = 3535161 < 16777215,$$

the desired contradiction. $\square$

Remark 6.45 (Exact absorption margin). The contradiction bound has margin 13242054. This payment does not claim safety at $e = 98233$ or provide an unsafe certificate. Its exact replay and independent audit are `experimental/verify_mca_full_lift_residue_zero_core_absorption_v1.py` and `experimental/audit_mca_full_lift_residue_zero_core_absorption_v1.py`.

Corollary 6.46 (Fixed-cutoff boundary-stack continuation). *Retain the pair-noncontained full-lift hypotheses and notation above. Fix a target slope budget B and $h_0 < H$, and let $P_{h_0}(e)$ be the independently truncated Johnson/mean-centered prefix through h_0 . For every $h_0 < h \leq H$, put*

$$A_h = 2h - e, \quad J_h = \left\lfloor \frac{e(A_h - c)}{A_h^2 - ec} \right\rfloor, \quad Q_h = \left\lfloor \frac{N - e - c}{m - h - c} \right\rfloor,$$

and suppose

$$2h > e, \quad A_h^2 > ec, \quad N - e > m - h > c$$

throughout this range. Define

$$(BS1) \quad F_e = P_{h_0}(e) + \sum_{h=h_0+1}^H \{1 + J_h(Q_h - 1)\}.$$

Then

$$(BS2) \quad |Z| \leq F_e + |\mathcal{T}|.$$

In particular, $F_e + N - m + 1 \leq B$ implies safety. Otherwise suppose $F_e < B$, and put

$$L_e = B - F_e + 1, \quad g_e = \left\lfloor \frac{L_e m - N}{L_e - 1} \right\rfloor, \quad u_e = g_e - c, \quad a_e = e - u_e + K.$$

If the punctured ordinary-Johnson cap

$$(BS3) \quad M_e = \left\lfloor \frac{(N - e)(m - a_e + 1 - c)}{(m - a_e + 1)^2 - (N - e)c} \right\rfloor$$

is defined and

$$(BS4) \quad eM_e + N - m + 1 \leq B,$$

then the support is safe.

Proof. The prefix handles deficits through h_0 . Fix an exact layer $h_0 < h \leq H$. Every explanation owns one slope. Relative to an anchor, its nonzero normalized codeword direction agrees with the gauged direction on at least $A_h = 2h - e$ coordinates. Distinct directions have intrinsic agreement sets meeting in at most c coordinates, so the constant-block Johnson count gives at most J_h classes. Each class and the anchor form one nonzero affine line. Its outside common core has size at most c , and outside-core packing gives Q_h members. Subtracting the repeated anchor prices the layer by $1 + J_h(Q_h - 1)$. The unpriced top union is the one synchronized line $\mathcal{T}$, proving (BS2).

If (BS2) does not pay directly and the family is unsafe, then $|\mathcal{T}| \geq L_e$ and total-core line packing forces core at least g_e . At most c core coordinates lie outside the gauged direction support, so at least u_e lie inside it. The two-anchor argument from [Corollary 6.44](#) absorbs every assigned pair of deficit at least a_e into the same line. The remaining explanations have outside agreement at least $m - a_e + 1$; (BS3) counts them, the conservative owner multiplier e prices their slopes, and (BS4) pays. $\square$

Remark 6.47 (Exact fixed-cutoff Mersenne interval). For Mersenne-31, take $h_0 = 65200$. Exact integer replay pays every support

$$98232 \leq e \leq 101155.$$

The direct branch applies through $e = 101149$; the final six supports use core absorption. At $e = 101155$,

$$\begin{aligned} F_e &= 16667033, & L_e &= 110183, & g_e &= 67446, \\ M_e &= 28, & eM_e + N - m + 1 &= 3813469. \end{aligned}$$

The final bound has slack 12963746. At adjacent $e = 101156$, the same fixed-cutoff charge is 16951223, above budget by 174008. This is a fixed-cutoff method wall, not an unsafe certificate. The Mersenne full-lift residual interval is now $101156 \leq e \leq 1044241$.

The endpoint verifier and full constant-memory C replay are `experimental/verify_mca_full_lift_fixed_cutoff_boundary_stack_v1.py` and `experimental/verify_mca_full_lift_fixed_cutoff_boundary_stack_v1.c`.

Corollary 6.48 (Residue-two repair at the fixed-cutoff wall). *Retain the pair-noncontained Mersenne-31 full-lift hypotheses. At*

$$e = 101156, \quad (s, q, H) = (33716, 2, 67439),$$

use the fixed-cutoff charge (BS1) with $h_0 = 65258$. Let $\mathcal{T}$ be the synchronized top union, and let $\mathcal{D}_1, \mathcal{D}_2$ be the exact layers of deficits $H, H - 1$. The resulting charges are

$$F = 16895280, \quad D_1 = 284224, \quad D_2 = 258385.$$

Then the complete slope family satisfies

$$|Z| \leq 16705799 < 16777215.$$

Consequently this support is safe, and the residual Mersenne full-lift interval is $101157 \leq e \leq 1044241$.

Proof. Every member of $\mathcal{T} \cup \mathcal{D}_1 \cup \mathcal{D}_2$ owns at most one slope. If $|\mathcal{T}| \geq 2$, two top anchors and any boundary member have at least

$$e - s - s - (s + 2) = K + q - 2 = K$$

common inside agreement coordinates. Thus the top union and both boundary layers lie on one affine codeword line. Removing both coarse boundary charges leaves

$$G = F - D_1 - D_2 = 16352671.$$

Unsafety would force at least $16777215 - G + 1 = 424545$ members on that line. Total-core packing then forces common core $67452 = m - 2$, with at least 67447 core coordinates inside the gauged direction support. The two-anchor absorption argument therefore synchronizes every assigned explanation of deficit at least 33715 onto the line. The remaining punctured ordinary list has agreement at least 33740 and Johnson cap 28, whence

$$|Z| \leq 101156 \cdot 28 + 981129 = 3813497.$$

Suppose next that $|\mathcal{T}| = 1$. If $|\mathcal{D}_1| \geq 2$, the top anchor and any two first-boundary members share at least

$$e - s - 2(s + 1) = K$$

inside coordinates, so $\mathcal{D}_1$ is one affine codeword line. Its outside agreement is $m - H = 15$, and outside-core packing gives $|\mathcal{D}_1| \leq 94742$. The two subcases give

$$F - D_1 + 94742 + 1 = 16705799, \quad F - D_1 + 2 = 16611058.$$

Finally suppose $\mathcal{T}$ is empty. If two first-boundary missed sets intersect, those two members and any third share at least

$$e - \{3(s + 1) - 1\} = K$$

inside coordinates. The layer is one affine line and contributes at most 94742, giving $F - D_1 + 94742 = 16705798$. Otherwise its missed sets are pairwise disjoint, so $|\mathcal{D}_1| \leq \lfloor e/(s + 1) \rfloor = 3$, giving $F - D_1 + 3 = 16611059$. These five exhaustive cases have maximum 16705799, as claimed. $\square$

Remark 6.49 (Scope and replay). At adjacent $e = 101157$, the residue resets to zero, so the preceding repair is inapplicable; no unsafe certificate is claimed. The source-bound replay and independent constant audit are `experimental/verify_mca_full_lift_fixed_cutoff_q2_anchor_repair_v1.py` and `experimental/audit_mca_full_lift_fixed_cutoff_q2_anchor_repair_v1.py`.

Corollary 6.50 (Boundary direction-class affine-line bank). *Retain the pair-noncontained full-lift hypotheses and put*

$$s = \left\lfloor \frac{e - K}{3} \right\rfloor, \quad H = e - s - 1, \quad U = \min(H, m), \quad c = K - 1.$$

Fix $h_0 < U$, and let $P_{h_0}(e)$ be the independently truncated Johnson/mean-centered prefix through h_0 . For $h_0 < h \leq U$, put

$$A_h = 2h - e, \quad J_h = \left\lfloor \frac{e(A_h - c)}{A_h^2 - ec} \right\rfloor,$$

and suppose

$$2h > e, \quad A_h > c, \quad A_h^2 > ec$$

throughout this range. Define

$$G_e = \mathbf{1}_{\{H < m\}} + \sum_{h=h_0+1}^U J_h, \quad C_e = P_{h_0}(e) + \sum_{h=h_0+1}^U (1 - J_h).$$

Then there are G_e affine explanation-line slots $\mathcal{L}_i$ such that

$$(LB1) \quad |Z| \leq C_e + \sum_{i=1}^{G_e} |\mathcal{L}_i|.$$

In particular, for target budget B , unsafety forces one slot to have size at least

$$(LB2) \quad \lambda_e = \left\lceil \frac{B - C_e + 1}{G_e} \right\rceil.$$

If $\lambda_e \geq 2$, put

$$g_e = \left\lceil \frac{\lambda_e m - N}{\lambda_e - 1} \right\rceil, \quad u_e = g_e - c, \quad a_e = e - u_e + K.$$

When the punctured ordinary-Johnson cap

$$M_e = \left\lfloor \frac{(N - e)(m - a_e + 1 - c)}{(m - a_e + 1)^2 - (N - e)c} \right\rfloor$$

is defined and $eM_e + N - m + 1 \leq B$, the support is safe.

Proof. Fix one anchor in an exact layer $h_0 < h \leq U$. Every selected explanation owns at most one slope because $2h > e$. Relative to the anchor, each nonzero normalized codeword direction agrees with the gauged direction on at least A_h inside coordinates. Distinct directions have agreement blocks meeting in at most c , so the constant-block Johnson count permits at most J_h direction classes.

Each class together with the anchor is one affine explanation line. Pad to exactly J_h classes with anchor-only slots. The exact layer then satisfies

$$|\mathcal{D}_h| = 1 - J_h + \sum_{j=1}^{J_h} |\mathcal{L}_{h,j}|.$$

If $H < m$, the cross-layer top union contributes one further affine line. Summing these identities and the prefix proves (LB1). Notice that this argument uses no outside-core denominator.

If the family is unsafe, pigeonhole gives (LB2). Total-core line packing forces a common core of size at least g_e , of which at least u_e coordinates lie inside the gauged direction support. Two anchors from the large slot synchronize every assigned explanation of deficit at least a_e onto that same line. The remaining explanations have outside agreement at least $m - a_e + 1$; the punctured ordinary-Johnson cap counts them, and the conservative owner multiplier e completes the payment. $\square$

Remark 6.51 (Exact Mersenne line-bank interval). For Mersenne-31, take $h_0 = 65272$. Exact replay pays every support

$$101157 \leq e \leq 124805.$$

At the endpoint,

$$\begin{aligned} P_{h_0} &= 1636955, & C_e &= 1604577, & G_e &= 34560, \\ \lambda_e &= 440, & g_e &= 65220, & M_e &= 126, \\ eM_e + N - m + 1 &= 16706559. \end{aligned}$$

The final margin is 70656. At adjacent $e = 124806$, the same legal compiler gives 16831491, above budget by 54276. This is a method wall, not an unsafe certificate. The Mersenne residual interval is now $124806 \leq e \leq 1044241$.

The source-bound endpoint replay, independent audit, and full constant-memory interval replay are `experimental/verify_mca_full_lift_boundary_line_bank_absorption_v1.py`, `experimental/audit_mca_full_lift_boundary_line_bank_absorption_v1.py`, and `experimental/verify_mca_full_lift_boundary_line_bank_absorption_v1.c`.

Corollary 6.52 (Recursive affine-line peeling and inside-core packing). *Assume the hypotheses of Corollary 6.50, suppose $H \geq m$, and put $Q = N - m + 1$. Fix a cutoff $b < m$ for which the prefix and every line-bank guard are valid. Let $\mathcal{Z}_r$ be the residual assigned slopes after r affine lines have been removed, with $\mathcal{Z}_0 = \mathcal{Z}$ and deficit ceiling $U_0 = m$, and charge the removed lines by rQ . Put*

$$T_r = B - rQ, \quad G_r = \sum_{h=b+1}^{U_r} J_h, \quad C_r = P_b(e) + (U_r - b) - G_r.$$

If the suffix-minimum prefix through U_r is defined and at most T_r , then the support is safe. Otherwise, whenever $G_r > 0$ and $T_r - C_r + 1 > 0$, an unsafe residual forces an affine line of size at least

$$(RP1) \quad \lambda_r = \left\lceil \frac{T_r - C_r + 1}{G_r} \right\rceil.$$

For $\lambda_r \geq 2$, define

$$(RP2) \quad g_r = \max\left(0, \left\lceil \frac{\lambda_r m - N}{\lambda_r - 1} \right\rceil\right), \quad u_r = \max(g_r - c, 0), \quad U_{r+1} = \min(U_r, e - u_r + K - 1).$$

Remove the whole selected line and repeat. The selected lines are distinct, and their inside-core sets I_i satisfy

$$|I_i| \geq u_i, \quad |I_i \cap I_j| \leq c.$$

Consequently, if at any stage

$$(RP3) \quad \sum_{i=0}^r u_i - \binom{r+1}{2} c > e,$$

the assumed unsafe family is impossible. Thus termination either at the weighted prefix or at (RP3) certifies safety.

Proof. Apply the exact-layer line-bank identity to the residual assigned family. Empty layers only improve the resulting inequality, so

$$|\mathcal{Z}_r| \leq C_r + \sum_{i=1}^{G_r} |\mathcal{L}_i|.$$

If the original family is unsafe after r line charges, then $|\mathcal{Z}_r| > T_r$, and pigeonhole gives (RP1). The guard $\lambda_r \geq 2$ makes the selected slot an actual affine line rather than a padded singleton. Total-core line packing gives g_r . At most c core coordinates lie outside the gauged direction support, and the common-core absorption argument places every residual explanation of deficit at least $e - u_r + K$ on the same line. Removing the whole line costs at most Q and gives (RP2).

Even when (RP2) does not lower the deficit ceiling, the selected line is absent from the next residual.

Write a selected line as $c_\gamma = a_i + \gamma b_i$. Its total common core consists of coordinates where $(r_0, r_1) = (a_i, b_i)$. Distinct selected lines have distinct codeword pairs. At least one of $a_i - a_j$ and $b_i - b_j$ is therefore a nonzero degree- $< K$ codeword, so two total cores, and hence two inside cores, meet in at most $c = K - 1$ coordinates. Since every inside core lies in the same e -coordinate support, first-order inclusion-exclusion gives the reverse of (RP3), proving the contradiction. $\square$

Remark 6.53 (Exact Mersenne recursive-peeling interval). For Mersenne-31, start from $b = 65304$. When its first line-bank layer fails a guard, take the least legal cutoff and add two guard layers. Exact replay checks every prefix and line-bank denominator under this printed rule. It pays all

$$124806 \leq e \leq 130198.$$

Of these 5393 supports, 3837 terminate at the weighted prefix and 1556 at (RP3), with at most five peeled lines. The first packing termination is $e = 128340$, where two inside cores give $62822 + 66574 - 5 = 129391 > 128340$. At the last support, the five inside-core lower bounds give

$$37718 + 33617 + 28204 + 20729 + 12942 - 10 \cdot 5 = 133160 > 130198.$$

At adjacent $e = 130199$, nine legal peels give packing lower bound $126052 \leq 130199$; the next residual target is 7947054, while the certified base charge is 8154082, so (RP1) no longer forces a line. This is a method wall, not an unsafe certificate. The Mersenne residual interval is now $130199 \leq e \leq 1044241$.

The endpoint replay, independent audit, and full constant-memory interval replay are `experimental/verify_mca_full_lift_recursive_line_peeling_core_packing_v1.py`, `experimental/audit_mca_full_lift_recursive_line_peeling_core_packing_v1.py`, and `experimental/verify_mca_full_lift_recursive_line_peeling_core_packing_v1.c`.

Corollary 6.54 (Joint-core charge for peeled explanation lines). *In the setting of Corollary 6.52, suppose r distinct parameterized affine explanation lines have been removed, and let g_i be their actual total common-core sizes. Put*

$$S_r = \min \left\{ r(m-1), e + \binom{r+1}{2} c \right\}, \quad q_r = \left\lfloor \frac{S_r}{m-1} \right\rfloor, \quad z_r = S_r - q_r(m-1),$$

and $f(g) = (N - g)/(m - g)$, $Q = N - m + 1$. Their total number of assigned slopes is at most

$$(JC1) \quad L_r = \begin{cases} rQ, & q_r = r, \\ \lfloor q_r Q + f(z_r) + (r - q_r - 1)f(0) \rfloor, & q_r < r. \end{cases}$$

Consequently the residual peeling target may be strengthened from $B - rQ$ to $B - L_r$.

Proof. Let I_i be the part of the i th total core inside the fixed e -coordinate gauged direction support. Distinct parameterized lines have distinct codeword pairs, so $|I_i \cap I_j| \leq c$. Each line direction is a nonzero degree- $< K$ codeword, so at most c further core coordinates lie outside that support. Therefore

$$\sum_i g_i \leq e + \binom{r}{2} c + rc = e + \binom{r+1}{2} c,$$

and $g_i \leq m - 1$ by pair noncontainment.

Off-core agreement sets on one line are disjoint, giving the rational cap $f(g_i)$. The function

$$f(g) = 1 + \frac{N - m}{m - g}$$

is increasing and convex. Moving core mass from the smaller of two interior arguments to the larger cannot decrease their sum. Repeating this exchange concentrates all available mass at $m - 1$, at most one remainder z_r , and zero. The total number of slopes is integral, so flooring the resulting rational sum gives (JC1). $\square$

Remark 6.55 (Exact Mersenne joint-charge interval). Using $B - L_r$ in the recursive line bank pays every support

$$130199 \leq e \leq 130219.$$

All 21 supports terminate by inside-core packing, using between four and thirteen lines. At the endpoint, the thresholds are 21 and then twelve copies of 18, giving

$$18393 + 12 \cdot 9736 - \binom{13}{2} \cdot 5 = 134835 > 130219.$$

At adjacent $e = 130220$, the first threshold is 20 and the next 42 are 16. Their packing lower bound is only

$$15811 + 42 \cdot 2041 - \binom{43}{2} \cdot 5 = 97018.$$

For $r = 43$, the joint core allowance is $134950 = 2(m - 1) + 44$, the exact joint charge is 1962895, and the next threshold drops to 13, whose forced-core lower bound is zero. Later thresholds cannot increase. This is a method wall, not an unsafe certificate; the Mersenne residual is now $130220 \leq e \leq 1044241$.

The endpoint replay, independent rational audit, and full interval replay are `experimental/verify_mca_full_lift_joint_core_charge_peeling_v1.py`, `experimental/audit_mca_full_lift_joint_core_charge_peeling_v1.py`, and `experimental/verify_mca_full_lift_joint_core_charge_peeling_v1.c`.

Corollary 6.56 (Lower-aware joint-core charge). *In the setting of Corollary 6.54, suppose the actual total cores also have certified lower bounds $0 \leq \ell_i \leq g_i$. Sort the ℓ_i decreasingly, start with $x_i = \ell_i$, and distribute*

$$S_r - \sum_i \ell_i$$

units by filling x_1 to $m - 1$, then x_2 , and so on. Whenever the lower bounds are feasible, the removed lines contain at most

$$(LA1) \quad L_r^{\text{low}} = \left\lfloor \sum_{i=1}^r \frac{N - x_i}{m - x_i} \right\rfloor$$

assigned slopes. Consequently the residual peeling target may be strengthened to $B - L_r^{\text{low}}$.

Proof. Since line labels do not occur in the objective, sort both the lower bounds and a feasible core vector decreasingly. This preserves the constraints: if $\ell_i \geq \ell_j$ but $g_i < g_j$, then swapping the two core values preserves both lower bounds.

For two nonterminal coordinates $x_i \geq x_j$, transfer mass from $x_j - \ell_j$ to x_i , stopping when $x_i = m - 1$ or $x_j = \ell_j$. Convexity of $f(g) = (N - g)/(m - g)$ gives

$$f(x_i + d) + f(x_j - d) \geq f(x_i) + f(x_j).$$

Iterating the exchange, and using monotonicity to spend all of the allowed mass, gives exactly the greedy vector in the statement. Equation (JC1) supplies the common sum budget S_r , so flooring the rational sum proves (LA1). $\square$

Remark 6.57 (Exact Mersenne lower-aware payment). At $e = 130220$ and $e = 130221$, the first 37 removed lines have forced total-core lower-bound runs $15816 \cdot 4, 2046 \cdot 33$. The maximizing vectors are respectively

$$18769, 15816 \cdot 3, 2046 \cdot 33 \quad \text{and} \quad 18770, 15816 \cdot 3, 2046 \cdot 33,$$

and both charges are 609. One final threshold 20 is forced. The 38 inside-core lower bounds then give

$$5 \cdot 15811 + 33 \cdot 2041 - \binom{38}{2} \cdot 5 = 142893 > e.$$

Thus both supports are safe.

At adjacent $e = 130222$, the compiler reaches 288 removed lines. The maximizing vector is $67453 \cdot 5, 1037, 0 \cdot 282$, its charge is 4910044, and the residual target 11867171 lies below the certified base 12148280. This is a method wall, not an unsafe certificate; the Mersenne residual is now $130222 \leq e \leq 1044241$.

The endpoint replay and independent rational audit are `experimental/verify_mca_full_lift_lower_aware_joint_core_charge_v1.py` and `experimental/audit_mca_full_lift_lower_aware_joint_core_charge_v1.py`.

Corollary 6.58 (Core-dichotomy capped charge). *Fix an integer absorption cutoff b . At every line selected by the recursive bank, split on its actual total core g .*

If $g \geq e + 10 - b$, every explanation of deficit at least $b + 1$ lies on that line, and the original family has size at most

$$(CD1) \quad P_e(b) + (N - m + 1),$$

where $P_e(b)$ is the exact weighted prefix through b .

In the complementary branch, every peeled core has the individual cap

$$(CD2) \quad g_i \leq G_e := e + 9 - b.$$

The lower-aware convex envelope remains valid after replacing its fill ceiling by G_e and its common budget by

$$(CD3) \quad S_r = \min \left\{ rG_e, e + \binom{r+1}{2}c \right\}.$$

Proof. A selected line has at least $g - c$ core coordinates inside the gauged direction support. An explanation with inside agreement size h and two line anchors shares at least $(g - c) + h - e$ of those coordinates. Restriction injectivity therefore puts it on the line whenever

$$h \geq e - g + c + K = e - g + 11.$$

The high-core threshold in the statement makes this at most $b + 1$, proving (CD1).

The integer complement of the high-core branch is (CD2). Under that cap, sort the forced core lower bounds and greedily spend (CD3), filling each coordinate only to G_e . The convex exchange in Corollary 6.56 applies unchanged and proves the capped joint charge. $\square$

Remark 6.59 (Exact Mersenne core-dichotomy payment). Take $b = 65450$. On $130222 \leq e \leq 130226$, the high-core bound (CD1) is at most $5161307 < 16777215$. In the complementary branch, $e = 130222, 130223$ force fourteen threshold-18 lines and

$$14 \cdot 9736 - \binom{14}{2} \cdot 5 = 135849 > e.$$

The next two supports force seventy threshold-16 lines and

$$70 \cdot 2041 - \binom{70}{2} \cdot 5 = 130795 > e.$$

Thus every support $130222 \leq e \leq 130225$ is safe.

At adjacent $e = 130226$, the first threshold is 14 and has zero forced core. After 14763 zero-lower-bound peels, the capped budget is $545032556 = 8412 \cdot 64785 + 61136$, the charge is 3199542, and the next pigeonhole threshold is one. This is a method wall, not an unsafe certificate; the Mersenne residual is now $130226 \leq e \leq 1044241$.

The endpoint replay and independent rational audit are `experimental/verify_mca_full_lift_core_dichotomy_capped_charge_v1.py` and `experimental/audit_mca_full_lift_core_dichotomy_capped_charge_v1.py`.

Corollary 6.60 (Exact-layer slot-core incidence). *Suppose a line slot selected by the recursive bank belongs to exact layer h and contains $\lambda \geq 2$ explanations. If u is their common core inside the gauged e -coordinate direction support, then*

$$(EL1) \quad u \geq \left\lceil \frac{\lambda h - e}{\lambda - 1} \right\rceil.$$

Consequently, in the low-core branch of Corollary 6.58, (EL1) may be used as the forced lower bound in the capped convex charge.

Proof. Write the parameterized affine explanation line as $c_\gamma = a + \gamma b$. At an inside coordinate outside the common core, the equation

$$r_0 + \gamma r_1 = a + \gamma b$$

holds for at most one line parameter γ . Every one of the λ members has at least h inside agreements, whereas a core coordinate is counted λ times. Incidence counting in the fixed inside support therefore gives

$$\lambda h \leq e + (\lambda - 1)u,$$

which is equivalent to (EL1). This is already an inside-core bound and does not spend the outside zero allowance. In the calibrated support interval the parent ceiling has $H \geq m$, so the optional cross-layer top slot is absent and every recursive-bank slot has one exact-layer owner. $\square$

Remark 6.61 (Exact Mersenne layer-aware payment). Retain absorption cutoff 65450. In the complementary capped-core branch, legal cutoffs 65516, ..., 65521 and (EL1) force three line cores of the following sizes:

e	b	λ	u
130226	65516	14	60540
130227	65516	5	49340
130228	65517	13	60126
130229	65517	4	43948
130230	65518	11	59048
130231	65518	4	43949
130232	65519	9	57431
130233, 130234	65520	7	54736
130235, 130236	65521	4	43951

The displayed λ is the forced minimum slot size and $h = b + 1$ is the minimum possible exact layer. In this range the right side of (EL1) is nondecreasing in both variables, so these substitutions are conservative. Distinct selected lines have pairwise inside-core intersection at most $c = 5$. Exact capped-charge replay therefore makes three lines violate inside-core packing on every support $130226 \leq e \leq 130236$; the smallest printed lower bound is

$$3 \cdot 43948 - \binom{3}{2} \cdot 5 = 131829 > 130229.$$

At adjacent $e = 130237$, cutoff 65521 forces only size-two affine explanation slots. Equation (EL1) gives $u = 807$, and the first-order packing expression has maximum

$$\max_s \left\{ 807s - 5 \binom{s}{2} \right\} = 65529 < e.$$

After 7583 such lower bounds the capped charge is 882245, the residual target is 15894970, and the next threshold is one. This is a method wall, not an unsafe certificate. The Mersenne full-lift residual is now $130237 \leq e \leq 1044241$; primitive shift-pair control is the next structural obstruction.

The endpoint replay and independent rational audit are `experimental/verify_mca_full_lift_exact_layer_slot_core_packing_v1.py` and `experimental/audit_mca_full_lift_exact_layer_slot_core_packing_v1.py`.

Corollary 6.62 (First-wall interpolation common-factor router). *At Mersenne full-lift support $e = 130237$, give (X, Y, Z) weights $(1, 5, 5)$ and let I_{264} be the vector space of weighted-degree at most 264 polynomials vanishing at all gauged inside received points $(x, r_0(x), r_1(x))$. Then*

$$(IF1) \quad \dim I_{264} \geq 938.$$

If the selected family is unsafe, the nonzero members of I_{264} have a common factor of positive (Y, Z) -degree over the algebraic closure of $\mathbb{F}(X)$.

Proof. The weight- $(1, 5, 5)$ monomial count is

$$\sum_{s=0}^{52} (s+1)(264-5s+1) = 131175.$$

Evaluation at the 130237 received points has rank at most 130237, which proves (IF1).

In the complementary capped-core branch, every selected line at cutoff 65521 has exact layer at least 65522 and at least two members. By Corollary 6.60, its inside core has size at least 807. Write the line as $c_\gamma = a + \gamma b$. Its inside core is precisely the set on which $(r_0, r_1) = (a, b)$.

For $Q \in I_{264}$, substitution gives a polynomial $Q(X, a(X), b(X))$ of degree at most 264. It has at least 807 distinct roots, hence vanishes identically. Thus every selected pair (a, b) is a common $\mathbb{F}(X)$ -rational zero of the interpolation kernel.

After 2704 removed lines, the capped convex charge is 132203 and the residual target is 16645012. With base 13961576 and group count 1933560, the next strict pigeonhole threshold remains two. Therefore an unsafe family forces 2705 distinct polynomial-pair cores.

Suppose instead that the kernel has no common factor of positive (Y, Z) -degree. Over the algebraic closure of $\mathbb{F}(X)$, two generic kernel members are then coprime. Their total (Y, Z) -degrees are at most $\lfloor 264/5 \rfloor = 52$. Affine Bezout bounds their common zero set by $52^2 = 2704$, contradicting the 2705 distinct selected pairs. $\square$

Remark 6.63 (Common-factor residual). The corollary pays the coprime interpolation branch but does not identify the residual factor with a split pencil. A further theorem must classify its ruled components or charge their line/core mass. The Mersenne support residual remains $130237 \leq e \leq 1044241$, with its first support reduced to this common-factor branch.

The exact calibration and independent audit are `experimental/verify_mca_full_lift_interpolation_common_factor_router_v1.py` and `experimental/audit_mca_full_lift_interpolation_common_factor_router_v1.py`.

Corollary 6.64 (Common-factor mass concentration). *In the unsafe branch of Corollary 6.62, let P be the full common interpolation factor and put $d = \deg_{Y,Z} P$. At least*

$$(FM1) \quad 7583 - (52 - d)^2 \geq 4982$$

distinct selected polynomial pairs lie on P . Moreover,

$$(FM2) \quad \#\{x \in E : P(x, r_0(x), r_1(x)) = 0\} \geq 126188.$$

Thus at most 4049 of the 130237 inside coordinates are exceptional to the factor relation.

Proof. After 7582 low-core lines have been removed, the capped charge is 881897, the target is 15895318, and the strict line-bank threshold is still two. Hence unsafety forces 7583 distinct polynomial-pair cores.

Divide every interpolation-kernel member by the full gcd P . The cofactor family has gcd one and total (Y, Z) -degree at most $52 - d$. Two generic cofactors are coprime, so affine Bezout permits at most $(52 - d)^2$ selected pairs outside $P = 0$. This proves (FM1), whose right side is minimized at $d = 1$.

For each of t captured pairs, its inside core U_i has size at least $u = 807$ and lies in the set on the left of (FM2). Distinct pair cores meet in at most five coordinates. If m_x is the core-incidence multiplicity, then

$$\sum_x m_x \geq tu, \quad \sum_x \binom{m_x}{2} \leq 5 \binom{t}{2}.$$

Cauchy's inequality therefore gives

$$\left| \bigcup_i U_i \right| \geq \left\lceil \frac{tu^2}{u + 5(t-1)} \right\rceil.$$

This is increasing in t . At $t = 4982$ it equals 126188, proving (FM2). $\square$

Remark 6.65. The set in (FM2) is not asserted to be a common core shared by all captured sections. Nor does the corollary assert that P is irreducible, rational, or a split pencil. Its output is a degree-at-most-52 relation holding on at least 96.89% of the inside support and carrying at least 4982 degree-five polynomial sections.

The exact calibration and all-degree audit are `experimental/verify_mca_full_lift_common_factor_mass_router_v1.py` and `experimental/audit_mca_full_lift_common_factor_mass_router_v1.py`.

Corollary 6.66 (Common-factor weighted-degree bound). *In the setting of Corollary 6.64, the primitive full gcd satisfies*

$$(WD1) \quad \text{wdeg}_{(1,5,5)} P \leq 217, \quad \deg_{Y,Z} P \leq 43.$$

Consequently, if $\deg_{Y,Z} P \geq 2$, then at least 5083 selected degree-five polynomial pairs lie on P , and

$$(WD2) \quad \#\{x \in E : P(x, r_0(x), r_1(x)) = 0\} \geq 126266.$$

Thus this higher-degree branch has at most 3971 exceptional inside coordinates.

Proof. Work over the algebraic closure. Gauss's lemma makes the primitive representative P a polynomial divisor of every member of I_{264} . If $w = \text{wdeg}_{(1,5,5)} P$, division by P embeds the at-least-938-dimensional kernel into weighted degree at most $264 - w$. The exact weighted-monomial count

$$M(D) = \sum_{s=0}^{\lfloor D/5 \rfloor} (s+1)(D - 5s + 1)$$

satisfies

$$M(46) = 935 < 938 \leq 990 = M(47).$$

Hence $264 - w \geq 47$, proving $w \leq 217$ and $\deg_{Y,Z} P \leq \lfloor 217/5 \rfloor = 43$.

In the higher-degree branch, the factor-mass bound is minimized at $d = 2$:

$$7583 - (52 - d)^2 \geq 7583 - 50^2 = 5083.$$

Applying the same core-incidence inequality as in Corollary 6.64 gives

$$\left\lceil \frac{5083 \cdot 807^2}{807 + 5(5083 - 1)} \right\rceil = 126266,$$

which proves (WD2). □

Remark 6.67. The full gcd in this corollary may be reducible. No irreducible-component classification or safety claim is implicit. The exact threshold replay and independent enumeration are `experimental/verify_mca_full_lift_common_factor_weighted_degree_bound_v1.py` and `experimental/audit_mca_full_lift_common_factor_weighted_degree_bound_v1.py`.

Corollary 6.68 (Base-field component descent). *Retain the higher-degree branch of Corollary 6.66. Over the deployed Mersenne field*

$$\mathbb{F} = \mathbb{F}_{p^4}, \quad p = 2^{31} - 1,$$

let P_{bf} be the product of the geometric irreducible components of P that are individually defined over $\mathbb{F}(X)$. At least 5079 selected polynomial pairs lie on P_{bf} , and one absolutely irreducible $\mathbb{F}(X)$ -component contains at least 132 selected pairs. Moreover,

$$(FD1) \quad \#\{x \in E : P_{\text{bf}}(x, r_0(x), r_1(x)) = 0\} \geq 126263.$$

Thus at most 3974 inside coordinates are exceptional to the base-field component union.

Proof. Factor the radical of P geometrically as $\prod_i R_i$, with $\delta_i = \deg_{Y,Z} R_i$ and $\sum_i \delta_i \leq d \leq 43$. Since $p > d$, a component field of definition cannot have nontrivial purely inseparable degree. If R_i is not defined over $\mathbb{F}(X)$, it therefore has a distinct conjugate R_i^σ . Every $\mathbb{F}(X)$ -rational point of R_i also lies on R_i^σ , so Bezout bounds their number by δ_i^2 .

Assign each selected pair outside the base-field component union to one non-base-field component containing it. At most

$$\sum_{R_i \text{ mnot over } \mathbb{F}(X)} \delta_i^2 \leq d^2$$

pairs are assigned. Hence the retained number is at least

$$7583 - (52 - d)^2 - d^2 = 4879 + 104d - 2d^2.$$

This concave expression has endpoint values 5079 and 5653 on $2 \leq d \leq 43$, proving the uniform 5079 bound. There are at most d components, and exact integer minimization gives

$$\min_{2 \leq d \leq 43} \left\lceil \frac{7583 - (52 - d)^2 - d^2}{d} \right\rceil = 132.$$

Finally, the core-incidence bound at $t = 5079$ gives

$$\left\lceil \frac{5079 \cdot 807^2}{807 + 5(5079 - 1)} \right\rceil = 126263,$$

which proves (FD1). $\square$

Remark 6.69. This is a base-field normalization, not an irreducibility or split-pencil theorem. The exact degree scan and independent degree-partition audit are `experimental/verify_mca_full_lift_common_factor_base_field_descent_v1.py` and `experimental/audit_mca_full_lift_common_factor_base_field_descent_v1.py`.

Corollary 6.70 (Linear factor gives a projective star). *In the setting of Corollary 6.64, suppose the full common interpolation factor has total (Y, Z) -degree one. Then, after projective rescaling, it has a primitive equation*

$$(LS1) \quad P(X, Y, Z) = AY + BZ + C(X), \quad A, B \in \mathbb{F}, \quad (A, B) \neq (0, 0), \quad \deg C \leq 5.$$

All captured affine explanation lines share one $\mathbb{F}$ -rational projective slope-codeword center. If $A \neq 0$, the center is finite:

$$\gamma_* = B/A, \text{quad}c_* = -C/A.$$

If $A = 0$, every captured line has the same direction codeword $-C/B$, the center at projective slope infinity.

Proof. Initially work over the algebraic closure and write $P = A(X)Y + B(X)Z + C(X)$ primitively. The factor-mass corollary supplies at least 4982 distinct sections $(a_i, b_i) \in \mathbb{F}[X]_{<6}^2$. Subtracting two section equations and removing $\gcd(A, B)$ gives

$$(LS2) \quad (a_i, b_i) = (a_0 + Bt_i, b_0 - At_i), \quad \deg t_i \leq s := 5 - \max(\deg A, \deg B).$$

Indeed, one section shows that $\gcd(A, B)$ also divides C , so primitivity makes it a unit. In particular $\deg A, \deg B \leq 5$.

Since A and B are coprime, the received pair induces a scalar word τ on every inside coordinate: use $(r_0 - a_0)/B$ where $B \neq 0$, and $-(r_1 - b_0)/A$ otherwise. Each t_i agrees with τ on its core of at least 807 coordinates. The ordinary constant-block Johnson cap for degree s is

$$J_s = \left\lfloor \frac{130237(807 - s)}{807^2 - 130237s} \right\rfloor.$$

For $s = 0, 1, 2, 3, 4$ its exact values are 161, 201, 268, 401, 802. If either A or B were nonconstant, then $s \leq 4$, contradicting $4982 > 802$. Hence A, B are constants and $\deg C \leq 5$.

If $A \neq 0$, two distinct captured $\mathbb{F}$ -rational pairs give $B/A = -(a_i - a_j)/(b_i - b_j) \in \mathbb{F}$, and then $-C/A = a_i + (B/A)b_i \in \mathbb{F}[X]_{<6}$. Every captured line passes through that finite center. If $A = 0$, the common direction $-C/B$ lies in $\mathbb{F}[X]_{<6}$, proving the projective-infinity case. $\square$

Remark 6.71. This identifies the degree-one branch with the primitive projective-star shape; it does not prove the star population bound. The complementary common-factor branch has (Y, Z) -degree at least two.

The exact Johnson replay and independent audit are `experimental/verify_mca_full_lift_linear_factor_projective_star_router_v1.py` and `experimental/audit_mca_full_lift_linear_factor_projective_star_router_v1.py`.

Theorem 6.72 (Full-explanation lifted-rank gauge dichotomy). *Let $C \leq \mathbb{F}^D$ have dimension K , let $r_\gamma = r_0 + \gamma r_1$, and let $Z \subseteq \mathbb{F}$ contain distinct slopes. For every $\gamma \in Z$, choose $c_\gamma \in C$, and assume that the c_γ have full affine rank K . Fix $\gamma_0 \in Z$ and put*

$$V = \text{span}\{(\gamma - \gamma_0, c_\gamma - c_{\gamma_0}) : \gamma \in Z\} \leq \mathbb{F} \oplus C.$$

Then $\dim V \in \{K, K+1\}$, and exactly one of the following holds.

- (1) *If $\dim V = K$, there is a unique nonzero linear functional $\ell : C \rightarrow \mathbb{F}$ such that $V = \{(\ell(u), u) : u \in C\}$. The gauge*

$$r_1 \mapsto r_1 - b, \quad c_\gamma \mapsto c_\gamma - \gamma b$$

drops the explanation affine rank to exactly $K-1$ precisely for $b \in C$ satisfying $\ell(b) = 1$.

Thus the rank-dropping gauges form an affine hyperplane in C .

- (2) *If $\dim V = K+1$, then $V = \mathbb{F} \oplus C$, and every $b \in C$ leaves the explanation affine rank equal to K .*

If, in addition, every chosen maximal agreement support is same-support pair-noncontained, then $r_1 \notin C$, and the selected error vectors $r_\gamma - c_\gamma$ have affine rank exactly $\dim V$.

Proof. Projection of V to C is onto because the explanations have full affine rank. Hence $K \leq \dim V \leq K+1$. If $\dim V = K$, projection is an isomorphism and V is the graph of a unique functional ℓ . It is nonzero because the slopes are not all equal. On anchored differences, the gauge is the map

$$T_b : V \longrightarrow C, \quad T_b(a, u) = u - ab.$$

For $V = \{(\ell(u), u) : u \in C\}$, this map is $u \mapsto u - \ell(u)b$. It has rank $K-1$ exactly when $\ell(b) = 1$, in which case its image is $\ker \ell$; otherwise it is invertible. If $\dim V = K+1$, then $V = \mathbb{F} \oplus C$, and T_b is surjective for every b .

Suppose now that $r_1 \in C$. On any selected agreement support, r_1 explains the direction and $c_\gamma - \gamma r_1 \in C$ explains the base, contradicting pair noncontainment. Thus $r_1 \notin C$. The linear map

$$(a, u) \mapsto ar_1 - u$$

is consequently injective on $\mathbb{F} \oplus C$, and it sends the anchored lifted differences to the anchored error-vector differences. Their affine rank is therefore $\dim V$. $\square$

Theorem 6.73 (Full-lift near-MDS extension reduction). *Retain the hypotheses of [Theorem 6.72](#), assume $\dim V = K+1$, write $N = |D|$, and put*

$$W = C + \langle r_1 \rangle, \quad e = \min_{b \in C} \text{wt}(r_1 - b).$$

Then $\dim W = K+1$, and the generalized Hamming weights of W are

$$(NM) \quad d_1(W) = e, \quad d_j(W) = N - K + j - 1 \quad (2 \leq j \leq K+1).$$

The selected errors $r_\gamma - c_\gamma$ lie in one affine coset of W , have full affine rank $K+1$, and determine their slopes uniquely. Moreover, if every selected maximal agreement support has size greater than K , restriction of W to each such support is injective.

Proof. Pair noncontainment gives $r_1 \notin C$, so W has dimension $K+1$. Interpolation on K coordinates gives $e \leq N - K$. Every nonzero word of W is either a nonzero RS word, of weight at least $N - K + 1$, or a nonzero scalar multiple of $r_1 - b$, of weight at least e ; a nearest b attains e . Thus $d_1(W) = e$.

Let $A \leq W$ have dimension $j \geq 2$. If $A \leq C$, the RS MDS hierarchy gives $|\text{supp } A| \geq N - K + j$. Otherwise $\dim(A \cap C) = j - 1$, and that RS subcode has at most $K - j + 1$ common zero coordinates. Hence

$$|\text{supp } A| \geq N - K + j - 1.$$

The generalized Singleton bound for the $[N, K + 1]$ code W is the reverse inequality, proving (NM).

The map $(a, u) \mapsto ar_1 - u$ is an isomorphism $\mathbb{F} \oplus C \rightarrow W$, so the full lifted rank is exactly the error affine rank and equal errors have equal slopes. Finally, a nonzero $ar_1 - u \in W$ vanishing on a selected support either contradicts the RS root bound when $a = 0$, or supplies simultaneous direction and base explanations when $a \neq 0$. Pair noncontainment excludes both cases. $\square$

Remark 6.74 (Exact top-rank split at the first shortened rows). The gauge preserves agreement supports, pair noncontainment, and $e = \min_{b \in C} \text{wt}(r_1 - b)$. Therefore the $\dim V = K$ branch may use [Theorem 6.11](#) at explanation rank $K - 1$. Combining this with the already separate low-support and recursive-direction gates gives the following exact routing table; q is explanation affine rank and $h = \dim V$, equivalently the error affine rank.

row	q	h	currently paid support range
KoalaBear	14	14	$e \leq 64047$ or $e \geq 992852$
KoalaBear	14	15	$e \leq 96150$ or $e \geq 1044239$
Mersenne-31	6	6	$e \leq 65454$ or $e \geq 1037876$
Mersenne-31	6	7	$e \leq 130236$ or $e \geq 1044242$

The proper-subspace/gauge improvement is only in the $h = K$ rows. The separate full-lift continuations extend the printed low walls in the $h = K + 1$ rows, but their remaining middle intervals stay unpaid. After the sparse-direction, centered-Gram, exact-layer, and interpolation continuations, the full-lift intervals are $96151 \leq e \leq 1044238$ for KoalaBear and $130237 \leq e \leq 1044241$ for Mersenne-31. The latter first support is not paid: [Corollary 6.62](#) routes it to the separately declared common-factor branch. [Theorem 6.73](#) shows that every generalized weight from rank two onward is already MDS-sharp. Even at the best endpoint $e = N - K$, the generic corrected compiler gives

$$743896698428332665 > 274980728111395087 \quad \text{and} \quad 219426634 > 16777215$$

on KoalaBear and Mersenne-31, respectively. Closing these cells therefore requires structure of the codimension-one RS extension beyond its weight hierarchy, or a sharper row-specific affine-list theorem. A finite-field audit and the exact adjacent arithmetic are in `experimental/verify_mca_full_explanation_lifted_rank_gauge_dichotomy_v1.py`.

Theorem 6.75 (Directional Johnson ray compiler). *Retain the hypotheses and notation of [Theorem 4.5](#), but allow $\nu \geq 0$. Choose a lift $b_1 \in \mathbb{F}^U$ of y_1 , put $C_U = \ker H_U$, and define the lift-invariant directional agreement*

$$(6.38) \quad \alpha_U(y_1) = \max_{c \in C_U} |\{x \in U : b_1(x) + c(x) = 0\}|.$$

If $h = N - t$ and

$$(6.39) \quad h^2 > N\alpha_U(y_1),$$

then

$$(6.40) \quad |Z| \leq \left\lfloor \frac{N(h - \alpha_U(y_1))}{h^2 - N\alpha_U(y_1)} \right\rfloor.$$

Moreover, with

$$d_U(y_1) = \min\{\text{wt}(e) : e \in \mathbb{F}^U, H_U e = y_1\},$$

one has

$$(6.41) \quad \alpha_U(y_1) = N - d_U(y_1) \geq \nu.$$

Equality holds exactly when the direction syndrome has no representation on fewer than R columns of U .

Proof. Changing the lift b_1 by a kernel word merely permutes the maximization in (6.38), so $\alpha_U(y_1)$ is intrinsic. Choose a lift b_0 of y_0 . For every slope write

$$e_\gamma = b_0 + \gamma b_1 + k_\gamma, \quad k_\gamma \in C_U,$$

and choose an h -set A_γ of zero coordinates of e_γ . On A_γ , the codeword $-k_\gamma$ agrees with $b_0 + \gamma b_1$. If $\gamma \neq \delta$, then on $A_\gamma \cap A_\delta$,

$$b_1 = \frac{-k_\gamma + k_\delta}{\gamma - \delta},$$

which is a codeword of C_U . Hence $|A_\gamma \cap A_\delta| \leq \alpha_U(y_1)$. The same Johnson incidence count as in Theorem 4.7, with $\alpha_U(y_1)$ in place of $\nu - 1$, gives (6.40).

Finally, the words $b_1 + c$, $c \in C_U$, are exactly all lifts of y_1 supported in U . Maximizing their number of zero coordinates is therefore equivalent to minimizing their weight, proving the equality in (6.41). Any R columns of H_U span syndrome space, so $d_U(y_1) \leq R$ and $\alpha_U(y_1) \geq N - R = \nu$. The equality criterion is immediate. $\square$

Corollary 6.76 (Rank-regular fixed-union ray bound). *In Theorem 6.75, suppose the direction syndrome is rank-regular on U , meaning $d_U(y_1) = R$. If $(R + \nu - t)^2 > (R + \nu)\nu$, then*

$$(6.42) \quad |Z| \leq \left\lfloor \frac{(R + \nu)(R - t)}{(R + \nu - t)^2 - (R + \nu)\nu} \right\rfloor.$$

Thus every failure of this stronger low-nullity payment is either a high-nullity chart or an explicit direction-rank-defect cell $d_U(y_1) < R$.

Proof. By (6.41), rank regularity is exactly $\alpha_U(y_1) = \nu$. Substitute this value into (6.40). $\square$

Corollary 6.77 (Quantitative direction-rank-defect payment). *Retain the notation of Theorem 6.75, and write*

$$d = R - t, \quad j = R - d_U(y_1), \quad N = R + \nu, \quad D_0 = (d + \nu)^2 - N\nu.$$

Thus $j \geq 0$ measures the direction-rank defect and $\alpha_U(y_1) = \nu + j$. Whenever $D_0 - Nj > 0$,

$$(6.43) \quad |Z| \leq \left\lfloor \frac{N(d - j)}{D_0 - Nj} \right\rfloor.$$

For an integer budget $B > 1$, define

$$(6.44) \quad J_B(\nu) = \min \left\{ d - 1, \left\lfloor \frac{D_0 - 1}{N} \right\rfloor, \left\lfloor \frac{BD_0 - Nd}{N(B - 1)} \right\rfloor \right\}.$$

If $0 \leq j \leq J_B(\nu)$, then the whole fixed-union ray chart costs at most B slopes. Hence the residual direction cell at a fixed nullity can be stated quantitatively as $j > J_B(\nu)$, rather than merely as $j > 0$.

Proof. By (6.41), $\alpha_U(y_1) = N - d_U(y_1) = \nu + j$. Substitution into (6.40) gives (6.43). The second term in (6.44) makes its denominator positive, and the first makes its numerator positive. The final term is exactly the rearrangement

$$N(d - j) \leq B(D_0 - Nj).$$

Thus the real fraction in (6.43) is at most B , and so is its floor. $\square$

Theorem 6.78 (Paving-basis direction-defect ray compiler). *Retain the fixed-union MCA notation of Theorems 4.5 and 6.75. Write*

$$N = R + \nu, \quad h = N - t, \quad d = R - t, \quad j = R - d_U(y_1).$$

Assume $0 \leq j < d$, equivalently that the direction word is outside the agreement ball on this chart. Then

$$(6.45) \quad |Z| \leq \left\lfloor \frac{(\nu + 1) \binom{R + \nu}{\nu + 1}}{(R - t - j) \binom{R + \nu - t}{\nu}} \right\rfloor.$$

This estimate is independent of the Johnson denominator and is strongest for small nullity, including large direction defect.

Proof. Choose a minimum-weight lift $q \in \mathbb{F}^U$ of the direction syndrome, so $\text{wt}(q) = R - j$, and choose a basis $k_1, \dots, k_\nu$ of the MDS kernel $C_U = \ker H_U$. In the parameter space with coordinates $(\gamma, \lambda_1, \dots, \lambda_\nu)$, agreement at coordinate x has normal

$$v_x = (q_x, k_1(x), \dots, k_\nu(x)).$$

Every set of at most ν such normals is independent, because its projection to the last ν coordinates is a set of columns of a generator matrix of the $[R + \nu, \nu, R + 1]$ MDS code C_U .

Every rank- ν flat of these normals contains at most $\nu + j$ coordinates. Indeed, such a flat is annihilated by a nonzero functional (δ, a) . If $\delta \neq 0$, the vanishing coordinates are zeros of a lift $\delta q + c$ of δy_1 , which has weight at least $R - j$. If $\delta = 0$, they are zeros of a nonzero MDS kernel word and number at most $\nu - 1$.

Fix a slope point and choose h incident coordinate hyperplanes. Put $r = \nu + 1$ and $\alpha = \nu + j$. Every $(r - 1)$ -subset of their normals is independent. A dependent r -subset is counted by each of its r subsets of size $r - 1$, while the closure of any such subset contains at most α coordinate normals. Hence the number of independent r -subsets among the h incident normals is at least

$$\binom{h}{r} - \frac{\alpha - r + 1}{r} \binom{h}{r - 1} = \frac{h - \alpha}{r} \binom{h}{r - 1} = \frac{d - j}{\nu + 1} \binom{R + \nu - t}{\nu}.$$

Globally there are at most $\binom{R + \nu}{\nu + 1}$ independent coordinate bases, and each basis cuts out at most one parameter point. Double counting proves (6.45). $\square$

Corollary 6.79 (Exact paving defect threshold). *For an integer budget $B \geq 0$, define*

$$(6.46) \quad J_B^{\text{pav}}(\nu) = \max \left\{ -1, d - 1 - \left\lfloor \frac{(\nu + 1) \binom{R + \nu}{\nu + 1}}{(B + 1) \binom{R + \nu - t}{\nu}} \right\rfloor \right\}.$$

Every separated fixed-union MCA chart with $0 \leq j \leq J_B^{\text{pav}}(\nu)$ costs at most B slopes. Combining this with Corollary 6.77, the exact certified defect range is

$$0 \leq j \leq \max\{J_B(\nu), J_B^{\text{pav}}(\nu)\}.$$

Proof. The floor in (6.45) is at most B exactly when its numerator is strictly smaller than $B + 1$ times its denominator. Solving this integer inequality for j gives (6.46). $\square$

7. MINIMUM-DISTANCE SHARPENING AND THE EXACT PRIMITIVE WEDGE

Retain the notation of Theorem 6.4. Let

$$\mu = d_1(C') = \min_{0 \neq c \in C'} |\text{supp}(c)|.$$

The generalized Singleton bound gives

$$\mu \geq n - K + 1 = R + 1, \quad R = n - K.$$

The general flag theorem uses only the ambient MDS lower bound for each flag-annihilator direction. Every such direction actually belongs to C' , so the true minimum support gives a sharper resource count.

Theorem 7.1 (Minimum-distance enhanced flag resources). *Retain the hypotheses and notation of Theorem 6.4. For a profile $\mathbf{b} = (b_1, \dots, b_r)$ define*

$$(7.1) \quad P_{\mathbf{b}}^{\sharp} = \prod_{i=1}^j \left(f_i - t + b_0 + \sum_{q=1}^r b_q \right),$$

$$(7.2) \quad Q_{\ell}^{\sharp}(\mathbf{b}) = \mu - t + b_0 + \sum_{q > \ell} b_q,$$

$$(7.3) \quad A_{\ell}(\mathbf{b}) = e_{\ell} - b_{\ell}, \quad B_{\ell}^{\sharp}(\mathbf{b}) = (Q_{\ell}^{\sharp}(\mathbf{b}) - e_{\ell} + b_{\ell})_+.$$

For $\sigma \in \{T, U\}^r$, put

$$(7.4) \quad C_\sigma^\sharp(\mathbf{b}) = P_\mathbf{b}^\sharp \prod_{\ell: \sigma_\ell = T} A_\ell(\mathbf{b}) \prod_{\ell: \sigma_\ell = U} B_\ell^\sharp(\mathbf{b}),$$

$$(7.5) \quad G_\sigma = d_0^j \prod_{\ell: \sigma_\ell = T} e_\ell \prod_{\ell: \sigma_\ell = U} (d_{\ell-1} - (j + \ell - 1)).$$

Then

$$(7.6) \quad \boxed{\sum_{\mathbf{b}} L_{\mathbf{b}} C_\sigma^\sharp(\mathbf{b}) \leq G_\sigma}$$

for every binary resource pattern. Consequently every nonnegative family x_σ satisfying

$$\sum_{\sigma} x_\sigma C_\sigma^\sharp(\mathbf{b}) \geq 1 \quad \text{for every profile } \mathbf{b}$$

gives the whole-chart bound

$$(7.7) \quad \boxed{L \leq \left\lfloor \sum_{\sigma} x_\sigma G_\sigma \right\rfloor}.$$

The tensor and all-extender corollaries remain valid after replacing $w + 1 + b_0$ by $\mu - t + b_0$ and replacing every old-support factor by (7.3).

Proof. Fix a flag level ℓ and an adapted coordinate basis through $V_{\ell-1}$. Its annihilator in V_ℓ is generated by a nonzero word

$$h_\ell \in V_\ell \setminus V_{\ell-1} \subseteq C'.$$

Therefore $\text{wt}(h_\ell) \geq \mu$. The word vanishes outside U_ℓ . The number of mismatches outside U_ℓ is

$$b_{\text{out}} = b_0 + \sum_{q > \ell} b_q,$$

and all of those coordinates are zeros of h_ℓ . Among the zero coordinates of h_ℓ , at most

$$(n - \mu) - b_{\text{out}}$$

are agreements. Since every listed point has at least m agreements, at least

$$m - ((n - \mu) - b_{\text{out}}) = \mu - t + b_{\text{out}} = Q_\ell^\sharp(\mathbf{b})$$

agreement coordinates are nonzeros of h_ℓ and hence extend the adapted basis. The $A_\ell = e_\ell - b_\ell$ agreeing coordinates in the new layer all extend, so at least $B_\ell^\sharp$ extending agreements lie in the old support. The global basis capacities are unchanged. The double count and finite dual are therefore identical to the proof of Theorem 6.4 with the sharpened factors. $\square$

Corollary 7.2 (Minimum-word all-extender bound). *Choose $V_0 = \langle h \rangle$, where h is a minimum-weight word of C' . If $s = \dim C'$, then*

$$(7.8) \quad \boxed{L \leq \left\lfloor \frac{\mu \prod_{\ell=1}^{s-1} (d_\ell - \ell)}{(\mu - t + b_0)^s} \right\rfloor \leq \left\lfloor \frac{\mu (n - s + 1)^{s-1}}{(\mu - t + b_0)^s} \right\rfloor}.$$

Writing $\mu = R + a$, $a \geq 1$, the soft form is

$$(7.9) \quad L \leq \left\lfloor \frac{(R + a)(n - s + 1)^{s-1}}{(w + a + b_0)^s} \right\rfloor.$$

Proof. A generalized-weight minimizer is support-saturated by Lemma 6.2, so the minimum-weight line admits a support-saturated flag. Apply the all-extender form of Theorem 7.1. The initial basis count is μ , and $\mu - t + b_0$ is a lower bound at every extension level. Since every remaining support layer is nonempty and $d_{s-1} \leq n$, one has $d_\ell - \ell \leq n - s + 1$. $\square$

Proposition 7.3 (Exact active minimum-support transitions). *At $b_0 = 0$, (7.9) pays every first-unpaid affine chart at and above the following support excesses:*

row	s	a_{paid}	bound at a_{paid}	bound at $a_{\text{paid}} - 1$
<i>KoalaBear MCA</i>	15	70832	274974324324469113	275003903412638889
<i>KoalaBear list</i>	15	70832	274974569967317804	275004149082130935
<i>Mersenne-31 MCA</i>	7	111529	16776653	16777295
<i>Mersenne-31 list</i>	7	111529	16776668	16777309

Thus a surviving zero-mismatch chart has $a \leq 70831$ on *KoalaBear* or $a \leq 111528$ on *Mersenne-31*. Positive b_0 only shrinks this residual.

Proposition 7.4 (Exact mismatch curve). *For selected support excesses, the least outside mismatch paid by (7.9) is*

	$a = 1$	$a = 1000$	$a = 10000$	$a = 50000$	$a = 70000$
<i>KoalaBear MCA/list</i>	70230	69240	60318	20659	825
<i>Mersenne-31 MCA/list</i>	108962	107987	99202	60141	40599

The curve is strictly decreasing in the actual direction-code support excess.

Remark 7.5 (Combination with the actual first generalized weight). Let the full support have size $d_s = R + \nu$. Replacing the ambient lower bound in the common-zero Johnson compiler by $d_1(C') = R + a$ gives

$$(7.10) \quad L \leq \left\lfloor \frac{(R + \nu)(w + a + b_0)}{(w + \nu + b_0)^2 - (R + \nu)(\nu - a)} \right\rfloor$$

when the denominator is positive. Combining (7.10) with the nonuniform form of (7.8) gives the exact residual wedge in (a, ν, b_0) .

Theorem 7.6 (Fixed-mismatch puncturing compiler). *Let $\mathcal{L} \subseteq \mathcal{A}$ be a family of codewords, each agreeing with a received word u on at least m coordinates. Suppose a fixed set $Z \subseteq D$, $|Z| = z$, is a mismatch set for every member of $\mathcal{L}$. Then $z \leq t$, and if the affine span of $\mathcal{L}$ has direction dimension at most s , then*

$$(7.11) \quad |\mathcal{L}| \leq \left\lfloor \frac{\binom{R-z+s}{s}}{\binom{w+s}{s}} \right\rfloor.$$

Proof. Every listed word has at most t mismatches, so $z \leq t$. Puncture the coordinates in Z . Since $z < t < R = n - K$, evaluation of degree- $< K$ polynomials remains injective on $D \setminus Z$, no agreement is deleted, and the affine direction dimension does not increase. Apply the recursive affine-span compiler to the punctured code of length $n - z$. $\square$

Corollary 7.7 (Active puncture thresholds). *At the first unpaid affine dimensions, a subfamily with a common fixed mismatch set is paid at the following first integers:*

row	z_{paid}	bound at z_{paid}	bound at $z_{\text{paid}} - 1$
<i>KoalaBear MCA</i>	67313	274978201175970494	274982404608707323
<i>KoalaBear list</i>	67314	274978201175970494	274982404608707323
<i>Mersenne-31 MCA</i>	322321	16777133	16777295
<i>Mersenne-31 list</i>	322322	16777133	16777295

Proposition 7.8 (Exchange defect dominates support excess). *In one boundary prefix fiber, let the associated codeword direction space have minimum support $R + a$. For distinct supports M, M' of size m , write*

$$e = |M \setminus M'| = |M' \setminus M| = w + r.$$

Then

$$(7.12) \quad r \geq a.$$

Proof. The prefix-rigidity identity of CAP25 v13.2 gives

$$\Lambda_M - \Lambda_{M'} = G(A - B), \quad \deg(A - B) \leq r - 1,$$

where G is the locator of $M \cap M'$ and has degree $m - e$. The corresponding nonzero codeword difference vanishes on the $m - e$ common support coordinates, so its weight is at most

$$n - (m - e) = t + e = R + r.$$

By the definition of a , its weight is at least $R + a$. Hence $r \geq a$. $\square$

Corollary 7.9 (Top-seam localization). *The top seam $e = w + 1$ can occur only in the exact-minimum-distance branch $a = 1$. Every chart with $a \geq 2$ has first possible exchange defect at least a .*

Theorem 7.10 (Exact primitive residual after minimum-distance payment). *After all unconditional first-match payments in Parts I–IV, an over-budget first-unpaid Q or list chart must lie in the following simultaneous wedge:*

- (R1) *its support excess a lies below the exact curve (7.9);*
- (R2) *its full-support nullity lies beyond both (7.10) and the minimum-word flag bound;*
- (R3) *it contains no fixed mismatch block meeting Theorem 7.6;*
- (R4) *every boundary neighbour has exchange defect at least a , and the top seam is restricted to $a = 1$;*
- (R5) *it is not in an earlier quotient, occupancy, extension, common-zero, saturation, or rank-defect cell.*

8. THE QUADRATIC MEAN-OVERLAP STAIRCASE

Lemma 8.1 (Large-overlap collapse). *Fix a received pair and an exact agreement $a = n - r \geq k + 1$. Choose one noncommon exact witness support S_z and explaining polynomial p_z for every slope z in its bad set Z . If two distinct chosen supports satisfy*

$$(MO1) \quad |S_z \cap S_w| \geq k + r,$$

then $|Z| \leq r + 1$.

Proof. Put

$$c_1 = \frac{p_z - p_w}{z - w}, \quad c_0 = p_z - zc_1.$$

On $S_z \cap S_w$, the two line equations imply that the received pair equals (c_0, c_1) . Let C_0 be the full common support of this codeword pair and write $c = |C_0| \geq k + r$. For every $u \in Z$, $|S_u \cap C_0| \geq a + c - n \geq k$, so the explaining polynomial p_u equals $c_0 + uc_1$.

Subtract this codeword pair from the received pair. At every coordinate outside C_0 at least one of the two residual values is nonzero, and a coordinate can cancel at at most one finite slope. A noncommon witness for u must contain at least $\max\{1, a - c\}$ such cancellations: the support-size condition gives $a - c$, and one is still necessary when $a \leq c$. Double counting therefore gives

$$(MO2) \quad |Z| \leq \left\lfloor \frac{n - c}{\max\{1, a - c\}} \right\rfloor \leq r + 1.$$

Indeed, for $c \geq a$ the numerator is at most r ; for $c = a - 1$ it is $r + 1$; and, with $x = a - c \geq 2$, the quotient is $(r + x)/x \leq r + 1$. $\square$

Theorem 8.2 (Quadratic mean-overlap theorem). *Let $C = \text{RS}_{\mathbb{F}}(D, k)$, put $R = n - k$, let $0 \leq r \leq R - 1$, and let $\emptyset \neq \Gamma \subseteq \mathbb{F}$. If*

$$(MO3) \quad (n - r)^2 \geq n(k + r), \quad \text{equivalently} \quad r^2 \geq n(3r - R),$$

then

$$(MO4) \quad B_{C, \Gamma}^{\text{MCA}}(n - r) = \min\{|\Gamma|, r + 1\}.$$

The same upper bound holds for every linear $[n, k]$ MDS code.

Proof. The universal tangent construction imported from [Cho26, Sec. 13] gives the lower bound. For the upper bound, fix a received pair and its challenge-restricted bad set Z . If $|Z| \leq r + 1$ there is nothing to prove. Otherwise choose $r + 2$ slopes and distinct exact witness supports $S_1, \dots, S_{r+2}$, each of size $a = n - r$. Suppose every pair had intersection at most $k + r - 1$, and put $d_x = \#\{i : x \in S_i\}$. Then

$$(MO5) \quad \sum_x \binom{d_x}{2} \leq \binom{r+2}{2} (k + r - 1),$$

whereas Cauchy–Schwarz and $a^2/n \geq k + r$ give

$$(MO6) \quad \sum_x \binom{d_x}{2} \geq \frac{(r+2)^2 a^2/n - (r+2)a}{2} \geq \frac{(r+2)^2 (k+r) - (r+2)a}{2}.$$

The last lower bound exceeds the upper bound in (MO5) by

$$\frac{r+2}{2} (k + 3r - n + 1).$$

When $3r > R$ this is positive. When $3r \leq R$, the exact deep theorem already gives the desired upper bound. Thus in the remaining case some pair satisfies (MO1), and Lemma 8.1 contradicts $|Z| \geq r + 2$. The exact-support reduction and large-overlap lemma use only the MDS information-set property, proving the stated extension. $\square$

Corollary 8.3 (Exact quadratic staircase). *Put*

$$(MO7) \quad r_{\text{quad}}(n, k) = \left\lfloor \frac{3n - \sqrt{n(5n + 4k)}}{2} \right\rfloor.$$

For every $0 \leq r \leq r_{\text{quad}}(n, k)$ and every nonempty challenge set Γ ,

$$B_{C, \Gamma}^{\text{MCA}}(n - r) = \min\{|\Gamma|, r + 1\}.$$

If $k/n \rightarrow \rho$, then $r_{\text{quad}}(n, k)/n \rightarrow \alpha_\rho = (3 - \sqrt{5 + 4\rho})/2$.

Proof. The smaller root of $r^2 - 3nr + n(n - k) = 0$ is the real number inside the floor in (MO7). Apply Theorem 8.2; the limiting formula follows by division by n . $\square$

Proposition 8.4 (Limit of pairwise-overlap information). *Let $k/n \rightarrow \rho \in (0, 1)$, $r/n \rightarrow \alpha \in (0, 1)$, and suppose*

$$(MO8) \quad \alpha^2 < 3\alpha - (1 - \rho).$$

For all sufficiently large n , there are $r + 2$ distinct r -subsets T_i such that, whenever $j = 3r - (n - k) \leq r$, $|T_i \cap T_\ell| \leq j - 1$ for every $i \neq \ell$. If $j > r$, that intersection condition is automatic. Hence pairwise support overlap alone cannot extend the exact theorem past the quadratic boundary; further progress must use the syndrome, determinant, or fiber incidence. This is a limitation of the method, not a construction of an MCA-bad line.

Proof. When $j \leq r$, choose $r + 2$ independent uniform r -subsets. The intersection of a fixed pair is hypergeometric with mean $r^2/n = (\alpha^2 + o(1))n$, whereas $j = (3\alpha - (1 - \rho) + o(1))n$. More explicitly,

$$\Pr(|T_1 \cap T_2| = s) = \frac{\binom{r}{s} \binom{n-r}{r-s}}{\binom{n}{r}}.$$

After writing $s = xn$, Stirling's formula gives the exponent

$$\alpha h(x/\alpha) + (1 - \alpha)h((\alpha - x)/(1 - \alpha)) - h(\alpha),$$

which is strictly concave and has its unique maximum zero at $x = \alpha^2$. The strict gap in (MO8), followed by summing at most $n + 1$ terms, therefore gives probability $e^{-\Omega(n)}$ that one fixed intersection reaches j . A union bound over $O(n^2)$ pairs is smaller than one. A successful family is automatically distinct, because equal members have intersection $r \geq j$. When $j > r$, choose any $r + 2$ distinct r -sets. $\square$

The following complementary theorem can be stronger when the overlap defect $j = 3r - R$ is small. Its only new input is a packing inequality; in particular it uses no smoothness or Fourier estimate.

Theorem 8.5 (Exact overlap–packing criterion). *Let $C = \text{RS}_{\mathbb{F}}(D, k)$, put $R = n - k$, and fix an integer $0 \leq r \leq R - 1$. Set $a = n - r$ and $j = 3r - R$. If either*

$$(OP0) \quad j \leq 0,$$

or

$$(OP) \quad 1 \leq j \leq r, \quad \binom{n}{j} \leq (r+1) \binom{r}{j},$$

then

$$(OP') \quad B_C^{\text{MCA}}(n - r) = r + 1.$$

The same upper bound holds for every linear $[n, k]$ MDS code; the matching lower bound holds whenever the field contains at least $r + 1$ distinct slope parameters.

Proof. Condition (OP0) is the imported deep-regime case [Cho26, Sec. 13], so assume (OP). Fix an arbitrary received pair (f, g) , let Z be its bad-slope set, and use the exact-agreement reduction imported from [Cho26, Sec. 13] to choose, for every $z \in Z$, an exact witness support S_z , $|S_z| = a$, and an explaining polynomial $p_z \in \mathbb{F}[X]_{<k}$. A single noncommon support cannot carry two distinct slopes, so the supports chosen for distinct elements of Z are distinct.

For the stated MDS extension, the same exact-support reduction is available: interpolate the unexplained component on any k -subset of the original witness support; because every k -subset is an information set, nonexplanation is detected at one further coordinate, and that $(k + 1)$ -set can be enlarged to size a .

First suppose that two of them have a large intersection:

$$(OP1) \quad |S_z \cap S_w| \geq n + k - a = k + r \quad (z \neq w).$$

Then Lemma 8.1 gives $|Z| \leq r + 1$.

It remains to consider the complementary case, in which every pair of chosen supports satisfies $|S_z \cap S_w| \leq n + k - a - 1$. Put $T_z = D \setminus S_z$, so $|T_z| = r$. For distinct slopes,

$$(OP3) \quad |T_z \cap T_w| = n - 2a + |S_z \cap S_w| \leq 3r - R - 1 = j - 1.$$

No j -subset of D can therefore lie in two different T_z 's. Counting pairs (z, J) with $J \in \binom{T_z}{j}$ yields

$$|Z| \binom{r}{j} \leq \binom{n}{j} \leq (r+1) \binom{r}{j},$$

and hence $|Z| \leq r + 1$. The universal tangent construction gives the reverse inequality. The proof used only linearity, the fact that agreement of two codewords on k coordinates forces equality, and the exact-support reduction; these are MDS properties. The coordinate tangent construction gives the stated MDS lower bound. $\square$

Proposition 8.6 (General overlap–packing upper bound). *For $1 \leq j = 3r - (n - k) \leq r$, let $A(n, r, j)$ be the largest size of a family of r -subsets of an n -set whose distinct members intersect in at most $j - 1$ points. Then*

$$(OP5) \quad B_{C, \Gamma}^{\text{MCA}}(n - r) \leq \min \{|\Gamma|, \max(r + 1, A(n, r, j))\}$$

and

$$(OP6) \quad A(n, r, j) \leq \left\lfloor \frac{\binom{n}{j}}{\binom{r}{j}} \right\rfloor.$$

Proof. For a fixed received pair, either two chosen exact witness supports have intersection at least $k + r$, in which case [Lemma 8.1](#) gives at most $r + 1$ slopes, or every pair of complementary r -sets intersects in at most $j - 1$ points. The latter family has size at most $A(n, r, j)$. Finally, no j -subset can lie in two members of such a packing, so counting contained j -subsets proves [\(OP6\)](#). $\square$

9. EXACT THRESHOLDS FOR EXPLICIT PRIZE ROWS

TABLE 1. Explicit maximal-dimension prize-admissible rows. In each row $k = 2^{40}$, $B = \lfloor p/2^{128} \rfloor = r_\rho(n) + 1$, and (s, a_0) is part of the Proth certificate in [Proposition 9.1](#).

rate	length	B	bits of p	(s, a_0)
1/2	2^{41}	389500552609	167	(92, 3)
1/4	2^{42}	1210584858040	169	(93, 13)
1/8	2^{43}	2879806199253	170	(95, 5)
1/16	2^{44}	6233898019554	171	(97, 5)

For completeness, the four field orders are

$$p_{1/2} = 132540169958804033333249306710494641010898987122689,$$

$$p_{1/4} = 411940680852499481698306614369841346700408394874881,$$

$$p_{1/8} = 979947269755402568812854322316630667196565607677953,$$

$$p_{1/16} = 2121285573237585848299875619011192262679065433997313.$$

Their odd Proth coefficients u , in the same order, are

$$u_{1/2} = 26766274163673319604503,$$

$$u_{1/4} = 41595378994516821279015,$$

$$u_{1/8} = 24737346889219389259839,$$

$$u_{1/16} = 13387194060291799253121.$$

For $F_{n,k}(r) = r^2 - n(3r - (n - k))$, the exact boundary checks are

	ρ	$F_{n,k}(B - 1)$	$F_{n,k}(B)$
(PC0)	1/2	5154112775168	-663955886271
	1/4	7590647904465	-3182321912768
	1/8	13908181940112	-6720484728007
	1/16	19335616403905	-20973145690236.

Proposition 9.1 (Exact arithmetic and primality certificates). *For the four rows in [Table 1](#), the displayed integer p satisfies*

$$(PC1) \quad p = u2^s + 1, \quad u \text{ odd}, \quad u < 2^s, \quad a_0^{(p-1)/2} \equiv -1 \pmod{p}.$$

Consequently every p is prime. Moreover

$$(PC2) \quad n \mid p - 1, \quad p < 2^{256}, \quad B2^{128} \leq p < (B + 1)2^{128},$$

where n and B are the corresponding entries of the table. They also satisfy $F_{n,k}(B - 1) \geq 0 > F_{n,k}(B)$. Since $F'_{n,k}(r) = 2r - 3n < 0$ on $[0, n]$, these two signs locate its smaller root and give $B - 1 = r_\rho(n) = r_{\text{quad}}(n, k)$. Thus $\mathbb{F}_p^$ contains a subgroup of order n , and each displayed code exists with the asserted slope budget.*

Proof. All identities and inequalities in [\(PC0\)](#)–[\(PC2\)](#) are direct integer calculations; the values of (s, a_0) , u , p , n , and B have all been printed above so that the calculation is independently reproducible. We recall why the certificate proves primality. If a prime ℓ divides p , then $a_0^{u2^{s-1}} \equiv -1 \pmod{\ell}$, whereas $a_0^{u2^s} \equiv 1 \pmod{\ell}$. The multiplicative order of a_0 modulo ℓ therefore has exact 2-part 2^s , so $2^s \mid \ell - 1$. Since $u < 2^s$, one has $2^s > \sqrt{p}$; hence a composite p would have a prime divisor both at most $\sqrt{p}$ and at least $2^s + 1$, a contradiction. This is Proth's theorem in the special form used here. Finally $s \geq \log_2 n$ gives $n \mid p - 1$, and cyclicity of $\mathbb{F}_p^*$ gives the required subgroup. $\square$

Theorem 9.2 (Exact first adjacent row over a separating field). *Let $C = \text{RS}_{\mathbb{F}}(D, k)$, put $R = n - k$ and $M = \binom{n}{R-1} = \binom{n}{k+1}$, and define*

$$Q_{\text{sep}}(M) = \max \left\{ M, \binom{M}{2} \right\}.$$

If $|\mathbb{F}| > Q_{\text{sep}}(M)$, then for the full-field challenge

$$(AD1) \quad B_C^{\text{MCA}}(k+1) = M.$$

If $R \geq 2$ and $b = \lfloor \varepsilon |\mathbb{F}| \rfloor$, the exact target threshold therefore satisfies

$$(AD2) \quad \begin{aligned} b \geq M &\implies a^* = k+1, \\ \binom{n}{R-2} \leq b < M &\implies a^* = k+2. \end{aligned}$$

Thus the first finite adjacent comparison has literal binomial constants; no $e^{o(n)}$ or unspecified polynomial factor occurs.

Proof. The upper bound in (AD1) is the exact support-atlas bound imported from [Cho26, Sec. 13]. For the reverse inequality use the Reed–Solomon parity check of redundancy R . For every $(R-1)$ -set $E \subseteq D$, its column span V_E is a hyperplane in $\mathbb{F}^R$, and distinct sets give distinct hyperplanes: otherwise the union of two distinct $(R-1)$ -sets would contain at least R columns in one $(R-1)$ -space, contradicting the MDS column property.

Every proper linear hyperplane in $\mathbb{F}^R$ has $|\mathbb{F}|^{R-1}$ points. Since $M < |\mathbb{F}|$, the union of the M spaces V_E has fewer than $|\mathbb{F}|^R$ points; choose y_1 outside it. For each E , choose a nonzero linear functional ℓ_E with kernel V_E . The line $y_0 + \gamma y_1$ meets V_E at

$$\gamma_E = -\ell_E(y_0)/\ell_E(y_1).$$

For $E \neq E'$, equality $\gamma_E = \gamma_{E'}$ cuts out a proper hyperplane in the variable y_0 , because the normalized functionals $\ell_E/\ell_E(y_1)$ and $\ell_{E'}/\ell_{E'}(y_1)$ are distinct. There are at most $\binom{M}{2} < |\mathbb{F}|$ collision hyperplanes; their union also has fewer than $|\mathbb{F}|^R$ points, so choose y_0 outside it. Since $y_1 \notin V_E$, the line is not contained in V_E , and every intersection is transverse. The M parameters γ_E are then finite and pairwise distinct. Syndrome surjectivity and Theorem 4.2 realize them as M bad slopes at agreement $k+1$, proving (AD1).

If $b \geq M$, the support upper bound makes the first allowed agreement safe. If $\binom{n}{R-2} \leq b < M$, equality (AD1) makes $k+1$ unsafe, whereas the exact support-atlas bound imported from [Cho26, Sec. 13] makes $k+2$ safe. Monotonicity proves (AD2). $\square$

10. ASYMPTOTIC CONSEQUENCES OF THE IMPORTED PREFIX FLOOR

Theorem 10.1 (Unconditional shallow-prefix RS–MCA exponent). *Let $C_n = \text{RS}_{\mathbb{F}_n}(D_n, k_n)$, where $D_n \subseteq \mathbb{B}_n \subseteq \mathbb{F}_n$, $|D_n| = n$, and $k_n/n \rightarrow \rho \in (0, 1)$. Let $a_n \geq k_n + 1$, put $w_n = a_n - k_n - 1$, and assume*

$$(AS1) \quad \frac{a_n}{n} \rightarrow \alpha \in (0, 1), \quad w_n \log |\mathbb{B}_n| = o(n),$$

$|\mathbb{F}_n| > n$, and, for some finite $\phi \geq 0$,

$$(AS2) \quad \frac{1}{n} \log |\mathbb{F}_n| \rightarrow \phi, \quad \frac{1}{n} \log (|\mathbb{F}_n| - n) \rightarrow \phi.$$

For the full-field challenge,

$$(AS3) \quad \frac{1}{n} \log B_{C_n}^{\text{MCA}}(a_n) = \min\{h(\alpha), \phi\} + o(1),$$

and therefore

$$(AS4) \quad \frac{1}{n} \log e_{\text{MCA}} \left(C_n, 1 - \frac{a_n}{n} \right) = -(\phi - h(\alpha))_+ + o(1).$$

Because $|\mathbb{B}_n| \geq n$, condition (AS1) forces $\alpha = \rho$. No smoothness, Fourier estimate, or auxiliary ray hypothesis is assumed.

Proof. Put $M_n = \binom{n}{a_n}$ and $L_n = \lceil M_n |\mathbb{B}_n|^{-w_n} \rceil$. Exact-agreement reduction and fixed-support slope uniqueness give the unconditional upper bound

$$(AS5) \quad B_{C_n}^{\text{MCA}}(a_n) \leq \min\{|\mathbb{F}_n|, M_n\}.$$

The exact locator-prefix list and collision-aware pole compiler imported from [Cho26, Secs. 14–18] give

$$(AS6) \quad B_{C_n}^{\text{MCA}}(a_n) \geq \left\lceil \frac{L_n(|\mathbb{F}_n| - n)}{|\mathbb{F}_n| - n + k_n(L_n - 1)} \right\rceil.$$

By Stirling's formula and (AS1), $\log L_n = h(\alpha)n + o(n)$. For positive x, y , $xy/(x + k_n y) \geq \frac{1}{2} \min\{y, x/k_n\}$; hence (AS2), $\log k_n = o(n)$, and the ceiling show that the right side of (AS6) has logarithm $n \min\{h(\alpha), \phi\} + o(n)$. The upper bound (AS5) has the same exponent, proving (AS3); subtracting $\log |\mathbb{F}_n|$ proves (AS4). Finally $w_n \log n = o(n)$, so $(a_n - k_n)/n \rightarrow 0$ and $\alpha = \rho$. $\square$

Corollary 10.2 (Unconditional large-challenge RS-MCA threshold). *Under the row hypotheses of Theorem 10.1, let the full-field target satisfy $\varepsilon_n = \exp(-\lambda n + o(n))$ for a finite $\lambda \geq 0$. If*

$$(AS7) \quad \lambda < \phi - h(\rho),$$

then, for every sufficiently large n ,

$$(AS8) \quad a_n^* = k_n + 1, \quad \delta_n^* = 1 - \frac{k_n + 1}{n} = 1 - \rho + o(1).$$

Conversely, if $\lambda > (\phi - h(\rho))_+$, every agreement sequence satisfying $(a_n - k_n - 1) \log |\mathbb{B}_n| = o(n)$ is unsafe. The critical boundary $\lambda = (\phi - h(\rho))_+$ is governed by the exact finite constants and is not decided by exponent asymptotics. This is a large scalar-extension regime. If $|\mathbb{F}_n|$ is polynomial in n , then $\phi = 0$ and (AS7) cannot hold; the corollary does not settle that Proximity-Prize parameter regime.

Proof. At $a = k_n + 1$, (AS5) is $B_{C_n}^{\text{MCA}}(a) \leq \binom{n}{k_n+1} = \exp(h(\rho)n + o(n))$. Under (AS7), $\varepsilon_n |\mathbb{F}_n| = \exp((\phi - \lambda)n + o(n))$ has strictly larger exponent, and taking the integer floor does not change the comparison. Thus the first permitted agreement is safe and (AS8) follows. The converse follows from (AS4) and the strict reverse exponential inequality. $\square$

Part 2. Exact profile refinements and the profile envelope

11. PROFILE ENVELOPE AND THE SEQUEL COMPILER

A realized first-match profile λ has a full support slice Ω_λ^0 , a coefficient field $\mathbb{B}_\lambda$, a proved boundary map $\Phi_\lambda : \Omega_\lambda^0 \rightarrow \mathbb{B}_\lambda^{R_\lambda}$, realized image size $L_\lambda = |\Phi_\lambda(\Omega_\lambda^0)|$, and image-normalized average

$$\bar{N}_\lambda = \frac{|\Omega_\lambda^0|}{L_\lambda}.$$

For a received line and agreement a , let $\Lambda(a)$ be the realized first-match profile set and define

$$(11.1) \quad \mathfrak{E}_n(a) = 1 + (n - a + 1) + \sup_{\text{received lines}} \sum_{\lambda \in \Lambda(a)} (1 + \bar{N}_\lambda).$$

The formal codomain size $|\mathbb{B}_\lambda|^{R_\lambda}$ may replace L_λ only after a full-image certificate $L_\lambda \geq e^{-o(n)} |\mathbb{B}_\lambda|^{R_\lambda}$ has been proved. Otherwise the leaf stays at realized-image scale or is routed to an earlier effective-image-collapse cell.

Definition 11.1 (Admissible sequel compiler). *A row sequence is admissible when: (i) every bad witness belongs to an ordered first-match atlas with $e^{o(n)}$ realized profiles; (ii) all quotient, planted, field-drop, tangent, rank, extension, and directly counted ray cells have proved distinct-slope payments; (iii) every remaining primitive profile has a fixed-density Boolean support slice, an exact boundary map, and image normalization; (iv) its nontrivial effective Fourier characters satisfy the certified minor/major payment of Definitions 13.2, 13.6 and 13.9, or an equivalent direct Sidon payment; and (v) every residual ray chart has either a direct slope bound or the ray compiler of Hypothesis 22.13. All profile fields and quotient scales are included in (11.1).*

Theorem 11.2 (Conditional profile-envelope compiler). *Every admissible row sequence satisfies*

$$(11.2) \quad B_{C_n}^{\text{MCA}}(a_n) \leq e^{o(n)} \mathfrak{E}_n(a_n).$$

For a finite row the same proof is literal after every $e^{o(n)}$ term is replaced by its exact integer constant and all first-match cells are summed once.

Proof. The first-match atlas pays the algebraic cells. The effective Fourier/Sidon step and [Theorem 21.2](#) pay primitive Q , while the ray compiler and [Proposition 22.14](#) convert the support bounds to distinct slopes. Summation over the realized profile set gives (11.2). The finite statement is the same disjoint first-match count with integer constants. $\square$

Definition 11.3 (Complete-fiber folding map). *Let $D \subseteq \mathbb{B}$ be finite. A polynomial map $\phi : D \rightarrow Q = \phi(D)$ is a c -fold complete-fiber folding map when every fiber has exactly c points. A support is complete at this scale when it is a union of full fibers. Multiplicative power maps and the Chebyshev twin-coset foldings of the foundation [[Cho26](#), Secs. 7 and 11] are the motivating examples.*

12. QUOTIENT-REMAINDER PREFIX NORMAL FORM

The quotient-prefix floors and descent mechanisms of CAP25 v13.2 are imported from [[Cho26](#), Secs. 3, 8, and 14]. The normal forms and obstruction theorems below refine that package for the sequel's Q compiler.

The quotient and remainder variables can be separated exactly at the locator-prefix level. This removes a spurious factor coming from summing over marked remainder sets, but it also exposes a genuine scaled quotient- coefficient-field term which prevents an unrestricted comparison with the identity profile.

Definition 12.1 (Quotient/remainder profile and field drop). *Let $\phi : D \rightarrow Q$ be a c -fold complete-fiber folding map in the sense of [Definition 11.3](#). A complete-fiber-plus-remainder support is*

$$S = \phi^{-1}(E) \sqcup R,$$

where $E \subseteq Q$ and $R \subseteq D \setminus \phi^{-1}(E)$. The set R is the support remainder: it records the selected points in fibers that are not declared complete. A quotient/remainder profile fixes ϕ , its degree c , the size $r = |R|$, the relevant coefficient field, and any planted remainder data, while E and the unfixed part of R range subject to the displayed disjointness. A fixed- R profile is the subprofile in which the remainder support itself is part of the label. The adjective Euclidean used below refers to separating the locator factor of degree $r < c$ from the composed quotient locator; it does not mean that an arbitrary support remainder has size less than c .

We call $\mathbb{B}_\phi \subseteq \mathbb{B}$ a scaled quotient coefficient field when, for some $\eta \in \mathbb{B}^\times$,

$$Q = \phi(D) \subseteq \eta \mathbb{B}_\phi.$$

When this field is recorded in a profile, it is taken minimal among the subfields with this property. Then the j -th elementary quotient coefficient lies in $\eta^j \mathbb{B}_\phi$. A field drop is the use of a proper such subfield. It is a positive-rate field drop along a sequence when

$$1 - \frac{\log |\mathbb{B}_\phi|}{\log |\mathbb{B}|}$$

stays bounded away from zero and a positive-density number on the entropy scale of quotient-prefix slots is live. The payment saved in each live slot is the factor $|\mathbb{B}|/|\mathbb{B}_\phi|$; this is the term quantified in [Proposition 12.12](#).

Definition 12.2 (Balanced quotient core). *A split locator is a monic squarefree polynomial all of whose roots lie in the evaluation domain D . After common locator factors, fixed quotient data, and planted remainder points have been removed, let A and B be two monic split locators of the same degree e . They form a depth- w balanced core if they have the same first w nonleading coefficients, equivalently*

$$\deg(A - B) \leq e - w - 1.$$

It is a balanced quotient core when the remaining locator factors are pullbacks through the same folding map ϕ . The word balanced records equality of degree and leading normalization; the core consists of the roots still allowed to move after common and planted factors are removed. If the resulting projective locator family is one-dimensional it is a split pencil $Q_0 + \lambda Q_1$; in higher dimension this definition supplies no distinct-ray bound, and the separate ray compiler is required.

For a monic polynomial

$$A(X) = X^d + a_1 X^{d-1} + \cdots + a_d$$

and an integer $0 \leq u \leq d$, write $\text{pref}_u(A) = (a_1, \dots, a_u)$. Equivalently, if $A^\vee(Z) := Z^d A(Z^{-1})$, then $\text{pref}_u(A)$ is the coefficient vector of $A^\vee(Z) \bmod Z^{u+1}$, with its constant coefficient omitted.

Theorem 12.3 (Exact quotient–remainder prefix normal form). *Let $\mathbb{B} \subseteq \mathbb{F}$ be finite fields, let $D \subseteq \mathbb{B}$ have n distinct elements, and let $\phi \in \mathbb{B}[X]$ be monic of degree $c \geq 2$. Assume that the restriction*

$$\phi : D \longrightarrow Q := \phi(D)$$

has complete fibers of size c ; in particular, $N := |Q| = n/c$. Fix integers m, r with $0 \leq r < c$ and $cm + r \leq n$. For

$$E \in \binom{Q}{m}, \quad R \in \binom{D \setminus \phi^{-1}(E)}{r},$$

put

$$\begin{aligned} V_E(Y) &:= \prod_{y \in E} (Y - y) = Y^m + \sum_{j=1}^m v_j(E) Y^{m-j}, \\ (QR1) \quad P_R(X) &:= \prod_{x \in R} (X - x) = X^r + \sum_{i=1}^r p_i(R) X^{r-i}, \\ L_{E,R}(X) &:= P_R(X) V_E(\phi(X)). \end{aligned}$$

Then $L_{E,R}$ is the monic locator of the support $\phi^{-1}(E) \sqcup R$, of degree $A = cm + r$.

Let $0 \leq w \leq A$ and put $d = \lfloor w/c \rfloor$.

- (i) *If $w \geq r$, every nonempty fiber of $(E, R) \mapsto \text{pref}_w(L_{E,R})$ determines R uniquely and, after R is recovered, is exactly one quotient–prefix fiber. More precisely, for every prefix value z there are unique data R_z and $\eta_z \in \mathbb{B}^d$, whenever the fiber is nonempty, such that*

$$\begin{aligned} (QR2) \quad & \{(E, R) : \text{pref}_w(L_{E,R}) = z\} \\ &= \{(E, R_z) : E \in \binom{Q \setminus \phi(R_z)}{m}, \text{pref}_d(V_E) = \eta_z\}. \end{aligned}$$

Thus no factor counting possible remainder sets occurs in a fixed global prefix fiber.

- (ii) *If $w < r$, then necessarily $w < c$, and the quotient set E is invisible to the first w coefficients. There is a fixed unitriangular bijection $T_{\phi, m, w} : \mathbb{B}^w \rightarrow \mathbb{B}^w$ such that*

$$(QR3) \quad \text{pref}_w(L_{E,R}) = T_{\phi, m, w}(\text{pref}_w(P_R))$$

for every admissible (E, R) . Consequently, with the convention $\binom{u}{v} = 0$ for $v > u$,

$$\begin{aligned} (QR4) \quad & \#\{(E, R) : \text{pref}_w(L_{E,R}) = z\} \\ &= \sum_{\substack{R \in \binom{P}{r} \\ \text{pref}_w(P_R) = T_{\phi, m, w}^{-1}(z)}} \binom{N - |\phi(R)|}{m}. \end{aligned}$$

Thus this case is a remainder–prefix problem, multiplied by the number of quotient sets still available after R is fixed.

Proof. For a monic degree- e polynomial A , set $A^\vee(Z) = Z^e A(Z^{-1})$. Also set

$$\Phi(Z) := Z^c \phi(Z^{-1}) = 1 + \phi_1 Z + \cdots + \phi_c Z^c.$$

Taking reciprocal polynomials in (QR1) gives the exact formal-power-series identity

$$(QR5) \quad L_{E,R}^\vee(Z) = P_R^\vee(Z) \left(\Phi(Z)^m + \sum_{j=1}^m v_j(E) Z^{cj} \Phi(Z)^{m-j} \right).$$

All factors on the right have constant coefficient one except for the displayed scalar coefficients $v_j(E)$.

Suppose first that $w \geq r$. Since $r < c$, no term containing a v_j occurs below degree c . For $1 \leq i \leq r$, the coefficient of Z^i on the right of (QR5) is $p_i(R)$ plus an expression depending only on $p_1(R), \dots, p_{i-1}(R)$, on m , and on Φ . The first r coefficients therefore recover $p_1(R), \dots, p_r(R)$ successively. They recover the monic polynomial P_R , hence its root set R , uniquely.

We may now divide $L_{E,R}^\vee$ by the unit $P_R^\vee$ modulo Z^{w+1} . In the remaining factor of (QR5), the coefficient of Z^{cj} is $v_j(E)$ plus an expression depending only on $v_1(E), \dots, v_{j-1}(E)$, on m , and on Φ . Hence the coefficients at degrees $c, 2c, \dots, dc$ recover $v_1(E), \dots, v_d(E)$ successively. Conversely, P_R and these d coefficients determine the right side of (QR5) modulo Z^{w+1} , because every term containing v_j with $j > d$ starts in degree $cj > w$. This proves (QR2); admissibility of E is exactly the condition $E \cap \phi(R) = \emptyset$.

If $w < r < c$, every term involving a v_j starts in degree at least $c > w$, so modulo Z^{w+1} the bracket in (QR5) is the fixed unit $\Phi(Z)^m$. Multiplication by this unit sends the first w coefficients of $P_R^\vee$ to the first w coefficients of $L_{E,R}^\vee$ by a unitriangular transformation. It is independent of E and invertible. For a fixed R , the admissible E are precisely the m -subsets of $Q \setminus \phi(R)$, of which there are $\binom{N-|\phi(R)|}{m}$. Summing over the indicated remainder-prefix fiber proves (QR4). $\square$

Proposition 12.4 (Complete support factorization and its limitation). *Under the complete-fiber hypotheses of Theorem 12.3, every support $S \subseteq D$ has the unique decomposition*

$$(QR5a) \quad E(S) = \{y \in Q : \phi^{-1}(y) \subseteq S\}, \quad R(S) = S \setminus \phi^{-1}(E(S)),$$

and $Q_S(X) = P_{R(S)}(X) V_{E(S)}(\phi(X))$. The remainder meets every nonfull ϕ -fiber in at most $c-1$ points, but its total size may be $\Theta(n)$. The reciprocal identity (QR5) remains valid for this arbitrary remainder. When $|R(S)| \geq c$, however, the quotient coefficients and the remainder coefficients begin to interlace at degree c , so the triangular recovery theorem above does not give a comparison without an additional partial-occupancy atlas.

Proof. A fiber is put into $E(S)$ exactly when all of its points lie in S ; all other selected points form $R(S)$. This proves existence, disjointness, the occupancy bound, and uniqueness. Multiplying the corresponding linear factors gives the locator identity. The reciprocal calculation used for (QR5) is purely multiplicative and does not require $\deg P_R < c$. If $\deg P_R \geq c$, both $p_c(R)$ and $v_1(E)$ occur in degree c , which proves the stated limitation of the triangular separation. $\square$

Theorem 12.5 (Exact partial-occupancy disintegration). *Let $D = D_0 \sqcup X$, where $|X| = b$, and let $\phi : D_0 \twoheadrightarrow Q$ have exactly $N = |Q|$ fibers, each of cardinality c . For $S \subseteq D$, put*

$$X_S = S \cap X, \quad E(S) = \{y \in Q : \phi^{-1}(y) \subseteq S\}, \quad R(S) = (S \cap D_0) \setminus \phi^{-1}(E(S)).$$

Write

$$t = |X_S|, \quad m = |E(S)|, \quad p = |\phi(R(S))|, \quad r = |R(S)|.$$

Then this datum is unique, every fiber meeting $R(S)$ contains between one and $c-1$ points of $R(S)$, and $|S| = t + cm + r$. If $\Omega_{t,m,p,r}$ denotes the family of supports with these four parameters, then

$$(PO1) \quad |\Omega_{t,m,p,r}| = \binom{b}{t} \binom{N}{p} \binom{N-p}{m} [x^r] ((1+x)^c - 1 - x^c)^p.$$

Consequently the nonempty families $\Omega_{t,m,p,r}$ partition $\binom{D}{a}$, and their exact add-back is

$$(PO2) \quad \sum_{\substack{t,m,p,r \\ t+cm+r=a}} |\Omega_{t,m,p,r}| = \binom{cN+b}{a}.$$

If ϕ is the restriction of a monic degree- c polynomial and every fiber $\phi^{-1}(y)$ is the complete simple root set in D_0 of $\phi(X) - y$, then their locators satisfy

$$Q_S = Q_{X_S} P_{R(S)} V_{E(S)}(\phi).$$

Proof. Full fibers, partial fibers, and empty fibers are mutually exclusive, so the displayed decomposition is canonical. Choose the t exceptional points, the p partial fibers, and the m full fibers in $\binom{b}{t} \binom{N}{p} \binom{N-p}{m}$ ways. On one partial fiber the weight enumerator of the allowed nonempty proper subsets is $(1+x)^c - 1 - x^c$; hence the coefficient in (PO1) counts the choices having total partial weight r . Summing (PO1) and extracting the coefficient of x^a gives

$$[x^a](1+x)^b (1+x^c + ((1+x)^c - 1 - x^c))^N = [x^a](1+x)^{b+cN} = \binom{b+cN}{a}.$$

Multiplying the linear factors belonging to the three disjoint parts of the support proves the locator identity. $\square$

Proposition 12.6 (Partial-occupancy Fourier factorization). *In the setting of Theorem 12.5, let G be a finite abelian group, let $g : D \rightarrow G$, and put $\Phi(S) = \sum_{x \in S} g(x)$. For a character $\chi \in \widehat{G}$ and a regular fiber $F_y = \phi^{-1}(y)$, define*

$$A_y(\chi) = \prod_{x \in F_y} \chi(g(x)), \quad P_y(\chi; z) = \prod_{x \in F_y} (1 + z\chi(g(x))) - 1 - z^c A_y(\chi).$$

Then

$$(PO3) \quad \sum_{S \in \Omega_{t,m,p,r}} \chi(\Phi(S)) = [u^t v^m \zeta^p z^r] \prod_{x \in X} (1 + u\chi(g(x))) \prod_{y \in Q} (1 + vA_y(\chi) + \zeta P_y(\chi; z)).$$

Consequently, for every $s \in G$,

$$(PO4) \quad |\{S \in \Omega_{t,m,p,r} : \Phi(S) = s\}| = \frac{1}{|G|} \sum_{\chi \in \widehat{G}} \overline{\chi(s)} \sum_{S \in \Omega_{t,m,p,r}} \chi(\Phi(S)).$$

At the trivial character, (PO3) is exactly the support count (PO1). Thus support add-back is unconditional, while a finite flatness theorem is exactly a bound for the nontrivial-character sum in (PO4).

Proof. For a fixed regular fiber the three terms in the second product encode, respectively, an empty fiber, a full fiber, and a nonempty proper subset. The variables v, ζ, z record the numbers m, p, r , while u records the exceptional weight t . Multiplicativity of χ turns the character of the support sum into the product of its selected-point characters, proving (PO3) by coefficient extraction. Formula (PO4) is character orthogonality on G . $\square$

Remark 12.7 (Effective normalization of a partial-occupancy slice). Fix a nonempty slice $\Omega_\lambda = \Omega_{t,m,p,r}$, choose $S_0 \in \Omega_\lambda$, and let

$$G_\lambda = \langle \Phi(S) - \Phi(S_0) : S \in \Omega_\lambda \rangle \leq G.$$

The attained boundary values lie in the affine coset $\Phi(S_0) + G_\lambda$, and the effective inversion is

$$(PO5) \quad |\{S \in \Omega_\lambda : \Phi(S) = s\}| = \frac{1}{|G_\lambda|} \sum_{\chi \in \widehat{G_\lambda}} \overline{\chi(s - \Phi(S_0))} \sum_{S \in \Omega_\lambda} \chi(\Phi(S) - \Phi(S_0)).$$

Every character of G_λ extends to the ambient finite abelian group, so the inner sum may be evaluated from (PO3) after multiplication by the fixed factor $\overline{\chi(\Phi(S_0))}$. Different extensions agree on support differences. Equivalently, ambient characters group by their restriction to G_λ , and the annihilator multiplicity cancels exactly against the ambient denominator. This identifies the correct Fourier denominator; it does not assert that the realized image fills the affine group.

Remark 12.8 (What partial-occupancy add-back does and does not prove). The exponentially many local proper-subset choices are already contained inside the coarse slices $\Omega_{t,m,p,r}$; they incur no separate profile-count loss in (PO2). The theorem does not bound the nontrivial Fourier

terms in (PO4), the size of a realized boundary image, or the projection of a witness cell to distinct slopes. Those are separate analytic and ray payments.

Theorem 12.9 (Canonical partial-occupancy atlas and bounded-boundary closure). *Fix a received line and exact agreement a . For every nonempty canonical slice $\Omega_\lambda = \Omega_{t,m,p,r}$ of Theorem 12.5, let C_λ be the witnesses whose support lies in Ω_λ , and order these cells arbitrarily. Then they form a witness-exhaustive first-match atlas and, with Z_λ° its first-match slope projections,*

$$(PO6) \quad |Z_\lambda^\circ| \leq |\Omega_\lambda|, \quad \sum_\lambda |Z_\lambda^\circ| \leq \sum_\lambda |\Omega_\lambda| = \binom{n}{a}.$$

More generally, give each slice a boundary map $\Phi_\lambda : \Omega_\lambda \rightarrow G_\lambda^{\text{amb}}$ into a finite abelian group of order A_λ , let N_λ be its active-coordinate count, let L_λ be its realized image size, and put $\bar{N}_\lambda = |\Omega_\lambda|/L_\lambda$. The exact finite ray multiplier and challenge-set bound are

$$(PO7) \quad |Z_\lambda^\circ| \leq L_\lambda \bar{N}_\lambda.$$

$$(PO7a) \quad B_{C,\Gamma}^{\text{MCA}}(a) \leq \min \left\{ |\Gamma|, \sup_r \sum_{\lambda \in \Lambda(r;a)} L_\lambda \bar{N}_\lambda \right\} \leq \min \left\{ |\Gamma|, \binom{n}{a} \right\}.$$

Define the coarse partial-occupancy envelope

$$\mathfrak{E}_n^{\text{PO}}(a) = 1 + \sup_r \sum_{\lambda \in \Lambda(r;a)} (1 + \bar{N}_\lambda).$$

If, uniformly over received lines and the polynomially many nonempty canonical slices,

$$(PO8) \quad \log A_\lambda = o(n),$$

then this atlas has a closed asymptotic profile-envelope compiler:

$$(PO9) \quad B_C^{\text{MCA}}(a) \leq e^{o(n)} \mathfrak{E}_n^{\text{PO}}(a).$$

Condition (PO8) discharges atlas exhaustivity, support add-back, and the direct alternative in (RC), whose exact multiplier is $L_\lambda \leq A_\lambda$. If, in addition,

$$(PO10) \quad \log A_\lambda = o(N_\lambda)$$

on every slice to which effective Fourier terminology is applied, then the effective inversion (PO5) may designate every nontrivial effective character major: its minor sum is empty and its exact major multiplier is at most $A_\lambda - 1$. Thus (PO10) verifies the slice analogues of MI and MA rather than assuming them. A concrete specialization is a polynomial-size coefficient field with uniformly bounded boundary depth and $N_\lambda = \Theta(n)$ on every live canonical folding profile. This is a coarse support-paid envelope, not an identity-scale or refined quotient-envelope comparison. It gives a safe target whenever the literal right side of (PO7a) is at most B^* ; a sharp frontier additionally uses the unsafe inequality in (SB2). Condition (PO8) alone does not locate a target frontier.

Proof. The support slices partition $\binom{D}{a}$ by (PO2), hence the corresponding witness cells partition the exact witness incidence. For each first-match slope choose one witness support in its cell. A fixed noncommon support carries at most one slope, so this choice is injective and proves the first inequality in (PO6). Exact add-back gives the equality there, and (PO7) is the same inequality rewritten using $|\Omega_\lambda| = L_\lambda \bar{N}_\lambda$.

The exact challenge bound follows by summing (PO7), taking the maximum over received lines, and using (PO2). Since $L_\lambda \leq A_\lambda$, condition (PO8) turns (PO7) into $e^{o(n)} \bar{N}_\lambda$. There are only polynomially many quadruples (t, m, p, r) , so their sum proves (PO9). If $A_\lambda \leq |\mathbb{B}|^{R_\lambda}$, a polynomial-size $\mathbb{B}$ and bounded R_λ imply (PO8). Finally, the effective difference group in (PO5) has at most A_λ characters. Under (PO10), designating all of its nontrivial characters major makes the minor sum empty and bounds the normalized major aggregate by $A_\lambda - 1 = e^{o(N_\lambda)}$. $\square$

Remark 12.10 (Exact dimension rounding). If the global locator degree is $A = cm + r$, the list-code dimension is K , and $w = A - K$, then the quotient depth in [Theorem 12.3](#) is

$$(QR5b) \quad d = \left\lfloor \frac{A - K}{c} \right\rfloor = m - \left\lceil \frac{K - r}{c} \right\rceil.$$

Thus the descended list dimension is exactly $K_c = \lceil (K - r)/c \rceil$. More explicitly, if $w = cu + s$ with $0 \leq s < c$, then

$$(QR5c) \quad |\mathbb{B}_\phi|^{-\lfloor w/c \rfloor} = |\mathbb{B}_\phi|^{s/c} |\mathbb{B}_\phi|^{-w/c}.$$

Thus the exact rounding multiplier is $|\mathbb{B}_\phi|^{s/c}$, lying between one and $|\mathbb{B}_\phi|^{(c-1)/c}$. Retaining the integer $\lfloor w/c \rfloor$ avoids even this comparison; no polynomial or $e^{o(n)}$ placeholder is needed in a finite certificate.

Remark 12.11 (Nonmonic maps and Chebyshev maps). If the leading coefficient of ϕ is $b \neq 0$, replace ϕ by $b^{-1}\phi$ and Q by $b^{-1}Q$. This changes quotient coefficients by fixed nonzero scalar factors and leaves every support and every fiber cardinality unchanged. The theorem therefore applies to the Chebyshev folding map T_c after this normalization. On a circle twin-coset domain on which T_c has complete fibers of size c , it gives the same statement verbatim. The fixed or ramified endpoints form an exceptional set X of cardinality b and are included exactly by the factor $\binom{b}{t}$ in (PO1); the crude total exceptional-point multiplier is at most 2^b , with $b \leq 2$ for involution-fixed endpoints. Coincident twin cosets are not a bounded planted set: they form the separate dihedral collapse profile of the circle edge-case audit imported from [\[Cho26, Sec. 11\]](#).

For a declared complete-fiber-plus-Euclidean-remainder profile, the exact normal-form theorem separates algebraic bookkeeping from prefix equidistribution only when the remainder degree is less than c . The exact disintegration of [Theorem 12.5](#) exhausts arbitrary partial-occupancy supports, but when $\deg P_R \geq c$ the quotient and remainder coefficients interlace. Thus that disintegration does not by itself prove prefix flatness on the resulting coarse slices. Let $\mathbb{B}_\phi \subseteq \mathbb{B}$ be a subfield and suppose $Q \subseteq \eta \mathbb{B}_\phi$ for some $\eta \in \mathbb{B}^\times$. Then $v_j(E) \in \eta^j \mathbb{B}_\phi$. Hence, in case (i), the natural average scale of the quotient-prefix fiber with fixed R is

$$(QR6) \quad \overline{N}_{\phi,R}(w) := \binom{N - |\phi(R)|}{m} |\mathbb{B}_\phi|^{-\lfloor w/c \rfloor}.$$

Pigeonholing proves a fiber of size at least $\lceil \overline{N}_{\phi,R}(w) \rceil$. An upper bound of the same order is not a consequence of the normal form: it is precisely the uniform prefix-flatness, or Sidon/Fourier, input on the quotient row.

Proposition 12.12 (Identity–quotient exponent comparison). *Consider a sequence to which [Theorem 12.3](#) applies with fixed $c \geq 2$ and fixed $0 \leq r < c$. Suppose*

$$(QR7) \quad \frac{a}{n} \longrightarrow \alpha \in (0, 1), \quad \log |\mathbb{B}| = o(n), \quad \log \binom{n}{a} - w \log |\mathbb{B}| = o(n),$$

where $a = cm + r$, and suppose $Q \subseteq \eta \mathbb{B}_c$ for a subfield $\mathbb{B}_c \subseteq \mathbb{B}$. Write $\overline{N}_{c,r}(w) = \overline{N}_{\phi,R}(w)$ for the natural scale in [\(QR6\)](#), with this fixed remainder profile. Then $w \rightarrow \infty$, so $w \geq r$ eventually, and

$$(QR8) \quad \log \overline{N}_{c,r}(w) = \frac{1}{c} \log \left(\binom{n}{a} |\mathbb{B}|^{-w} \right) + \frac{w}{c} \log \frac{|\mathbb{B}|}{|\mathbb{B}_c|} + o(n),$$

where replacing $N - |\phi(R)|$ by N changes only the $o(n)$ term. Equivalently, if $\lambda_c = \log |\mathbb{B}_c| / \log |\mathbb{B}|$ has a limit, then at the identity crossing

$$(QR9) \quad \frac{1}{n} \log \overline{N}_{c,r}(w) = \frac{h(\alpha)}{c} (1 - \lambda_c) + o(1),$$

using the entropy convention fixed in the Introduction. Thus the natural average scale of this bounded-degree quotient profile is comparable with the identity scale at exponential accuracy exactly when $\log |\mathbb{B}_c| / \log |\mathbb{B}| = 1 - o(1)$. A proper field drop of fixed relative degree produces a positive exponential quotient term.

Proof. Stirling's formula, $a/n \rightarrow \alpha$, and fixed c, r give

$$\log \binom{n/c + O(1)}{(a-r)/c} = \frac{1}{c} \log \binom{n}{a} + o(n).$$

Also $\lfloor w/c \rfloor \log |\mathbb{B}_c| = (w/c) \log |\mathbb{B}_c| + o(n)$, because the rounding error is at most $\log |\mathbb{B}_c| \leq \log |\mathbb{B}| = o(n)$. This proves (QR8). Since $\log \binom{n}{a} = nh(\alpha) + o(n)$, (QR7) gives $w \log |\mathbb{B}| = nh(\alpha) + o(n)$. In particular $w \rightarrow \infty$ because $\log |\mathbb{B}| = o(n)$. Substitution in (QR8) proves (QR9). $\square$

Remark 12.13 (Unbounded quotient degrees). If $c = c_n \rightarrow \infty$, the complete-fiber choice contributes at most $2^{n/c} = e^{o(n)}$. When $w \geq r$, the remainder is recovered exactly and no further remainder multiplicity occurs. When $w < r < c$, the exact formula (QR4) reduces the profile to an ordinary prefix problem for the remainder support, multiplied by $e^{o(n)}$ quotient choices. Thus only bounded quotient degrees can create a new positive-rate field-drop term; large degrees either cost $e^{o(n)}$ or return to an identity/remainder prefix problem. This does not prove flatness of the remaining prefix problem.

13. EFFECTIVE-IMAGE FOURIER NORMALIZATION

There is, however, an exact normalization that must be performed before declaring any character major. Assume $1 \leq m \leq |T| - 1$, as on every fixed-density primitive leaf, put $p = \text{char } \mathbb{B}$, and fix $S_0 \in \binom{T}{m}$ and $t_0 \in T$. Then

$$(EF0) \quad \text{Span}_{\mathbb{F}_p} \{ \Psi(S) - \Psi(S_0) : S \in \binom{T}{m} \} = \text{Span}_{\mathbb{F}_p} \{ g(t) - g(t_0) : t \in T \}.$$

Define the *effective Fourier span*

$$(EF1) \quad V_g = \text{Span}_{\mathbb{F}_p} \{ g(t) - g(t_0) : t \in T \} \subseteq \mathbb{B}^R, \quad A_{\text{eff}} = |V_g|.$$

The fixed-weight sum map takes values in the affine space $mg(t_0) + V_g$. Thus Fourier inversion belongs on the additive group V_g , not automatically on the ambient group $\mathbb{B}^R$.

Lemma 13.1 (Effective-span Fourier inversion). *For $z \in V_g$, let*

$$N_g(z) = \left| \left\{ S \in \binom{T}{m} : \sum_{t \in S} g(t) = mg(t_0) + z \right\} \right|.$$

Then

$$(EF2) \quad N_g(z) = \frac{1}{A_{\text{eff}}} \sum_{\chi \in \widehat{V}_g} \overline{\chi(z)} e_m(\chi(g(t) - g(t_0)) : t \in T).$$

In particular,

$$(EF3) \quad \max_z N_g(z) \leq \frac{1}{A_{\text{eff}}} \left(\binom{|T|}{m} + \sum_{1 \neq \chi \in \widehat{V}_g} |e_m(\chi(g(t) - g(t_0)) : t \in T)| \right).$$

Every character of V_g extends to a trace character of $\mathbb{B}^R$, and a nontrivial effective character is not constant on all the generators $g(t) - g(t_0)$. Ambient characters that are trivial on V_g occur with multiplicity $|\mathbb{B}|^R / A_{\text{eff}}$, which cancels exactly against the ambient Fourier normalization.

Proof. The inclusion from left to right in (EF0) follows by expanding a difference of two support sums. Conversely, for distinct $t, u \in T$, choose an $(m-1)$ -set $A \subseteq T \setminus \{t, u\}$; such a set exists because $1 \leq m \leq |T| - 1$. Then $\Psi(A \cup \{t\}) - \Psi(A \cup \{u\}) = g(t) - g(u)$, so the two spans agree. The first formula is character orthogonality on the finite additive group V_g , applied after translating by $mg(t_0)$; the second follows by the triangle inequality. The trace pairing on the $\mathbb{F}_p$ -space $\mathbb{B}^R$ is nondegenerate, so every character of its subspace V_g extends to the ambient space. A character trivial on each displayed generator is trivial on their span. Finally, the annihilator $V_g^\perp$ has size $|\mathbb{B}|^R / A_{\text{eff}}$; grouping ambient characters by their restriction to V_g proves the cancellation claim. $\square$

Definition 13.2 (Certified effective major/minor partition). *Let $\text{res} : \widehat{\mathbb{B}}^R \rightarrow \widehat{V}_g$ be restriction. A certified effective Fourier partition is a disjoint union*

$$\widehat{V}_g \setminus \{1\} = \mathfrak{m}_{\text{eff}} \sqcup \mathfrak{M}_{\text{eff}}$$

such that every $\chi \in \mathfrak{m}_{\text{eff}}$ is supplied with an ambient lift whose rational phase satisfies the character-sum hypotheses used to pay it. Characters without such a certified lift may be placed in $\mathfrak{M}_{\text{eff}}$, but their aggregate must then be paid by (MA). This designation is data of the chart; it is not a property of an unspecified ambient lift.

Definition 13.3 (Effective Fourier payment). *The full slice has an effective Fourier payment with constant $\kappa \geq 1$ if*

$$(EFP) \quad \sum_{1 \neq \chi \in \widehat{V}_g} |e_m(\chi(g(t) - g(t_0)) : t \in T)| \leq (\kappa - 1) \binom{|T|}{m}.$$

By Lemma 13.1, this implies the exact max-fiber bound

$$(EF4) \quad \max_z N_g(z) \leq \kappa \binom{|T|}{m} / A_{\text{eff}}.$$

It also implies that the realized image contains at least A_{eff}/κ points. The ambient instances of (PF) and (MA) are one sufficient way to prove (EFP), but (EFP) is phase-dependent and does not pay an ambient annihilator twice.

Proof of the image-size assertion. The fibers sum to $\binom{|T|}{m}$, and each nonempty fiber is at most $\kappa \binom{|T|}{m} / A_{\text{eff}}$ by (EF4). Hence at least A_{eff}/κ fibers are nonempty. $\square$

Lemma 13.4 (Constant and Artin–Schreier modes are annihilator modes). *Suppose an ambient trace character has phase constant on T . Then its restriction to V_g is trivial. In particular, if an ambient rational phase is Artin–Schreier of the form $G^p - G + c$ and G is defined at the $\mathbb{B}$ -rational points of T , then its character is an annihilator mode and disappears from the nontrivial sum in (EF2).*

Proof. The quotient of the character values at $g(t)$ and $g(t_0)$ is one for every t , so the character is trivial on the generators of V_g . For an Artin–Schreier phase, $\text{Tr}_{\mathbb{B}/\mathbb{F}_p}(G(t)^p - G(t)) = 0$, hence its additive-character value is constant on all points where it is defined. $\square$

Remark 13.5 (Effective normalization does not prove the remaining aggregate). The effective-span reduction removes exactly the constant annihilator multiplicity. A nonconstant character may still factor through a quotient map or have a phase-dependent estimate too weak for (EFP). Such modes require a recursive quotient estimate, the aggregate (MA), or a direct full-slice bound. Thus effective normalization repairs the Fourier denominator but is not a proof of unrestricted deep-prefix flatness.

Definition 13.6 (Major-character aggregate). *Let A_* be the additive target on which Fourier inversion is being performed, let $h : T \rightarrow A_*$, and let $\mathfrak{M}_* \subseteq \widehat{A}_* \setminus \{1\}$ be the designated major characters. The chart satisfies (MA) on A_* if*

$$(MA) \quad \sum_{\chi \in \mathfrak{M}_*} |e_m(\chi(h(t)) : t \in T)| \leq e^{o(|T|)} \binom{|T|}{m}.$$

Unless the ambient target is explicitly named, $A_ = V_g$ and $h(t) = g(t) - g(t_0)$. For the ambient target $A_* = \mathbb{B}^R$, the same condition is equivalently*

$$|\mathbb{B}|^{-R} \sum_{\alpha \in \mathfrak{M}_{\text{amb}}} |e_m(\psi(\alpha \cdot g(t)) : t \in T)| \leq e^{o(|T|)} \overline{N}_0.$$

The condition is vacuous when the designated major set is empty. It can be proved by a recursive quotient census or replaced by a direct full-slice max-fiber bound; first-match terminology alone does not imply it.

Lemma 13.7 (Sparse effective major set). *On the effective group V_g , let $\mathfrak{M}_{\text{eff}}$ be any set of nontrivial characters designated major. Its exact normalized Fourier loss is at most $|\mathfrak{M}_{\text{eff}}|$. Consequently $|\mathfrak{M}_{\text{eff}}| = e^{o(|T|)}$ proves the effective form of (MA); for a finite row, the literal cardinality is a valid major-character constant.*

Proof. Every fixed-weight Fourier coefficient is a sum of $\binom{|T|}{m}$ complex numbers of modulus one, and hence has absolute value at most that binomial coefficient. Sum this bound over the designated set. $\square$

Theorem 13.8 (Exact low-boundary-entropy closure). *Let $\Omega = \binom{T}{m}$, let W be a finite additive $\mathbb{F}_p$ -space, and let $\Psi : \Omega \rightarrow W$ be a fixed-weight additive boundary map. Fix $S_0 \in \Omega$, put $z_0 = \Psi(S_0)$, and define the effective difference span*

$$V = \text{Span}_{\mathbb{F}_p} \{ \Psi(S) - \Psi(S_0) : S \in \Omega \} \subseteq W.$$

Then $\Psi(\Omega) \subseteq z_0 + V$. Put

$$A = |V|, \quad L = |\Psi(\Omega)|, \quad M = \binom{|T|}{m}, \quad \bar{N} = M/L.$$

Designate every nontrivial character of V as major and no character as minor. Then the literal finite-row constants are

$$(SE1) \quad C_{\min} = 0, \quad C_{\text{maj}} \leq A - 1, \quad \max_z |\Psi^{-1}(z)| \leq L\bar{N} = M.$$

For every exact-agreement witness cell whose support projection is contained in Ω , its first-match slope set satisfies

$$(SE2) \quad |Z| \leq L\bar{N} = M.$$

Consequently, if $\log A = o(|T|)$, effective (MI), effective (MA), image-normalized Q , and the direct alternative in (RC) all hold with subexponential loss. Taking the support-restricted incidence over Ω as one ordered cell is exhaustive for that restricted incidence; it is exhaustive for the whole exact-agreement incidence when $\Omega = \binom{D}{a}$, or after an exhaustive partition into such slices. For a locator prefix with $|T| = \Theta(n)$, the concrete condition $(a - k - 1) \log |\mathbb{B}| = o(n)$ implies this hypothesis because $A \leq |\mathbb{B}|^{a-k-1}$.

Proof. Fourier inversion is performed on the effective group V . The minor sum is empty. Every fixed-weight Fourier coefficient has modulus at most M , so the normalized major loss is at most $A - 1$. The max-fiber estimate in (SE1) is the trivial bound by the whole support slice, written as the exact image-normalized multiplier L . By exact-agreement reduction, every bad slope in the cell has a noncommon m -support, and a fixed noncommon support carries at most one slope. Choosing one support for each slope is therefore injective and proves (SE2). Since $L \leq A$, all three finite multipliers $A - 1, L, L$ are $e^{o(|T|)}$ when $\log A = o(|T|)$. The one-cell atlas is exhaustive for its support-restricted incidence by construction; the stated whole-incidence alternatives give global exhaustivity. If it is refined by a partial-occupancy partition, (PO1)–(PO2) add the same support set back exactly and do not change the conclusion. $\square$

Definition 13.9 (Aggregate minor payment). *After effective-span normalization, let $\mathfrak{m}_{\text{eff}}$ be the nontrivial characters not designated major. The chart satisfies (MI) if*

$$(MI) \quad \sum_{\chi \in \mathfrak{m}_{\text{eff}}} |e_m(\chi(g(t) - g(t_0)) : t \in T)| \leq e^{o(|T|)} \binom{|T|}{m}.$$

For a finite row the corresponding exact loss is this sum divided by $\binom{|T|}{m}$. The pointwise condition (PF) is only one sufficient way to prove (MI); phase-dependent estimates may be strictly stronger.

Proposition 13.10 (Effective MI plus MA). *Partition the nontrivial effective characters as $\widehat{V}_g \setminus \{1\} = \mathfrak{m}_{\text{eff}} \sqcup \mathfrak{M}_{\text{eff}}$. Suppose (MI) and the effective form of (MA) in Definition 13.6 hold.*

Then the full slice and every first-match residual satisfy

$$(EF5) \quad \max_z |\{S : \Psi(S) = z\}| \leq e^{o(|T|)} \binom{|T|}{m} / A_{\text{eff}}.$$

The realized image has size at least $e^{-o(|T|)} A_{\text{eff}}$, so the same assertion holds at image normalization. For a finite row the exact multiplier is

$$(EF6) \quad 1 + \binom{|T|}{m}^{-1} \sum_{1 \neq \chi \in \widehat{V}_g} |e_m(\chi(g(t) - g(t_0)) : t \in T)|.$$

Proof. Add the minor and major sums and apply [Lemma 13.1](#) and [Definition 13.3](#). A first-match residual only deletes supports from a full-slice fiber. $\square$

Corollary 13.11 (Exact finite minor and major constants). *Put $M = \binom{|T|}{m}$ and, for an arbitrary partition of the nontrivial effective dual, define*

$$C_{\min} = M^{-1} \sum_{\chi \in \mathfrak{m}_{\text{eff}}} |E_m(\chi)|, \quad C_{\text{maj}} = M^{-1} \sum_{\chi \in \mathfrak{M}_{\text{eff}}} |E_m(\chi)|,$$

where $E_m(\chi) = e_m(\chi(g(t) - g(t_0)) : t \in T)$. Then

$$(EF7) \quad \max_z N_g(z) \leq \frac{M}{A_{\text{eff}}} (1 + C_{\min} + C_{\text{maj}}).$$

If every minor character has power sums bounded by Λ through order m , then

$$(EF8) \quad C_{\min} \leq \frac{|\mathfrak{m}_{\text{eff}}|}{M} \binom{\lceil \Lambda \rceil + m - 1}{m}.$$

Thus (EF7), with the two actual finite sums, records every finite Fourier loss in an adjacent-row certificate exactly.

Proof. Formula (EF7) is (EF3) after splitting the nontrivial characters. The cycle-index coefficient bound gives $|E_m(\chi)| \leq \binom{\lceil \Lambda \rceil + m - 1}{m}$, proving (EF8). $\square$

Definition 13.12 (Pointwise minor-arc range). *A primitive weighted prefix chart satisfies the pointwise minor-arc range if the characteristic is larger than m , every certified minor lift has a rational phase of total pole-degree at most Δ_{pole} and satisfies the nondegeneracy hypotheses of [Proposition 18.3](#), and, with*

$$\Lambda = C_0(\Delta_{\text{pole}} + 1) \sqrt{|\mathbb{B}|} + |P|$$

for the absolute Weil constant C_0 and planted exceptional set P , one has

$$(PF) \quad R \log |\mathbb{B}| + \log \binom{\Lambda + m - 1}{m} - \log \binom{|T|}{m} \leq o(|T|).$$

For circle twin-cosets the same definition is used with $2C_0(\Delta_{\text{pole}} + 1) \sqrt{|\mathbb{B}|}$ in place of $C_0(\Delta_{\text{pole}} + 1) \sqrt{|\mathbb{B}|}$, accounting for the two inverse-coset lifts on the torus before the involution quotient.

Part 3. Primitive Q and the residual compiler

14. THE PRIMITIVE PREFIX LEAF

After first-match removal of paid cells, a residual prefix-boundary chart has a fixed set of active coordinates T , $|T| = N$, and a fixed support size m . The full profile slice is

$$\Omega_{T,m} := \{x \in \{0, 1\}^T : \sum_{t \in T} x_t = m\},$$

and the primitive residual is a subset $\Omega^\circ \subseteq \Omega_{T,m}$. The superscript circle records a first-match residual; no Zariski openness is asserted. The normalization is taken from the full profile slice, not from the possibly much smaller residual mass. It is nevertheless the *realized image*, rather than the formal codomain, that supplies the default number of target fibers.

The three words in *primitive first-match residual* have separate content. “First match” means that a slope is removed as soon as one of its witnesses occurs in an earlier ordered cell; “residual” means that only witnesses surviving those removals remain; and “primitive” means that the row-specific quotient, planted, field-descent, rank, and ray-saturation certificates named by the atlas have all been applied. Primitivity does not mean merely trivial setwise stabilizer, and it does not assume the Fourier, Sidon, max-fiber, or ray estimates that will be imposed below.

The prefix syndrome map is a linear map $\Phi : \mathbb{Z}^T \rightarrow \mathbb{B}^R$, restricted first to the full slice and then to Ω° . In the unprofiled locator chart, when $R < \text{char } \mathbb{B}$, it is equivalent through Newton’s triangular change of coordinates to

$$(14.1) \quad \Phi(x) = \sum_{t \in T} x_t(t, t^2, \dots, t^R).$$

Equivalently, the augmented columns $\tilde{v}_t = (1, t, \dots, t^R)$ have constant first coordinate $\sum_t x_t = m$, and the other R coordinates are Φ . A residual quotient or rational chart is allowed to replace these columns by $v_t = \rho(t)(1, t, \dots, t^{R-1})$ only when that genuine weighted Vandermonde representation, with $\rho(t) \neq 0$, has been proved for the chart. It is not a consequence merely of constructibility or of deleting a Jacobian locus. The number R is the number of independent prefix equations; in the unprofiled MCA boundary leaf $R = w = a - K$.

Put $M = |\Omega_{T,m}|$, and write

$$(14.2) \quad \mathcal{S} = \Phi(\Omega_{T,m}), \quad L = |\mathcal{S}|, \quad A = |\mathbb{B}|^R, \quad \overline{N}^{\text{img}} = M/L, \quad \overline{N}^{\text{amb}} = M/A,$$

and for $s \in \mathcal{S}$ put $F_s = \Omega^\circ \cap \Phi^{-1}(s)$ and $f_s = |F_s|$. Unless an ambient full-image certificate has been proved, we set $\overline{N} = \overline{N}^{\text{img}}$. The certificate is

$$(FI) \quad L \geq e^{-o(N)} A.$$

The set $\mathcal{S}$ is the *effective image*: it consists of the prefix targets actually attained by the full profile slice. The formal ambient target space $\mathbb{B}^R$ can be much larger. Condition (FI), the *full-image certificate*, says that these two target counts differ by at most a subexponential factor. An *effective-image collapse* is the failure of this comparison; it must be retained as an effective-image rank-collapse profile, or handled throughout with image normalization. Under (FI), the two scales differ by only $e^{o(N)}$, so either may be used at exponential accuracy. If the certificate fails, effective-image collapse must be retained as an effective-image rank-collapse profile, or all estimates on this leaf must stay at image scale. Deleting earlier cells can only decrease every fiber.

The image-scale frontier condition is $\log M - \log L = o(N)$. The familiar ambient equation $\log \binom{N}{m} - R \log |\mathbb{B}| = o(N)$ is interchangeable with it only under (FI). This distinction is immaterial for a surjective prefix map but essential for a collapsed primitive image.

Lemma 14.1 (Exact image-ambient moment conversion). *For $q \geq 1$, extend f_s by zero from $\mathcal{S}$ to $\mathbb{B}^R$ and define*

$$\Gamma_q^{\text{img}} = \frac{1}{L} \sum_{s \in \mathcal{S}} \left(\frac{f_s}{\overline{N}^{\text{img}}} \right)^q, \quad \Gamma_q^{\text{amb}} = \frac{1}{A} \sum_{s \in \mathbb{B}^R} \left(\frac{f_s}{\overline{N}^{\text{amb}}} \right)^q.$$

Then

$$(14.3) \quad \Gamma_q^{\text{amb}} = \left(\frac{A}{L} \right)^{q-1} \Gamma_q^{\text{img}}.$$

Thus an ambient-normalized moment bound implies the corresponding image-normalized bound. The converse is valid with subexponential loss at logarithmic order only under (FI).

Proof. Since $\overline{N}^{\text{amb}} = M/A$ and $\overline{N}^{\text{img}} = M/L$, zero extension and one rearrangement give

$$\Gamma_q^{\text{amb}} = \frac{1}{A} \sum_{s \in \mathcal{S}} \left(\frac{A f_s}{M} \right)^q = \left(\frac{A}{L} \right)^{q-1} \frac{1}{L} \sum_{s \in \mathcal{S}} \left(\frac{L f_s}{M} \right)^q.$$

Because $L \leq A$, the asserted one-way implication follows. Under (FI), $(q-1) \log(A/L) = o(Nq)$, which proves the reverse transfer at the scale used here. $\square$

Remark 14.2 (Full-slice flatness certifies the image size). If a Fourier argument proves $\max_s |\Omega_{T,m} \cap \Phi^{-1}(s)| \leq e^{o(N)} \overline{N}^{\text{amb}}$, then averaging over the nonempty image fibers gives $A/L \leq e^{o(N)}$. Hence such an ambient flatness theorem itself proves (FI); image collapse cannot be silently assumed away before that estimate has been established.

Definition 14.3 (Primitive Q). *The Q branch is the primitive max-fiber problem: it asks whether one effective prefix target contains substantially more residual supports than the image-normalized average. A primitive prefix leaf satisfies Q with loss $e^{o(N)}$ if*

$$\max_{s \in S} f_s \leq e^{o(N)} \overline{N}^{\text{img}}.$$

A sequence of primitive leaves satisfies logarithmic-moment Q if, at every logarithmic order $q = q_N \rightarrow \infty$, $q \leq (\log N)^C$, used by the argument,

$$\Gamma_q^{\text{ord}} := \Gamma_q^{\text{img}} = L^{-1} \sum_{s \in S} \left(\frac{f_s}{\overline{N}^{\text{img}}} \right)^q \leq e^{o(Nq)}.$$

Here a logarithmic moment order means $q = q_N \rightarrow \infty$ with $q \leq (\log N)^C$ for a fixed C . The displayed normalized q -th moment is a Rényi-type probe of the fiber distribution: letting q grow slowly makes it approximate the largest normalized fiber without introducing an exponential moment-order cost. An ambient formulation is permitted only when it is proved directly or when (FI) has already been certified.

The next lemma is the primitive logarithmic moment equivalence used in the earlier RS–MCA reductions. It is included because it is the exact interface between moment estimates and max-fiber Q.

Lemma 14.4 (Logarithmic moments and primitive Q). *Assume $q = q_N \rightarrow \infty$ and $\log L/q = o(N)$. Then $\Gamma_q^{\text{ord}} \leq e^{o(Nq)}$ implies primitive Q. Conversely, primitive Q implies $\Gamma_q^{\text{ord}} \leq e^{o(Nq)}$.*

Proof. Let $M_* = \max_s f_s / \overline{N}^{\text{img}}$. Since $M_*^q \leq \sum_s (f_s / \overline{N}^{\text{img}})^q = L \Gamma_q^{\text{ord}}$, the moment bound gives $\log M_* \leq (\log L + \log \Gamma_q^{\text{ord}})/q = o(N)$. Conversely, $\sum_s f_s \leq |\Omega_{T,m}| = L \overline{N}^{\text{img}}$, so $\Gamma_q^{\text{ord}} \leq M_*^{q-1}$. Primitive Q gives $M_* = e^{o(N)}$, hence $\Gamma_q^{\text{ord}} = e^{o(Nq)}$. $\square$

The genuine Vandermonde condition removes the ordinary rank-defect alternative.

Lemma 14.5 (Augmented Vandermonde independence). *Let $t_1, \dots, t_r \in \mathbb{B}$ be distinct and $r \leq R+1$. Then the columns $\tilde{v}_{t_i} = (1, t_i, \dots, t_i^R)$ are linearly independent. The same holds for $\rho(t_i)(1, t_i, \dots, t_i^{R-1})$ when $r \leq R$ and every $\rho(t_i) \neq 0$.*

Proof. In either case the leading $r \times r$ minor, after removing the nonzero column weights when present, is $(t_i^j)_{0 \leq j < r, 1 \leq i \leq r}$. Its determinant is $\prod_{i < j} (t_j - t_i) \neq 0$. $\square$

In the high-energy argument below the Vandermonde lemma is not used: Boolean-cube growth gives the contradiction directly. The lemma certifies only the exact locator chart, or a residual chart for which the displayed weighted form has separately been verified; it does not classify arbitrary rank defects as paid cells.

15. PREFIX COORDINATES AND THE VANDERMONDE MAP

This section explains why the primitive residual has the linear Vandermonde form used in Section 14. Let $S \subseteq D$ have size m , and write its monic locator as

$$Q_S(X) = \prod_{t \in S} (X - t) = X^m + q_1(S)X^{m-1} + \dots + q_m(S).$$

The first R boundary coordinates may be taken either as elementary symmetric coefficients $q_1, \dots, q_R$ or, when $R < \text{char } \mathbb{B}$, as power sums $p_j(S) = \sum_{t \in S} t^j$, $1 \leq j \leq R$. Newton identities give triangular polynomial changes of variables between these two coordinate systems; the diagonal entries are $1, 2, \dots, R$, so the change is invertible in characteristic greater than R .

For the unprofiled locator chart, fixing the support size makes the zeroth power sum constant and Newton's identities identify the first R locator coefficients with the power sums $\sum_t x_t t^j$, $1 \leq j \leq R$. Thus (14.1) is exact. For a general residual chart one first fixes the planted roots, quotient data, common factors, and all coordinates charged to earlier cells. The fixed part contributes a constant syndrome s_0 . To use the primitive theorem one must then exhibit the remaining map explicitly, either as the same power-sum map, as a genuine common-weight Vandermonde map, or as a map for which the required Sidon estimate has been proved directly. Coordinate-dependent rational weights do not become a Vandermonde system merely because they are nonzero. Here "common-weight" means that one scalar $\rho(t)$ multiplies every monomial coordinate of the column at t , rather than using a different weight for each coordinate. This verification is part of the primitive-map requirement in Definition 11.1.

The fixed-density condition is also intrinsic. In a prefix leaf the total support size is fixed by the agreement threshold, and all paid planted blocks have fixed size. Thus the number of remaining choices is fixed. If a profile leaves several colors of coordinates, one obtains a product of fixed-density slices; the proof applies to each color after encoding the colors in a bounded alphabet. For the main statement we present the Boolean case because it is the MCA support case and because the Boolean-cube difference bound applies directly.

Lemma 15.1 (Newton-coordinate fiber equivalence). *Assume $R < \text{char } \mathbb{B}$. For supports of fixed size, the fibers of the first R elementary locator coefficients are exactly the fibers of the first R power sums. Consequently the Q max-fiber problem may be stated using either coordinate system.*

Proof. Newton identities give, for $1 \leq j \leq R$,

$$p_j + q_1 p_{j-1} + \cdots + q_{j-1} p_1 + j q_j = 0.$$

Since j is invertible in $\mathbb{B}$, q_j is determined by $p_1, \dots, p_j$, and conversely p_j is determined by $q_1, \dots, q_j$. The transformations are triangular inverse polynomial maps on the first R coordinates. $\square$

The *description entropy* of a profile family is the logarithm of the number of allowed auxiliary/profile descriptions. It measures the census of cells and is distinct from the algebraic description complexity of one cell, which measures the number and degrees of its defining equations.

Lemma 15.2 (Primitive profiles have subexponential multiplicity). *If the cell ledger fixes only bounded-complexity quotient, planted, tangent, extension, and split-pencil profiles, and the total description entropy of all data outside the moving support is $o(n)$ bits, then the number of primitive profiles is $e^{o(n)}$.*

Proof. There are at most $2^{o(n)} = e^{o(n)}$ descriptions by hypothesis. In particular, a planted or common block of size $o(n)$ has $\sum_{j=o(n)} \binom{n}{j} = e^{o(n)}$ possible supports, but recording $s = o(n)$ freely chosen field values costs $s \log_2 |\mathbb{B}|$ bits and is subexponential only when that product is $o(n)$. Positive-density data requires its own profile census and entropy payment; it cannot be discarded as a bounded-complexity label. $\square$

Remark 15.3. The subexponential profile lemma is the asymptotic counterpart of the finite ledger census. It is the reason a row proof must distinguish between a genuine cell payment and a mere change of variables. A change of variables with exponentially many charts would consume the whole frontier budget.

16. THE SIDON SPLIT

The Q moment can be large for two opposite reasons. A fiber may contain many repeated additive differences, which is the high-energy situation seen by inverse additive combinatorics, or it may be large while its differences almost never repeat, which is the Sidon-like situation. The split below separates these mechanisms before either one is estimated.

Every support is identified with its incidence vector in the torsion-free group $\mathbb{Z}^T$. In particular, all sums and differences in this section are integer operations, even when the Reed–Solomon coefficient field has characteristic two. For a finite set $F \subseteq \{0, 1\}^T$, define its additive energy

$$E(F) = |\{(a, b, c, d) \in F^4 : a - b = c - d\}|.$$

Let X, X' be independent uniform points of F . Then $E(F) = |F|^4 \sum_z \mathbb{P}(X - X' = z)^2$. It is convenient to normalize by

$$\Delta(F) = \frac{E(F)}{|F|^3}.$$

The quantity $\Delta(F)$ is an additive closure probability:

$$\Delta(F) = \mathbb{P}_{a,b,c \in F}(a - b + c \in F).$$

Indeed, once a, b, c are chosen, the fourth point in an energy quadruple is forced to be $d = a - b + c$.

Thus additive energy counts repeated differences, and $\Delta(F)$ is their normalized density. Values $\Delta(F) = e^{-o(N)}$ indicate substantial additive closure at exponential scale, whereas $\Delta(F) \leq e^{-\sigma N}$ is the low-energy, Sidon-like regime. The logarithmic moments below are the slowly growing normalized moments from [Definition 14.3](#); they measure how much of the Q obstruction is carried by fibers of either type. Throughout these sections, *positive rate* means a fixed positive exponential rate along a subsequence: for a fiber it means an excess of the form $e^{cN-o(N)}$, and for a logarithmic moment it means an excess of the form $e^{cNq-o(Nq)}$, for some constant $c > 0$.

For the fiber F_s , write $\Delta_s = E(F_s)/f_s^3$, with $\Delta_s = 0$ if $f_s = 0$.

Definition 16.1 (Sidon-heavy logarithmic moment). *For $\sigma > 0$, the Sidon-heavy part of the normalized moment is*

$$\Gamma_{q,\sigma}^{\text{sid}} = L^{-1} \sum_{\Delta_s \leq e^{-\sigma N}} \left(\frac{f_s}{N} \right)^q.$$

A sequence has no positive-rate Sidon-heavy obstruction if $\Gamma_{q,\sigma}^{\text{sid}} \leq e^{o(Nq)}$ for every fixed $\sigma > 0$ and every logarithmic q . Here “Sidon” refers to the exponentially small normalized energy Δ_s , while “heavy” refers to the positive-exponential contribution of those fibers to the normalized logarithmic moment; it does not assert that every low-energy fiber is individually large.

The following dichotomy is the correct replacement for the naive assertion that collision excess directly creates small doubling.

Proposition 16.2 (Ordinary moment split). *Fix $\eta > 0$ and $\sigma > 0$. If $\Gamma_q^{\text{ord}} \geq e^{\eta Nq}$ and $\Gamma_{q,\sigma}^{\text{sid}} \leq \frac{1}{2}e^{\eta Nq}$, then there is a syndrome s such that*

$$f_s \geq e^{\eta N - o(N)} \overline{N}, \quad \Delta_s > e^{-\sigma N}.$$

If no positive-rate Sidon-heavy obstruction is present, then any positive-rate moment failure gives, after taking $\sigma \rightarrow 0$ slowly, a fiber with $f_s \geq e^{\eta N - o(N)} \overline{N}$ and $E(F_s) \geq f_s^3 e^{-o(N)}$.

Proof. The high-energy part of the moment is at least $\frac{1}{2}e^{\eta Nq}$. Since there are L syndromes, some high-energy syndrome satisfies

$$\left(\frac{f_s}{N} \right)^q \geq \frac{L}{2} e^{\eta Nq} \cdot L^{-1} = \frac{1}{2} e^{\eta Nq}.$$

Thus $f_s \geq e^{\eta N - o(N)} \overline{N}$, and by construction $\Delta_s > e^{-\sigma N}$. If $\Gamma_{q,\sigma}^{\text{sid}} \leq e^{o(Nq)}$ for every fixed σ , apply the first part along a sequence $\sigma_j \downarrow 0$ and diagonalize; the normalization $\Delta_s = E(F_s)/f_s^3$ gives the energy conclusion. $\square$

The proof also explains the no-go point. Suppose a fiber is heavy but $\Delta_s \leq e^{-\sigma N}$. A closure-rich diagram using tests $a - b + c \in F_s$ pays a factor Δ_s for each nonredundant closure. Such diagrams cannot account for the ordinary contribution f_s^q of q independent choices in F_s . Therefore Sidon-heavy fibers must be removed by a separate cell; they cannot be forced into the high-energy inverse branch.

Definition 16.3 (Sidon moment payment). *A primitive prefix leaf has a Sidon moment payment if $\Gamma_{q,\sigma}^{\text{sid}} = e^{o(Nq)}$ for every fixed $\sigma > 0$ and every logarithmic moment order used by the primitive-Q argument, uniformly over profiles. This is an analytic support-fiber statement. It is not equivalent to a first-match distinct-slope payment; the latter still requires (RC) or a direct ray bound.*

The equivalence in the definition follows from [Lemma 14.4](#) restricted to the low-energy set of syndromes. Operationally, one proves payment either directly by bounding these fibers or indirectly by proving Fourier flatness on the primitive complement of algebraic major arcs.

17. WHY THE SIDON CELL IS NECESSARY

It is tempting to try to prove primitive Q by expanding the logarithmic moment into diagrams and then showing that every primitive diagram contains many additive closure tests. That strategy is correct only for a closure-rich submoment. It cannot capture the ordinary moment on a heavy low-energy fiber. The obstruction is simple enough to record as a lemma.

Let $F = F_s$ be a fiber of size f . A closure test is an equation of the form $a - b + c \in F$, with $a, b, c \in F$ already chosen. Its success probability is $\Delta(F) = E(F)/f^3$. A closure-rich diagram with ℓ nonredundant closure tests and r free choices contributes, up to a subexponential diagram factor,

$$f^r \Delta(F)^\ell.$$

After the free coordinates have been chosen, in the image normalization of a primitive prefix leaf, the remaining $q - r$ coordinates are supplied only at the random scale $\bar{N}$. Thus its contribution is at most

$$e^{o(Nq)} \bar{N}^{q-r} f^r \Delta(F)^\ell.$$

Lemma 17.1 (No-go for ordinary moment capture). *Assume that $f \geq e^{\eta N \bar{N}}$, $\Delta(F) \leq e^{-\sigma N}$, and that a family of diagrams has $r = o(q)$ free fiber choices and $\ell \geq cq$ nonredundant closure tests. Then the total contribution of these diagrams is smaller than the ordinary fiber contribution f^q by a factor*

$$\exp(-(\eta + c\sigma)Nq + o(Nq))$$

up to changing η by $o(1)$.

Proof. For one diagram the ratio to f^q is at most

$$e^{o(Nq)} \left(\frac{\bar{N}}{f}\right)^{q-r} \Delta(F)^\ell \leq e^{o(Nq)} \exp(-\eta N(q-r) - c\sigma Nq).$$

Since $r = o(q)$, this is $\exp(-(\eta + c\sigma)Nq + o(Nq))$. The number of bounded-complexity logarithmic diagrams is $e^{O(q \log q)} = e^{o(Nq)}$, so summing over them does not change the exponential rate. $\square$

The lemma has an important conceptual consequence. A large Sidon-like fiber is not a hidden approximate group. It is a different phenomenon: many independent choices land in the same syndrome, but their pairwise differences almost never repeat. Additive-combinatorics inverse tools are designed for the opposite situation, where repeated differences are abundant. Therefore the first-match ledger must contain a Sidon/Fourier cell before the primitive inverse argument begins.

The valid implication is exactly [Proposition 16.2](#). Ordinary moment mass splits into a low-energy part, which needs a separately proved Sidon moment estimate, and a high-energy part, which is dispatched by [Proposition 21.1](#).

This mechanism does not replace a prefix-flatness or Sidon-payment theorem; it identifies exactly where one is needed. The first-match order is: algebraic major arcs first, then a separately certified Sidon/Fourier cell, and only then the high-energy primitive inverse step. Naming a large low-energy fiber a cell does not pay it.

18. FOURIER FLATNESS AS A SIDON PAYMENT

This section records the general weighted form of the Fourier/Sidon payment. It is phrased for an arbitrary finite field $\mathbb{B}$, because the primitive residual charts that come from the moduli ledger need not have unweighted power-sum coordinates. Let T be a finite set, $|T| = N$, let $g : T \rightarrow \mathbb{B}^R$ be any function, and set

$$\Omega_{T,m} = \{x \in \{0,1\}^T : \sum_{t \in T} x_t = m\}, \quad \Psi(x) = \sum_{t \in T} x_t g(t).$$

In applications $g(t) = (\rho_0(t), \rho_1(t)t, \dots, \rho_{R-1}(t)t^{R-1})$ with nonzero rational weights ρ_i on the primitive chart.

Fourier inversion has three pieces. The zero character produces the average-fiber main term. On the effective dual, the chart-specific partition of [Definition 13.2](#) designates certified minor characters and the complementary major characters. The ambient algebraic exceptional locus of [Definition 18.5](#) helps construct that partition but does not canonically determine it, because an effective character has many ambient lifts. The pointwise condition (PF) controls every certified minor lift, while the aggregate condition (MA) controls the sum of all designated major characters. Neither estimate substitutes for the other.

The condition (PF') below is the raw power-sum flatness inequality when one has a uniform bound for every nonzero character. Condition (PF) is its row-specific Weil specialization for minor characters after substituting the pole-degree bound; the omitted major characters are then handled by (MA).

Theorem 18.1 (Prefix flatness from power-sum control). *Fix a nontrivial additive character ψ of $\mathbb{B}$. For $\alpha \in \mathbb{B}^R \setminus \{0\}$ and $1 \leq j \leq m$, put*

$$P_j(\alpha) = \sum_{t \in T} \psi(j\alpha \cdot g(t)).$$

Assume $|P_j(\alpha)| \leq \Lambda$ for every $\alpha \neq 0$ and every $1 \leq j \leq m$. Then, for every $z \in \mathbb{B}^R$,

$$|\Psi^{-1}(z)| \leq |\mathbb{B}|^{-R} \binom{N}{m} + \binom{\Lambda + m - 1}{m}.$$

Consequently, if

$$(PF') \quad R \log |\mathbb{B}| + \log \binom{\Lambda + m - 1}{m} - \log \binom{N}{m} \leq o(N),$$

then $\max_z |\Psi^{-1}(z)| \leq e^{o(N)} |\mathbb{B}|^{-R} \binom{N}{m}$.

Proof. By orthogonality,

$$|\Psi^{-1}(z)| = |\mathbb{B}|^{-R} \sum_{\alpha \in \mathbb{B}^R} \psi(-\alpha \cdot z) e_m(\psi(\alpha \cdot g(t)) : t \in T),$$

where e_m denotes the m -th elementary symmetric coefficient. The term $\alpha = 0$ is the main term $|\mathbb{B}|^{-R} \binom{N}{m}$. For $\alpha \neq 0$, write $a_t = \psi(\alpha \cdot g(t))$. The generating identity $\sum_{r \geq 0} e_r(a_t) u^r = \exp(\sum_{j \geq 1} (-1)^{j-1} (\sum_t a_t^j) u^j / j)$ and the hypothesis $|\sum_t a_t^j| \leq \Lambda$ for $j \leq m$ give $|e_m(a_t : t \in T)| \leq [u^m](1-u)^{-\Lambda} = \binom{\Lambda+m-1}{m}$. Summing the nontrivial characters and using the outside factor $|\mathbb{B}|^{-R}$ gives the first bound. The displayed condition (PF') makes the error term subexponential relative to the random-fiber main term. $\square$

Remark 18.2 (Small characteristic cycles). If some integers $j \leq m$ vanish in the characteristic, the formal Li–Wan cycle index [\[LW10\]](#) separates the cycles divisible by $p = \text{char } \mathbb{B}$; the corresponding coefficient majorant is $[u^m](1-u)^{-\Lambda}(1-u^p)^{-(N-\Lambda)/p}$. This bookkeeping identity does not supply the missing character-sum bounds for those cycles, exclude Artin–Schreier phases, or prove a uniform Sidon payment. The pointwise theorem above is therefore deliberately stated in the large-characteristic range. A small-characteristic leaf needs a direct cycle-by-cycle certificate or a separate image-normalized Q/Sidon theorem.

Proposition 18.3 (Weighted Weil minor arcs). *Let T be either a multiplicative coset in $\mathbb{B}^\times$ or the x -coordinate of a torus twin coset. Suppose that every certified minor lift $\alpha \in \mathbb{B}^R$ has phase $R_\alpha(t) = \alpha \cdot g(t)$ which is a nonconstant rational function on the corresponding genus-zero curve, is not of Artin–Schreier form, and has total pole-degree at most Δ_{pole} . If $m < \text{char } \mathbb{B}$, then the stated power-sum bound holds for every certified minor lift, with*

$$\Lambda = C_0(\Delta_{\text{pole}} + 1)\sqrt{|\mathbb{B}|} + |P|$$

for multiplicative cosets, and with $2C_0(\Delta_{\text{pole}} + 1)\sqrt{|\mathbb{B}|} + |P|$ for circle twin-cosets. Here P is the planted exceptional set and C_0 is an absolute Weil constant.

Proof. For a multiplicative coset $\theta H \subseteq \mathbb{B}^\times$, write the coset indicator as the average of the multiplicative characters of $\mathbb{B}^\times/H$. Each nontrivial power sum becomes an average of mixed multiplicative-additive character sums $\sum_x \chi(x)\psi(jR_\alpha(\theta x))$. Since $1 \leq j < m < \text{char } \mathbb{B}$, multiplication by j does not make the rational function constant or Artin–Schreier. Weil’s bound for one-variable mixed character sums [Wei48] gives $C_0(\Delta_{\text{pole}} + 1)\sqrt{|\mathbb{B}|}$. Removing planted exceptional points changes the sum by at most $|P|$. The circle case is the same after pulling back by $x = (u + u^{-1})/2$ to the norm-one torus; the two torus branches give the factor two. $\square$

Theorem 18.4 (Unconditional shallow-prefix Fourier flatness). *Let $p \rightarrow \infty$ through odd primes. For each p , let T_p be either a multiplicative coset in $\mathbb{F}_p^\times$ or a Chebyshev twin-coset x -domain over $\mathbb{F}_p$, in the latter case with $p \equiv 3 \pmod{4}$, and put $N_p = |T_p|$. Fix $0 < \alpha < 1/2$, choose integers m_p, R_p with*

$$(SFM1) \quad 1 \leq R_p, \quad \alpha \leq \frac{m_p}{N_p} \leq 1 - \alpha, \quad g_p(t) = (t, t^2, \dots, t^{R_p}), \quad R_p\sqrt{p} = o(N_p),$$

and let $\Psi_p(S) = \sum_{t \in S} g_p(t)$ on $\binom{T_p}{m_p}$. Then $V_{g_p} = \mathbb{F}_p^{R_p}$, and the certified effective partition may take every nontrivial effective character to be minor and no character to be major. For some $c_\alpha > 0$,

$$(SFM2) \quad \sum_{1 \neq \chi \in \widehat{V_{g_p}}} |e_{m_p}(\chi(g_p(t) - g_p(t_0)) : t \in T_p)| \leq e^{-c_\alpha N_p} \binom{N_p}{m_p}.$$

Consequently effective (MI) holds, (MA) is vacuous, and

$$(SFM3) \quad \max_z |\Psi_p^{-1}(z)| \leq (1 + e^{-c_\alpha N_p}) p^{-R_p} \binom{N_p}{m_p}.$$

Indeed every $z \in \mathbb{F}_p^{R_p}$ is attained for all sufficiently large rows, and uniformly in z ,

$$(SFM4) \quad |\Psi_p^{-1}(z)| = p^{-R_p} \binom{N_p}{m_p} (1 + O(e^{-c_\alpha N_p})).$$

Thus these smooth and circle prefix charts have unconditional effective Fourier payment and image coverage in the shallow range (SFM1).

Proof. Condition (SFM1) implies $R_p < p$, $R_p < N_p$, and $R_p \log p = o(N_p)$. Moreover $N_p \leq p + 1$, so the fixed upper density bound gives $m_p < p$ for every sufficiently large p . For a nonzero ambient parameter, let $d = \max\{j : a_j \neq 0\}$. Then $1 \leq d \leq R_p < p$, and $P(X) = \sum_{j=1}^{R_p} a_j X^j$ has a pole of order d at infinity. Every pole of a nonconstant Artin–Schreier function $G^p - G + c$ has order divisible by p , so P is non-Artin–Schreier. In the circle case $P((u + u^{-1})/2)$ has poles of the same order $d < p$ at the two ends of the torus compactification and is likewise non-Artin–Schreier.

For a multiplicative coset, the character expansion of its indicator and the mixed Weil bound give, uniformly for $1 \leq j \leq m_p < p$,

$$\left| \sum_{t \in T_p} \psi(jP(t)) \right| = O(R_p\sqrt{p}) = o(N_p).$$

For a circle domain, pullback by $x = (u + u^{-1})/2$ gives the same conclusion with an inessential factor two. Put $\Lambda_p = O(R_p\sqrt{p}) = o(N_p)$. The cycle-index estimate therefore bounds every nontrivial fixed-weight Fourier coefficient by

$$\binom{\lceil \Lambda_p \rceil + m_p - 1}{m_p} = e^{o(N_p)}.$$

If a nonzero ambient character annihilated V_{g_p} , its character value would be constant on T_p , making the power sum have modulus N_p , contrary to the preceding $o(N_p)$ bound. Thus the annihilator is zero and $V_{g_p} = \mathbb{F}_p^{R_p}$. There are $p^{R_p} = e^{o(N_p)}$ effective characters, whereas Stirling's formula and the density interval give $\binom{N_p}{m_p} \geq e^{(h(\alpha)+o(1))N_p}$. Summing the coefficient bounds proves (SFM2), after decreasing the positive constant c_α . All nontrivial effective characters have just received a certified minor lift, so the major set is empty. Finally Fourier inversion shows that the absolute nonzero-character error in each fiber is $e^{o(N_p)}$, whereas its main term is $p^{-R_p} \binom{N_p}{m_p} = e^{(h(m_p/N_p)+o(1))N_p}$. This proves (SFM3)–(SFM4) and eventual surjectivity. $\square$

Definition 18.5 (Ambient algebraic exceptional locus). *An ambient phase parameter is algebraically exceptional if its weighted phase is constant on a routed quotient fiber, is Artin–Schreier degenerate, has a pole or zero supported on a positive-density moving divisor, or factors through a proper quotient, Chebyshev map, field-descent locus, tangent rank drop, differential-locator defect, or ray-saturation collapse. On an effective chart this locus is used only to construct the certified partition of Definition 13.2; it is not itself a numerical estimate. In particular, an Artin–Schreier phase defined on the whole active set restricts to the trivial effective character by Lemma 13.4.*

Proposition 18.6 (Major-arc algebraic interface). *Each condition in Definition 18.5 is a constructible condition on the dual phase parameters and points to a quotient, Chebyshev, field-descent, planted-divisor, differential, or rank-loss profile. This classification does not imply that the character vanishes after primal first-match deletion and does not prove (MA). A flatness argument must also bound (MA), recursively or directly.*

Proof. Constancy on a quotient fiber, Chebyshev factorization, Artin–Schreier degeneracy, a positive-density pole divisor, and a rank drop are respectively coefficient, Frobenius, resultant, or determinantal conditions. They therefore identify constructible dual loci on which a recursive estimate can be attempted. Fourier inversion sums characters, not primal witnesses, so no support-level payment follows from this fact. $\square$

Proposition 18.7 (Frontier Weil separation). *Let $N = |T|$, let $m = \theta N + o(N)$ with $0 < \theta < 1$, and suppose a primitive smooth or circle chart has minor-arc power-sum parameter $\Lambda = o(N)$. If the ambient random-fiber scale $\bar{N} = |\mathbb{B}|^{-R} \binom{N}{m}$ satisfies $\log \bar{N} = o(N)$, or more generally $\log \bar{N} \geq -o(N)$, then the prefix-flat inequality (PF') holds. In particular, minor arcs alone cannot create a positive-rate max-fiber or Sidon-heavy moment at this near-zero ambient entropy scale.*

Proof. The elementary estimate $\binom{\Lambda+m-1}{m} = \binom{\Lambda+m-1}{\Lambda-1} \leq (e(m+\Lambda)/\Lambda)^\Lambda$ gives $\log \binom{\Lambda+m-1}{m} = o(N)$, since $\Lambda = o(N)$. The relative minor-arc error exponent is

$$\log \binom{\Lambda+m-1}{m} - \log \bar{N} = R \log |\mathbb{B}| + \log \binom{\Lambda+m-1}{m} - \log \binom{N}{m}.$$

The assumed frontier inequality makes this at most $o(N)$; it may be negative by a linear amount, which is only stronger. Thus the minor-arc contribution to every full-profile fiber is at most $e^{o(N)} \bar{N}$. This statement concerns only minor characters; a full flatness conclusion also requires (MA). $\square$

Theorem 18.8 (Sidon payment in the pointwise flat range). *Let $\Omega^\circ \subseteq \Omega_{T,m}$ be a primitive first-match residual of a smooth or circle chart, with weighted prefix map Ψ and ambient scale $\bar{N}^{\text{amb}} = |\mathbb{B}|^{-R} \binom{N}{m}$. Put $A_{\text{amb}} = |\mathbb{B}|^R$ and $L = |\Psi(\Omega_{T,m})|$. Suppose the pointwise minor-arc*

condition (PF) and the ambient instance of the aggregate condition (MA) hold. Then, for every fixed $\sigma > 0$ and every logarithmic $q \leq (\log N)^C$, for a fixed C ,

$$\Gamma_{q,\sigma}^{\text{sid}}(\Omega^\circ) \leq e^{o(Nq)},$$

where Γ^{sid} uses the image normalization of [Definition 16.1](#). Equivalently, no such pointwise-flat residual leaf supports positive-rate low-energy heavy fibers.

Proof. It is enough to bound every residual fiber, and a residual is contained in the full slice. Fourier inversion on that slice splits into the zero, minor, and major characters. The zero character contributes $\overline{N}^{\text{amb}}$. The cycle-index estimate and [Proposition 18.3](#), together with (PF), bound the complete minor contribution by $e^{o(N)}\overline{N}^{\text{amb}}$. Condition (MA) bounds the complete major contribution by the same quantity. Thus every full-slice, and hence every residual, fiber is $e^{o(N)}\overline{N}^{\text{amb}}$. This proves the full-image certificate by [Remark 14.2](#), and [Lemma 14.1](#) transfers the ambient moment bound to image normalization. $\square$

Remark 18.9 (Role of the Sidon-heavy cell). The theorem implements “pay structure, flatten Sidon, inverse high energy” only in the range where both (PF) and (MA) are verified. At entropy depth a phase can have degree up to R , and the raw Weil parameter is of order $R\sqrt{|\mathbb{B}|}$; this is not automatically $o(N)$ when $R \asymp N/\log|\mathbb{B}|$. Outside the pointwise-flat range a different Sidon payment remains necessary.

19. EXTERNAL ADDITIVE COMBINATORICS

Only two external additive-combinatorics inputs are used in the primitive high-energy proof.

Theorem 19.1 (Balog–Szemerédi–Gowers, energy form). *There is an absolute constant C such that if A is a finite subset of an abelian group and $E(A) \geq |A|^3/K$, then there is $A' \subseteq A$ with $|A'| \geq K^{-C}|A|$ and $|A' - A'| \leq K^C|A'|$.*

This is the usual energy form of the Balog–Szemerédi–Gowers theorem; polynomial dependence on K follows from Gowers’s refinement of the Balog–Szemerédi theorem [[BS94](#), [Gow98](#)]. The difference-set form follows from the sumset form by replacing A by $-A$ and using standard Ruzsa calculus [[TV06](#)], or directly by applying the graph form to differences.

Theorem 19.2 (Boolean-cube difference growth). *If $U \subseteq \{0, 1\}^d$, then*

$$|U - U| \geq |U|^{3/2}.$$

This is the Boolean-cube specialization of a broader sumset inequality of Green–Matolcsi–Ruzsa–Shakan–Zhelezov [[GMR+22](#)]. It also follows from the induced-doubling identity of Matolcsi–Ruzsa–Shakan–Zhelezov [[MRSZ22](#)]. Only the displayed Boolean statement is used below.

Remark 19.3. Entropy inverse theorems of Tao and the entropy-doubling theory of Green–Manners–Tao are compatible with this proof and are often useful after one has produced small-doubling trades [[Tao10](#), [GMT23](#)]. They are not needed as hypotheses here. The Boolean-cube growth theorem gives a direct contradiction to a large small-doubling subset of a primitive Boolean support fiber.

20. ENTROPY INVERSE THEORY AND ITS ROLE

The terminology “Tao–Gowers method” can refer to two related but distinct steps. The first is BSG, used directly in [Proposition 21.1](#). The second is Freiman-type structural theory, meaning inverse theorems that model a set with small sumset or difference set by a generalized arithmetic progression or a coset progression. It is not needed for the Boolean contradiction, but it can suggest candidate geometric strata in more general bounded-alphabet residuals; those strata still require direct payment proofs.

A *bounded integer trade* is a vector $Y \in \mathbb{Z}^T$ whose coordinates are bounded independently of N , with $\Phi(Y) = 0$ and, on a fixed-weight slice, $\sum_t Y_t = 0$. If Y is random, Y' is an independent

copy, and $H(Y - Y') \leq H(Y) + o(N)$, then entropy BSG and entropy Freiman theory produce, after conditioning and losing $o(N)$ entropy, a set model with doubling $e^{o(N)}$. Tao's entropy inverse theorem gives this in *coset-progression* language—a subgroup coset plus a generalized arithmetic progression—for Shannon entropy [Tao10]; Green–Manners–Tao recast the theory in terms of entropic doubling and Ruzsa distance [GMT23]; Goh formulates an entropic analogue of additive energy and the entropic BSG connection [Goh24]. In a torsion model one may also use polynomial Freiman–Ruzsa technology after reducing bounded integer trades modulo a large prime with no wraparound.

In the present paper the Boolean-cube difference theorem makes this structural compression unnecessary for the primitive Q step. Once BSG finds $A' \subseteq F_s \subseteq \{0, 1\}^T$ with $|A' - A'| \leq e^{o(N)}|A'|$, Boolean-cube difference growth already forces $|A'| = e^{o(N)}$. This contradicts the exponential size supplied by moment excess. Therefore there is no separate entropy–Freiman hypothesis in the proof.

For non-Boolean residual trades, entropy inverse theory can suggest coset-progression incidence strata or rank-defect tests. This is a useful heuristic for constructing a row-specific atlas, not an exhaustive dichotomy: each suggested stratum still needs the explicit first-match census and slope-projection payment required by Definition 11.1, and arbitrary weighted charts are not covered by Lemma 14.5. No such heuristic is used in the proof of the Boolean primitive theorem.

21. PRIMITIVE Q AFTER THE SIDON CUT

The key primitive theorem is now short and self-contained modulo Theorems 19.1 and 19.2. It also addresses the main obstruction: ordinary Rényi mass can be Sidon-heavy, but after that cell is removed it must be high-energy, and high-energy is impossible in a large Boolean fiber.

Proposition 21.1 (High-energy Boolean fibers are small). *Let $F \subseteq \{0, 1\}^T$, $|T| = N$. If $E(F) \geq |F|^3/K$, then $|F| \leq K^{O(1)}$. Consequently an exponential set $|F| \geq e^{cN-o(N)}$ cannot have $E(F) \geq |F|^3/e^{o(N)}$.*

Proof. Apply Theorem 19.1. There is $F' \subseteq F$ with $|F'| \geq K^{-C}|F|$ and $|F' - F'| \leq K^C|F'|$. Since $F' \subseteq \{0, 1\}^T$, Theorem 19.2 gives $|F' - F'| \geq |F'|^{3/2}$. Hence $|F'|^{1/2} \leq K^C$, so $|F| \leq K^{3C}$ after increasing C . Taking $K = e^{o(N)}$ rules out $|F| \geq e^{cN-o(N)}$. $\square$

Theorem 21.2 (Primitive Q after a Sidon moment payment). *Let $\Omega_N^0 \subseteq \{0, 1\}^{T_N}$ be fixed-density full profile slices, with $|T_N| = N$ and $m_N/N \in [\alpha, 1 - \alpha]$, and let $\Omega_N^0 \subseteq \Omega_N^0$ be primitive first-match residuals. Suppose there is a logarithmic moment order $q_N \rightarrow \infty$ with $\log L_N/q_N = o(N)$, where $L_N = |\Phi_N(\Omega_N^0)|$ and $\bar{N}_N = |\Omega_N^0|/L_N$, for which $\Gamma_{q_N, \sigma}^{\text{sid}} \leq e^{o(Nq_N)}$ for every fixed $\sigma > 0$. Then the leaves satisfy primitive Q:*

$$\max_s |\Omega_N^0 \cap \Phi_N^{-1}(s)| \leq e^{o(N)} \bar{N}_N.$$

Proof. Assume primitive Q fails at positive exponential rate. The lower half of the moment sandwich gives $\Gamma_{q_N}^{\text{ord}} \geq e^{\eta N q_N - o(Nq_N)}$ along a subsequence. The Sidon moment hypothesis and Proposition 16.2 then give a fiber F_s with $f_s \geq e^{\eta N - o(N)} \bar{N}_N$ and $E(F_s) \geq f_s^3 e^{-o(N)}$. Since $\bar{N}_N = |\Omega_N^0|/L_N \geq 1$, this fiber has size $e^{\eta N - o(N)}$, contradicting Proposition 21.1 with $K = e^{o(N)}$. Thus no positive-rate max-fiber failure occurs. $\square$

22. Q, SP, AND THE MCA COMPILER

This section uses the prefix rigidity, exact second moment, and slope-elimination results of [Cho26, Secs. 19–21 and Apps. E–F] as imported inputs. The new content is the occupancy/ray compiler and the precise hypotheses under which Q controls MCA slopes.

The *Q branch* is the maximum-prefix-fiber problem. The *SP branch* (support-pair, or shift-pair, branch) is its second-moment problem: it counts ordered pairs of supports in one prefix fiber. These names label two counting interfaces and carry no assertion that either bound has already been proved. Neither count is automatically the MCA numerator: a bad slope begins with one support, while converting support pairs to projective slope directions requires a proved

multiplicity or incidence bound. We first record the exact moment identities and then state that separate ray compiler.

Definition 22.1 (Q and SP ledgers). *For a first-match residual profile λ , let Ω_λ^0 be its full profile slice, let $\Omega_\lambda^\circ \subseteq \Omega_\lambda^0$ be the residual, and put*

$$\mathcal{S}_\lambda = \Phi_\lambda(\Omega_\lambda^0), \quad L_\lambda = |\mathcal{S}_\lambda|, \quad \bar{N}_\lambda = |\Omega_\lambda^0|/L_\lambda, \quad Q_\lambda(s) = |\Omega_\lambda^\circ \cap \Phi_\lambda^{-1}(s)|.$$

The Q ledger is discharged if $\max_s Q_\lambda(s) \leq e^{o(n)} \bar{N}_\lambda$ for every primitive profile. The support-pair SP statistic is $\sum_s Q_\lambda(s)^2$. A distinct-ray ledger is discharged only by a direct slope bound or by (RC) below. If the ambient target count $|\mathbb{B}_\lambda|^{R_\lambda}$ is used in place of L_λ , the ledger must also record the full-image certificate (FI) or a direct ambient theorem.

The compiler uses two elementary inequalities repeatedly. They are included to make explicit how the logarithmic moment, max-fiber Q, and shift-pair second moment interact.

Imported moment inequality. For fiber sizes $N(s)$, mean $\bar{N}$, ratios $R(s) = N(s)/\bar{N}$, and $\Gamma_r = L^{-1} \sum_s R(s)^r$, CAP25 v13.2 [Cho26, Sec. 19 and Apps. E–F] gives the elementary bounds

$$(M) \quad \max_s R(s) \leq (L\Gamma_r)^{1/r}, \quad \Gamma_r \leq (\max_s R(s))^{r-1}.$$

Corollary 22.2 (Finite moment floor). *Suppose a finite row has remaining Q bit margin Δ , meaning that one needs $\max_s R(s) \leq 2^\Delta$. A moment proof of order r can certify this only if*

$$\log_2 \Gamma_r + \log_2 L \leq r\Delta,$$

where L is the number of realized targets in the selected normalization. In an ambient prefix proof this specializes to $\log_2 \Gamma_r^{\text{amb}} + R \log_2 |\mathbb{B}| \leq r\Delta$; transferring an image-normalized moment to that statement additionally costs $(r-1) \log_2 (|\mathbb{B}|^R/L)$.

Proof. The moment sandwich gives $\log_2 \max_s R(s) \leq (\log_2 L + \log_2 \Gamma_r)/r$. To make this at most Δ one needs the displayed inequality. $\square$

The finite corollary explains why a small adjacent margin cannot be closed by a fixed-order moment when $R \log |\mathbb{B}|$ is large.

Definition 22.3 (Explanation image and retained-support occupancy). *Fix $a \geq k+1$, a received pair $r = (r_0, r_1)$, and a cell $\mathcal{C} \subseteq \mathcal{W}_a(r)$. Its explanation image and slope image are*

$$\mathcal{R}(\mathcal{C}) = \{(\gamma, h) : \exists S \ (\gamma, S, h) \in \mathcal{C}\}, \quad Z(\mathcal{C}) = \{\gamma : \exists h \ (\gamma, h) \in \mathcal{R}(\mathcal{C})\}.$$

For $\rho = (\gamma, h) \in \mathcal{R}(\mathcal{C})$, its retained-support occupancy is

$$\nu_{\mathcal{C}}(\rho) = |\{S \in \binom{D}{a} : (\gamma, S, h) \in \mathcal{C}\}|.$$

Thus the witness projection factors canonically as

$$\mathcal{C} \longrightarrow \mathcal{R}(\mathcal{C}) \longrightarrow Z(\mathcal{C}), \quad (\gamma, S, h) \longmapsto (\gamma, h) \longmapsto \gamma.$$

The first map forgets an exact-agreement support; the second forgets the explaining polynomial. Distinct explanation states may have the same slope, so neither map is silently treated as a bijection. A ray compiler means a proved estimate for the final slope image, obtained from these actual fibers or from a direct incidence bound.

Lemma 22.4 (Exact occupancy compiler). *With the notation above,*

$$(RC_{\text{occ}}) \quad |\mathcal{C}| = \sum_{\rho \in \mathcal{R}(\mathcal{C})} \nu_{\mathcal{C}}(\rho), \quad |\mathcal{R}(\mathcal{C})| = \sum_{w \in \mathcal{C}} \frac{1}{\nu_{\mathcal{C}}(\pi_{\text{exp}}(w))}, \quad |Z(\mathcal{C})| \leq |\mathcal{R}(\mathcal{C})|.$$

Here $\pi_{\text{exp}}(\gamma, S, h) = (\gamma, h)$. In particular, if every retained explanation state has occupancy at least $H \geq 1$, then

$$(RC1) \quad |Z(\mathcal{C})| \leq \left\lfloor \frac{|\mathcal{C}|}{H} \right\rfloor.$$

Also, because one noncommon exact- a support carries at most one slope and the explaining polynomial on $a \geq k + 1$ points is unique,

$$(RC2) \quad |Z(\mathcal{C})| \leq |\{S : \exists \gamma, h \ (\gamma, S, h) \in \mathcal{C}\}|.$$

All these statements are exact finite-row identities or inequalities.

Proof. Partition $\mathcal{C}$ by the fibers of π_{exp} . This gives the first identity; summing the reciprocal of the fiber size over each fiber gives the second. Projection $(\gamma, h) \mapsto \gamma$ can only decrease cardinality. If every first fiber has size at least H , (RC1) follows. Finally, if the same noncommon support carried two distinct slopes, slope elimination would simultaneously explain the received pair on that support. For fixed support and slope, uniqueness of interpolation gives at most one h . This proves (RC2). $\square$

Lemma 22.5 (Max-occupancy add-back). *Let $\mathcal{C} \subseteq \mathcal{W}_a(r)$ be partitioned into profile cells $\mathcal{C} = \coprod_{\lambda \in \Lambda} \mathcal{C}_\lambda$, and put $P = |\Lambda|$. For every explanation state $\rho = (\gamma, h) \in \mathcal{R}(\mathcal{C})$, let*

$$M(\rho) = |\{S : (\gamma, S, h) \in \mathcal{C}\}|.$$

Assign ρ to a profile $\lambda(\rho)$ on which $\nu_{\mathcal{C}_\lambda}(\rho)$ is maximal, and let $\mathcal{R}_\lambda^\star$ be the assigned states. Then

$$(RC_{\max 1}) \quad \nu_{\mathcal{C}_\lambda}(\rho) \geq \lceil M(\rho)/P \rceil \quad (\rho \in \mathcal{R}_\lambda^\star).$$

Consequently, if $M(\rho) \geq H$ for every retained state, then

$$(RC_{\max 2}) \quad |Z(\mathcal{C})| \leq \sum_{\lambda \in \Lambda} |\mathcal{R}_\lambda^\star| \leq \sum_{\lambda \in \Lambda} \left\lfloor \frac{|\mathcal{C}_\lambda^\star|}{\lceil H/P \rceil} \right\rfloor,$$

where $\mathcal{C}_\lambda^\star$ consists of the witnesses in $\mathcal{C}_\lambda$ lying above assigned states. Call $\mathcal{C}$ explanation-state complete if, whenever $(\gamma, h) \in \mathcal{R}(\mathcal{C})$, every exact- a noncommon witness $(\gamma, S, h) \in \mathcal{W}_a(r)$ also belongs to $\mathcal{C}$. If $\mathcal{C}$ is explanation-state complete, every retained state has exact agreement $b \geq a$, and the received pair is not simultaneously explained on any a -subset of that agreement set, then one may take $H = \binom{b}{a}$.

Proof. The P profile occupancies of a fixed state sum to $M(\rho)$, so their maximum is at least $\lceil M(\rho)/P \rceil$. Apply Lemma 22.4 on each assigned cell and sum. Under the final hypotheses, explanation-state completeness retains every a -subset of the b -point agreement set, and every such subset is noncommon. Hence $M(\rho) = \binom{b}{a}$. $\square$

Definition 22.6 (Support interpolation functionals). *Let $T \subseteq D$, $|T| = m \geq k$, put $w = m - k$, and let $Q_T(X) = \prod_{x \in T} (X - x)$. For a word u , define*

$$(22.1) \quad \mathcal{A}_T(u) = \left(\sum_{x \in T} \frac{u(x)x^r}{Q'_T(x)} \right)_{0 \leq r < w} \in \mathbb{F}^w.$$

Imported slope interface. CAP25 v13.2 [Cho26, Sec. 19 and Apps. E–F] proves that a word u is explained by $\text{RS}_{\mathbb{F}}(D, k)$ on T exactly when $\mathcal{A}_T(u) = 0$. Hence

$$(SE) \quad \mathcal{A}_T(u_0) + \gamma \mathcal{A}_T(u_1) = 0,$$

and every noncommon support carries at most one finite slope. We use this only as an imported incidence input to the occupancy compiler above.

Imported support-pair identity. For every residual prefix profile λ , the foundational paper [Cho26, Sec. 19 and Apps. E–F] identifies the ordered same-prefix support-pair count as

$$(SP_0) \quad \sum_s Q_\lambda(s)^2 \leq (\max_s Q_\lambda(s)) \sum_s Q_\lambda(s).$$

Q therefore controls a support-pair statistic. It does not justify dividing that statistic by $|\Omega_\lambda^0|$: a ray can have a single representative at exact agreement.

Proposition 22.7 (Q and SP do not determine rays). *Let $\Phi : \Omega \rightarrow \Sigma$ be any support-profile map with $\Omega \neq \emptyset$, put $Q(s) = |\Phi^{-1}(s)|$, and $SP = \sum_s Q(s)^2$. Regard Ω as an abstract set of candidate noncommon supports, and suppose the only retained slope-incidence axiom is that each support carries at most one slope. For every nonempty finite slope set Γ and every $1 \leq M \leq \min\{|\Omega|, |\Gamma|\}$, within the class of abstract incidence structures obeying these retained data there is a partial support-to-slope map having exactly M distinct values. There is also a one-value map. Hence no ray bound can be deduced from Q , SP , and that axiom alone; an RS-specific bound must use additional algebraic incidence information.*

Proof. Choose M distinct supports and M distinct slopes and map them bijectively, leaving all other supports unlabeled. This changes neither Φ , its fibers, nor SP , and it respects one-slope-per-support. Mapping the same chosen supports to one slope gives identical Q and SP data with a one-ray image. The missing datum is the image incidence, not another moment of the prefix fibers. $\square$

Proposition 22.8 (Exact pair-to-ray multiplicity requirement). *Let $\mathcal{P} = \{(A, B) \in \Omega^2 : \Phi(A) = \Phi(B)\}$, so $|\mathcal{P}| = SP$. Let Z be a bad-slope set and $I \subseteq Z \times \mathcal{P}$ an incidence relation, and let $H, J \geq 1$. If every $\gamma \in Z$ is incident with at least H pairs and every pair is incident with at most J slopes, then*

$$(22.2) \quad |Z| \leq \frac{J SP}{H} \leq \frac{J(\max_s Q(s))|\Omega|}{H}.$$

Consequently this pair-count argument yields a ray estimate at the same profile scale provided $J|\Omega|/H = e^{o(n)}$; alternatively one may prove a direct transverse-secant estimate. Merely proving that every ray has one representing pair is insufficient.

Proof. Double counting gives $H|Z| \leq |I| \leq J|\mathcal{P}| = J SP$. The second inequality follows from $SP \leq (\max_s Q(s)) \sum_s Q(s) = (\max_s Q(s))|\Omega|$. $\square$

Corollary 22.9 (Exact finite pair-to-ray constant). *For a residual profile λ , put $m_\lambda = \sum_s Q_\lambda(s)$, $L_\lambda = |\Phi_\lambda(\Omega_\lambda^0)|$, and $\bar{N}_\lambda = |\Omega_\lambda^0|/L_\lambda$. For an integer $q \geq 2$, define*

$$\Gamma_{q,\lambda} = L_\lambda^{-1} \sum_s \left(\frac{Q_\lambda(s)}{\bar{N}_\lambda} \right)^q.$$

If the pair-to-ray incidence has lower ray degree H_λ and upper pair degree J_λ , then

$$(RC_{\text{pair}}) \quad |Z_\lambda^o| \leq \left\lfloor \frac{J_\lambda m_\lambda}{H_\lambda} \bar{N}_\lambda (L_\lambda \Gamma_{q,\lambda})^{1/q} \right\rfloor.$$

Thus the finite compiler uses the actual residual mass and literal incidence degrees; it contains no unnamed polynomial or asymptotic loss.

Proof. The moment sandwich gives $\max_s Q_\lambda(s) \leq \bar{N}_\lambda (L_\lambda \Gamma_{q,\lambda})^{1/q}$. Hence $\sum_s Q_\lambda(s)^2 \leq m_\lambda \max_s Q_\lambda(s)$. Apply [Proposition 22.8](#) and take the integer floor. $\square$

Proposition 22.10 (Curve-degree ray compiler). *Let V be a finite-dimensional space of locator polynomials and, for $x \in D$, let $H_x = \{[L] \in \mathbb{P}(V) : L(x) = 0\}$ be its evaluation hyperplane. Suppose a residual slope set Z admits an injective assignment into a disjoint union of projective integral curves*

$$Z \hookrightarrow \prod_{i=1}^s X_i(\mathbb{F}),$$

where “integral curve” means an irreducible reduced one-dimensional projective variety, together with locator morphisms $\psi_i : X_i \rightarrow \mathbb{P}(V)$. Put

$$\delta_i = \deg \psi_i^* \mathcal{O}_{\mathbb{P}(V)}(1),$$

the degree on X_i of the pullback of a projective hyperplane. For every $\gamma \in Z$, if its assigned point is $y \in X_i(\mathbb{F})$, require $\psi_i(y) = [L_{\gamma,y}]$, where $L_{\gamma,y}$ is a locator from an actual residual witness to γ . For each i , put

$$G_i = \{x \in D : \psi_i(X_i) \subseteq H_x\}, \quad D_i^{\text{mov}} = D \setminus G_i.$$

Assume the locator attached to every chosen representative on X_i has at least $h_i \geq 1$ roots in D_i^{mov} . Then

$$(RC_{\text{curve}}) \quad |Z| \leq \left\lfloor \sum_{i=1}^s \frac{|D_i^{\text{mov}}| \delta_i}{h_i} \right\rfloor \leq \left\lfloor n \sum_{i=1}^s \frac{\delta_i}{h_i} \right\rfloor.$$

In particular, a cover of subexponential total locator degree gives a subexponential ray budget when the h_i have no exponential loss. The standard degree-one pencil $[s : t] \mapsto [sA + tB]$ is the special case $X = \mathbb{P}^1$ and $\delta = 1$.

Proof. Choose one representative point for each slope. For fixed i and $x \in D_i^{\text{mov}}$, the pullback of H_x is a nonzero effective divisor of degree δ_i , and hence contains at most δ_i distinct rational points. Each chosen point on X_i supplies at least h_i moving-root incidences. Thus the number assigned to X_i is at most $|D_i^{\text{mov}}| \delta_i / h_i$. Sum over i and take the integer floor. $\square$

Corollary 22.11 (Automatic exact-support locator-curve compiler). *Fix a received line and exact agreement $a \geq k + 1$. For each slope in a residual set Z , choose one noncommon witness support $S_\gamma \in \binom{D}{a}$ and its locator Q_{S_γ} . Suppose the locator points are covered by $\psi_i(X_i(\mathbb{F}))$, where $\psi_i : X_i \rightarrow \mathbb{P}(V)$ are nonconstant morphisms from projective integral curves. Put*

$$\delta_i = \deg \psi_i^* \mathcal{O}(1), \quad G_i = \{x \in D : \psi_i(X_i) \subseteq H_x\}, \quad g_i = |G_i|.$$

Then

$$(RC_{\text{loc}}) \quad |Z| \leq \#\{i : g_i = a\} + \left\lfloor \sum_{i: g_i < a} \frac{(n - g_i) \delta_i}{a - g_i} \right\rfloor.$$

No separate slope-to-curve injectivity hypothesis is required. A component with $g_i = a$ contains at most one selected degree- a locator, because its common roots determine the monic support locator uniquely; this accounts for the first term. Components with $g_i > a$ contain no selected degree- a locator and are discarded.

Proof. Two distinct slopes cannot use the same chosen noncommon support. Since a monic squarefree locator determines its root set, their projective locator points are distinct as well. Assign each locator to the first covering component and choose one rational preimage; this is the injective assignment required by Proposition 22.10. On component i , every selected degree- a locator has the g_i common roots and at least $a - g_i$ roots outside them. If $g_i = a$, there is only one such monic locator. On the other components apply (RC_{curve}) with $h_i = a - g_i$. $\square$

Corollary 22.12 (Automatic exact-support one-pencil compiler). *In the setting above, suppose all selected locators lie in the projective pencil spanned by A, B , and let $g < a$ be the number of common roots of A, B in D . Then*

$$(RC_{\text{pen}}) \quad |Z| \leq \lfloor (n - g) / (a - g) \rfloor.$$

The same argument applies to residual locators of degree e whose roots lie in a fixed active domain $T \subseteq D$, provided the fixed profile data together with the residual locator determine the exact support. If $N = |T|$ and g_T is the number of common roots of A, B in T , then

$$|Z| \leq \lfloor (N - g_T) / (e - g_T) \rfloor.$$

Proof. The exact-support locator assignment is injective by the preceding proof. Apply Lemma 23.2, or the curve corollary with $X = \mathbb{P}^1$ and $\delta = 1$. $\square$

By Theorem 4.2, the genuinely missing invariant can be taken to be $\Theta_t(y_0, y_1)$. Equivalently one may construct a higher-dimensional split-pencil *eliminant*, meaning a polynomial in the slope parameter obtained after eliminating locator and support variables, whose degree bounds that transverse secant incidence. By Lemma 4.3, only dependent support unions can support more than $t + 1$ rays. Those unions are no longer untreated: Theorems 4.5 and 6.75 bound every fixed union, and Proposition 23.16 gives the exact active low-nullity ranges. What remains is an uncontrolled family of many such unions, a high-nullity chart, or a direction whose rank defect

exceeds the quantitative threshold of [Corollary 6.77](#); a higher-dimensional balanced core must still be routed into one of those certified cells or paid by its own eliminant.

Here an *MDS-circuit stratum* is a chart on which the relevant union of generalized Vandermonde parity-check columns has a minimal linear dependence. Equivalently, deleting any one column restores independence. The name identifies the only dependence pattern left by the MDS property; it does not by itself place all witnesses in one fixed circuit or provide a count for an unrestricted union of such strata.

The next hypothesis packages exactly this missing conversion. Its number H_λ is a lower bound on how many certified support pairs represent one bad ray, while J_λ is an upper bound on how many bad rays one pair can represent. The ratio in (RC) is what makes the resulting double count stay at the profile scale.

Hypothesis 22.13 (Ray compiler). *For every received line and primitive residual profile λ , let Z_λ° be its actual first-match set of unpaid bad slopes and put*

$$\mathcal{P}_\lambda = \{(A, B) \in (\Omega_\lambda^\circ)^2 : \Phi_\lambda(A) = \Phi_\lambda(B)\}.$$

Put $m_\lambda = |\Omega_\lambda^\circ| = \sum_s Q_\lambda(s)$. The row satisfies (RC) on this profile if either it has the direct bound $|Z_\lambda^\circ| \leq e^{o(n)}(1 + \bar{N}_\lambda)$, or there are an incidence $I_\lambda \subseteq Z_\lambda^\circ \times \mathcal{P}_\lambda$ and numbers $H_\lambda, J_\lambda \geq 1$ such that

$$(RC) \quad \deg_{I_\lambda}(\gamma) \geq H_\lambda, \quad \deg_{I_\lambda}(A, B) \leq J_\lambda, \quad \frac{J_\lambda m_\lambda}{H_\lambda} = e^{o(n)}.$$

All bounds are uniform in the received line and realized profile. The incidence must arise from the RS witness geometry; its existence is not a consequence of Q or SP. In particular, one representing pair per ray does not supply the required lower degree.

Proposition 22.14 (Q plus RC implies rays). *If primitive Q is discharged and (RC) holds on every residual profile, then all primitive ray ledgers are discharged at total cost $e^{o(n)}\mathfrak{E}_n$.*

Proof. The boundary prefix fibers are bounded by Q. For every remaining unpaid slope choose one noncommon support, as permitted by the imported slope interface (SE); (RC) bounds their distinct slope set by $e^{o(n)}\bar{N}_\lambda$. Sum over the subexponential profile set and use (11.1). $\square$

The proof of [Theorem 11.2](#) is now complete: the Sidon payment and [Theorem 21.2](#) pay primitive Q, while [Proposition 22.14](#) and (RC) convert the remaining support bounds into distinct-ray bounds.

Part 4. Finite deployment: Q, one-parameter BC, and the exact ledger

23. FINITE MAX-FIBER CALIBRATION

We import the four unsafe endpoints and challenge budgets from CAP25 v13.2 [[Cho26](#), Cor. 14.3 and Def. 21.1]. Only the adjacent-row quantities consumed by the new upper compiler are printed here.

TABLE 2. Imported active frontier rows and the adjacent safe-side calibration.

row	a_0	$a_+ = a_0 + 1$	$w = a_+ - K$	B^*
KoalaBear MCA	1116047	1116048	67471	274980728111395087
KoalaBear list	1116046	1116047	67471	274980728111395087
Mersenne-31 MCA	1116023	1116024	67447	16777215
Mersenne-31 list	1116022	1116023	67447	16777215

Here $n = 2^{21}$, $K = 2^{20} + 1$ on the MCA route, and $K = 2^{20}$ on the list route. Put

$$\bar{F} = \binom{n}{a_+} |\mathbb{B}|^{-w}.$$

The exact ceilings of this average are, in table order,

$$(23.1) \quad 57198030366, \quad 65065153468, \quad 1752700, \quad 1993678.$$

Hence the full-budget Q overhead ceilings are respectively

$$(23.2) \quad 4807520.9295, \quad 4226236.5253, \quad 9.5722, \quad 8.4152.$$

Any payment made to another cell lowers these values.

The exact moment interface imported from the foundation is

$$(23.3) \quad \max_z R(z) \leq (|\mathbb{B}|^w \Gamma_r)^{1/r}, \quad R(z) = \frac{|\text{Fib}_w(z)|}{\binom{n}{a_+} |\mathbb{B}|^{-w}}.$$

The exact second-moment decomposition is written

$$(23.4) \quad \sum_z N_w(z)^2 = \binom{n}{a_+} + \sum_e P_e.$$

Theorem 23.1 (Finite moment criterion for Q). *If $\Gamma_r \leq e^A$, then the normalized maximum fiber satisfies*

$$R_Q^{\max} \leq \exp\left(\frac{A + w \log |\mathbb{B}|}{r}\right).$$

Thus a finite row with remaining bit margin Δ can be proved by this route only if

$$(23.5) \quad \log_2 \Gamma_r + w \log_2 |\mathbb{B}| \leq r\Delta.$$

Proof. Take logarithms in (23.3). □

Lemma 23.2 (Pencil payment). *Let $L_{[s:t]} = sA + tB$ be a projective polynomial pencil and let T be a finite set. If every counted member has at least h roots in T that are not common roots of A and B , then at most $\lfloor |T|/h \rfloor$ projective parameters are counted.*

Proof. A point $x \in T$ that is not a common root solves the homogeneous linear equation $sA(x) + tB(x) = 0$ for exactly one projective parameter. Count moving-root incidences. □

Definition 23.3 (projective locator pencil and its common D -part). *Let $D \subseteq \mathbb{F}$ be finite, put $\Lambda_D = \prod_{x \in D} (X - x)$, and let $A, B \in \mathbb{F}[X]$ be not both zero. The projective locator pencil generated by A, B is*

$$L_{[s:t]}(X) = sA(X) + tB(X), \quad [s : t] \in \mathbb{P}^1(\mathbb{F}).$$

Its fixed D -part is

$$G_D(A, B) = \gcd(A, B, \Lambda_D), \quad g_D(A, B) = \deg G_D(A, B).$$

A root of $G_D(A, B)$ is a common root of every member of the pencil; these are precisely the roots assigned to the common-GCD branch in the first-match ledger.

Theorem 23.4 (moving-root bound for one-parameter split pencils). *Let $D \subseteq \mathbb{F}$ have size n , and let $L_{[s:t]} = sA + tB$ be a projective locator pencil as in Definition 23.3. Put $G = G_D(A, B)$ and $g = \deg G$. Let $\mathcal{Z} \subseteq \mathbb{P}^1(\mathbb{F})$, and suppose that for every $\lambda = [s : t] \in \mathcal{Z}$, the polynomial L_λ/G has at least $h \geq 1$ distinct roots in $D \setminus Z(G)$. Then*

$$|\mathcal{Z}| \leq \left\lfloor \frac{n - g}{h} \right\rfloor.$$

In particular, if every counted member has degree j , has fixed D -part exactly G , and is D -split, then

$$|\mathcal{Z}| \leq \left\lfloor \frac{n - g}{j - g} \right\rfloor, \quad j > g.$$

The case $j = g$ is not a moving split-pencil cell: all D -roots are fixed and the member belongs to the common-GCD branch.

Proof. Count incidences

$$I = \{(\lambda, x) : \lambda \in \mathcal{Z}, x \in D \setminus Z(G), L_\lambda(x) = 0\}.$$

For a fixed moving root $x \in D \setminus Z(G)$, the pair $(A(x), B(x))$ is not $(0, 0)$. Hence the homogeneous linear equation

$$sA(x) + tB(x) = 0$$

has exactly one solution $[s : t] \in \mathbb{P}^1(\mathbb{F})$. Thus each moving domain point contributes to at most one incidence, and $|I| \leq n - g$. On the other hand, every $\lambda \in \mathcal{Z}$ contributes at least h incidences by hypothesis, so $|I| \geq h|\mathcal{Z}|$. This proves the first bound.

If a counted member is D -split of degree j and has fixed D -part exactly G , then exactly $j - g$ of its D -roots are moving roots in $D \setminus Z(G)$. Applying the first bound with $h = j - g$ gives the displayed inequality. If $j = g$, no moving root remains; the split condition has already been paid by the common-GCD branch rather than by a primitive moving pencil. $\square$

Corollary 23.5 (BC is settled on primitive one-parameter locator pencils). *Consider an MCA split-pencil chart at agreement m , and write $\omega = n - m$. Suppose that after the first-match tangent, common-support, quotient, extension, degree-drop, and common-GCD branches have been removed, the residual candidate error locators in this chart are represented by a projective locator pencil $L_\lambda = sA + tB$, with the slope parameter injecting into λ , and that every residual bad slope in the chart has at least one degree- ω D -split locator in the pencil. Then the number of residual bad finite slopes in this chart is at most*

$$\left\lfloor \frac{n - g}{\omega - g} \right\rfloor, \quad g = g_D(A, B) < \omega.$$

In the primitive case $g = 0$, this is $\lfloor n/(n - m) \rfloor$. For the two active MCA adjacent rows in Table 2,

row	$m = a_+$	$\omega = n - m$	$\lfloor n/\omega \rfloor$
KoalaBear MCA	1116048	981104	2
Mersenne-31 MCA	1116024	981128	2

so every primitive one-parameter BC pencil contributes at most two finite bad slopes.

Proof. The first-match removals ensure that the chart is counted only through moving roots of the displayed pencil. A residual bad slope supplies at least one degree- ω split locator, and the slope-to-pencil parameter is injective by hypothesis; hence the slope count is bounded by the number of split parameters. Apply Theorem 23.4 with $j = \omega$. The displayed row values use $n = 2^{21}$, $a_+ = 1116048$ for KoalaBear MCA, and $a_+ = 1116024$ for Mersenne-31 MCA. $\square$

Definition 23.6 (saturated codeword rays and support census). *Let $C = \text{RS}[\mathbb{F}, D, K]$, let $U : D \rightarrow \mathbb{F}$, and put*

$$S_c(U) = \{x \in D : U(x) = c(x)\}, \quad s_c(U) = |S_c(U)|$$

for $c \in C$. The saturated ray set at agreement m is

$$\text{Ray}(U; m) = \{c \in C : s_c(U) \geq m\}.$$

Let $\omega = n - m$, and define the support-locator census

$$\begin{aligned} \text{Cen}(U; m) = \#\{(W, N) \in \mathbb{F}[X]^2 : W \text{ monic, } \deg W = \omega, W \mid \ell_D, \\ WU = N \text{ on } D, W \mid N, \deg(N/W) < K\}. \end{aligned}$$

Equivalently, $W = \ell_{D \setminus T}$ for an agreement support $T \subseteq D$ of size m .

Theorem 23.7 (exact saturation identity). *With $C = \text{RS}[\mathbb{F}, D, K]$ and $\text{Cen}(U; m)$ as in Definition 23.6, for every word $U : D \rightarrow \mathbb{F}$,*

$$\text{Cen}(U; m) = \sum_{c \in C} \binom{s_c(U)}{m}.$$

Moreover, the support locators lying above a fixed codeword ray c are exactly

$$(W, N) = H(G_c, G_c c), \quad G_c = \ell_{D \setminus S_c(U)}, \quad H \mid \ell_{S_c(U)}, \quad \deg H = s_c(U) - m.$$

Consequently the fiber size of the forgetful map from support locators to saturated codeword rays is $\binom{s_c(U)}{m}$ above c .

Proof. Let (W, N) be counted by $\text{Cen}(U; m)$. Since W is monic, divides ℓ_D , and has degree $n - m$, there is a unique m -set $T \subseteq D$ such that $W = \ell_{D \setminus T}$. Since $W \mid N$ and $\deg(N/W) < K$, write $N = Wc$ for a unique codeword $c \in C$. On every $x \in T$, the value $W(x)$ is nonzero; the identity $WU = N = Wc$ on D therefore gives $U(x) = c(x)$. Thus $T \subseteq S_c(U)$.

Conversely, if $c \in C$ and $T \subseteq S_c(U)$ has $|T| = m$, then $W = \ell_{D \setminus T}$ and $N = Wc$ satisfy all defining conditions of $\text{Cen}(U; m)$. For fixed c , the number of such supports is $\binom{s_c(U)}{m}$, and summing over c gives the first identity.

For the fiber statement, write $G_c = \ell_{D \setminus S_c(U)}$. If $T \subseteq S_c(U)$, then

$$\ell_{D \setminus T} = \ell_{D \setminus S_c(U)} \ell_{S_c(U) \setminus T} = G_c H,$$

where $H = \ell_{S_c(U) \setminus T}$, hence $H \mid \ell_{S_c(U)}$ and $\deg H = s_c(U) - m$. The converse is the same construction reversed, because every monic divisor H of $\ell_{S_c(U)}$ of degree $s_c(U) - m$ is the locator of $S_c(U) \setminus T$ for a unique m -set $T \subseteq S_c(U)$. $\square$

Corollary 23.8 (raw support BC is not the MCA object). *In the setting of Theorem 23.7, if U agrees with a single codeword c on $n - d$ positions and with no other codeword on m or more positions, then*

$$|\text{Ray}(U; m)| = 1, \quad \text{Cen}(U; m) = \binom{n - d}{m}.$$

Thus an exponentially large raw support-locator family may represent one MCA ray. A split-pencil proof for MCA must therefore count saturated primitive line-rays after quotient, prefix, tangent, and common-support branches have been paid; a raw support census is a list/moment object, not the final MCA numerator.

Proof. The ray statement is the hypothesis. The support-census identity is Theorem 23.7 with one nonzero summand $\binom{s_c(U)}{m} = \binom{n-d}{m}$. The final sentence is just the same equality interpreted in MCA terms: MCA counts a bad slope once, whereas $\text{Cen}(U; m)$ counts all size- m supports below the same saturated agreement ray. $\square$

Definition 23.9 (line rays). *In the setting of Definition 23.6, for an affine line $U_Z = u + Zv$ and a set of finite slopes $E \subseteq \mathbb{F}$, define*

$$\text{LineRay}_E(u, v; m) = \{(z, c) \in E \times C : |\{x \in D : u(x) + zv(x) = c(x)\}| \geq m\}.$$

Write $s_{z,c} = |\{x \in D : u(x) + zv(x) = c(x)\}|$.

Proposition 23.10 (line-ray saturation identity). *With notation as in Definition 23.9, for every $E \subseteq \mathbb{F}$,*

$$\sum_{z \in E} \text{Cen}(U_z; m) = \sum_{(z,c) \in \text{LineRay}_E(u,v;m)} \binom{s_{z,c}}{m}.$$

Consequently

$$|\{z \in E : \exists c (z, c) \in \text{LineRay}_E(u, v; m)\}| \leq |\text{LineRay}_E(u, v; m)| \leq \sum_{z \in E} \text{Cen}(U_z; m).$$

The two inequalities are the exact deduplication losses between raw support locators, saturated codeword rays, and slope numerators.

Proof. Apply Theorem 23.7 separately to the word $U_z = u + zv$ for each $z \in E$, and then sum over z . The first inequality forgets the codeword coordinate of a line ray, and the second follows because each line ray has $s_{z,c} \geq m$, hence contributes at least one support to the right-hand side of the identity. Corollary 23.8 shows that the second inequality can be exponentially loose for a single slope and a single codeword ray, so an MCA upper ledger must either count slopes directly or control this saturation factor. $\square$

Theorem 23.11 (Q discharges the shift-pair ledger). *Let $N_w(z) = |\text{Fib}_w(z)|$, let*

$$\overline{F} = \binom{n}{m} |\mathbb{B}|^{-w},$$

and let P_e denote the full off-diagonal shift-pair stratum in the imported decomposition (23.4), namely

$$P_e = \sum_{\substack{R \subseteq D \\ |R|=m-e}} \text{sp}_w(e; D \setminus R).$$

If Q holds in max-fiber form with constant κ ,

$$\max_z N_w(z) \leq \kappa \overline{F},$$

then necessarily $\kappa \overline{F} \geq 1$ whenever $\binom{D}{m}$ is nonempty, and

$$\sum_{e=w+1}^{\min(m, n-m)} P_e \leq \left(\kappa - \frac{1}{\overline{F}} \right) \binom{n}{m} \overline{F}.$$

Every quotient-removed or primitive shift-pair subledger satisfies the same upper bound. Hence SP is not an independent residual input in any closing theorem that already proves Q as a max-fiber theorem.

Proof. The imported exact second-moment identity (23.4) says

$$\sum_z N_w(z)^2 = \binom{n}{m} + \sum_{e=w+1}^{\min(m, n-m)} P_e.$$

On the other hand,

$$\sum_z N_w(z)^2 \leq (\max_z N_w(z)) \sum_z N_w(z) \leq \kappa \overline{F} \binom{n}{m},$$

because $\sum_z N_w(z) = \binom{n}{m}$. Since at least one fiber is nonempty when $\binom{D}{m} \neq \emptyset$, the hypothesis also implies $\kappa \overline{F} \geq 1$, so the displayed right-hand side is nonnegative in the nonvacuous case. Subtracting the diagonal term $\binom{n}{m}$ gives the displayed inequality. Any primitive or quotient-removed ledger is a sub-sum of the full off-diagonal ledger, so it is bounded a fortiori. $\square$

Theorem 23.12 (SP: unconditional primitive shift-pair bound). *Use the notation of the imported exact second-moment ledger. For $w+1 \leq e \leq \min(m, n-m)$, put*

$$T_e(n, m) = \binom{n}{m-e} \binom{n-m+e}{e} \binom{n-m}{e}.$$

Then the total off-diagonal contribution at distance e satisfies

$$P_e^{\text{quot}} + P_e^{\text{prim}} \leq T_e(n, m).$$

If $D \subseteq \mathbb{B}$, then the top stratum $e = w+1$ also satisfies the sharper constant-shift bound

$$P_{w+1}^{\text{quot}} + P_{w+1}^{\text{prim}} \leq (|\mathbb{B}| - 1) \binom{n}{m-w-1} \binom{n-m+w+1}{w+1}.$$

Consequently

$$\Gamma_2 \leq \frac{1}{\overline{F}} + \frac{\sum_{e=w+1}^{\min(m, n-m)} T_e(n, m)}{\binom{n}{m} \overline{F}},$$

with the option of replacing the $e = w+1$ term by the sharper displayed top-stratum bound.

Proof. Fix e . In the imported exact second-moment decomposition (23.4), an off-diagonal pair is encoded by a common part $R \subseteq D$ of size $m - e$, then by two disjoint degree- e root sets inside $D \setminus R$. There are $\binom{n}{m-e}$ choices for R . Once R is fixed, the ambient set $D \setminus R$ has size $n - m + e$; the first root set has at most $\binom{n-m+e}{e}$ choices and the second, disjoint from it, has at most $\binom{n-m}{e}$ choices. This ignores the depth equation and the quotient/primitive first-match assignment, so it is an upper bound for $P_e^{\text{quot}} + P_e^{\text{prim}}$.

When $e = w + 1$, the imported decomposition (23.4) shows that every top-stratum shift pair has the form $A - B = c \in \mathbb{F}^\times$. If both polynomials split over $D \subseteq \mathbb{B}$, then both have coefficients in $\mathbb{B}$, hence $c \in \mathbb{B}^\times$. For each fixed R , choose A by choosing its e -element root set in $D \setminus R$, and choose $c \in \mathbb{B}^\times$; the polynomial $B = A - c$ is then determined. Requiring B to split over $D \setminus R$, to be disjoint from A , and then assigning the pair to either the quotient or primitive ledger can only reduce the count. This gives the top-stratum bound. The displayed Γ_2 inequality is the imported exact second-moment ledger with the total quotient-plus-primitive off-diagonal contribution in every stratum replaced by the just-proved upper bound. $\square$

Proposition 23.13 (size of the proper Q bound at the active adjacent rows). *At the four active adjacent agreements, the imported Q packing estimate [Cho26, Secs. 19–21] gives the following one-term upper bounds for the remaining normalized prefix overhead:*

row	w	$\log_2 R_Q^{\text{pack}}$	allowed by the bound	spare margin
KoalaBear MCA	67471	1660789.024		22.1969
KoalaBear list	67471	1660789.017		22.0109
Mersenne-31 MCA	67447	1660926.120		3.2589
Mersenne-31 list	67447	1660926.114		3.0730.

Thus the unconditional packing proof is a real theorem but is far too weak to supply the adjacent safe ledger.

Proof. For each row, $n = 2097152$, $K = k + 1 = 1048577$ on the MCA route, $K = k = 1048576$ on the list route, and $w = a_+ - K$. The base-field orders are $2^{31} - 2^{24} + 1$ for KoalaBear and $2^{31} - 1$ for Mersenne-31. The imported Q packing inequality [Cho26, Secs. 19–21] gives

$$\log_2 R_Q^{\text{pack}} \leq w \log_2 |\mathbb{B}| - \log_2 \binom{a_+}{\lfloor w/2 \rfloor} - \log_2 \binom{n - a_+}{\lfloor w/2 \rfloor}.$$

Substituting the exact field orders and

$$a_+ \in \{1116048, 1116047, 1116024, 1116023\}$$

gives the displayed values. For these four rows the Johnson-ball summand

$$J_i = \binom{a_+}{i} \binom{n - a_+}{i}$$

is increasing on $0 \leq i \leq t$. Indeed

$$\frac{J_i}{J_{i-1}} = \frac{(a_+ - i + 1)(n - a_+ - i + 1)}{i^2},$$

and at $i = t$ this ratio is already > 900 in all four rows, while the ratio is larger for smaller i . Hence $V_t(n, a_+) \leq (t + 1)J_t$, so replacing the one-term denominator by the full ball can lower the displayed logarithmic upper bound by less than $\log_2(t + 1) < 16$ bits. This is negligible compared with the gap between the displayed million-bit overhead and the row margins. $\square$

Definition 23.14 (row-sharp Q atom bound). *Fix a deployed adjacent row, let a_+ be its adjacent agreement, let K be the list dimension used on that route, and put $w = a_+ - K$. For a first-match residual prefix-boundary family $\mathcal{P}_Q(z) \subseteq \text{Fib}_w(z)$, define*

$$R_Q^{\text{max}} = |\mathbb{B}|^w \max_z \frac{|\mathcal{P}_Q(z)|}{\binom{n}{a_+}}.$$

The row-sharp Q atom bound with bit margin Δ_Q is

$$R_Q^{\max} \leq 2^{\Delta_Q}.$$

Equivalently, if the entire row budget is hypothetically assigned to Q , the sufficient integer form is

$$\max_z |\mathcal{P}_Q(z)| \leq B^*.$$

Any non- Q cell paid in the same first-match ledger only decreases the available Δ_Q .

Proposition 23.15 (exact finite Q atom target at the four adjacent rows). *At the active adjacent agreements, the full-budget Q atom target is as follows:*

row	w	$\left\lceil \binom{n}{a_+} \mathbb{B} ^{-w} \right\rceil$	B^*	$B^* / \left\lceil \binom{n}{a_+} \mathbb{B} ^{-w} \right\rceil$
KoalaBear MCA	67471	57198030366	274980728111395087	4807520.9295
KoalaBear list	67471	65065153468	274980728111395087	4226236.5253
Mersenne-31 MCA	67447	1752700	16777215	9.5722
Mersenne-31 list	67447	1993678	16777215	8.4152

In bits, these are exactly the spare margins already printed in the row table:

$$22.1969, \quad 22.0109, \quad 3.2589, \quad 3.0730.$$

Thus the binding Q problem is the Mersenne-31 list row: after all first-match payments, the remaining primitive prefix-boundary atom must be at most about 8.42 times its average size, and less if any other cell consumes budget.

Proof. For each row, substitute $n = 2^{21}$, $k = 2^{20}$, the displayed adjacent agreement a_+ , and $K = k + 1$ on the MCA route and $K = k$ on the list route. The average prefix fiber is $\binom{n}{a_+} |\mathbb{B}|^{-w}$, and the integer lower-route replay uses its ceiling. The budgets are

$$B_{KB}^* = \left\lfloor \frac{(2^{31} - 2^{24} + 1)^6}{2^{128}} \right\rfloor, \quad B_{M31}^* = \left\lfloor \frac{(2^{31} - 1)^4}{2^{100}} \right\rfloor.$$

The displayed ratios are the exact integer quotients evaluated as decimals. Taking base-two logarithms of $B^* / (\binom{n}{a_+} |\mathbb{B}|^{-w})$ gives the four printed bit margins. $\square$

Proposition 23.16 (Exact fixed-union payments at the active rows). *Let $t = n - a_+$.*

(a) Fixed-word Q and list charts. *Let $K = k + 1$ on an MCA route and $K = k$ on a list route, put $R_Q = n - K$, and let $\mathcal{A}$ be a subfamily of one boundary- Q prefix fiber, or of one arbitrary-word list, whose error supports are all contained in a fixed union U of size $R_Q + \nu$. Then [Theorem 4.7](#) gives a full-budget payment for every integer nullity in the ranges below:*

	row	R_Q	$R_Q - t$	$0 \leq \nu \leq$	$\max_{\nu \text{ in range}} \mathcal{A} $
(FU-Q/list)	KoalaBear MCA	1048575	67471	4983	94008
	KoalaBear list	1048576	67471	4983	94632
	Mersenne-31 MCA	1048575	67447	4980	3555853
	Mersenne-31 list	1048576	67447	4980	4735771

Every number in the last column is below the corresponding challenge budget B^* of [Proposition 23.15](#). In particular, a global Q counterexample at the binding Mersenne-31 list row cannot be contained in one support union of nullity at most 4980.

(b) Fixed-line MCA charts. *For the actual MCA code, $R = n - k = 1048576$. The unconditional multiplicity compiler [Theorem 4.5](#) pays the following fixed unions within the full row budget:*

	row	$0 < \nu \leq$	$\max_{\nu \text{ in range}} Z $	B^*
(FU-MCA)	KoalaBear MCA	2	2847909263951	274980728111395087
	Mersenne-31 MCA	1	8150817	16777215

If the direction is rank-regular on its union, so that $d_U(y_1) = R$, the directional compiler [Corollary 6.76](#) pays the much larger ranges

	row	$0 \leq \nu \leq$	$\max_{\nu \text{ in range}} Z $
(FU-MCA-reg)	KoalaBear MCA	4982	94008
	Mersenne-31 MCA	4979	3555853

The direction-defect threshold of [Corollary 6.77](#) pays more than the rank-regular case. With the full row budget, selected exact thresholds are

(FU-KB-def)	KoalaBear ν	0	1000	4000	4500	4979	4980	4982
	$J_{B^*}(\nu)$	4341	3466	853	418	3	2	0

and

(FU-M31-def)	Mersenne-31 ν	0	1000	4000	4500	4979
	$J_{B^*}(\nu)$	4338	3463	849	415	0

Thus an unpaid fixed-union MCA chart is either above the displayed nullity range or has defect $j > J_{B^*}(\nu)$; it need not be declared exceptional merely because $j > 0$. At the Mersenne row, two disjoint one-circuit charts cost at most $2 \cdot 8150817 = 16301634 < 16777215$, leaving 475581 units of the full budget.

Proof. For a boundary-Q fiber, the imported exact witness-list dictionary [[Cho26](#), Secs. 18–21] associates to every support a distinct codeword of $\text{RS}[\mathbb{F}, D, K]$ agreeing with one received word on that support. Hence a subfamily whose error supports lie in U is a subset of the affine syndrome fiber $\mathcal{L}_U(y, t)$ in [Theorem 4.7](#). The same statement is tautological for an arbitrary-word list. Substituting $N = R_Q + \nu$ and $h = N - t$ into (4.9), and evaluating every integer in the displayed ranges, gives (FU-Q/list). At the next nullity the Johnson denominator is nonpositive in every row: $\nu = 4984$ for KoalaBear and $\nu = 4981$ for Mersenne-31.

For MCA, substitute $R = 1048576$ and the two deployed values of t into (4.5). The largest full-budget nullities are 2 and 1, giving (FU-MCA). Under rank regularity, substitute the same row data into (6.42); its denominator is positive through $\nu = 4982$ and 4979, respectively, and the displayed maxima occur at the final integer in each range. Finally substitute the row budgets into (6.44) to obtain Eqs. (FU-KB-def) and (FU-M31-def). All comparisons are exact integer comparisons and are reproduced by the release verifier. $\square$

Proposition 23.17 (Exact recursive affine-span payments at the active rows). *In the boundary-Q witness list, or in an arbitrary-word list for the same route, suppose the associated codewords lie in one affine subspace of $\text{RS}[\mathbb{F}, D, K]$ of dimension s . The recursive compiler [Theorem 4.8](#) gives the following full-budget payments:*

row	$0 \leq s \leq$	$\max_{s \text{ in range}} \mathcal{L}_A $	compiler value at the next s
KoalaBear MCA	14	47876303026096432	743896698428332665
KoalaBear list	14	47876942243704279	743907339955962922
Mersenne-31 MCA	6	14115447	219426634
Mersenne-31 list	6	14115528	219428099

The middle column is below the corresponding full challenge budget in every row. The final column only records that this compiler no longer fits at dimension 15 or 7; it is not a lower bound on an actual list. Consequently a Q counterexample surviving the full-budget payments must have affine codeword-span dimension at least 15 on a KoalaBear row and at least 7 on a Mersenne-31 row.

Proof. Apply [Theorem 4.8](#) with $m = a_+$, with $K = k + 1$ on an MCA boundary route and $K = k$ on a list route. Evaluate the single rational expression in (4.10) by exact integer products. The value is increasing over each displayed range, so the maximum occurs at its last dimension. Comparison with the exact budgets in [Proposition 23.15](#) gives the claim; the next-dimension values mark the limit of the recursive compiler. $\square$

Remark 23.18 (Retraction of the former active affine-span MCA payments). The former KoalaBear and Mersenne-31 MCA affine-span payments used the refuted ordered-basis denominator of [Proposition 6.9](#). Their falling- and rising-product evaluations were arithmetically correct evaluations of that expression, but they are not upper bounds for MCA slope families. They therefore imply no affine-rank lower wall for a surviving MCA chart. The ordinary affine-span list payments immediately above are unaffected. The next proposition gives a distinct conditional wall from the repaired support-transverse theorem, with its deleted stratum and local margin stated explicitly.

Proposition 23.19 (Conditional post-deletion affine-error router on KoalaBear). *Use the KoalaBear parameters*

$$(n, k, m, w, B_*) = (2097152, 1048576, 1116048, 67472, 274980728111395087).$$

On one received line, let Z_{bad} be the finite support-wise MCA-bad slopes and define the intrinsic near-rational stratum

$$N = \{\gamma \in Z_{\text{bad}} : \min_{c \in C} \text{wt}(r_\gamma - c) \leq w\}, \quad Z = Z_{\text{bad}} \setminus N.$$

Suppose $|N| \leq 2w = 134944$, as supplied under the printed KoalaBear hypotheses by the separately pinned two-anchor near-rational theorem. Its guards hold here because $w \geq 1$ and $3w = 202416 \leq 1048576 = n - k$. For every $\gamma \in Z$, choose one exact same-support-noncontained m -record, and let a be the affine rank of the resulting error words. If $a \leq 9$, then

$$|Z_{\text{bad}}| \leq 110390969172173096 + 134944 = 110390969172308040 < B_*,$$

with slack 164589758939087047.

For $|Z| \geq 2$, apply [Theorem 6.13](#) and recompute the margin θ of (6.14) in the translated explanation direction space. At error ranks $a = 10, 11, 12$, the line is also paid whenever respectively

$$\theta \geq 13, \quad \theta \geq 388, \quad \theta \geq 12050.$$

Consequently an over-budget line forces $a \geq 10$ for every selected post-deletion family. At ranks 10, 11, 12, every such gauge emits an actual selected support and a codeword direction agreeing with the original r_1 on all but at most respectively 12, 387, 12049 support coordinates.

Proof. The exact-record selection is legitimate. Given a bad witness (S, c) with $|S| \geq m$, suppose every m -subset of S were pair-contained. Adjacent m -subsets overlap in $m - 1 \geq K$ coordinates, so Reed–Solomon uniqueness makes their explaining pairs identical. Connectedness of the m -subset Johnson graph would then glue one pair explaining all of S , contradicting badness. Hence some exact m -subwitness is pair-noncontained.

If $|Z| \leq 1$, the displayed bound is immediate. Otherwise [Theorem 6.13](#) gives a translated explanation family of affine dimension $s = a - 1$. For $2 \leq a \leq 9$, apply [Theorem 6.12](#) with $\theta \geq 1$. The largest cap occurs at $s = 8$ and is 110390969172173096.

If $a = 1$, all translated explanations coincide. Same-support noncontainment forces each selected support to contain a coordinate where the translated line direction is nonzero; at such a coordinate the agreement equation determines at most one finite slope. Hence $|Z| \leq n$, which is smaller than the displayed cap. The case $a = 0$ has at most one slope, since the injective map in the proof of [Theorem 6.13](#) has nonzero slope projection.

For $a = 10, 11, 12$, substitute $s = 9, 10, 11$ in (6.15). Exact integer arithmetic shows that the least paying margins after the $2w$ add-back are 13, 388, 12050. At the preceding integers the totals are respectively

$$285894151677963688, \quad 275500176064828033, \quad 274992929018868606,$$

all strictly above B_* . Negating the paying condition gives the three same-record direction-exception terminals. Explicitly, if $b_0 \in C$ is the gauge codeword from [Theorem 6.13](#) and failure of the relevant margin produces b_1 in the translated explanation direction space, then $r'_1 = r_1 - b_0$ agrees with b_1 off the printed exception set. Thus the inverse gauge transports the terminal to the original record as $r_1 = b_0 + b_1$ on the identical support outside that same set. $\square$

Theorem 23.20 (Sub-square-root common-support interleaving collapse). *Let $C \leq \mathbb{F}^D$ be an $\mathbb{F}$ -linear code, put $q = |\mathbb{F}|$, and fix an agreement threshold A . Let L be the largest ordinary list size at agreement A , over all received words. For $\mu \geq 1$, let L_μ be the corresponding largest list of μ -tuples in C^μ agreeing with one received μ -tuple on one common support of size at least A . If $1 \leq L < q$, then*

$$L \leq L_\mu \leq \left\lfloor \frac{L(q-1)}{q-L} \right\rfloor.$$

In particular, if $L^2 < q$, then $L_\mu = L$ for every $\mu \geq 1$.

Proof. The case $\mu = 1$ is the definition of L , so assume $\mu \geq 2$. The diagonal embedding gives $L_\mu \geq L$. For the reverse bound fix a common-support list $X \subseteq C^\mu$, $N = |X|$, and for $\alpha = (\alpha_2, \dots, \alpha_\mu) \in \mathbb{F}^{\mu-1}$ project by

$$\Phi_\alpha(c_1, \dots, c_\mu) = c_1 + \sum_{i=2}^{\mu} \alpha_i c_i.$$

The projected set is an ordinary list at the identical agreement threshold, so it has at most L images. Cauchy–Schwarz gives at least $(N^2/L - N)/2$ unordered colliding pairs for every α . A fixed distinct pair collides for at most $q^{\mu-2}$ choices of α : if some difference after the first component is nonzero, fix all other coefficients and solve for that coefficient; if all those differences vanish, the nonzero first difference prevents collision. Summing over all $q^{\mu-1}$ projections yields

$$q^{\mu-1}(N^2/L - N) \leq N(N-1)q^{\mu-2},$$

and hence $N \leq L(q-1)/(q-L)$. Finally

$$\frac{L(q-1)}{q-L} - L = \frac{L(L-1)}{q-L} < 1$$

when $L^2 < q$; integrality and the diagonal lower bound give equality. $\square$

Theorem 23.21 (Margin/interleaving split for a gauged MCA family). *Let $C = \text{RS}[\mathbb{F}, D, K]$ have length n , put $m = K + w \leq n$, and let Z be a finite family of distinct slopes with exact same-support-pair-noncontained m -records. Suppose their explanations lie in one affine flat $c_0 + C' \subseteq C$, where $s = \dim C'$. For each $\gamma \in Z$, define*

$$\theta_\gamma = \min \left\{ w + 1, \min_{b \in C'} |\{x \in S_\gamma : r_1(x) \neq b(x)\}| \right\}.$$

Fix an integer T with $2 \leq T \leq w$, put

$$A = m - T + 1 > K, \quad M_s(T) = \left\lfloor \frac{\binom{n-K+s}{s}}{\binom{w-T+1+s}{s}} \right\rfloor,$$

and assume $M_s(T)^2 < |\mathbb{F}|$. Define

$$H_0(T) = \left\lfloor \frac{n}{T} \right\rfloor$$

and, for $1 \leq j \leq s$,

$$H_j(T) = \left\lfloor \max \left\{ \frac{n^{j+1}}{mT(w+1)^{j-1}}, \frac{(n-K+j)^{j+1}}{T(w+1)^j} \right\} \right\rfloor.$$

Then

$$(23.6) \quad |Z| \leq \max_{0 \leq j \leq s} H_j(T) + (n-A)M_s(T).$$

Proof. Partition $Z = Z_{\text{hi}} \sqcup Z_{\text{lo}}$ according as $\theta_\gamma \geq T$ or $\theta_\gamma \leq T-1$. If the affine explanation span of Z_{hi} has direction dimension $j \geq 1$, its direction space is contained in C' ; minimizing over that smaller space cannot decrease the margins. Thus [Theorem 6.12](#) gives $|Z_{\text{hi}}| \leq H_j(T)$. If $j = 0$, every explanation equals one word. Each support contains at least T coordinates where the line direction is nonzero, and at each coordinate the agreement equation determines at most one slope. Incidence counting gives $|Z_{\text{hi}}| \leq H_0(T)$.

For every $\gamma \in Z_{\text{lo}}$, choose $b_\gamma \in C'$ with at most $T - 1$ mismatches on S_γ , and put $a_\gamma = h_\gamma - \gamma b_\gamma \in c_0 + C'$. The pair (a_γ, b_γ) agrees with (r_0, r_1) on at least A common coordinates. Translating the first component by c_0 places all such pairs in the two-fold common-support interleaving of C' . The ordinary affine-span list compiler [Theorem 4.8](#) bounds each projected ordinary list by $M_s(T)$. The hypothesis $M_s(T)^2 < |\mathbb{F}|$ and [Theorem 23.20](#) therefore bound the number of distinct pairs by $M_s(T)$.

For one fixed pair (a, b) , let $H_{a,b} = \{x : r_0(x) = a(x), r_1(x) = b(x)\}$; then $|H_{a,b}| \geq A$. Pair noncontainment gives every owning slope an actual $x \in S_\gamma \setminus H_{a,b}$. At such an x ,

$$\gamma = -\frac{r_0(x) - a(x)}{r_1(x) - b(x)}.$$

Choosing one exception coordinate per slope is injective, so the pair owns at most $n - |H_{a,b}| \leq n - A$ slopes. Adding the disjoint high and low bounds proves (23.6). $\square$

Proposition 23.22 (KoalaBear affine-error rank-ten payment). *Retain the intrinsic near-rational deletion and exact post-deletion record selection of [Proposition 23.19](#). If the selected error words have affine rank $a = 10$, and the received line is over the deployed sextic field*

$$|\mathbb{F}| = 2130706433^6,$$

then

$$|Z_{\text{bad}}| \leq 2w + 5143522968716559 + 56727790457914040 = 61871313426765543 < B_*.$$

The slack is 213109414684629544. Thus the direct KoalaBear post-near affine-error-rank-ten branch is paid. The same printed split does not pay affine error ranks 11 or 12.

Proof. The reversible gauge [Theorem 6.13](#) converts error rank $a = 10$ into an explanation flat of dimension $s = 9$, preserving the slopes, errors, exact supports, and same-support pair noncontainment. Use [Theorem 23.21](#) with $T = 667$. Exact integer arithmetic gives

$$\begin{aligned} A &= 1115382, & M_9(667) &= 57781140652, & M_9(667)^2 &< 2130706433^6 = |\mathbb{F}|, \\ \max_{0 \leq j \leq 9} H_j(667) &= 5143522968716559, & (n - A)M_9(667) &= 56727790457914040. \end{aligned}$$

Adding the disjoint near charge $2w = 134944$ proves the displayed bound. An exhaustive exact scan of $2 \leq T \leq w$ shows that $T = 16$ is the first paying threshold and $T = 667$ is the unique minimizer for $s = 9$. For $s = 10$, corresponding to error rank 11, the minimum of this formula is $1040506078215897711 > B_*$, attained at $T = 876$; the minima at $s = 11, 12$ are larger still. $\square$

Theorem 23.23 (Nonuniform support-margin resource). *Retain the hypotheses and notation of [Theorem 6.12](#), but define the individual support margins*

$$\theta_\gamma = \min \left\{ w + 1, \min_{b \in C'} |\{x \in S_\gamma : r_1(x) \neq b(x)\}| \right\}.$$

Then

$$(23.7) \quad \sum_{\gamma \in Z} \theta_\gamma \leq C_s, \quad C_s = \left\lfloor \max \left\{ \frac{n^{s+1}}{m(w+1)^{s-1}}, \frac{(n-K+s)^{s+1}}{(w+1)^s} \right\} \right\rfloor.$$

Proof. Use the ordered-normal-basis proof of [Theorem 6.12](#), before replacing the individual final-extension counts by their minimum. With the proof's global zero-normal parameters (z, g) , the parameter point belonging to γ owns at least

$$(m - g)\theta_\gamma(w + 1)^{\overline{s-1}}$$

ordered independent coordinate-normal tuples. One independent $(s + 1)$ -tuple determines at most one parameter point. Summing first gives

$$(m - g)(w + 1)^{\overline{s-1}} \sum_{\gamma} \theta_\gamma \leq (n - z)^{\overline{s+1}}.$$

Here (z, g) are global: a zero normal's affine incidence equation is either identically satisfied on the parameter flat or never satisfied. Now $0 \leq g \leq z \leq K - s$. For fixed z , the ratio is largest

at $g = z$, and the successive-ratio calculation in the cited proof puts its maximum at $z = 0$ or $z = K - s$. These are the two terms in (23.7). $\square$

Theorem 23.24 (Weighted minimizing-pair core route cut). *Retain the hypotheses of Theorem 23.23. Fix $1 \leq \tau \leq w - 1$, and put*

$$Q_s(\tau) = \left\lfloor \frac{\binom{n-K+s}{s}}{\binom{w-\tau+s}{s}} \right\rfloor, \quad Q_s(\tau)^2 < |\mathbb{F}|.$$

For every γ with $\theta_\gamma \leq \tau$, fix one minimizing $b_\gamma \in C'$, and put

$$a_\gamma = c_\gamma - \gamma b_\gamma, \quad \lambda_\gamma = \tau + 1 - \theta_\gamma,$$

and group records by their ordered pair (a_γ, b_γ) . If $|Z| \geq E$ and $(\tau + 1)E > C_s$, some fixed pair has total weight at least

$$(23.8) \quad L_s(\tau, E) = \left\lceil \frac{\max\{0, (\tau + 1)E - C_s\}}{Q_s(\tau)} \right\rceil.$$

For a fixed pair (a, b) , write

$$H_{a,b} = \{x \in D : r_0(x) = a(x), r_1(x) = b(x)\}, \quad \delta = \max\{1, m - |H_{a,b}|\}.$$

The pair owns at most

$$(23.9) \quad c_\delta = \left\lfloor \frac{n - m + \delta}{\delta} \right\rfloor$$

records, and its weight is at most $c_\delta(\tau + 1 - \delta)$.

Proof. The positive-part identity and (23.7) give

$$\sum_{\theta_\gamma \leq \tau} (\tau + 1 - \theta_\gamma) = \sum_{\gamma \in Z} (\tau + 1 - \theta_\gamma)_+ \geq (\tau + 1)|Z| - C_s.$$

The selected pair (a_γ, b_γ) agrees with the received pair on at least $m - \tau > K$ coordinates. The ordinary affine-span list theorem and the sub-square interleaving collapse Theorem 23.20 bound the number of distinct ordered pairs by $Q_s(\tau)$. Pigeonholing proves (23.8); there is no bipartite factor two because the interleaving theorem counts complete ordered pairs.

Fix one pair and an owning record. On S_γ , the identity

$$r_0 - a = -\gamma(r_1 - b)$$

shows that $S_\gamma \cap H_{a,b} = \{x \in S_\gamma : r_1(x) = b(x)\}$. Since the margin cap is inactive, $\theta_\gamma = |S_\gamma \setminus H_{a,b}| \geq \delta$. For distinct finite slopes these exception sets are disjoint: at a point outside $H_{a,b}$, the displayed equation determines γ uniquely. Their union lies in $D \setminus H_{a,b}$. Division by δ , with the harmless overestimate $n - |H_{a,b}| \leq n - m + 1$ when $|H_{a,b}| \geq m$, proves (23.9). Each edge weight is at most $\tau + 1 - \delta$, proving the final assertion. $\square$

Proposition 23.25 (KoalaBear rank-eleven pair-core terminal). *Retain the intrinsic near deletion, exact record selection, and reversible gauge of Proposition 23.19. Suppose the selected error words have affine rank 11, so the gauged explanations lie in an affine ten-flat over the deployed field $|\mathbb{F}| = 2130706433^6$. If the received line has more than B_* bad slopes, then for each row below there exists a fixed actual minimizing pair satisfying that row. The pairs supplied by the two independently optimized cutoffs need not coincide.*

	τ	pair weight	owned slopes	$\delta \leq$
(PC11)	6486	743449148	114624	8
	1795	360132809	200632	4

Parallel records are allowed; the owned-slope column is not a distinct-neighbor degree.

Moreover, the complete one-cutoff upper bound obtained from the printed high-margin theorem, all cumulative interleaved pair caps, and the exact core-deficiency multiplicities has unique minimum

$$813929118931913384 = 134944 + 5401690553097387 + 808527428378681053 > B_*$$

at cutoff $J = 19737$. Thus those inputs do not pay rank eleven. The remaining theorem must couple different fixed pairs or route the dense parallel pair cores to a chronology-correct owner.

Proof. An over-budget line leaves at least $E = B_* - 2w + 1$ post-near slopes. Apply [Theorem 23.24](#) with $s = 10$. Exact integer optimization of (23.8) gives the two rows of (PC11); comparison with $c_\delta(\tau + 1 - \delta)$ gives the displayed deficiency walls.

For the final assertion, first fix one minimizing b_γ , and hence one pair type (a_γ, b_γ) , for every selected record, once and for all cutoffs. Let $G(t)$ count distinct selected pair types whose deficiency is at most t . Such a pair has common agreement at least $m - t$, so interleaving gives $G(t) \leq Q_{10}(t)$. Since c_δ decreases, summation by parts bounds the low records through J by

$$Q_{10}(1)c_1 + \sum_{\delta=2}^J (Q_{10}(\delta) - Q_{10}(\delta - 1))c_\delta.$$

The records with margin at least $J + 1$ are bounded by the maximum of the support-transverse theorem over affine subranks $0 \leq j \leq 10$. Scanning all $1 \leq J \leq w - 1$, subject to the exact sub-square guard, gives the unique displayed minimum. The release certificate recomputes every term, the adjacent cutoffs, and hostile mutations. $\square$

Remark 23.26 (Scope of the rank-eleven route cut). The proposition moves no active-v4 atom and does not close KoalaBear. A small constant-code star shows that the fixed-pair multiplicity and parallel edges are genuine under the local hypotheses, so no simple-graph expansion or smaller per-pair multiplier may be inferred. The missing result is an actual-record cross-pair compatibility or owner theorem, not another scalar threshold optimization.

Proposition 23.27 (Two-anchor reserve repricing). *Suppose the near-rational first-match set has size at most $2w$, the residual exception set has size at most 31, and a large rational owner has size $g > 2m - K$. If the active S/A/E assembly is used after those two sets, then its contained-owner input must be required at the target*

$$(23.10) \quad B_{\text{owner}}^{(2w)}(g) \leq B_* - (2w + 31) - (n - g).$$

With this target the four charges sum exactly to B_* . On the deployed KoalaBear row, the targets at $g = 2m - K + 1$ and $g = n$ are respectively

$$274980728110346481, \quad 274980728111260112.$$

On the Mersenne-31 arithmetic stress row they are respectively

$$15728609, \quad 16642288.$$

The corresponding full-owner targets are at least 4807520 and 9 times the proved exact average ceilings.

Proof. Write N for the near-rational set and E for the residual exception set. Since $2w$ equals 134944 on KoalaBear and 134896 on the Mersenne-31 stress row, neither bound can be absorbed into the 31-slope bound without an additional containment theorem. Charging the two sets in first-match order leaves at most $n - g$ crossing slopes, so the remaining target is forced by

$$2w + 31 + (n - g) + (B_* - (2w + 31) - (n - g)) = B_*.$$

Substitution at the two endpoints gives the displayed integers. Exact division of the full-owner targets by the average ceilings gives quotients 4807520 and 9, with positive remainders. This proves arithmetic viability, not the contained-owner bound in (23.10). $\square$

Theorem 23.28 (Degree-guarded effective-shift witness adapter). *Let D be an n -point Reed–Solomon evaluation domain, let $k < m \leq n$, put $\omega = n - m$, and for a word U define*

$$\mathcal{M}_U = \{(W, N) : W(x)U(x) = N(x) \text{ for every } x \in D\}.$$

For a nonzero vector define

$$s_k(W, N) = \max\{\deg W, \deg N - (k - 1)\}, \quad s_{k+1}(W, N) = \max\{\deg W, \deg N - k\}.$$

Then

$$s_{k+1}(W, N) \leq s_k(W, N) \leq s_{k+1}(W, N) + 1,$$

and the two lattice minima obey the same inequalities.

Suppose W is the monic split squarefree locator of $D \setminus T$, where $|T| = m$, and $N = Wc$. Under the effective envelope $s_{k+1}(W, N) \leq \omega$, the following are equivalent:

$$(23.11) \quad \deg c < k, \quad \deg N \leq \omega + k - 1, \quad s_k(W, N) \leq \omega.$$

They hold exactly when the lattice vector reconstructs a degree-below- k explanation of U on the identical support T . Conversely every such explanation has a unique guarded representation.

For a received line $U = u + \gamma v$, interpolate $u|_T$ and $v|_T$ to their unique degree-below- m polynomials. The pair is simultaneously code-explained on T if and only if both interpolants have degree below k . Thus a guarded explanation together with failure of this pair test is an actual support-wise MCA-bad witness on the same line, slope, and support.

Proof. Put $a = \deg W$ and $b = \deg N - k$, with the usual convention $\deg 0 = -\infty$. Then

$$s_{k+1} = \max\{a, b\}, \quad s_k = \max\{a, b + 1\},$$

which proves the one-step comparison, including after taking minima. Since $\deg W = \omega$ and $N = Wc$, the effective cap permits $\deg c \leq k$, while each condition in (23.11) removes exactly the degree- k coefficient.

At every point of T , W is nonzero, so the lattice identity gives $U = c$. Conversely an explanation h on T gives

$$W = \prod_{x \in D \setminus T} (X - x), \quad N = Wh.$$

The monic locator and h are unique because $m \geq k$. Finally, unique degree-below- m interpolation on T turns simultaneous pair explainability into the two stated degree tests. $\square$

Proposition 23.29 (Transport mutation and typed pole-line control). *On the deployed KoalaBear row let $D = \langle \zeta \rangle$ have order 2^{21} , let E be the first 67473 subgroup powers, and let S be the following 1116048 powers.*

First, for

$$u = \mathbf{1}_E, \quad v = X^k, \quad \gamma = 0,$$

the identical slope/support record is support-wise MCA-bad for $RS[\mathbb{F}, D, k]$ but not for $RS[\mathbb{F}, D, k+1]$. Hence silent dimension transport does not preserve badness or first-owner semantics.

Second, over the deployed sextic extension, let α be a root of $X^6 + X + 6$ and set

$$v(x) = -\frac{1}{x - \alpha}, \quad u(x) = \mathbf{1}_E(x) + \frac{\alpha}{x - \alpha}, \quad \gamma = \alpha.$$

This is an actual support-wise MCA-bad witness on S , it passes (23.11) with quotient zero, and its lattice minimum is exactly 67473 under both shifts. Its Q , BC , and new-owner fields remain unassigned.

Proof. For the first record, zero explains u on S . A degree-below- k polynomial agreeing with X^k on $m > k$ points would make a nonzero degree- k polynomial have more than k roots. At dimension $k + 1$, the pair $(0, X^k)$ is a simultaneous explanation on the same support.

For the pole line, direct cancellation gives $u + \alpha v = \mathbf{1}_E$, again explained by zero on S . If g of degree below $k + 1$ agreed with v there, then $(X - \alpha)g + 1$ would have $m > k + 1$ roots and degree at most $k + 1$, which is impossible. The complement locator with numerator zero therefore passes the guarded adapter.

The prefix locator gives shifted degree 67473. If a nonzero lattice vector had shifted degree at most 67472, its numerator would vanish on all $2097152 - 67473 = 2029679$ points outside E , while its degree would be at most 1116048 under either shift. Thus the numerator is zero. The denominator would then vanish at all 67473 points of E despite degree at most 67472, another contradiction. No owner inference occurs in this argument. $\square$

Remark 23.30 (Acceptance contract for the rank-eleven owner route). The pair-core terminal of [Proposition 23.25](#) forces a possibly fixed pair of deficiency at most 4 owning at least 200632 slopes. A chronology-correct owner theorem for that terminal may use the preceding results as an interface: if it reuses S/A/E it must prove the repriced bound (23.10); if it computes certificates in the effective $k + 1$ envelope it must impose (23.11) and exact-support pair noncontainment; and it must preserve the actual line, support, slope, and owner chronology. The preceding results prove witness soundness and arithmetic consistency only. They do not construct the owner, pay rank eleven, move an active-v4 atom, or close KoalaBear.

Proposition 23.31 (Exact rank-flat and common-zero reductions at the active rows). *Consider the first affine dimensions not paid by [Proposition 23.17](#): $s = 15$ on KoalaBear and $s = 7$ on Mersenne-31. If such an affine Q /list chart exceeds the full row budget, [Corollary 4.10](#) forces a proper direction subcode with the following coarse support ceiling:*

	row	s	$\Theta_{s,B^*}^{\text{coarse}}$	$d_j \leq$	$d_j - (n - K) \leq$
(RF-active)	KoalaBear MCA	15	152252	1133356	84781
	KoalaBear list	15	152252	1133357	84781
	Mersenne-31 MCA	7	232423	1213551	164976
	Mersenne-31 list	7	232423	1213552	164976

Thus an over-budget first-unpaid affine chart is not generalized-weight uniform: it contains a proper direction subcode supported on a union much smaller than the full domain.

There is also an exact common-zero refinement beyond the former fixed-union boundary. Write the common-zero set of the affine direction space as $z = K - \nu$, and let b be the number of those coordinates on which the common affine value disagrees with the received word. At the first nullity not paid by [Proposition 23.16](#), the common-zero Johnson compiler gives

	row	ν	$\min b$	bound	residual b
(CZ-active)	KoalaBear MCA	4984	2	537111	$b \leq 1$
	KoalaBear list	4984	2	558125	$b \leq 1$
	Mersenne-31 MCA	4981	7	590532	$b \leq 6$
	Mersenne-31 list	4981	7	616024	$b \leq 6$

At nullity 5000, the corresponding paid mismatch thresholds are $b \geq 102$ on the two KoalaBear rows and $b \geq 126$ on the two Mersenne-31 rows.

Proof. For (RF-active), use the coarse form of [Corollary 4.10](#). The integer roots are

$$\left\lfloor \left(\frac{n^{15}}{(B_{\text{KB}}^* + 1)(67471 + 15)} \right)^{1/14} \right\rfloor = 152252$$

and

$$\left\lfloor \left(\frac{n^7}{(B_{\text{M31}}^* + 1)(67447 + 7)} \right)^{1/6} \right\rfloor = 232423.$$

Adding the exact radii $t = n - a_+$ gives the support and nullity ceilings in the table.

For (CZ-active), substitute the row data into (4.20) and choose the least nonnegative integer b for which the denominator is positive and the integer floor is at most the row budget. All displayed comparisons, including the nullity- 5000 thresholds, are exact and are reproduced by the verifier. $\square$

Proposition 23.32 (Exact codimension-one recursion at the active rows). *Consider a first-unpaid affine Q /list chart: dimension 15 on a KoalaBear route and dimension 7 on a Mersenne-31 route. Suppose the low-support direction subcode forced by [Proposition 23.31](#) has support closure of codimension one in the chart direction space. Write $j = 14$ or $j = 6$ for the dimension of that support-saturated hyperplane, d for its support size, e for the support increment to the whole chart, and b_0 for the common mismatch count outside the whole support. Then the exact*

full-budget payments are

(CR-active)

row	j	$d - (n - K) \leq$	e paid from	b_0 paid from	largest residual nullity
KoalaBear MCA	14	84781	8565	9813	93345
KoalaBear list	14	84781	8565	9813	93345
Mersenne-31 MCA	6	164976	73266	40074	238241
Mersenne-31 list	6	164976	73267	40074	238242

Here “ e paid from” means that every chart with at least that support increment is below the full row budget, and similarly for b_0 . The largest residual nullity is $d + e - (n - K)$ after both payments have failed. More explicitly, every surviving codimension-one branch lies in

(CR-KB)	$e \leq 8564$,	$b_0 \leq 9812$	on either KoalaBear route,
(CR-MCA)	$e \leq 73265$,	$b_0 \leq 40073$	on the Mersenne-31 MCA route,
(CR-list)	$e \leq 73266$,	$b_0 \leq 40073$	on the Mersenne-31 list route.

The corresponding certified maxima at the first paid integers are

	row	C_e	C_{b_0}	B^*
(CR-budget)	KoalaBear MCA	274959863953460332	274954310643125130	274980728111395087
	KoalaBear list	274963447501329540	274957707108580869	274980728111395087
	Mersenne-31 MCA	16777189	16776155	16777215
	Mersenne-31 list	16777174	16776251	16777215

where C_e is the closed bound at the first paid support increment and C_{b_0} is the all-extender bound at the first paid mismatch count. The immediately preceding integer fails for the displayed MDS-soft certificate in every row; this records the sharp integer transition of this compiler, not a lower bound on an actual list.

Proof. The support closure has the same support as the direction subcode supplied by [Proposition 23.31](#); hence its support size is in the exact ranges

row	d range
KoalaBear MCA	$1048589 \leq d \leq 1133356$
KoalaBear list	$1048590 \leq d \leq 1133357$
Mersenne-31 MCA	$1048581 \leq d \leq 1213551$
Mersenne-31 list	$1048582 \leq d \leq 1213552$.

For the MDS-soft profile Π of [Corollary 5.3](#), the largest initial factor is $A = d - t + b_0$ and $Q = w + 1 + b_0$. Uniformly over the displayed ranges,

$$(j-1)Q - A \geq \begin{cases} 13 \cdot 67472 - 152252 = 724884, & \text{KoalaBear,} \\ 5 \cdot 67448 - 232423 = 104817, & \text{Mersenne-31.} \end{cases}$$

The final term grows by $(j-2)b_0$, so [Lemma 5.2](#) proves profile interpolation for every e and b_0 in the active ranges.

The closed bound (5.11) decreases with e and with b_0 , while the all-extender bound (5.12) decreases with b_0 . It therefore suffices to use $b_0 = 0$ for the support-increment payment and to test the displayed mismatch threshold for the all-extender payment. Exact rational comparison over every integer d in the four finite ranges gives the maxima in (CR-budget). At the preceding support increments the maxima are respectively

$$274997419229066860, \quad 275001003273987385, \quad 16777303, \quad 16777288,$$

and at the preceding mismatch counts they are

$$275005816112956831, \quad 275009213214651183, \quad 16777247, \quad 16777343.$$

Each exceeds its row budget. Finally combine the support ceiling with $e \leq e_{\text{paid}} - 1$ to obtain the residual nullities in (CR-active). The release verifier performs all comparisons with exact integers. $\square$

Proposition 23.33 (Active higher-codimension flag payment). *Consider a first-unpaid affine Q /list chart: dimension 15 on a KoalaBear route and dimension 7 on a Mersenne-31 route. If the chart is over the full row budget, choose a generalized-weight minimizer supplied by Corollary 4.10. By Lemmas 6.2 and 6.3, its support closure admits the exact coset-free flag certificate of Theorem 6.4 in every codimension.*

For support-closure codimension at least two, the coarse all-extender certificate already pays every chart with

(23.12)

$$b_0 \geq 70230 \quad \text{on either KoalaBear route,} \quad b_0 \geq 108962 \quad \text{on either Mersenne-31 route.}$$

At the first paid integers, the worst exact bounds are

	row	b_0	certified maximum	B^*
(HF-active)	KoalaBear MCA	70230	274974976450914526	274980728111395087
	KoalaBear list	70230	274975238687487221	274980728111395087
	Mersenne-31 MCA	108962	16776934	16777215
	Mersenne-31 list	108962	16776950	16777215

The preceding mismatch integer fails the same uniform certificate, with values respectively

$$275004931457252515, \quad 275005193722392530, \quad 16777600, \quad 16777616.$$

These are transitions of the all-extender flag certificate, not lower bounds on actual lists. Below (23.12), the exact binary-pattern linear program (6.8) and the closed tensor certificate (6.12) remain available for the actual support-layer profile.

Proof. The rank-flat conclusion supplies a generalized-weight minimizer of some dimension $j < s$. It is support-saturated by Lemma 6.2, and therefore has a support-saturated flag to the whole chart by Lemma 6.3. For codimension at least two, use (6.9), replace the flag-support product by $(n - j)^{\overline{s-j}}$, and use the exact support ceilings in (RF-active). The resulting bound is

$$\frac{d^j (n - j)^{\overline{s-j}}}{(d - t + b_0) \prod_{i=1}^{j-1} (w + i + b_0) (w + 1 + b_0)^{s-j}}.$$

For fixed j its maximum over the allowed support interval occurs at an endpoint: the ratio of the d -dependent factor at $d + 1$ and d changes sign at most once, since the numerator of the comparison is

$$(j - 1)d - j(t - b_0) + j - 1.$$

Checking the two endpoints for $1 \leq j \leq 13$ on KoalaBear and $1 \leq j \leq 5$ on Mersenne-31 gives the exact transitions displayed above. The independent verifier performs these integer comparisons without floating point. $\square$

Proposition 23.34 (Exact paving-basis MCA payments at the active rows). *Let $j = R - d_U(y_1)$ be the direction defect of a separated fixed-union MCA chart. With the full row budget, the paving threshold of Corollary 6.79 is*

(PAV-KB)	KoalaBear ν	0-9	10	11	12	13
	$J_{B^*}^{\text{pav}}(\nu)$	67471	67468	67423	66715	55724

and

(PAV-M31)	Mersenne-31 ν	0-1	2	3	4	5
	$J_{B^*}^{\text{pav}}(\nu)$	67447	67432	67213	63797	10700.

Consequently every separated KoalaBear direction defect is paid through nullity 9. At Mersenne nullity 2, only the top fifteen possible defects $67433 \leq j \leq 67447$ remain outside this compiler. At every nullity the certified range is the maximum of the paving threshold and the directional-Johnson threshold in Eqs. (FU-KB-def) and (FU-M31-def).

Proof. Substitute $R = 1048576$, the deployed radii, and the exact row budgets into (6.46). At the last paid defects in the nontrivial columns, the paving bounds are respectively

ν	$J_{B^*}^{\text{pav}}(\nu)$	bound at $J_{B^*}^{\text{pav}}(\nu)$
10	67468	215264294199835875
11	67423	273052400452931367
12	66715	274631447525088415
13	55724	274966795861149003

and

ν	$J_{B^*}^{\text{pav}}(\nu)$	bound at $J_{B^*}^{\text{pav}}(\nu)$
2	67432	15838845
3	67213	16764415
4	63797	16774561
5	10700	16776956

The next defect exceeds the budget in every displayed nonterminal column. All calculations use exact binomial integers. $\square$

Proposition 23.35 (moment-order floor for a finite Q proof). *Let*

$$\tau = \frac{\sum_z |\mathcal{P}_Q(z)|}{\binom{n}{m}} \in (0, 1]$$

be the mass of the pruned residual relative to the full support slice. A proof based only on the moment criterion of Theorem 23.1 must satisfy the necessary condition

$$r(\Delta_Q - \log_2 \tau) \geq w \log_2 |\mathbb{B}|.$$

In the full-mass case $\tau = 1$, this reduces to

$$r \geq \left\lceil \frac{w \log_2 |\mathbb{B}|}{\Delta_Q} \right\rceil.$$

For the active rows the corresponding full-mass floors are

row	w	Δ_Q bits	minimum r
KoalaBear MCA	67471	22.1969	94196
KoalaBear list	67471	22.0109	94991
Mersenne-31 MCA	67447	3.2589	641593
Mersenne-31 list	67447	3.0730	680397

Thus fixed moments cannot prove the finite adjacent rows in the full-mass model. After first-match pruning, the displayed row floors remain applicable only when $\tau = 1$; otherwise the residual mass must be carried explicitly, or one must use an off-diagonal or falling-factorial moment or isolate sparse-heavy fibers as a separate cell.

Proof. Normalize the residual masses by the full slice and put $\mu(z) = |\mathcal{P}_Q(z)| / \binom{n}{m}$, so $\sum_z \mu(z) = \tau$. Convexity over at most $|\mathbb{B}|^w$ prefix values gives

$$|\mathbb{B}|^{w(r-1)} \sum_z \mu(z)^r \geq \tau^r.$$

In the notation of Theorem 23.1, this is $\Gamma_r \geq \tau^r$. The finite criterion requires

$$\log_2 \Gamma_r + w \log_2 |\mathbb{B}| \leq r \Delta_Q.$$

Substituting the lower bound and rearranging gives the stated necessary condition. Setting $\tau = 1$ yields the table. The numerical entries use the exact row prime in $\log_2 |\mathbb{B}|$. $\square$

Proposition 23.36 (the BC moving-root theorem does not pay boundary Q). *In the boundary-Q profile of the imported boundary-Q dictionary [Cho26, Secs. 18–21], the unknown divisor is*

$$Q_z + A, \quad \deg A \leq \omega - w - 1, \quad \omega = n - a_+.$$

Thus the coefficient parameter space has dimension $\omega - w$. At the active rows this dimension is

row	$\omega = n - a_+$	w	$\omega - w$
KoalaBear MCA	981104	67471	913633
KoalaBear list	981105	67471	913634
Mersenne-31 MCA	981128	67447	913681
Mersenne-31 list	981129	67447	913682

Therefore [Theorem 23.4](#) and [Corollary 23.5](#) cannot be applied directly to boundary Q : those theorems count split parameters on a projective one-parameter locator pencil, while Q is a high-dimensional affine divisor slice. To route Q through BC one would need an additional theorem decomposing this affine slice into paid one-parameter pencils with a first-match bound on the number of relevant pencils. No such theorem is present in this manuscript.

Proof. The imported boundary- Q dictionary [[Cho26](#), Secs. 18–21] states exactly that a prefix fiber is counted by the divisibility condition $Q_z + A \mid X^n - \beta$ with $\deg A \leq \omega - w - 1$. A polynomial of degree at most $\omega - w - 1$ has $\omega - w$ free coefficients. Substituting the four values of a_+ , K , and $w = a_+ - K$ gives the table. Since $\omega - w$ is larger than nine hundred thousand in all four cases, the family is not a one-parameter pencil. The moving-root theorem is still valid for every line inside this affine space, but a line-by-line decomposition without a bound on the number of lines gives no row budget. $\square$

Part 5. Order-32 line rank and rational-atom classification

The rest of the paper changes the MCA completion mechanism. Instead of adding many local cell budgets, it seeks a maximum-type classification: every large bad-slope set is either one global affine codeword line, one coherent rational atom, or one primitive spread family. This part proves the exact rank dictionary, the dyadic quotient branch, low-degree and low-complexity atom extraction, and atom coherence.

24. RANK-THEORETIC SETUP

Let $D \subseteq \mathbb{F}$ consist of n distinct evaluation points and let

$$C = \text{RS}[\mathbb{F}, D, k] = \{(h(x))_{x \in D} : h \in \mathbb{F}[X], \deg h < k\}.$$

Fix distinct finite slopes $\gamma_1, \dots, \gamma_L$ and a received line $r_\gamma = r_0 + \gamma r_1$. Suppose $h_i \in \mathbb{F}[X]_{<k}$ agrees with r_{γ_i} on a support S_i of size m ; shrinking supports to exact size m does not affect the rank calculation. Put

$$I_x = \{i : x \in S_i\}.$$

The tuple is globally affine if $h_i = A + \gamma_i B$ for some $A, B \in \mathbb{F}[X]_{<k}$.

25. THE RANK-TWO REDUCED INTERSECTION MATRIX

25.1. Local relation spaces. For $I \subseteq [L]$, define

$$\mathcal{K}_\gamma(I) := \left\{ c \in \mathbb{F}^L : \text{supp}(c) \subseteq I, \sum_{i=1}^L c_i = 0, \sum_{i=1}^L c_i \gamma_i = 0 \right\}.$$

Since the slopes are distinct,

$$\dim \mathcal{K}_\gamma(I) = (|I| - 2)_+.$$

Let $v_k(x) = (1, x, \dots, x^{k-1}) \in \mathbb{F}^k$.

Definition 25.1 (rank-two reduced intersection matrix). *Choose any basis $\mathcal{B}_x$ of $\mathcal{K}_\gamma(I_x)$ for every $x \in D$. The matrix*

$$\mathcal{M}^{(2)}(D, \mathcal{S}, \gamma)$$

has Lk columns, grouped into L blocks of length k , and one row

$$c \otimes v_k(x)$$

for every $x \in D$ and $c \in \mathcal{B}_x$.

Changing the local bases left-multiplies the corresponding row blocks by invertible matrices, so the kernel and rank are intrinsic.

Theorem 25.2 (exact kernel dictionary). *For $h_i \in \mathbb{F}[X]_{<k}$, let $\mathbf{h}$ be the concatenated coefficient vector. Then*

$$\mathbf{h} \in \ker \mathcal{M}^{(2)}(D, \mathcal{S}, \gamma)$$

if and only if, for every $x \in D$, there exist scalars $a_x, b_x \in \mathbb{F}$ such that

$$h_i(x) = a_x + \gamma_i b_x \quad (i \in I_x).$$

In particular, every family of codeword explanations of one received line belongs to the kernel.

Proof. For fixed x , the local equations are

$$\sum_{i \in I_x} c_i h_i(x) = 0 \quad (c \in \mathcal{K}_\gamma(I_x)).$$

The orthogonal complement of $\mathcal{K}_\gamma(I_x)$ in $\mathbb{F}^{I_x}$ is exactly

$$\text{Span}\{\mathbf{1}_{I_x}, (\gamma_i)_{i \in I_x}\}.$$

Thus the local value vector is affine in the slope exactly when all local rows annihilate it. Applying this independently at each coordinate proves the claim. For a received line, take $a_x = r_0(x)$ and $b_x = r_1(x)$. $\square$

Proposition 25.3 (row count and trivial kernel). *The matrix has*

$$\sum_{x \in D} (r_x - 2)_+ \geq Lm - 2n$$

rows. It always contains the $2k$ -dimensional kernel

$$h_i = A + \gamma_i B, \quad A, B \in \mathbb{F}[X]_{<k}.$$

Consequently its maximum useful rank is $(L - 2)k$, and the equation surplus over the nontrivial unknown dimension is

$$L(m - k) - 2(n - k).$$

Proof. Since $(r - 2)_+ \geq r - 2$,

$$\sum_x (r_x - 2)_+ \geq \sum_x r_x - 2n = \sum_i |S_i| - 2n = Lm - 2n.$$

A globally affine tuple satisfies every local relation. Distinct slopes make the map $(A, B) \mapsto (A + \gamma_i B)_i$ injective, so this kernel has dimension $2k$. Subtracting $(L - 2)k$ from the row lower bound gives the displayed surplus. $\square$

26. WHY THE DEPLOYED CRITICAL ORDER IS EXACTLY 32

The actual Reed–Solomon dimension on both active MCA rows is

$$n = 2^{21} = 2097152, \quad k = 2^{20} = 1048576.$$

The adjacent safe agreements are

$$m_{\text{KB}} = 1116048, \quad m_{\text{M31}} = 1116024.$$

Proposition 26.1 (critical order). *For either deployed row,*

$$L_* := \left\lfloor \frac{2(n - k)}{m - k} \right\rfloor + 1 = 32.$$

More precisely,

row	$m - k$	$31(m - k) - 2(n - k)$	$32(m - k) - 2(n - k)$
KoalaBear MCA	67472	−5520	61952
Mersenne-31 MCA	67448	−6264	61184

Thus order 31 can remain underdetermined for dimension reasons, whereas order 32 has a positive row surplus.

Proof. Direct substitution into [Proposition 25.3](#). $\square$

27. WHAT FULL RELATIVE RANK WOULD BUY

Theorem 27.1 (one global code-line block). *Let Z be a set of bad slopes whose explaining codewords lie on one global codeword line,*

$$h_\gamma = A + \gamma B.$$

Assume the received line is support-wise MCA-bad at agreement m . Then

$$\boxed{|Z| \leq n - m + 1.}$$

Proof. Let

$$C_0 := \{x \in D : r_0(x) = A(x), r_1(x) = B(x)\}, \quad c := |C_0|.$$

MCA-badness gives $c < m$. Outside C_0 , the equation

$$r_0(x) - A(x) + \gamma(r_1(x) - B(x)) = 0$$

has at most one finite solution γ . Counting agreements over all $\gamma \in Z$ therefore yields

$$|Z|m \leq |Z|c + (n - c).$$

Hence

$$|Z| \leq \frac{n - c}{m - c} \leq n - m + 1,$$

where the last expression is the maximum over integers $0 \leq c \leq m - 1$. $\square$

Corollary 27.2 (conditional generic closure). *Suppose every 32-subfamily of a set of bad explanations has full relative rank, meaning that the kernel of its rank-two matrix is exactly the global affine kernel. Then all explanations lie on one global codeword line and*

$$|Z| \leq n - m + 1.$$

At the deployed rows the bounds are

$$981105 \quad \text{and} \quad 981129,$$

respectively.

Proof. Any 32 explanations are globally affine. If there are more than 32, fix two slopes and place any third slope in a 32-subset containing the fixed pair; the unique codeword line through the fixed pair contains the third. Apply [Theorem 27.1](#). $\square$

The Mersenne-31 challenge numerator is 16777215, so the generic order-32 branch would fit with more than fifteen million units of margin. The difficult issue is therefore not the generic branch. It is classification and exclusive payment of specialized rank defects.

28. A NECESSARY PARTITION INEQUALITY

The following obstruction identifies the correct rank-two analogue of the partition conditions in higher-order MDS theory.

For a partition $\mathcal{P}$ of $[L]$, put

$$c(\mathcal{P}) := \sum_{P \in \mathcal{P}} \min\{2, |P|\},$$

and for $I \subseteq [L]$ define

$$w_{\mathcal{P}}(I) := \sum_{P \in \mathcal{P}} \min\{2, |I \cap P|\} - \min\{2, |I|\}.$$

Proposition 28.1 (partition obstruction). *If*

$$\sum_{x \in D} w_{\mathcal{P}}(I_x) < k(c(\mathcal{P}) - 2)$$

for some partition $\mathcal{P}$, then the rank-two matrix has kernel dimension strictly larger than $2k$. Consequently full relative rank requires

$$\boxed{\sum_{x \in D} w_{\mathcal{P}}(I_x) \geq k(c(\mathcal{P}) - 2) \quad \text{for every partition } \mathcal{P}.}$$

Proof. Consider tuples that are independently affine on each block of $\mathcal{P}$: a singleton block contributes one free degree- $< k$ polynomial, and a block of size at least two contributes two. This space has dimension $kc(\mathcal{P})$. At coordinate x , forcing the blockwise-affine local values on I_x to lie on one affine slope line imposes at most $w_{\mathcal{P}}(I_x)$ scalar conditions. Hence the intersection with the rank-two kernel has dimension at least

$$kc(\mathcal{P}) - \sum_x w_{\mathcal{P}}(I_x).$$

Under the strict inequality this exceeds $2k$. □

For generic evaluation points, higher-order MDS and row-span constrained local-property machinery make closely related partition inequalities sufficient for the expected rank and for curve-decoding phenomena [BGM23, GGSW26]. The deployed domains are not generic. The point of the next sections is to identify what specializations must be removed before a sufficiency theorem can be true. Recent work explicitly notes that identifying deterministic Reed–Solomon evaluation sets with the random-evaluation curve-decoding behavior remains open [GGSW26].

29. EXACT DYADIC QUOTIENT DESCENT

Theorem 29.1 (pair descent on a multiplicative dyadic domain). *Assume $\text{char } \mathbb{F} \neq 2$ and*

$$D = D_+ \sqcup (-D_+).$$

Suppose the incidence pattern is pair-invariant:

$$I_x = I_{-x} \quad (x \in D_+).$$

Write uniquely

$$h_i(X) = u_i(X^2) + Xv_i(X^2),$$

where

$$\deg u_i < \left\lceil \frac{k}{2} \right\rceil, \quad \deg v_i < \left\lfloor \frac{k}{2} \right\rfloor.$$

Then the rank-two kernel on D is naturally the direct sum of two rank-two kernels on the squared domain

$$D^{(2)} := \{x^2 : x \in D_+\},$$

with the common incidence sets $I_x = I_{-x}$ and degree bounds $\lceil k/2 \rceil$ and $\lfloor k/2 \rfloor$. In particular, every nontrivial pair-invariant defect descends to a nontrivial defect on the half-size quotient in at least one parity component.

Proof. Fix $x \in D_+$ and $c \in \mathcal{K}_{\gamma}(I_x) = \mathcal{K}_{\gamma}(I_{-x})$. The two local equations are

$$\sum_i c_i(u_i(x^2) + xv_i(x^2)) = 0, \quad \sum_i c_i(u_i(x^2) - xv_i(x^2)) = 0.$$

Because $2x \neq 0$, adding and subtracting gives

$$\sum_i c_i u_i(x^2) = 0, \quad \sum_i c_i v_i(x^2) = 0.$$

The converse follows by reversing the calculation. The even/odd polynomial decomposition is a vector-space isomorphism, so the full kernel is the direct sum claimed. The global affine kernels descend componentwise; hence any excess kernel dimension appears in at least one component. □

Iterating [Theorem 29.1](#) proves exact descent for every fully periodic dyadic incidence atom. This is the rigorous quotient branch of the proposed specialization theorem. What fails is the assertion that every non-descending atom must be a simple pole.

30. RATIONAL ATOMS AND PROFILE DEPTH

Let

$$\Lambda_i(X) := \prod_{x \in S_i} (X - x)$$

be the monic support locator of degree m .

Definition 30.1 (scalar-locator rational atom). *A family is a scalar-locator rational atom if there exist*

$$Q, A, B \in \mathbb{F}[X], \quad c_i = c_0 + c_1 \gamma_i,$$

with $Q(x) \neq 0$ for all $x \in D$, such that

$$\boxed{Qh_i + c_i \Lambda_i = A + \gamma_i B \quad (i \in [L]).}$$

Proposition 30.2 (rational atoms are MCA explanations). *Every atom of [Definition 30.1](#) is explained by the rational received line*

$$r_\gamma(x) = \frac{A(x) + \gamma B(x)}{Q(x)}.$$

The codeword h_i agrees with r_{γ_i} on all roots of Λ_i .

Proof. At $x \in S_i$, the locator term vanishes and division by $Q(x)$ is valid. □

Proposition 30.3 (converse under the scalar degree profile). *Put $m = k + d$. Suppose*

$$r_\gamma(x) = \frac{A(x) + \gamma B(x)}{Q(x)}$$

on D , where Q has no root in D , $\deg Q = s \leq d$, and $\deg A, \deg B \leq m$. If $h_i \in \mathbb{F}[X]_{<k}$ agrees with r_{γ_i} on the m roots of Λ_i , then

$$A + \gamma_i B - Qh_i = c_i \Lambda_i$$

for a scalar c_i , and c_i is affine in γ_i .

Proof. The left side has degree at most m and vanishes at the m distinct roots of Λ_i , so it is a scalar multiple of Λ_i . Since $\deg(Qh_i) \leq k + s - 1 \leq m - 1$, its degree- m coefficient is zero. Comparing degree- m coefficients gives

$$c_i = [X^m]A + \gamma_i [X^m]B.$$

□

Corollary 30.4 (scaled-locator prefix depth). *Under the hypotheses of [Proposition 30.3](#), the coefficients of*

$$c_i \Lambda_i$$

in degrees

$$k + s, k + s + 1, \dots, m$$

are affine functions of γ_i . Apart from the leading scaling coefficient, this is a depth- $(d - s)$ affine prefix constraint.

Proof. The product Qh_i has degree at most $k + s - 1$, so the displayed high coefficients of $c_i \Lambda_i$ equal those of $A + \gamma_i B$. □

Thus a simple pole, $s = 1$, is only one endpoint of a profile continuum. At the opposite endpoint $s = d$, only the leading scaling is forced, while the locator residues modulo Q must lie on a projective affine line. Any correct rank-defect classification must account for the entire continuum or prove that the already-declared first-match cells remove all but selected endpoints.

31. AN EXACT DEGREE-TWO RANK DEFECT OVER $\mathbb{F}_{17}$

We now give a completely explicit obstruction to an unqualified “quotient or simple pole” theorem.

Theorem 31.1 (degree-two support-wise MCA atom). *Let*

$$\mathbb{F} = \mathbb{F}_{17}, \quad D = \mu_8 = \{1, 9, 13, 15, 16, 8, 4, 2\}, \quad k = 4, \quad m = 6.$$

Take slopes

$$(\gamma_0, \gamma_1, \gamma_2, \gamma_3, \gamma_4) = (0, 1, 4, 7, 2)$$

and supports

i	S_i
0	$\{1, 15, 16, 8, 4, 2\}$
1	$\{1, 9, 13, 15, 16, 8\}$
2	$\{1, 13, 15, 16, 4, 2\}$
3	$\{1, 9, 15, 16, 4, 2\}$
4	$\{1, 9, 13, 15, 8, 4\}$.

Let

$$\begin{aligned} h_0 &= h_1 = 0, \\ h_2 &= 9 + 13X + 8X^2 + 4X^3, \\ h_3 &= 3 + 10X + 14X^2 + 7X^3, \\ h_4 &= 11 + 8X + 9X^2 + 6X^3. \end{aligned}$$

Define a received line by the following values, in the displayed order of D :

x	1	9	13	15	16	8	4	2
$r_0(x)$	0	16	6	0	0	0	0	0
$r_1(x)$	0	1	11	0	0	0	5	12.

Then:

- (i) h_i agrees with $r_0 + \gamma_i r_1$ on S_i for every i ;
- (ii) the received line is support-wise MCA-bad at agreement $m = 6$;
- (iii) the rank-two matrix has rank 11, whereas full relative rank is 12;
- (iv) the defect is not first-level quotient-periodic;
- (v) it has no scalar-locator rational representation with $\deg Q \leq 1$, but it has one with $\deg Q = 2$.

Proof. Direct substitution proves (i). For (ii), observe that r_0 is zero at exactly the six points S_0 . On any other six-point set, a degree- < 4 polynomial agreeing with r_0 would have at least four zero roots and also take a prescribed nonzero value, which is impossible; with five zero roots the same conclusion is immediate. Hence the only possible six-point common explanation set for the first coefficient is S_0 . On S_0 , the word r_1 is zero at 1, 15, 16, 8 and nonzero at 4, 2, so no degree- < 4 polynomial can agree with it there. Thus no pair of codeword polynomials simultaneously explains the received line on six coordinates.

Part (i) and [Theorem 25.2](#) give a nontrivial kernel vector. It is not globally affine because $h_0 = h_1 = 0$ at slopes 0 and 1, while $h_2 \neq 0$. Exact row reduction gives rank 11; one 11×11 pivot minor has determinant $9 \in \mathbb{F}_{17}^\times$. This proves (iii).

For (iv), the incidence sets at 1 and $-1 = 16$ are different: the former contains all five slopes, while the latter omits slope γ_4 . Thus the hypothesis of [Theorem 29.1](#) fails already at the first dyadic level.

For (v), let

$$\begin{aligned} Q &= 12 + 3X + 5X^2, \\ A &= 13 + 10X + 7X^2 + 13X^3 + 5X^4 + 11X^5 + 9X^6, \\ B &= 7 + 5X + 6X^2 + 14X^3 + 3X^4 + 15X^5 + X^6, \end{aligned}$$

and $c_i = 9 + \gamma_i$. The polynomial Q is nonzero on every point of D , and exact multiplication gives

$$Qh_i + c_i\Lambda_i = A + \gamma_i B \quad (i = 0, \dots, 4).$$

For a denominator bound s , the identities of [Definition 30.1](#) form a homogeneous coefficient system in Q, A, B, c_0, c_1 . At $s = 0$ and $s = 1$ the matrices have full column ranks 17 and 18, respectively. At $s = 2$ the 35×19 system has rank 18 and its one-dimensional kernel is generated by the displayed certificate. Hence degree two is minimal within this atom model. $\square$

Remark 31.2 (the common-support qualifier is essential). The example has the common support factor

$$H = (X - 1)(X - 15) = X^2 + X + 15.$$

It divides every locator and also $A, B, h_0, \dots, h_4$. Therefore [Theorem 31.1](#) refutes an *unqualified* quotient-or-pole theorem, but it does not refute a theorem stated only after the common-support first-match branch has been removed. Instead, it proves that the qualifier cannot be omitted and that rational denominator profiles can be hidden inside common-support reductions.

The incidence pattern also satisfies every partition inequality of [Proposition 28.1](#), with minimum slack 1. Thus the defect is a genuine specialization phenomenon rather than a failure of the obvious partition count. This finite claim is included in the verifier.

32. MAXIMAL SUPPORTS AND THE SLOPE-DEGREE BARRIER

For each explanation, take its *maximal* agreement support

$$\widehat{S}_i := \{x \in D : h_i(x) = r_0(x) + \gamma_i r_1(x)\}, \quad |\widehat{S}_i| \geq m.$$

Put

$$\widehat{I}_x := \{i : x \in \widehat{S}_i\}, \quad C_\cap := \bigcap_{i=1}^L \widehat{S}_i, \quad c := |C_\cap|.$$

Let $R_0, R_1 \in \mathbb{F}[X]_{<n}$ be the interpolation polynomials of r_0, r_1 on D . Interpolating the codeword polynomials coefficientwise in the slope gives the unique polynomial

$$H(X, Z) \in \mathbb{F}[X, Z], \quad \deg_X H < k, \quad \deg_Z H < L,$$

with

$$H(X, \gamma_i) = h_i(X) \quad (i \in [L]).$$

Define the slope-error polynomial

$$E(X, Z) := H(X, Z) - R_0(X) - ZR_1(X).$$

Lemma 32.1 (zero slope fibers are exactly common agreements). *For $x \in D$,*

$$E(x, Z) \equiv 0 \iff x \in C_\cap.$$

Proof. If $x \in C_\cap$, then $E(x, \gamma_i) = 0$ for all L distinct slopes. Since $\deg_Z E(x, Z) < L$, the polynomial is zero. Conversely, if $E(x, Z)$ is zero, then evaluating at every γ_i gives

$$h_i(x) = R_0(x) + \gamma_i R_1(x) = r_0(x) + \gamma_i r_1(x),$$

so x belongs to every maximal support. $\square$

Theorem 32.2 (maximal-support slope-degree barrier). *Assume the explanations are not globally affine: there do not exist $A, B \in \mathbb{F}[X]_{<k}$ with*

$$h_i = A + \gamma_i B \quad (i \in [L]).$$

Then $E \neq 0$. Writing

$$r := \deg_Z E,$$

one has $c < k$ and

$$\boxed{Lm \leq Lc + r(n - c).}$$

Equivalently,

$$r \geq \left\lceil \frac{L(m-c)}{n-c} \right\rceil.$$

In particular, after the common-support branch is removed ($c = 0$),

$$r \geq \left\lceil \frac{Lm}{n} \right\rceil.$$

Proof. If $c \geq k$, choose two slopes and let $A, B \in \mathbb{F}[X]_{<k}$ be the unique codeword line through their polynomials. For every i , the polynomial $h_i - A - \gamma_i B$ has degree $< k$ and vanishes on $C_\cap$, hence is zero. This contradicts non-affinity, so $c < k$.

For $x \notin C_\cap$, Lemma 32.1 says $E(x, Z)$ is nonzero. It has the $|\hat{I}_x|$ distinct roots γ_i indexed by agreements at x , whence

$$|\hat{I}_x| \leq r.$$

Counting maximal-support incidences gives

$$Lm \leq \sum_i |\hat{S}_i| = \sum_{x \in C_\cap} L + \sum_{x \notin C_\cap} |\hat{I}_x| \leq Lc + r(n-c).$$

It remains only to justify that $E \neq 0$. If it were zero, then H would be affine in Z and all $h_i = H(X, \gamma_i)$ would lie on one global codeword line. $\square$

Corollary 32.3 (the active primitive degree is at least eighteen). *For $L = 32$ and either deployed adjacent row, a non-affine explanation with no common agreement coordinate satisfies*

$$18 \leq r \leq 31.$$

Indeed,

$$\left\lceil \frac{32m_{\text{KB}}}{n} \right\rceil = \left\lceil \frac{32m_{\text{M31}}}{n} \right\rceil = 18.$$

Thus every slope-degree ≤ 17 explanation is already global-affine or common-support. This removes the entire low-degree branch before any determinant specialization argument begins.

33. LOCAL INCIDENCE COMPLEXITY

For the rational extraction theorem, choose exact m -subsupports

$$S_i \subseteq \hat{S}_i, \quad |S_i| = m,$$

and set

$$I_x := \{i : x \in S_i\}, \quad r_x := |I_x|.$$

Let N_a count coordinates of incidence multiplicity a , let

$$N_{\geq 2} := |\{x : r_x \geq 2\}|, \quad U := \bigcup_i S_i, \quad u := |U|.$$

Definition 33.1 (two-cover incidence complexity). *Define*

$$\chi(\mathcal{S}) := \sum_{x \in D} \min\{2, r_x\} = N_1 + 2N_{\geq 2}.$$

Equivalently,

$$\chi(\mathcal{S}) = 2u - N_1 = 2n - 2N_0 - N_1.$$

Proposition 33.2 (exact relation to the rank-two row count). *For exact m -supports,*

$$\sum_{x \in D} (r_x - 2)_+ = Lm - \chi(\mathcal{S}).$$

Proof. At multiplicity zero or one, both sides contribute zero after subtracting $\min\{2, r_x\}$. At multiplicity at least two, the difference is $r_x - 2$. Sum over x and use $\sum_x r_x = Lm$. $\square$

The quantity χ is the number of scalar conditions needed to force one common rational received line on the occupied coordinates: two conditions at a multiply occupied coordinate and one at a singleton.

34. SUPPORT-COLLAPSED RATIONAL INTERPOLATION

Let

$$\Lambda_i(X) := \prod_{x \in S_i} (X - x)$$

be the monic locator of S_i .

Definition 34.1 (scalar-locator certificate). *A scalar-locator certificate of denominator degree at most s is a tuple*

$$(Q, A, B, c_0, c_1)$$

with

$$\deg Q \leq s, \quad \deg A, \deg B \leq m,$$

not all zero, such that, with $c_i = c_0 + c_1 \gamma_i$,

$$\boxed{Qh_i + c_i \Lambda_i = A + \gamma_i B \quad (i \in [L]).}$$

It is a rational atom when $Q \neq 0$ and $(c_0, c_1) \neq (0, 0)$, a pure locator pencil when $Q = 0$, and globally affine when $(c_0, c_1) = (0, 0)$.

When Q has no root on D , a rational atom is explained by the received line $(A + \gamma B)/Q$. A root of Q on D is retained as a separate denominator-root first-match branch rather than being silently divided away.

Theorem 34.2 (low-complexity atom extraction). *Assume $L \geq 2$. Let*

$$d = m - k$$

and fix an integer $0 \leq s \leq d$. If

$$\boxed{\chi(\mathcal{S}) < 2m + s + 3,}$$

then the explanations admit a scalar-locator certificate of denominator degree at most s .

More precisely, one of the following holds:

- (i) the codeword explanations are globally affine;
- (ii) they lie in a pure locator pencil;
- (iii) they form a rational atom, possibly with a denominator-root degeneracy on D .

Proof. Use as unknowns the coefficients of

$$Q \in \mathbb{F}[X]_{\leq s}, \quad A, B \in \mathbb{F}[X]_{\leq m}.$$

There are

$$(s + 1) + 2(m + 1) = 2m + s + 3$$

unknowns.

At a coordinate x with $r_x \geq 2$, impose the two homogeneous linear equations

$$A(x) = Q(x)r_0(x), \quad B(x) = Q(x)r_1(x).$$

At a singleton coordinate, with $I_x = \{i\}$, impose

$$A(x) + \gamma_i B(x) = Q(x)h_i(x).$$

No equation is needed at an unoccupied coordinate. The total number of equations is at most $\chi(\mathcal{S})$. Under the strict inequality, the homogeneous system has a nonzero solution (Q, A, B) .

For every i and every $x \in S_i$, the imposed equations give

$$A(x) + \gamma_i B(x) = Q(x)h_i(x).$$

The polynomial

$$P_i := A + \gamma_i B - Qh_i$$

has degree at most m because

$$\deg(Qh_i) \leq s + k - 1 \leq m - 1.$$

It vanishes on the m distinct roots of Λ_i , so

$$P_i = c_i \Lambda_i$$

for a scalar c_i . Comparing coefficients of X^m gives

$$c_i = [X^m]A + \gamma_i[X^m]B,$$

which is affine in γ_i . This is the required certificate.

If $Q = 0$, nontriviality and two distinct slopes rule out $c_0 = c_1 = 0$, so the certificate is a pure locator pencil. If $Q \neq 0$ but $c_0 = c_1 = 0$, then

$$Qh_i = A + \gamma_i B.$$

Using two slopes shows Q divides both A and B , and cancellation gives a global affine codeword line. The remaining case is a rational atom. $\square$

Remark 34.3 (relation to rational-function interpolation). The proof is a support-collapsed Berlekamp–Welch argument: it interpolates only the two received-line coefficients on multiply occupied coordinates, rather than imposing a separate condition for every support incidence. This is compatible with the rational-function and curve-decoding viewpoint in recent proximity-gap work [BCHKS25, GG25], but the theorem above is elementary and finite; no random-evaluation or asymptotic result is imported.

Corollary 34.4 (full denominator profile). *Taking $s = d = m - k$, every explanation with*

$$\boxed{\chi(\mathcal{S}) < 3m - k + 3}$$

is globally affine, a pure locator pencil, or a scalar-locator rational atom of denominator degree at most $m - k$.

Corollary 34.5 (support-union form). *If*

$$2u - N_1 < 3m - k + 3,$$

then the same alternatives hold. In particular, it suffices that

$$u \leq \left\lfloor \frac{3m - k + 2}{2} \right\rfloor.$$

Corollary 34.6 (exact deployed extraction thresholds). *For the active 32-tuples, put*

$$\Xi := 3m - k + 3.$$

Then

row	Ξ	$2n - \Xi$	automatic union ceiling	union nullity ceiling
KoalaBear MCA	2299571	1894733	1149785	101209
Mersenne-31 MCA	2299499	1894805	1149749	101173

where the last column subtracts $n - k = 1048576$ from the union ceiling.

Equivalently, extraction is automatic if

$$2N_0 + N_1 \geq 1894734$$

on KoalaBear, or if

$$2N_0 + N_1 \geq 1894806$$

on Mersenne-31.

These are classification thresholds, not challenge-numerator payments. Their use is to show that a large low-complexity support family belongs to one atom instead of being charged as many unrelated cells.

35. THE DEGREE-TWO COUNTEREXAMPLE LIES AT THE EXTRACTION THRESHOLD

The degree-two $\mathbb{F}_{17}$ example of [Theorem 31.1](#) has

$$n = 8, \quad k = 4, \quad m = 6, \quad L = 5.$$

Its five listed supports are maximal. Their incidence multiplicities are

$$(5, 3, 3, 5, 4, 3, 4, 3),$$

so

$$\chi = 16 \quad \text{while} \quad 3m - k + 3 = 17.$$

Thus [Theorem 34.2](#) forces a denominator of degree at most

$$d = m - k = 2,$$

which is exactly the minimal degree established by [Theorem 31.1](#).

The common support has size $c = 2$. The coefficientwise slope interpolation of the five codewords has degree $r = 4$, and

$$\left\lceil \frac{5(6-2)}{8-2} \right\rceil = 4.$$

Hence the example also attains equality in [Theorem 32.2](#). The counterexample is therefore not an unexplained accident: it sits simultaneously on the common-support slope-degree boundary and the low-complexity rational-extraction boundary.

36. COLLISION PENCILS FOR TWO ATOM CERTIFICATES

Suppose the same indexed explanations

$$(\gamma_i, h_i, \Lambda_i), \quad i \in I,$$

satisfy two nontrivial scalar-locator certificates

$$Qh_i + (c_0 + c_1\gamma_i)\Lambda_i = A + \gamma_i B,$$

$$Q'h_i + (c'_0 + c'_1\gamma_i)\Lambda_i = A' + \gamma_i B'.$$

Assume $|I| \geq 2$, $Q, Q' \neq 0$, and both scalar pairs are nonzero. Define

$$P_0 := Q'c_0 - Qc'_0, \quad P_1 := Q'c_1 - Qc'_1,$$

$$R_0 := Q'A - QA', \quad R_1 := Q'B - QB'.$$

Eliminating h_i gives the collision identity

$$(P_0 + \gamma_i P_1)\Lambda_i = R_0 + \gamma_i R_1.$$

Let

$$G := \{x \in D : R_0(x) = R_1(x) = 0\}, \quad g := |G|,$$

and

$$H := \{x \in G : P_0(x) = P_1(x) = 0\}, \quad h := |H|.$$

Theorem 36.1 (atom collision rigidity). *If $R_0 = R_1 = 0$ and the certificates share at least two distinct slopes, then the certificates are projectively identical.*

Otherwise, if $g < m$, then

$$|I| \leq \left\lfloor \frac{n-g}{m-g} \right\rfloor.$$

Moreover, every coordinate in $G \setminus H$ belongs to at least $|I| - 1$ of the supports S_i .

If both denominators have degree at most d , then two distinct certificates satisfy

$$h \leq d.$$

Proof. First suppose $R_0 = R_1 = 0$. Then

$$(P_0 + \gamma_i P_1) \Lambda_i = 0$$

as a polynomial. Since $\Lambda_i \neq 0$, two distinct slopes imply $P_0 = P_1 = 0$. Thus

$$Q'(c_0, c_1) = Q(c'_0, c'_1).$$

Choose a nonzero component of (c_0, c_1) . Because $Q, Q' \neq 0$, this identity gives a scalar $\lambda \in \mathbb{F}^\times$ with

$$Q' = \lambda Q, \quad (c'_0, c'_1) = \lambda(c_0, c_1).$$

The equations $R_0 = R_1 = 0$ then give

$$A' = \lambda A, \quad B' = \lambda B.$$

So the certificates are projectively identical.

Now assume $(R_0, R_1) \neq (0, 0)$. At every root of Λ_i ,

$$R_0(x) + \gamma_i R_1(x) = 0.$$

After removing the at most g roots in G , each S_i contributes at least $m - g$ incidences. At a coordinate outside G , the affine equation

$$R_0(x) + \gamma R_1(x) = 0$$

has at most one slope solution. Hence

$$|I|(m - g) \leq n - g,$$

which proves the bound.

For $x \in G \setminus H$, the polynomial

$$P_0(x) + \gamma P_1(x)$$

is a nonzero affine function of γ and therefore vanishes at at most one of the shared slopes. The collision identity then forces $\Lambda_i(x) = 0$ for all other i , proving the near-sunflower assertion.

Finally, if the certificates are distinct, then $(P_0, P_1) \neq (0, 0)$ by the first part. Both have degree at most d , so their common root set in D has size at most d . $\square$

Corollary 36.2 (primitive atom uniqueness). *If $G = \emptyset$, two distinct atom certificates share at most*

$$\left\lfloor \frac{n}{m} \right\rfloor = 1$$

explanation at either deployed row. Thus two primitive certificates sharing two slopes are projectively identical.

37. AFFINE LINES VERSUS RATIONAL ATOMS

A similar argument prevents a primitive rational atom from masquerading as a global affine codeword line.

Theorem 37.1 (affine-atom collision). *Suppose $h_i = A_0 + \gamma_i B_0$ for $i \in I$, and the same explanations satisfy a nontrivial atom certificate*

$$Qh_i + (c_0 + c_1 \gamma_i) \Lambda_i = A + \gamma_i B, \quad (c_0, c_1) \neq (0, 0).$$

Put

$$E_0 := A - QA_0, \quad E_1 := B - QB_0,$$

and let

$$G_{\text{aff}} := \{x \in D : E_0(x) = E_1(x) = 0\}, \quad g_{\text{aff}} := |G_{\text{aff}}|.$$

If $|I| \geq 2$, then $(E_0, E_1) \neq (0, 0)$, and if $g_{\text{aff}} < m$,

$$\boxed{|I| \leq \left\lfloor \frac{n - g_{\text{aff}}}{m - g_{\text{aff}}} \right\rfloor}.$$

Every point of G_{aff} belongs to at least $|I| - 1$ supports. In particular, a primitive atom and a global affine line share at most one explanation.

Proof. Substitution gives

$$(c_0 + c_1\gamma_i)\Lambda_i = E_0 + \gamma_i E_1.$$

If $E_0 = E_1 = 0$, then $c_0 + c_1\gamma_i = 0$ for every i , since $\Lambda_i \neq 0$. A nonzero affine function of the slope has at most one root, contradicting $|I| \geq 2$.

The root-incidence count for the pencil $E_0 + \gamma E_1$ is identical to the proof of [Theorem 36.1](#). At a common fixed root, all supports except possibly the unique slope with $c_0 + c_1\gamma = 0$ contain the coordinate. $\square$

38. COHERENCE OR A QUANTIFIED NEAR-SUNFLOWER

The graph whose vertices are the 32-subsets of a fixed slope set and whose edges join subsets with 31 common slopes is connected. Collision rigidity can therefore propagate a local classification globally.

Theorem 38.1 (order-32 coherence). *Let Z be a set of at least 32 bad slopes. Suppose every 32-subset of the corresponding explanations is classified as either*

- (a) *one global affine codeword line, or*
- (b) *one nontrivial scalar-locator rational atom of denominator degree at most $d = m - k$.*

Assume all cross-fixed sets arising between distinct local classifications have already been removed by the first-match rule. Then the entire set Z lies on one global affine codeword line or in one projective rational atom certificate.

Proof. Adjacent 32-subsets share 31 explanations. Two affine codeword lines sharing two distinct slopes are identical. By [Corollary 36.2](#), two primitive atom certificates sharing two explanations are identical. By [Theorem 37.1](#), a primitive atom and an affine line cannot share two explanations. Therefore adjacent vertices of the 32-subset graph receive the same type and the same structure. Connectedness propagates that structure to every 32-subset and hence to every explanation in Z . $\square$

Without the primitive hypothesis, failure of coherence is not unstructured.

Corollary 38.2 (exact near-sunflower alternative on a 31-overlap). *Suppose two distinct local classifications share $q = 31$ explanations. Then at least one classification is an atom, and their cross-fixed set has size at least*

$$g_{31} := \left\lceil \frac{31m - n}{30} \right\rceil.$$

For two rational atoms, at least $g_{31} - d$ of those coordinates belong to at least 30 of the 31 supports. For an affine-line/atom collision, all g_{31} fixed coordinates have that property.

At the deployed rows,

row	g_{31}	$d = m - k$	$g_{31} - d$
KoalaBear MCA	1083345	67472	1015873
Mersenne-31 MCA	1083320	67448	1015872

Thus every failure of local atom coherence produces a 30-of-31 near-sunflower on more than one million domain coordinates.

Proof. Rearrange the collision inequality

$$31(m - g) \leq n - g$$

to obtain

$$g \geq \left\lceil \frac{31m - n}{30} \right\rceil.$$

For two atoms, [Theorem 36.1](#) gives a collision core H of size at most d ; every point outside it lies in at least thirty supports. The affine-atom statement follows from [Theorem 37.1](#). $\square$

For reference, if two classifications share all 32 explanations, the corresponding fixed-root thresholds are

row	$g_{32} = \lceil (32m - n)/31 \rceil$	$g_{32} - d$
KoalaBear MCA	1084400	1016928
Mersenne-31 MCA	1084375	1016927

39. THE PROVED PART OF THE RELATIVE ORDER-32 THEOREM

The results can now be stated as one relative classification theorem for the actual explanation vector, which is weaker than classifying every vector in the rank-two kernel but is sufficient for MCA.

Theorem 39.1 (partial relative order-32 explanation theorem). *Fix either deployed adjacent MCA row and 32 distinct bad slopes. After maximalizing supports, removing their common-support branch, and selecting exact m -subsupports, any non-global-affine explanation outside the pure-locator, denominator-root, and scalar-locator rational-profile cells must satisfy both*

$$18 \leq \deg_Z E \leq 31$$

and

$$\chi(\mathcal{S}) \geq 3m - k + 3.$$

Numerically, the second inequality is

$$\chi(\mathcal{S}) \geq 2299571 \quad \text{on KoalaBear,}$$

and

$$\chi(\mathcal{S}) \geq 2299499 \quad \text{on Mersenne-31.}$$

Moreover, if every residual 32-subset is classified as affine or rational, then either all bad slopes belong to one coherent structure or a 31-overlap enters the quantified near-sunflower branch of [Corollary 38.2](#).

Proof. The slope-degree statement is [Corollary 32.3](#). The incidence-complexity statement is the contrapositive of [Corollary 34.4](#). The coherence alternative follows from [Theorem 38.1](#) and [Corollary 38.2](#). $\square$

This proves the relative theorem in two substantial regimes:

- low slope degree ($r \leq 17$) is impossible outside global-affine/common-support structure;
- low two-cover complexity ($\chi < 3m - k + 3$) automatically produces a rational or pure-locator atom.

It also discharges the separate atom-coherence obligation once the near-sunflower first-match cell is applied.

Part 6. Primitive spread abundance and the exact MCA endgame

The rank and atom theorems reduce a large bad-slope set to one coherent global structure, except for a quantified exceptional set. This part proves that universal spread exclusion is false, identifies the correct spread-abundance target, localizes rational atoms to a canonical owner, and bounds broad classes of spread-core corrections.

40. A CRITICAL SPREAD CORE EXISTS

A universal theorem excluding primitive spread critical tuples on dyadic domains would be stronger than the deployed abundance statement. The following exact example shows that such a universal prohibition is false.

Proposition 40.1 (exact dyadic critical spread core). *Let*

$$\mathbb{F} = \mathbb{F}_{17}, \quad D = \mathbb{F}_{17}^\times = \langle 3 \rangle, \quad n = 16, \quad k = 8, \quad m = 10.$$

The critical order is

$$L_* = \left\lfloor \frac{2(n-k)}{m-k} \right\rfloor + 1 = 9.$$

There exist nine distinct slopes, nine codeword explanations, and exact size-10 witness supports with all of the following properties.

- (i) *The received pair is support-wise MCA-bad at agreement 10.*
- (ii) *The maximal agreement supports have empty common intersection.*
- (iii) *The rank-two reduced intersection matrix has rank 55, whereas full relative rank is $(9-2)8 = 56$.*
- (iv) *The two-cover complexity is 32, above the extraction threshold of [Corollary 34.4](#) $3m - k + 3 = 25$.*
- (v) *The incidence pattern is not periodic at any nontrivial level of the dyadic tower.*
- (vi) *There is no scalar-locator certificate with denominator degree 0, 1, or $2 = m - k$.*
- (vii) *The coefficientwise slope interpolation has degree 8, above the common-support-free lower barrier $\lceil 9m/n \rceil = 6$.*
- (viii) *The direction code contains a word of weight $9 = n - k + 1$, so the minimum-support excess is exactly one.*

Moreover, for this same received line every slope in $\mathbb{F}_{17}$ has a degree- < 8 codeword agreeing on at least ten coordinates.

Exact certificate. Order the domain as

$$(1, 3, 9, 10, 13, 5, 15, 11, 16, 14, 8, 7, 4, 12, 2, 6).$$

Use slopes $\gamma_i = i$ for $0 \leq i \leq 8$ and the following exact size-10 supports, written by domain indices:

i	S_i
0	0, 3, 4, 5, 7, 9, 10, 11, 12, 13
1	1, 3, 5, 6, 7, 10, 11, 12, 13, 14
2	1, 4, 6, 7, 9, 10, 11, 12, 13, 15
3	0, 1, 2, 4, 5, 6, 7, 9, 14, 15
4	0, 2, 3, 4, 5, 6, 7, 9, 12, 13
5	0, 3, 4, 5, 6, 7, 8, 9, 10, 14
6	1, 2, 3, 5, 6, 8, 9, 11, 13, 15
7	0, 3, 4, 5, 6, 7, 11, 12, 14, 15
8	2, 3, 4, 5, 6, 7, 8, 9, 10, 14

The codeword coefficient vectors, constant coefficient first, are

i	1	X	X^2	X^3	X^4	X^5	X^6	X^7
0	1	13	16	16	9	16	7	12
1	13	9	14	12	6	4	2	13
2	2	2	2	4	2	12	0	1
3	13	3	3	4	0	13	16	5
4	1	4	11	4	2	16	9	16
5	16	13	15	13	2	15	5	7
6	3	3	0	13	0	8	12	4
7	3	2	6	8	10	13	10	12
8	1	12	2	14	9	10	9	4

and the received coefficient words on the ordered domain are

$$r_0 = (5, 16, 0, 0, 0, 0, 0, 5, 7, 8, 0, 8, 10, 0, 6),$$

$$r_1 = (6, 14, 15, 0, 0, 0, 2, 13, 14, 12, 6, 3, 15, 13, 6, 15).$$

Direct substitution verifies every claimed support. Maximalizing the nine supports gives sizes 10 or 11 and empty total intersection.

For the rank statement, form the rank-two matrix of [Definition 25.1](#). It has 58 rows and 72 columns. The displayed explanation vector supplies one kernel direction independent of the 16-dimensional global-affine kernel, so its rank is at most 55; exact Gaussian elimination over $\mathbb{F}_{17}$ gives rank 55. The scalar-locator systems at denominator degrees 0, 1, 2 have column counts and ranks

$$(25, 25), \quad (26, 26), \quad (27, 27),$$

respectively. Exact interpolation gives slope degree 8, and direct evaluation of all pairwise codeword differences gives minimum weight 9.

Finally, exhaustive interpolation over the $\binom{16}{10}$ candidate common supports finds no ten-coordinate support on which both r_0 and r_1 are degree- < 8 evaluations. Enumerating the $\binom{16}{8}$ interpolation bases at each of the seventeen slopes finds at least one agreement-10 codeword for every slope. The accompanying verifier repeats all calculations in exact arithmetic. $\square$

Remark 40.2 (what the example does and does not refute). The example refutes the *universal* prohibition of primitive spread critical tuples. It does not refute a row-sharp theorem for the two deployed length- 2^{21} domains. It also shows why the minimum-distance refinement is relevant but not by itself decisive: the hard spread core lies exactly in the support-excess-1 branch, the same branch in which the boundary top seam survives by [Corollary 7.9](#).

41. THE DEPLOYED SPREAD-ABUNDANCE TARGET

Definition 41.1 (first-match primitive spread core). *A set of 32 explanations is a first-match primitive spread core if it satisfies the residual conditions of [Theorem 39.1](#): no global affine codeword line, common-support factor, exact dyadic quotient descent, scalar-locator certificate of denominator degree at most $m - k$, denominator-root, extension, tangent, field-drop, or quantified near-sunflower atom; slope degree lies between 18 and 31; and two-cover complexity is at least $3m - k + 3$.*

A universal nonexistence theorem is false by [Proposition 40.1](#). The finite MCA theorem requires the weaker and more natural abundance statement.

Problem 41.2 (deployed spread-abundance theorem). Intermediate status. *This is the line-level abundance consequence required from the final spread-component routing problem [Problem 54.1](#); it is not an additional terminal input to the completion certificate.*

For either deployed domain, prove that if a support-wise MCA-bad received line contains one first-match primitive spread core, then its total number of bad challenge slopes is at most the row budget B^ .*

A stronger useful form would bound the number of slopes extending a fixed core by a quantity below $B^ - 32$, with every failure routed to one canonical quotient, extension, rational-owner, clone, or near-sunflower certificate.*

Remark 41.3. The distinction is essential. The finite theorem does not need local higher-order-MDS rank at every critical tuple. It needs a global statement that a line carrying one genuine spread defect cannot carry too many further slopes. The exact example in [Proposition 40.1](#) shows that any proof must use the deployed domain and field geometry, not only the abstract critical-order dimension count.

42. RATIONAL ATOMS HAVE A CANONICAL OWNER

Consider one coherent rational atom certificate

$$(42.1) \quad Qh_i + (c_0 + c_1\gamma_i)\Lambda_i = A + \gamma_i B, \quad i \in I,$$

where Λ_i is the locator of a witness support S_i , $|S_i| \geq m$, and Q has no root on D . Put

$$\tilde{r}_0(x) = \frac{A(x)}{Q(x)}, \quad \tilde{r}_1(x) = \frac{B(x)}{Q(x)}.$$

The line

$$\tilde{r}_\gamma = \tilde{r}_0 + \gamma \tilde{r}_1$$

is the *owner line* of the atom.

Definition 42.1 (fixed owner set). *Define*

$$G = \{x \in D : r_0(x) = \tilde{r}_0(x), r_1(x) = \tilde{r}_1(x)\}, \quad g = |G|.$$

Theorem 42.2 (exclusive rational-owner localization). *Retain (42.1) and assume every scalar $c_0 + c_1\gamma_i$ is nonzero; zero-scalar slopes are a separate global-affine degeneracy and contribute at most one slope. Then:*

- (a) *for every $x \in D \setminus G$, at most one support S_i contains x ;*
- (b) *if $g < m$, then*

$$|I| \leq \left\lfloor \frac{n-g}{m-g} \right\rfloor \leq n-m+1;$$

- (c) *if $g \geq m$, at most $n-g$ supports meet $D \setminus G$, while the supports contained in G form a support-wise MCA-bad instance for the punctured code*

$$C|_G \cong \text{RS}[\mathbb{F}, G, k].$$

Consequently

$$|I| \leq n-g + B_{\text{MCA}}(\text{RS}[\mathbb{F}, G, k], m).$$

Proof. On S_i , the locator vanishes, and (42.1) gives

$$h_i(x) = \tilde{r}_0(x) + \gamma_i \tilde{r}_1(x).$$

Since S_i is also an agreement support for the original received line,

$$(r_0 - \tilde{r}_0)(x) + \gamma_i(r_1 - \tilde{r}_1)(x) = 0.$$

For $x \notin G$ this is a nonzero affine equation in γ_i , hence has at most one solution. This proves (a).

If $g < m$, every support uses at least $m-g$ coordinates outside G . By (a), those outside incidences are disjoint across slopes, so

$$|I|(m-g) \leq n-g.$$

The quotient $(n-g)/(m-g)$ is increasing for $g < m$, and at $g = m-1$ equals $n-m+1$. This proves (b).

Now assume $g \geq m$. Distinct supports meeting $D \setminus G$ use disjoint nonempty outside-coordinate sets, so there are at most $n-g$ of them. Every remaining support lies in G and witnesses the punctured owner line there. Since $g \geq m > k$, evaluation on G is injective and the punctured code is Reed–Solomon of dimension k . If the punctured pair admitted a common size- m explaining support, the same support would explain the original pair on D , contradicting support-wise MCA-badness. The contained slopes are therefore bounded by the punctured MCA numerator. $\square$

Corollary 42.3 (all small-owner atoms are paid). *For either deployed row, every coherent rational atom with*

$$g \leq 2m - k$$

has at most n bad slopes. Explicitly, the paid owner ranges are

row	$2m - k$	atom bound
KoalaBear MCA	1183520	2097152
Mersenne-31 MCA	1183472	2097152

Both fit their challenge budgets; on Mersenne-31 the exact residual margin is

$$16777215 - 2097152 = 14680063.$$

Proof. The case $g < m$ is [Theorem 42.2\(b\)](#). For $m \leq g \leq 2m - k$, the punctured error count is $g - m$ and

$$2(g - m) \leq g - k.$$

CAP25 v13.2 proves, by its half-distance pincer and unique-decoding correlated-agreement input, that the punctured MCA numerator is at most g in this range [\[Cho26\]](#). Apply [Theorem 42.2\(c\)](#):

$$|I| \leq (n - g) + g = n.$$

□

43. A SELF-CONTAINED MINIMUM-DISTANCE OWNER PAYMENT

The corollary uses the half-distance proximity theorem of CAP25 v13.2. The next statement is elementary and retains the actual minimum support of the punctured direction code.

Theorem 43.1 (triple-collapse from direction minimum support). *Assume $g \geq m$ and restrict to the atom slopes whose supports lie in G . Let*

$$C'_G = \text{Span}\{h_i - h_j : i, j \in I\} \subseteq \text{RS}[\mathbb{F}, G, k]$$

be their direction code. If it is nonzero, write

$$d_1(C'_G) = g - k + a, \quad a \geq 1.$$

If

$$\boxed{3m - 2g \geq k - a + 1,}$$

then all contained codeword explanations lie on one global codeword line. Hence the full atom, including supports meeting $D \setminus G$, satisfies

$$\boxed{|I| \leq n - m + 1.}$$

Equivalently, this payment applies whenever

$$g \leq \left\lfloor \frac{3m - k + a - 1}{2} \right\rfloor.$$

Proof. For three distinct slopes define the affine deviation

$$p_{ijk} = (\gamma_j - \gamma_k)h_i + (\gamma_k - \gamma_i)h_j + (\gamma_i - \gamma_j)h_k.$$

It belongs to C'_G . On $S_i \cap S_j \cap S_k$, all three codewords equal the owner line at their slopes, so p_{ijk} vanishes. Inclusion-exclusion gives

$$|S_i \cap S_j \cap S_k| \geq 3m - 2g.$$

A nonzero word of C'_G has at most

$$g - d_1(C'_G) = k - a$$

zeros. Under the displayed hypothesis, p_{ijk} must therefore be zero for every triple.

Fix two slopes. The vanishing of every triple deviation shows that every h_i lies on the same global codeword line $H_0 + \gamma_i H_1$. Let $C_0 \subseteq G$ be the coordinate set on which this codeword line equals the owner line coefficientwise. Support-wise MCA-badness gives $|C_0| < m$. Outside C_0 , equality of the two affine lines determines at most one slope. Thus the contained slopes satisfy

$$|I_{\text{in}}|(m - |C_0|) \leq g - |C_0|, \quad |I_{\text{in}}| \leq g - m + 1.$$

There are at most $n - g$ crossing supports by [Theorem 42.2](#), and hence

$$|I| \leq (n - g) + (g - m + 1) = n - m + 1.$$

□

Corollary 43.2 (exact support-excess-one thresholds). *At $a = 1$, [Theorem 43.1](#) pays owners through*

row	$\lfloor (3m - k)/2 \rfloor$
<i>KoalaBear MCA</i>	1149784
<i>Mersenne-31 MCA</i>	1149748

Larger direction support excess increases the threshold by one owner coordinate for every two units of excess.

Remark 43.3 (connection with the minimum-distance sharpening of [Section 7](#)). For the boundary prefix dictionary, [Proposition 7.8](#) proves that a shift pair of exchange defect r can occur only when $r \geq a$, where a is the direction-code support excess; the top seam $r = 1$ is therefore confined to $a = 1$. The counterexample in [Proposition 40.1](#) also has $a = 1$. Thus the difficult seam is now consistently localized: large a is paid by the sharpened flag and owner bounds, while the exact-minimum-distance branch remains the genuine image-counting problem.

44. OWNER PENCILS AND A SEVEN-POINT RICH-INCIDENCE BOUND

The exclusive atom theorem also requires control of competing owner certificates. [Theorem 36.1](#) gives uniqueness once two slopes are shared. The following theorem gives a complementary geometric bound when owners vary in a projective pencil.

Let

$$(Q_\tau, A_\tau, B_\tau) = (Q_0, A_0, B_0) + \tau(Q_1, A_1, B_1), \quad \tau \in \mathbb{P}^1(\mathbb{F}),$$

be a projective pencil of owner triples. At a coordinate $x \in D$, set

$$F_x(\gamma, \tau) = A_\tau(x) + \gamma B_\tau(x) - Q_\tau(x)(r_0(x) + \gamma r_1(x)).$$

This has bidegree at most $(1, 1)$ on $\mathbb{P}_\gamma^1 \times \mathbb{P}_\tau^1$.

Theorem 44.1 (component-free owner-pencil rich points). *Assume no F_x is identically zero and no two coordinate curves $F_x = 0$ share an irreducible component. Then the number of parameter pairs (γ, τ) for which at least m coordinate equations vanish is at most*

$$\left\lfloor \frac{2\binom{n}{2}}{\binom{m}{2}} \right\rfloor = \left\lfloor \frac{2n(n-1)}{m(m-1)} \right\rfloor.$$

For both deployed MCA rows this number is exactly 7.

Proof. Two bidegree- $(1, 1)$ curves with no common component have intersection number at most

$$(1, 1) \cdot (1, 1) = 2$$

on $\mathbb{P}^1 \times \mathbb{P}^1$. A parameter pair lying on at least m coordinate curves contributes at least $\binom{m}{2}$ incidences with unordered coordinate-curve pairs. Summing first over parameter pairs and then over the $\binom{n}{2}$ curve pairs gives

$$N_{\text{rich}} \binom{m}{2} \leq 2 \binom{n}{2}.$$

The deployed integer evaluations are direct. □

Remark 44.2 (the shared-component branches are explicit). A nonzero bidegree- $(1, 1)$ polynomial is either irreducible or factors into one vertical and one horizontal linear factor. Hence a shared component is one of:

- (i) a fixed-slope vertical component;
- (ii) a fixed-owner horizontal component;
- (iii) a common irreducible $(1, 1)$ component, in which case the two coordinate equations are proportional and form a coordinate-clone cell.

Identically zero F_x are fixed owner coordinates. Thus every failure of [Theorem 44.1](#) is already a named structural branch rather than an unpriced multiplicity.

45. THE EXACT LARGE-OWNER ATOM TARGET

After [Corollary 42.3](#), only

$$g \geq 2m - k + 1$$

remains for a coherent rational atom. If at most 31 exceptional slopes are reserved for the local-to-global first-match transition, [Theorem 42.2](#) shows that the required bound for the owner-contained slopes is

$$(45.1) \quad B_{\text{owner}}(g) \leq B^* - 31 - (n - g).$$

At the first unproved owner size and at $g = n$, the targets are

row	$g_{\min} = 2m - k + 1$	target at $g_{\min}$	target at $g = n$
KoalaBear MCA	1183521	274980728110481425	274980728111395056
Mersenne-31 MCA	1183473	15863505	16777184

The Mersenne-31 row is binding.

For the full-owner boundary-prefix atom, use the CAP25 convention $K = k + 1$ and $w = m - K$. The average prefix-fiber sizes are

$$\bar{F} = \binom{n}{m} |\mathbb{B}|^{-w}.$$

Their high-precision values and the full-owner multiplicative targets are

row	w	$\log_2 \bar{F}$	$(B^* - 31)/\bar{F}$
KoalaBear MCA	67471	35.7352464170	4807520.929566
Mersenne-31 MCA	67447	20.7411470346	9.57219783037

Thus the exact remaining Mersenne theorem is a factor-9.5722 maximum-image or maximum-fiber statement before any additional cell consumes budget. Owner localization has removed every smaller owner and all multi-owner component-free pencils, but it does not prove this full-owner prefix bound.

Problem 45.1 (exclusive large-owner atom image bound). *For every first-match coherent rational atom on either deployed domain with $g > 2m - k$, prove (45.1). In the full-owner simple-pole branch it suffices to prove the row-sharp challenge-slope image bound*

$$|\{U_z(\alpha) - \Lambda_M(\alpha) : M \in \text{Fib}_w(z)\} \cap \Gamma| \leq B^* - 31.$$

For Mersenne-31, the stronger support-fiber estimate need only lose a factor below 9.57219783037 from the exact average.

46. A CONDITIONAL EXACT ADJACENT-ROW THEOREM

We can now state the final certificate without reverting to a sum of many unrelated local budgets.

Definition 46.1 (exception routing input). *Say that exception routing holds if, for every bad received line with more than 31 slopes, there is a set E of at most 31 slopes such that every remaining local order-32 classification is primitive and [Theorem 38.1](#) applies; every denominator-root, extension, coordinate-clone, imperfect-quotient, and quantified near-sunflower failure is either included in E or assigned to the spread-abundance or large-owner input with the same received-line owner retained.*

Theorem 46.2 (exact conditional adjacent-row closure). *Fix one of the two deployed MCA rows. Assume the following three inputs.*

- (S): The deployed spread-abundance theorem, [Problem 41.2](#).
- (A): The exclusive large-owner atom image bound, [Problem 45.1](#).
- (E): Exception routing in the sense of [Definition 46.1](#).

Then

$$\boxed{B_{\text{MCA}}(C, m) \leq B^*}.$$

Together with the unconditional CAP25 v13.2 witness at agreement $m - 1$, this proves that the row's exact adjacent MCA threshold is m .

Proof. Let Z be the set of bad challenge slopes. If $|Z| \leq 31$, the claim is immediate. If Z contains a primitive spread core, assumption (S) pays the whole line.

Otherwise remove the exceptional set E supplied by exception routing. Every 32-subset of $Z \setminus E$ is classified by [Theorem 39.1](#) as a global affine codeword line or one primitive rational atom. The overlap graph of 32-subsets is connected, and [Theorem 36.1](#) forces one coherent global structure.

If that structure is a global affine codeword line, the standard coordinate-injection argument gives at most $n - m + 1$ slopes. Adding $|E| \leq 31$ remains far below either budget.

Suppose instead that it is one rational atom with owner-set size g . For $g < m$, [Theorem 42.2](#) gives at most $n - m + 1$ atom slopes. For $m \leq g \leq 2m - k$, [Corollary 42.3](#) gives at most n . Again the additional 31 slopes fit both budgets. Finally, for $g > 2m - k$, there are at most $n - g$ crossing slopes and assumption (A), namely [\(45.1\)](#), pays the contained slopes; adding the reserved 31 gives exactly at most B^* .

CAP25 v13.2 already proves an unsafe witness at $m - 1$, so safety at m identifies the adjacent threshold. $\square$

47. CORE INTERPOLANT AND CORRECTION SETUP

Let

$$C = \text{RS}[\mathbb{F}, D, k], \quad |D| = n,$$

and let

$$r_\gamma = r_0 + \gamma r_1 \in \mathbb{F}^D$$

be support-wise MCA-bad at agreement $m > k$. Put

$$d = m - k, \quad R = n - k, \quad t = n - m.$$

Fix one first-match primitive spread core in the sense of [Problem 41.2](#)

$$(\gamma_i, h_i, S_i)_{i=1}^{32},$$

where the slopes are distinct, $h_i \in \mathbb{F}[X]_{<k}$, and S_i is the *maximal* agreement support of h_i with r_{γ_i} . The common-support branch has already been removed, so

$$(47.1) \quad \bigcap_{i=1}^{32} S_i = \emptyset.$$

There is a unique coefficientwise interpolant

$$(47.2) \quad H(X, Z) = \sum_{j=0}^{31} H_j(X) Z^j, \quad H(X, \gamma_i) = h_i(X), \quad \deg H_j < k.$$

For $x \in D$ define

$$(47.3) \quad E_x(Z) = H(x, Z) - r_0(x) - Z r_1(x)$$

and

$$q_x = |\{i : x \in S_i\}|.$$

Every core agreement contributes a root of E_x , so

$$(47.4) \quad \sum_{x \in D} q_x = \sum_{i=1}^{32} |S_i| \geq 32m.$$

Moreover, [\(47.1\)](#) and maximality imply

$$(47.5) \quad E_x \neq 0 \quad \text{for every } x \in D.$$

Indeed, $E_x \equiv 0$ would put x in every maximal core support.

The deployed values are

row	m	d	$n - m$	B^*
KoalaBear MCA	1116048	67472	981104	274980728111395087
Mersenne-31 MCA	1116024	67448	981128	16777215

with $n = 2^{21}$ and $k = 2^{20}$.

48. ONLY FIFTY-EIGHT SLOPES CAN USE THE CORE INTERPOLANT

Theorem 48.1 (core-compatible abundance bound). *Let Z_0 be the set of slopes $\gamma \in \mathbb{F}$ for which the codeword $H(X, \gamma)$ agrees with r_γ on at least m coordinates. Then*

$$(48.1) \quad |Z_0| \leq \left\lfloor \frac{31n}{m} \right\rfloor = 58.$$

More precisely, outside the 32 core slopes there are at most

$$(48.2) \quad \left\lfloor \frac{31n - 32m}{m} \right\rfloor = 26$$

additional core-compatible slopes on either deployed row.

Proof. By (47.5), every E_x has at most 31 roots. A slope in Z_0 contributes at least m coordinate roots, hence

$$|Z_0|m \leq \sum_{x \in D} \deg E_x \leq 31n,$$

which proves (48.1).

For a non-core slope, every agreement coordinate consumes a root distinct from the q_x core roots already present at x . The remaining root capacity is at most

$$\sum_x (31 - q_x) \leq 31n - 32m$$

by (47.4). Dividing by m gives (48.2). $\square$

Remark 48.2. The theorem uses maximal supports only to exclude identically zero coordinate error polynomials. It is independent of quotient, rational-owner, or minimum-distance estimates. Thus any over-budget spread line must obtain essentially all of its extra slopes from nonzero codeword corrections to H .

49. AN UNCONDITIONAL THEOREM FOR ONE COMPLETE CORRECTION RAY

Fix a nonzero $P \in \mathbb{F}[X]_{<k}$. Consider all parameter pairs (γ, c) for which

$$(49.1) \quad h_{\gamma,c}(X) = H(X, \gamma) + cP(X)$$

agrees with r_γ on at least m coordinates. At $x \in D$ the agreement equation is

$$(49.2) \quad E_x(\gamma) + cP(x) = 0.$$

No component-freeness assumption will be imposed.

For $P(x) \neq 0$, put

$$f_x(Z) = \frac{E_x(Z)}{P(x)} \in \mathbb{F}[Z]_{\leq 31}.$$

Coordinates with the same f_x form a *clone class*. Coordinates with $P(x) = 0$ form the vertical class.

Lemma 49.1 (clone-class dichotomy). *Let $C_f = \{x : P(x) \neq 0, f_x = f\}$. Then exactly one of the following holds.*

- (i) $|C_f| \leq k - 1$;

(ii) *the codeword curve*

$$G_f(X, Z) = H(X, Z) - f(Z)P(X)$$

has slope degree at most one. In this case the graph $c = -f(\gamma)$ is a global affine codeword line.

At most two clone classes of type (ii) can have size at least k .

Proof. For every $x \in C_f$,

$$G_f(x, Z) = r_0(x) + Zr_1(x)$$

as a polynomial identity in Z . If G_f has a nonzero coefficient of Z^j for some $j \geq 2$, that coefficient is a nonzero degree- $< k$ polynomial vanishing on C_f , so $|C_f| \leq k - 1$. Otherwise $G_f = A_f + ZB_f$ is a global affine codeword line.

Clone classes are disjoint subsets of D . A class of type (ii) not covered by (i) has at least k coordinates. Since $n = 2k$, there can be at most two. $\square$

Lemma 49.2 (global affine-line block). *For any fixed global codeword line $A + \gamma B$, the number of slopes for which it agrees with r_γ on at least m coordinates is at most*

$$(49.3) \quad n - m + 1.$$

Proof. Let C_0 be the coordinate set on which $A = r_0$ and $B = r_1$. Support-wise MCA-badness gives $|C_0| < m$. Outside C_0 , agreement of the two affine lines determines at most one slope. Hence

$$N(m - |C_0|) \leq n - |C_0|,$$

and the right quotient is maximized at $|C_0| = m - 1$. $\square$

Theorem 49.3 (complete correction-ray bound). *For either deployed row, the number of rich parameter pairs (γ, c) satisfying (49.1) is at most*

$$(49.4) \quad 2(n - m + 1) + \left\lfloor \frac{31 \binom{n}{2}}{(k - 1)(m - k + 1)} \right\rfloor.$$

The final fraction equals 963 on both rows. Thus one full correction ray contributes at most

row	ray bound	$B^* - \text{ray bound}$
KoalaBear MCA	1963173	274980728109431914
Mersenne-31 MCA	1963221	14813994

In particular, a single correction ray can never be the residual spread-abundance obstruction.

Proof. Remove all rich points lying on a type-(ii) clone graph whose class has at least k coordinates. By Lemmas 49.1 and 49.2, there are at most two such graphs and each contributes at most $n - m + 1$ slopes.

Consider a remaining rich point and choose exactly m of its agreement coordinates. Partition them into the vertical class and the clone classes meeting the point. The vertical class has size at most $k - 1$, because a nonzero degree- $< k$ polynomial has at most $k - 1$ zeros. Every clone class present at the point has size at most $k - 1$: type-(i) classes have this bound, while every larger type-(ii) graph was removed.

Since $m < 2(k - 1)$ on both rows, a partition of m elements into parts of size at most $k - 1$ has at least

$$(49.5) \quad (k - 1)(m - k + 1)$$

heterogeneous unordered pairs. This is the minimum obtained from the two-part partition $(k - 1, m - k + 1)$.

A heterogeneous pair contributes to at most 31 rich parameter points. For two graph coordinates from different clone classes, the condition is

$$f_x(\gamma) = f_y(\gamma),$$

a nonzero polynomial equation of degree at most 31, after which c is determined. For one vertical and one graph coordinate, $E_x(\gamma) = 0$ has at most 31 roots and again the graph equation determines c . We never count vertical-vertical pairs or pairs inside one clone class.

Double counting rich-point/heterogeneous-pair incidences gives

$$N_{\text{res}}(k-1)(m-k+1) \leq 31 \binom{n}{2}.$$

Adding the at most two affine-line blocks proves (49.4). Counting parameter pairs only overestimates the number of distinct slopes. $\square$

Remark 49.4 (minimum-support refinement). Suppose every nonzero codeword in the span of P and the high slope coefficients $H_2, \dots, H_{31}$ has support at least $R+a$. Then $k-1$ in the proof can be replaced by $k-a$. Writing $q = k-a$, the exact cross-pair denominator is

$$\Xi(m, q) = \binom{m}{2} - u \binom{q}{2} - \binom{r}{2}, \quad m = uq + r, \quad 0 \leq r < q,$$

and there are at most $\lfloor n/(q+1) \rfloor$ large affine clone classes. In the active residual wedge $m \leq 2q$, so $\Xi(m, q) = q(m-q)$. The exact-minimum-distance case $a = 1$ is the worst case in the range left by the minimum-distance sharpening of Section 7; the same section also proves that the boundary top seam occurs only at $a = 1$.

50. HIGHER-DIMENSIONAL CORRECTION SPACES

Let $W \leq C$ have dimension s , and consider all rich points

$$(\gamma, P) \in \mathbb{F} \times W$$

for which $H(X, \gamma) + P(X)$ agrees with r_γ on at least m coordinates. Choose a basis $P_1, \dots, P_s$ of W and write $P = \sum c_j P_j$. The coordinate equation

$$(50.1) \quad F_x(Z, \mathbf{c}) = E_x(Z) + \sum_{j=1}^s c_j P_j(x) = 0$$

has bidegree at most $(31, 1)$ on $\mathbb{P}_Z^1 \times \mathbb{P}_{\mathbf{c}}^s$.

50.1. The proper-intersection compiler.

Definition 50.1 (proper correction space). *The correction space W is proper for the core if every $s+1$ coordinate hypersurfaces from (50.1) have empty or zero-dimensional common intersection in $\mathbb{P}^1 \times \mathbb{P}^s$.*

Theorem 50.2 (proper-intersection correction compiler). *If W is proper for the core, then the number N_W of rich parameter points satisfies*

$$(50.2) \quad N_W \leq \left\lfloor 31(s+1) \frac{\binom{n}{s+1}}{\binom{m}{s+1}} \right\rfloor.$$

At the deployed rows this pays every proper correction space through

row	$s \leq$	bound at final s	bound at $s+1$
KoalaBear MCA	50	148068539552473273	283695391248936790
Mersenne-31 MCA	15	11989622	23938324

The next values only mark the limit of this compiler.

Proof. The divisor class of each projective hypersurface is at most

$$31H_Z + H_W.$$

In the Chow ring of $\mathbb{P}^1 \times \mathbb{P}^s$,

$$(31H_Z + H_W)^{s+1} = 31(s+1)H_Z H_W^s,$$

because $H_Z^2 = 0$ and $H_W^{s+1} = 0$. Hence every proper intersection of $s+1$ coordinate hypersurfaces has degree at most $31(s+1)$.

A rich point lies on at least m hypersurfaces and therefore contributes at least $\binom{m}{s+1}$ incidences with coordinate $(s+1)$ -subsets. Summing first over rich points and then over the $\binom{n}{s+1}$ coordinate subsets proves (50.2). $\square$

Lemma 50.3 (what a nonproper tuple means). *Let T be a set of $s+1$ coordinates.*

- (i) *If the evaluation functionals of T on W have rank below s , then a nonzero word of W vanishes on a subset spanning the same evaluation flat.*
- (ii) *If their rank is s and the common intersection is positive-dimensional, choose any evaluation basis $B \subset T$. There is a unique polynomial curve $P_B(Z) \in W \otimes \mathbb{F}[Z]_{\leq 31}$ solving the equations on B , and every equation in T vanishes identically on that curve.*

Thus every failure of properness is either an evaluation rank-flat or an exact polynomial clone component.

Proof. Part (i) is the annihilator statement for the column matroid of the evaluation map. For (ii), the constant evaluation matrix on B is invertible, so solving its equations gives $P_B(Z)$ coefficientwise and with degree at most 31. Substitution into the remaining equation gives a univariate polynomial. A positive-dimensional intersection forces that polynomial to vanish identically. $\square$

50.2. A clone-tolerant basis-curve compiler. Let

$$V = \text{Span}(W, H_2, \dots, H_{31}) \leq C$$

and write

$$(50.3) \quad d_1(V) = R + a, \quad a \geq 1.$$

We say that W *absorbs the high core* if $H_j \in W$ for every $j \geq 2$.

Theorem 50.4 (clone-tolerant basis-curve bound). *Assume $s \geq 1$ and that W does not absorb the high core. Put*

$$(50.4) \quad M_a = \left\lfloor \frac{31(R+a)}{d+a} \right\rfloor.$$

Then

$$(50.5) \quad \boxed{N_W \leq \left\lfloor M_a \frac{n^s}{(d+a)^s} \right\rfloor}.$$

No component-freeness assumption is required.

Proof. Fix a rich point and an exact size- m agreement support S . For every j -dimensional subcode $U \leq W$, strict increase of generalized Hamming weights and (50.3) give

$$|\text{supp}_D(U)| \geq R + a + j - 1.$$

At most $t = n - m$ of these support coordinates lie outside S , so the restricted code $W|_S$ has

$$d_j(W|_S) \geq d + a + j - 1.$$

Greedily choosing independent evaluation columns therefore gives at least

$$(50.6) \quad \frac{(d+a)^s}{s!}$$

unordered evaluation bases $B \subset S$.

For each such B , solve the basis equations to obtain a unique polynomial correction curve

$$P_B(Z) \in W \otimes \mathbb{F}[Z]_{\leq 31}.$$

Let C_B be the coordinates whose equations vanish identically after this substitution, and put $c_B = |C_B|$. Since W does not absorb all high core coefficients, the codeword curve $H + P_B$ has

a nonzero coefficient of Z^j for some $j \geq 2$. That coefficient lies in V , so it has at most $k - a$ zeros. Hence

$$(50.7) \quad c_B \leq k - a.$$

Every rich point on the curve $P = P_B(\gamma)$ agrees automatically on C_B and must obtain at least $m - c_B$ further agreements outside C_B . Each outside coordinate supplies a nonzero polynomial equation of degree at most 31. Root-capacity counting on this one curve gives

$$N(B)(m - c_B) \leq 31(n - c_B).$$

The quotient $(n - c)/(m - c)$ is increasing for $c < m$, and (50.7) yields

$$N(B) \leq \left\lfloor \frac{31(R + a)}{d + a} \right\rfloor = M_a.$$

Each rich point is incident to at least (50.6) basis curves. Globally there are at most $\binom{n}{s}$ evaluation bases, and each basis curve contains at most M_a rich points. Therefore

$$N_W \frac{(d + a)^{\bar{s}}}{s!} \leq M_a \binom{n}{s} = M_a \frac{n^{\bar{s}}}{s!},$$

which is (50.5). □

Corollary 50.5 (exact deployed transitions). *At exact minimum support $a = 1$, Theorem 50.4 pays through*

row	$s \leq$	bound at final s	bound at $s + 1$
KoalaBear MCA	9	13013823503882165	404431535289439486
Mersenne-31 MCA	3	14457337	449492838

The actual minimum-support excess enlarges the range. The first exact paid transitions beyond those dimensions are

row	s	a_{paid}	bound at a_{paid}	bound at $a_{\text{paid}} - 1$
KoalaBear MCA	10	2434	274948660791683460	274987992550612247
	11	25508	274953206296382843	274985734988552690
	12	50875	274977972290567960	275005854396152321
Mersenne-31 MCA	4	64258	16777099	16777608

All comparisons are exact integer comparisons with the row budgets.

Remark 50.6 (relation to the minimum-distance wedge). Proposition 7.3 proves that the first-unpaid charts already satisfy $a \leq 70831$ on KoalaBear and $a \leq 111528$ on Mersenne-31 at zero outside mismatch, with positive mismatch shrinking the residual further. It also confines the boundary top seam to $a = 1$. Thus Corollary 50.5 is not merely an abstract monotonicity statement: it removes correction dimensions 10–12 on substantial parts of the KoalaBear wedge and dimension 4 on a substantial part of the Mersenne wedge, while the genuine top-seam endpoint remains the exact-minimum-distance case.

51. A RANK-ELEVEN DENSE-LOCATOR COMPONENT ROUTE

We record a local theorem packet for the post-near rank-eleven route. Its input is the following pinned setup, supplied by the public prize-DAG source listed in the companion certificate. Let

$$R = 1048576, \quad d = 67472, \quad n' = R + K', \quad m' = d + K', \quad 10 \leq K' \leq 1048576.$$

There is a fixed ten-dimensional correction space V' , a degree-at-most 31 interpolant $H(X, Z)$, and one affine pair (a'_0, b'_0) . The deviation

$$D_H(X, Z) = H(X, Z) - a'_0(X) - Zb'_0(X)$$

vanishes at eighteen distinct dense-pair anchor slopes, while ten other values of D_H span V' . Every coefficient of D_H lies in V' . The post-near family has distinct slopes and, if the original line is unsafe, contains at least

$$N_{\min} = 274980728111260126$$

non-dense records after the exact 134944 near charge and the eighteen dense anchors. These hypotheses are the scope of this section; no statement here silently promotes an arbitrary rank-eleven family into this setup.

Proposition 51.1 (Dense-root high-coefficient saturation). *Let*

$$q(Z) = \prod_{i=1}^{18} (Z - \gamma_i)$$

be the monic locator of the dense anchor slopes. Then

$$D_H = qG, \quad \deg_Z G \leq 13, \quad \text{Span}\{H_j : j \geq 2\} = V'.$$

Consequently, if a correction space $W \leq V'$ absorbs every coefficient H_j with $j \geq 2$, then $W = V'$ and $\dim W = 10$.

Proof. The eighteen vector-valued zeroes give the coefficientwise divisibility $q \mid D_H$. Write $q = Z^{18} + q_{17}Z^{17} + \dots + q_0$ and $G = \sum_{j=0}^{13} G_j Z^j$. Descending from degree 31 to degree 18, the coefficients of D_H recover $G_{13}, G_{12}, \dots, G_0$ by a triangular system with diagonal one. Hence

$$\text{Span}\{(D_H)_{18}, \dots, (D_H)_{31}\} = \text{Span}\{G_0, \dots, G_{13}\}.$$

The ten nonzero anchor values span V' , while every coefficient of D_H lies in V' , so both displayed spans equal V' . Subtracting the affine pair changes only coefficients zero and one, hence $(D_H)_j = H_j$ for $j \geq 2$. If W absorbs those coefficients, then $V' \leq W \leq V'$. $\square$

For a non-dense record of slope γ , put

$$R_\gamma = \frac{h'_\gamma - a'_0 - \gamma b'_0}{q(\gamma)} \in V'.$$

This is well-defined because the record slopes are distinct from the roots of q . At a coordinate x , rich agreement is the equation

$$(51.1) \quad (a'_0 - r'_0)(x) + Z(b'_0 - r'_1)(x) + q(Z)R(x) = 0.$$

Proposition 51.2 (Dense-locator component-incidence bound). *For every eleven-coordinate tuple, the isolated solutions of (51.1) in $\mathbf{P}_Z^1 \times \mathbf{P}_R^{10}$, counted with local multiplicity, have total multiplicity at most 198. Uniformly in K' ,*

$$A_{\text{iso}}(K') = \left\lceil 198 \frac{\binom{n'}{11}}{\binom{m'}{11}} \right\rceil \leq 2526815879272440.$$

Thus at least 990810934 parts per billion of all pairs consisting of a non-dense record and an eleven-subset of its support lie on a positive-dimensional component through that record.

At such a point, the component is either a full-evaluation-rank affine-owner component or a rank-deficient kernel component. At least one of these two lanes carries 495405467 parts per billion of all incidences.

Proof. After bihomogenization, each coordinate equation has divisor class at most $18H_Z + H_R$. Generic independent perturbations in these basepoint-free classes make eleven equations intersect properly. Every isolated point of the original intersection persists with at least its local intersection multiplicity, while the total generic intersection number is

$$(18H_Z + H_R)^{11} = 18 \cdot 11H_Z H_R^{10} = 198H_Z H_R^{10}.$$

Summing this bound over all eleven-coordinate tuples and dividing by the $\binom{m'}{11}$ tuples contributed by each record gives $A_{\text{iso}}(K')$. The binomial ratio is the product, for $0 \leq i \leq 10$, of

$$\frac{1048576 + K' - i}{67472 + K' - i}.$$

Every factor decreases with K' , so the maximum occurs at $K' = 10$ and the displayed integer follows by exact evaluation. Dividing it by $N_{\min}$ gives an isolated-incidence ceiling of 9189066 parts per billion, leaving the stated component floor.

For the classification, inspect evaluation of V' on the eleven-tuple. At rank ten, ten independent coordinate equations determine one rational correction curve; an extra equation vanishes identically precisely on an affine-owner component. At rank at most nine, the compatible fiber contains a positive-dimensional evaluation kernel. These alternatives are exhaustive, and halving the total component floor gives the one-lane floor. $\square$

Corollary 51.3 (Component-star owner-pencil trichotomy). *At least 540546700 parts per billion of the non-dense records have at least 98 percent of their eleven-subsets on a positive-dimensional component. In particular, there are at least*

$$148639925144138894$$

such records. For each one there is a ten-subset B of its support with at least

$$E(K') = \left\lceil \frac{98(m' - 10)}{100} \right\rceil$$

component extensions, and exactly one of the following holds:

- (i) $\text{rank}(\text{ev}_B) = 10$, and one affine owner has support deficiency at most 22320;
- (ii) $\text{rank}(\text{ev}_B) = 9$, and at least 45153 full-rank extensions lie on one owner pencil with direction $\beta(-\gamma u, u)$, where $\ker(\text{ev}_B) = \mathbb{F}u$;
- (iii) $\text{rank}(\text{ev}_B) \leq 8$, and the evaluation kernel has dimension at least two.

Proof. If α is the fraction of records above the 98-percent threshold, then Proposition 51.2 gives

$$\frac{990810934}{10^9} \leq \alpha + (1 - \alpha) \frac{98}{100},$$

so $\alpha \geq 540546700/10^9$. Multiplication by $N_{\min}$ and integer rounding gives the absolute floor. Double-counting pairs (B, x) inside one support gives a ten-subset with $E(K')$ extensions.

At full rank those equations determine one owner. At rank nine, a kernel word u already has ten roots and has at most $K' - 11$ further roots, leaving at least

$$E(K') - (K' - 11) \geq 45153$$

full-rank extensions. Any two corresponding owners differ by $\beta(-\gamma u, u)$. Rank at most eight gives kernel dimension at least two by rank-nullity. The owner deficiency in the full-rank case is at most $m' - 10 - E(K') \leq 22320$. $\square$

Proposition 51.4 (Record-intrinsic full-rank star owner). *In the full-rank branch of Corollary 51.3, the owner of support deficiency at most 22320 is unique for its record. In particular it is independent of the ten-subset used to expose it.*

Proof. Suppose two owner pairs have within-support cores A_p, A_q , each of size at least $m' - 22320$. Inclusion-exclusion inside the exact support gives

$$|A_p \cap A_q| \geq m' - 2 \cdot 22320 = K' + 22832.$$

If the pairs are distinct, at least one component of their difference is a nonzero RS polynomial of degree below K' . Every coordinate in $A_p \cap A_q$ is a root of that component, so the RS root bound gives $|A_p \cap A_q| \leq K' - 1$, a contradiction. The gap is uniformly 22833. $\square$

Proposition 51.5 (One fixed rank-nine split-pencil cell). *Fix a rank-nine ten-subset B and let g records be routed to this cell. All owners agreeing with the received pair on B form an*

affine plane Π_B with directions $(u, 0)$ and $(0, u)$. Each record slope γ defines a line of direction $(-\gamma u, u)$ in Π_B . If t_p is the number of record lines through owner point p , then

$$\sum_p \binom{t_p}{2} = \binom{g}{2}.$$

Moreover, if J_B is the common owner-plane core on $Z(u)$ and $C_p = J_B \dot{\cup} P_p$, then the petals P_p are pairwise disjoint and

$$45153g \leq \sum_p t_p |P_p| \leq 981105(2097152 - 10) = 2057516501910.$$

Consequently

$$g \leq \lfloor 2057516501910 / 45153 \rfloor = 45567658.$$

Proof. All agreeing owner pairs differ by $(\alpha u, \beta u)$. Ownership at slope γ cuts out one affine line with direction $(-\gamma u, u)$, and distinct slopes give distinct directions. Every pair of record lines has one intersection point, proving the block-design identity.

On $Z(u)$ every point of Π_B has the same pair value. Away from $Z(u)$, evaluation identifies the parameter plane bijectively with the pair-value plane, so at most one owner agrees at a coordinate. This proves the common core and disjoint-petal decomposition. Every record supplies at least 45153 full-rank extension incidences in petals on its line.

For a fixed owner with core size $c_p < m'$, the outside-core agreement sets of its distinct slopes are disjoint. Hence

$$t_p(m' - c_p) \leq n' - c_p, \quad t_p \leq n' - m' + 1 = 981105.$$

Finally $|J_B| \geq 10$, so $\sum_p |P_p| \leq 2097152 - 10$, giving the displayed weighted cap and its integer floor quotient by 45153. $\square$

Remark 51.6 (Integer rounding repair). The public source node recorded the valid weaker bound 45567659 by taking a ceiling. Since g is an integer and $45153g \leq 2057516501910$, the sharp consequence is the floor 45567658 printed above. Both values and the rounding direction are retained in the companion certificate.

Proposition 51.7 (Split-pencil pair-core dichotomy). *Lift one fixed residual rank-nine owner plane to the original KoalaBear row. Let J be the coordinates where every owner in the plane equals the received pair, let $j = |J| \geq 10$, and let g be the number of record lines in the plane. Then either*

$$g \leq 1434405,$$

or the entire owner plane has a common received-pair core of size

$$j \geq 134944.$$

Proof. Write C_p for the pair core of owner p and $P_p = C_p \setminus J$. Proposition 51.5 gives pairwise disjoint petals P_p and $\sum_p \binom{t_p}{2} = \binom{g}{2}$.

If two distinct record lines meet at p , their selected original supports, each of size $m = 1116048$, intersect in at least

$$2m - n = 2(1116048) - 2097152 = 134944$$

coordinates. The two distinct slope equalities recover both received columns there, so the intersection lies in C_p .

Assume $j \leq 134943$. Put $D = n - m = 981104$, $x = m - |C_p|$, and $y = t_p - 1$ at a point with $t_p \geq 2$. Pair noncontainment gives $x \geq 1$, while fixed-owner exception disjointness gives $yx \leq D$. Since $m - j \geq D + 1$ and

$$D + 1 - x - y = (D - yx) + (y - 1)(x - 1) \geq 0,$$

we have $|P_p| = |C_p| - j \geq t_p - 1$ and $t_p \leq D + 1$. Consequently

$$\begin{aligned} g(g-1) &= \sum_p t_p(t_p-1) \\ &\leq 981105 \sum_p |P_p| \\ &\leq 981105(2097152 - 10) = 2057516501910. \end{aligned}$$

Now

$$1434405 \cdot 1434404 \leq 2057516501910 < 1434406 \cdot 1434405,$$

so $g \leq 1434405$. Thus any larger plane lies in the alternative $j \geq 134944$. $\square$

Proposition 51.8 (Fixed nine-subset component target). *One of the affine-owner and rank-deficient component lanes admits a fixed nine-coordinate subset B carrying at least*

$$2578110$$

distinct records in that same lane. Moreover, one of the following holds.

- (N1) *every retained record has a rank-deficient component extension whose kernel lies in the same nonzero $\ker(\text{ev}_B)$;*
- (N2) *$\text{rank}(\text{ev}_B) = 9$, and after lifting to the original row the resulting owner plane has a common received-pair core of size at least 134944;*
- (N3) *$\text{rank}(\text{ev}_B) = 8$, the retained owners lie in one affine U^2 flat with $\dim U = 2$, and the retained errors have affine rank at most three.*

Proof. Choose the heavier component lane. It carries at least $\lambda = 495405467/10^9$ of all record/eleven-subset incidences. Mark all $\binom{11}{9} = 55$ nine-subsets of every such eleven-subset. A fixed record/nine-subset pair is marked at most $\binom{m'-9}{2}$ times, while

$$55 \binom{m'}{11} = \binom{m'}{9} \binom{m'-9}{2}.$$

Averaging over the domain nine-subsets gives one B carrying at least

$$\left\lceil \lambda 274980728111260126 \frac{\binom{m'}{9}}{\binom{m'}{9}} \right\rceil \geq 2578110.$$

The nine factors in the binomial ratio increase with K' , and exact evaluation at $K' = 10$ gives the displayed endpoint.

In the rank-deficient lane, restriction from an extending eleven-subset puts every component kernel inside the fixed nonzero $\ker(\text{ev}_B)$, proving (N1). In the affine-owner lane, an evaluation-rank-ten eleven-subset loses at most two ranks on restriction to B , so the rank on B is eight or nine.

At rank nine the owners form one affine plane. The nine-coordinate version of Proposition 51.7 has ordered-pair resource

$$981105(2097152 - 9) = 2057517483015 < 1434406 \cdot 1434405,$$

so its low-core cap remains 1434405. Since $2578110 > 1434405$, the reversible original-row lift forces the shared core in (N2).

At rank eight put $U = \ker(\text{ev}_B)$, so $\dim U = 2$, and anchor one retained owner (A_*, B_*) . Every retained explanation has the form

$$h_\gamma = A_* + \gamma B_* + v_\gamma, \quad v_\gamma \in U.$$

Hence anchored error differences lie in $\text{span}(U, r_1 - B_*)$, of dimension at most three. Exact locator multiplication in the reverse common-core lift is injective, proving (N3). $\square$

Proposition 51.9 (A fixed-chart local-cap fence). *The local rank-nine conclusions of Proposition 51.8, considered without the weighted component incidences of their ancestors, admit an official-row family of*

$$4070408 > 2578110$$

distinct pair-noncontained exact-support slopes. The family has common owner-plane core $K - 1$, every record has a rank-ten eleven-subset on an affine-owner component, and its errors have affine rank at most two.

Proof. Choose $J \subset D$ with $|J| = K - 1$ and put

$$u(X) = \prod_{x \in J} (X - x).$$

Fix a nine-subset $B \subset J$. Nine interpolation polynomials of degree at most eight together with u span a ten-dimensional RS space whose evaluation on B has rank nine and kernel Fu .

Outside J , the coordinate and rich-support weights are

$$N = n - (K - 1) = 1048577, \quad L = m - (K - 1) = 67473.$$

Give eight owner points $P_i = (c + iM, 0)$ weight $L - 1 = 67472$, where $M = 508801$, and give the remaining M coordinates weight one at $Q_j = (j, 1)$, $0 \leq j < M$. The values

$$\gamma_{i,j} = c + iM - j$$

are distinct because their eight difference intervals are consecutive and disjoint. There are $8M = 4070408$ of them. Since the KoalaBear base prime is greater than $18 \cdot 4070408$, the common translate c may also avoid any prescribed eighteen slopes.

Label an outside coordinate by (α_x, β_x) and define

$$r_0(x) = \alpha_x u(x), \quad r_1(x) = \beta_x u(x),$$

with both columns zero on J . At slope $\gamma_{i,j}$ use the codeword $h_{i,j} = (c + iM)u$. The line $\alpha + \gamma_{i,j}\beta = c + iM$ contains exactly P_i and Q_j among the assigned owner points. Its support therefore has size $(K - 1) + (L - 1) + 1 = m$.

A containing polynomial pair would have its second component equal to zero on J and the heavy fibre, altogether $m - 1 > K - 1$ coordinates, and hence zero identically by the RS root bound. At Q_j it would instead have to equal the nonzero value of u , a contradiction. Two coordinates in the heavy fibre adjoined to B give evaluation rank ten and the fixed-owner component through P_i . Finally

$$e_{i,j} = r_0 + \gamma_{i,j}r_1 - (c + iM)u,$$

so all anchored error differences lie in $\text{span}(r_1, u)$. □

Proposition 51.10 (Weighted high-row exclusion of the rank-nine fixed target). *In the fixed target of Proposition 51.8, alternative (N2) is empty for every residual dimension $20618 \leq K' \leq 1048576$. Before deduplication, one fixed nine-subset carries marked component-extension weight at least*

$$(W9) \quad \left\lceil \frac{495405467}{10^9} 274980728111260126 \frac{\binom{m'}{9} \binom{m'-9}{2}}{\binom{n'}{9}} \right\rceil.$$

Every rank-nine chart has marked weight at most

$$(C9) \quad 981105(m' - 10)n'.$$

At the exact first weighted crossing the adjacent integer bounds are

K'	demand	cap
20617	92386821615379573	92394042904582935
20618	92397581841774591	92395178310909600.

Consequently the fixed target reduces to (N1) or (N3) only from $K' = 20618$ onward. The rank-nine branch remains open below that row.

Proof. The marking argument in the proof of Proposition 51.8 produces $55I$ marked triples (γ, T, B) . Average over B without dividing by $\binom{m'-9}{2}$ and use

$$55 \binom{m'}{11} = \binom{m'}{9} \binom{m'-9}{2}$$

to obtain (W9). At $K' = 10$ its right side is 5868470021012020; dividing by the maximum extension multiplicity recovers the 2578110 distinct records used above.

Now fix a rank-nine chart and let u span its evaluation kernel on B . Write $C_p = J \sqcup P_p$ for the pair core of an owner point in the affine owner plane. The petals P_p are pairwise disjoint. A rank-ten extension $T = B \cup \{x, y\}$ determines one owner point p , both added coordinates lie in C_p , and at least one lies outside $Z(u)$, hence in P_p . Therefore one owner point contributes at most

$$|P_p| |C_p \setminus B| \leq |P_p| (m' - 10)$$

extensions per owned record. Pair noncontainment gives $|C_p| < m'$, while fixed-owner exception disjointness gives at most $n' - m' + 1 = 981105$ owned records. Summing the disjoint petals proves (C9).

Exact cross-multiplication is negative at $K' = 20617$ and positive at $K' = 20618$; after rounding the first strict gap is 2403530864991. After cancelling $\binom{m'-9}{2} = (m' - 9)(m' - 10)/2$, the ratio of the unrounded lower bound to (C9) is a positive constant times

$$\frac{\binom{m'}{9} m' - 9}{\binom{n'}{9} n'}.$$

Every factor increases with K' because $n' - m' = 981104$. The boundary contradiction therefore persists through the deployed endpoint.

The former low-row argument applied the original-row intersection floor $2m - n = 134944$ after reversing the shortening and then compared that lifted core with residual m' . This mixes rows: the reverse lift already inserts $1048576 - K'$ deleted locator coordinates into every owner core. Hence an original-row core of size 134944 need not contain any residual core of that size. No conclusion for $10 \leq K' \leq 20617$ is claimed here. $\square$

Proposition 51.11 (Residual petal-pair refinement of the rank-nine target). *In the fixed target of Proposition 51.8, alternative (N2) is empty for every*

$$15635 \leq K' \leq 20617.$$

More precisely, let J be the common core of one residual rank-nine owner plane and put $j = |J|$. Then $9 \leq j \leq K' - 1$, and its marked component-extension weight satisfies

$$(RP9) \quad W_B \leq \left\lfloor \frac{981105(n' - j)(m' + j - 20)}{2} \right\rfloor.$$

The exact adjacent integer bounds at the first crossing are

K'	demand	cap
15634	50777401704768572	50779283449126807
15635	50783693985583057	50780312213264392.

Together with Proposition 51.10, this removes the rank-nine target for every $K' \geq 15635$. It remains open for $10 \leq K' \leq 15634$.

Proof. Let u span the evaluation kernel on the fixed nine-set B . Owners agreeing on B form one affine plane. If J is the set of coordinates where the complete plane agrees with the received pair, then $B \subseteq J \subseteq Z(u)$. Since u is nonzero of degree below K' ,

$$(1) \quad 9 \leq j \leq K' - 1.$$

Write the core of owner p as $C_p = J \sqcup P_p$ and put $s_p = |P_p|$. The petals P_p are pairwise disjoint, and support-wise pair noncontainment gives $s_p \leq m' - j - 1$.

A marked extension $B \cup \{x, y\}$ has at least one coordinate outside $Z(u)$, hence in P_p for its owner. The number of available unordered pairs for p is therefore at most

$$q_p = s_p(j - 9) + \binom{s_p}{2}.$$

Owners with $s_p = 0$ contribute no marked extension. For every remaining owner,

$$(2) \quad \frac{q_p}{s_p} \leq j - 9 + \frac{m' - j - 2}{2} = \frac{m' + j - 20}{2}.$$

At most 981105 selected records have the same owner. Summing (2) and using $\sum_p s_p \leq n' - j$ proves (RP9); the floor is valid because W_B is integral.

The unfloored envelope $F(j) = (n' - j)(m' + j - 20)$ has forward difference

$$F(j + 1) - F(j) = 981104 - 2j + 19 > 0$$

through the claimed interval. Its maximum is therefore at $j = K' - 1$. Exact cross-multiplication against (W9) is negative at $K' = 15634$ and positive at $K' = 15635$. The ratio of the lower bound to the unfloored cap is a positive constant times nine factors $(m' - i)/(n' - i)$, $0 \leq i \leq 8$, and

$$\frac{(K' + 67463)(K' + 67462)}{2K' + 67451}.$$

The nine factors increase because $n' - m' = 981104$. After cancelling the positive common factor $K' + 67463$, the last factor has forward cross-difference $2(K' - 11) > 0$. Thus the contradiction persists through $K' = 20617$. No original-row 134944 core floor is used. $\square$

Proposition 51.12 (Canonical kernel-basis capacity). *In the setup of Proposition 51.2, the rank-deficient component lane cannot carry its dominant-lane incidence floor for any residual dimension*

$$10 \leq K' \leq 4598.$$

More precisely, for $1 \leq d \leq 9$, put $r = 10 - d$ and let $M_d(K')$ be the support-local record cap after cancelling r coordinates; for $d = 9$, take $M_9 = 61871313426630599$. The complete kernel-lane capacity is at most

$$(CK) \quad \text{Cap}(K') = \sum_{d=1}^9 \binom{n'}{10-d} M_d(K') \binom{K'-10}{d+1}.$$

For every $10 \leq K' \leq 4598$, this is strictly below

$$\frac{495405467}{10^9} 274980728111260126 \binom{m'}{11}.$$

At $K' = 4598$, the integer demand exceeds the capacity by

$$219272330501201744129177266158988707697316238048878827197685.$$

At $K' = 4599$ the comparison reverses, so no larger interval is claimed.

Proof. Fix a rank-deficient eleven-subset T , let $r = \text{rank}(\text{ev}_T)$, and choose one canonical rank basis $B \subset T$, $|B| = r$. Rank zero is absent because the ten-dimensional residual correction space has empty global common zero set. Put $H = \ker(\text{ev}_B)$, so $\dim H = d = 10 - r$.

Solving the coordinate equations on B determines the quotient component modulo H . If two selected slopes use the same B , every remaining compatibility condition is affine-linear in the slope and has two roots, hence is an identity. Thus, after one common affine codeword-pair translation, all records assigned to B have explanations in H . Every word of H and both translated received columns vanish on B , so exact locator cancellation preserves slopes, supports, and pair noncontainment. Support-local transversality bounds the resulting family by $M_d(K')$; at $d = 9$, use the uniform post-near margin/interleaving cap printed above.

Since T and B have the same evaluation rank, the $d + 1$ coordinates of $T \setminus B$ are common zeros of H . A d -dimensional subspace of degree- $< K'$ polynomials has at most $K' - d$ common

zeros. The $r = 10 - d$ basis coordinates are already among them, leaving at most $K' - 10$ choices. Hence one canonical basis carries at most

$$M_d(K') \binom{K' - 10}{d + 1}$$

incidences. There are at most $\binom{n'}{r}$ bases of rank r , and canonical assignment counts each tuple once, proving (CK).

For $d \leq 8$, the exact cap used in the replay is

$$M_d(K') = \left\lfloor \max \left\{ \frac{(1048576 + J)_{\underline{d+1}}}{(67472 + J)(67473)_{\underline{d-1}}}, \frac{(1048576 + d)_{\underline{d+1}}}{(67473)_{\underline{d}}} \right\} \right\rfloor, \quad J = K' - (10 - d).$$

Exact integer evaluation of this finite sum on all 4589 rows gives the displayed strict interval and endpoint gap. At $K' = 4599$, capacity exceeds the integer demand by 95457494746881463288875361950515757435711627164872173764503, which records the honest method wall. $\square$

Corollary 51.13 (All-bases kernel-capacity compression). *In the setup of Proposition 51.2, the rank-deficient component lane cannot carry its dominant-lane incidence floor for any residual dimension*

$$10 \leq K' \leq 11641.$$

More precisely, its capacity is at most

$$\text{Cap}_{\text{all}}(K') = \sum_{d=1}^9 \left\lfloor \frac{\binom{n'}{10-d} M_d(K') \binom{K'-10}{d+1}}{d+2} \right\rfloor.$$

At $K' = 11641$, the integer demand exceeds this capacity by

$$17769453550459149385453824948016076737082337523706893862084.$$

At $K' = 11642$, the capacity exceeds demand by

$$187031323586740190878769118921060658362307444191332937452616,$$

so no larger interval is claimed.

Proof. The evaluation matroid on a rank- $(10 - d)$ eleven-set T is loopless: a loop would be a coordinate at which the complete residual correction space vanishes, contrary to its empty global common zero set. Fix a basis B of T . For every $e \in T \setminus B$, the fundamental circuit $C(e, B)$ contains some $b_e \in B$, and $(B \setminus \{b_e\}) \cup \{e\}$ is a basis. These $d + 1$ bases are pairwise distinct and differ from B . Hence T has at least $d + 2$ rank bases.

Decorate each (γ, T) incidence by every one of its bases. For a fixed basis, the common quotient, cancellation, support-local record cap, and $\binom{K'-10}{d+1}$ extension bound in the proof of Proposition 51.12 apply unchanged; none used the canonical tie-break. Thus the numerator of each summand in the displayed capacity bounds the decorated incidences, while every undecorated incidence is represented at least $d + 2$ times. Division and the integer floor prove the capacity formula.

Exact integer replay of all 11632 rows gives the displayed interval and endpoint gap. The next-row calculation gives the stated reversal. The factor $d + 2$ is sharp from looplessness alone: take $9 - d$ coloops and one parallel class of size $d + 2$. $\square$

Corollary 51.14 (Record-support/ambient kernel-capacity hybrid). *For one residual record with exact support S , $|S| = m'$, the number of rank- $(10 - d)$ eleven-subsets of S is at most*

$$P_d(K') = \left\lfloor \frac{\binom{m'}{10-d} \binom{K'-10}{d+1}}{d+2} \right\rfloor.$$

Combining this with the ambient all-bases capacity $A_d(K')$ from Corollary 51.13, the kernel lane cannot carry its dominant-lane incidence floor for

$$10 \leq K' \leq 11772.$$

At $K' = 11772$, the integer demand exceeds the hybrid capacity by

$$76504076505592948633027913576880724493595282142849410185084.$$

At $K' = 11773$, the hybrid capacity exceeds demand by

$$139343682529231472322825521514042608524569163680782450618944,$$

so no larger interval is claimed.

Proof. For one record, decorate each rank- $(10 - d)$ eleven-subset of S by all its bases. A fixed basis $B \subset S$ has at most $\binom{K'-10}{d+1}$ same-rank extensions, and every eleven-set has at least $d + 2$ bases. There are at most $\binom{m'}{10-d}$ candidate bases, which proves the displayed per-record bound without using $M_d(K')$.

Let $R_{\text{act}} \geq N_{\text{min}} = 274980728111260126$ be the actual residual record count and let I_d be the rank- $(10 - d)$ incidence count. The two bounds give

$$I_d \leq \min\{A_d(K'), R_{\text{act}}P_d(K')\}.$$

After division by R_{act} , each term $\min\{A_d/R_{\text{act}}, P_d\}$ is nonincreasing in R_{act} . Hence it suffices to compare

$$\sum_{d=1}^9 \min\{A_d(K'), N_{\text{min}}P_d(K')\}$$

with the integer demand at N_{min} . Exact replay of all 11763 rows gives the stated interval and endpoint gap. At both boundary rows the ambient branch is selected for $d = 1, 2$ and the record-support branch for $d = 3, \dots, 9$. The next-row calculation gives the displayed reversal. $\square$

Proposition 51.15 (Shared kernel nine-shadow resource). *For one residual record with exact support S , let $I_d(S)$ count the rank- $(10 - d)$ eleven-subsets of S . Then*

$$\sum_{d=1}^9 \frac{\binom{d+2}{2}}{\binom{K'-d-9}{2}} I_d(S) \leq \binom{m'}{9},$$

where a zero extension coefficient forces the corresponding incidence count to vanish.

Proof. Fix a counted corank- d eleven-set T and put $r = 10 - d$. Its evaluation matroid M is loopless. The dual M^* has rank $c = d + 1$ and no coloops, and a nine-subset spans M exactly when its complementary pair is independent in M^* .

Ignore the loops of M^* and partition its remaining elements into parallel classes. If there are at least $c + 1$ classes, representatives of any $c + 1$ classes give at least $\binom{c+1}{2}$ independent pairs. If there are exactly c classes, each has size at least two, since a singleton class would be a coloop. The cross-class pairs then number at least $4\binom{c}{2} \geq \binom{c+1}{2}$. Thus T has at least $\binom{d+2}{2}$ spanning nine-subsets.

Fix such a rank- r nine-subset $U \subset S$. Its kernel has dimension d , so generalized MDS bounds its common-zero set, and hence the closure of U , by $K' - d$ coordinates. Apart from the nine coordinates of U , there are at most $K' - d - 9$ choices. Thus U has at most $\binom{K'-d-9}{2}$ same-rank eleven-set extensions. Double counting spanning pairs (T, U) in each rank and summing over the rank partition of the $\binom{m'}{9}$ nine-subsets proves the proposition. $\square$

Corollary 51.16 (One-shadow kernel-capacity cut). *Intersecting the inequality in Proposition 51.15 with the individual ambient/record caps excludes the dominant kernel lane for*

$$10 \leq K' \leq 15445.$$

At $K' = 15445$, the integer demand exceeds capacity by

$$178044655461817065880792270525721984196903835342334290540589.$$

At $K' = 15446$, capacity exceeds demand by

$$124087038578417364551353992932097013573495323735890481286577,$$

so no larger interval is claimed by this one-resource method.

Proof. After division by the actual record count, corank d has upper bound $u_d = \min\{A_d/N_{\min}, P_d\}$. Its weight in that inequality is $w_d = \binom{d+2}{2}/\binom{K'-d-9}{2}$, which increases with d . The exact capacity is therefore the fractional-knapsack optimum obtained by filling coranks in order. Exact rational replay checks all 15436 rows. At both boundary rows corank one is full, corank two is partial, and all higher coranks vanish. The displayed next-row reversal is exact, not a rounding artifact. $\square$

Proposition 51.17 (Full nine-shadow containment resource). *With $I_d(S)$ as above, put*

$$E_0 = \binom{m' - 9}{2}, \quad E_1 = \binom{K' - 10}{2}.$$

Then

$$\left(52 + 3 \frac{E_0}{E_1}\right) I_1(S) + 55 \sum_{d=2}^9 I_d(S) \leq E_0 \binom{m'}{9}.$$

When $E_1 = 0$, the corank-one term is absent because $I_1(S) = 0$.

Proof. Let J_1 count rank-nine nine-subsets of S . Every kernel eleven-set contains exactly $\binom{11}{9} = 55$ nine-subsets. A rank-nine nine-subset can only extend inside its rank-nine closure, giving at most E_1 extensions; every lower-rank nine-subset has at most the unrestricted support-pair count E_0 . Hence, with $I = \sum_d I_d$,

$$55I \leq E_1 J_1 + E_0 \left(\binom{m'}{9} - J_1 \right).$$

The corank-one case of Proposition 51.15 gives $3I_1 \leq E_1 J_1$. Since $E_0 > E_1$, substituting this lower bound for J_1 and rearranging gives the proposition. $\square$

Corollary 51.18 (Two-resource nine-shadow kernel-capacity cut). *The individual ambient/record caps together with Propositions 51.15 and 51.17 exclude the dominant kernel lane for*

$$10 \leq K' \leq 15670.$$

At $K' = 15670$, the integer demand exceeds capacity by

$$60244744187647715538325354175068999745872308513185869854532.$$

At $K' = 15671$, capacity exceeds demand by

$$291105561463347587484268984669020036510369238771859813045635,$$

so no larger interval is claimed.

Proof. The primary exact optimizer enumerates every cap breakpoint and every intersection of the rank-preserving and full-containment resource bounds. It checks all 15661 rows. At both boundary rows only coranks one and two are positive, both resource inequalities bind, and both individual caps are slack.

For an independent certificate, write w_1, w_2 for the first two weights in Proposition 51.15 and $v_1 = 52 + 3E_0/E_1$. Put

$$D = v_1 w_2 - 55 w_1, \quad \lambda = \frac{v_1 - 55}{D}, \quad \mu = \frac{w_2 - w_1}{D}.$$

Exact replay verifies $\lambda, \mu \geq 0$, $\lambda w_1 + \mu v_1 = \lambda w_2 + 55\mu = 1$, and $\lambda w_d + 55\mu \geq 1$ for $d \geq 3$. The resulting dual bound is exact at both boundary rows and gives the same signs and capacities as the primal replay. $\square$

Proposition 51.19 (Rank-eight nine-shadow extension deficit). *Let U be a rank-eight nine-subset of one exact support S . At least*

$$L_2 = \binom{67474}{2} = 2276336601$$

unordered pairs in $S \setminus U$ raise its rank to ten. Consequently, with

$$E_2 = \binom{K' - 11}{2},$$

the full-containment resource of Proposition 51.17 sharpens to

$$(R8S) \quad \left(52 + 3\frac{E_0}{E_1}\right) I_1(S) + \left(55 + 6\frac{L_2}{E_2}\right) I_2(S) + 55 \sum_{d=3}^9 I_d(S) \leq E_0 \binom{m'}{9}.$$

A zero extension coefficient again forces the corresponding incidence count to vanish.

Proof. Put $C = \text{cl}_S(U)$, $c = |C|$, and $X = S \setminus C$. The generalized-MDS closure bound used in Proposition 51.15 gives $c \leq K' - 2$, so $q = |X| \geq 67474$. Contracting by U leaves a rank-two matroid on X . Every parallel class $P \subset X$ has

$$|P| \leq K' - 1 - c,$$

because $C \cup P$ is contained in a rank-nine closure, whose generalized-MDS cap is $K' - 1$. Hence every $x \in X$ has at least

$$q - (K' - 1 - c) = m' - K' + 1 = 67473$$

partners outside its parallel class. Dividing the ordered count by two gives at least $q \cdot 67473/2 \geq \binom{67474}{2}$ rank-raising pairs.

Let J_2 count rank-eight nine-subsets of S . Such a subset has at most $E_0 - L_2$ kernel extensions. Refining the first inequality in the proof of Proposition 51.17 therefore gives

$$55 \sum_d I_d \leq E_0 \binom{m'}{9} - (E_0 - E_1)J_1 - L_2 J_2.$$

The shared-shadow resource gives $3I_1 \leq E_1 J_1$ and $6I_2 \leq E_2 J_2$. Substitution and rearrangement give (R8S). $\square$

Corollary 51.20 (Rank-eight-deficit nine-shadow capacity cut). *The individual ambient/record caps, the shared rank-preserving resource, and (R8S) exclude the dominant kernel lane for*

$$10 \leq K' \leq 17608.$$

At $K' = 17608$, integer demand exceeds floored capacity by

$$126547040539829546354916747965612889135249249684319416999204.$$

At $K' = 17609$, capacity exceeds demand by

$$165662859003771823867021831078593815988062146919602894849014,$$

so no larger interval is claimed by this method.

Proof. The primary replay enumerates all nonnegative vertices of the two-resource dual and checks all 17599 closed rows exactly. At the endpoint and wall, coranks one and three attain their individual caps, coranks two and four are resource-tight, coranks five through nine vanish, and both shared resources bind. An independently written replay reconstructs the primal optimizer from the printed fifteen-interval active-set ledger and verifies feasibility, complementary slackness, and strong duality on every row. It obtains the same rational optima and the displayed opposite boundary signs. $\square$

Proposition 51.21 (Two-step nine-shadow hierarchy). *For one exact support S , let $I_d(S)$ count rank- $(10 - d)$ eleven-subsets. For every $3 \leq d \leq 9$, put*

$$s_d = \binom{d+2}{2}, \quad L_d = \binom{67472+d}{2}, \quad E_d = \binom{K'-d-9}{2}, \quad Q_d = \binom{11-d}{2}.$$

Then

$$(TS) \quad \frac{s_d L_d}{E_d} I_d(S) \leq Q_d I_{d-2}(S).$$

A zero source-extension coefficient forces the corresponding source incidence count to vanish.

Proof. Let $J_d(S)$ count rank- $(10-d)$ nine-subsets. The shared-shadow theorem, Proposition 51.15, gives

$$(1) \quad s_d I_d(S) \leq E_d J_d(S).$$

Fix a nine-subset U counted by $J_d(S)$ and put $C = \text{cl}_S(U)$, $c = |C|$, and $X = S \setminus C$. The generalized-MDS closure cap gives $c \leq K' - d$, hence $q = |X| \geq 67472 + d$. After contracting U , every parallel class in X has size at most $K' - d + 1 - c$: adjoining such a class to C gives a rank- $(11 - d)$ closure. Every point of X therefore has at least $67471 + d$ partners outside its parallel class. At least

$$\frac{q(67471 + d)}{2} \geq \binom{67472 + d}{2} = L_d$$

unordered pairs raise the rank of U by two.

Now fix a resulting rank- $(12 - d)$ eleven-subset T and write $D = T \setminus U$. The dual-rank identity

$$r_T(T \setminus D) = r_T(T) - |D| + r_{T^*}(D)$$

shows that U has rank $10 - d$ exactly when the complementary pair D consists of two coloops of the evaluation matroid on T . This matroid is loopless, has rank $12 - d$, and has eleven elements, so it has at most $11 - d$ coloops: $12 - d$ coloops would make every remaining element a loop. Thus T contains at most Q_d source shadows, and double counting gives

$$(2) \quad L_d J_d(S) \leq Q_d I_{d-2}(S).$$

Combining (1) and (2) proves (TS). $\square$

Corollary 51.22 (Two-step nine-shadow kernel-capacity cut). *The individual ambient/record caps, the full-containment resource, and the seven inequalities (TS) exclude the dominant kernel lane for*

$$10 \leq K' \leq 18101.$$

At $K' = 18101$, integer demand exceeds floored capacity by

$$33462159928103132226516704640419847248244116666500998762314.$$

At $K' = 18102$, capacity exceeds demand by

$$275016496133605602641019628236447268989861205055439981187167,$$

so no larger interval is claimed by this method.

Proof. On every newly checked row $17609 \leq K' \leq 18102$, the exact optimizer has one active-set pattern: the corank-one cap, full containment, and all seven two-step inequalities bind; every corank is positive; the rank-preserving nine-shadow resource and all other individual caps are slack. The primary replay constructs the odd and even primal chains and solves the nine exact dual equations by rational Gaussian elimination. An independently written replay instead derives the hierarchy multipliers by backward odd/even recurrences. Both verify feasibility and strong duality on all 494 rows and obtain the displayed opposite boundary signs. $\square$

Proposition 51.23 (Multistep kernel-shadow hierarchy). *For one exact support S , let $I_d(S)$ count rank- $(10 - d)$ eleven-subsets. For every $2 \leq t < d \leq 9$, put*

$$s_{d,t} = \binom{d+2}{t}, \quad L_{d,t} = \binom{67472+d}{t}, \quad E_{d,t} = \binom{K'-d-11+t}{t}, \quad Q_{d,t} = \binom{9-d+t}{t}.$$

Then

$$(MS) \quad \frac{s_{d,t} L_{d,t}}{E_{d,t}} I_d(S) \leq Q_{d,t} I_{d-t}(S).$$

There are 28 such inequalities; the seven rows with $t = 2$ are Proposition 51.21.

Proof. Let N be the dual of the evaluation matroid on a source eleven-set. It is coloopless of rank $d + 1$, because the evaluation matroid is loopless. Write f_j for the number of independent j -sets of N . Every independent j -set has at least $d + 2 - j$ one-point extensions. Otherwise at most $d + 1 - j$ elements lie outside its closure. Rank forces equality, and deleting any one of those elements leaves rank at most d , making each such element a coloop, a contradiction. Double counting extensions therefore gives

$$(j + 1)f_{j+1} \geq (d + 2 - j)f_j.$$

Iteration from $f_0 = 1$ yields $f_t \geq \binom{d+2}{t}$. Taking complements gives at least $s_{d,t}$ spanning $(11 - t)$ -shadows per source. A fixed shadow has at most $E_{d,t}$ same-rank source extensions by the generalized-MDS closure cap.

Starting from such a shadow, after j independent support points have been adjoined its closure has size at most $K' - d + j$. Thus at least $67472 + d - j$ support points raise rank at the next step. Dividing the resulting ordered count by $t!$ gives at least $L_{d,t}$ rank-raising t -sets.

Conversely, fix a resulting rank- $(10 - d + t)$ target eleven-set. By the dual-rank identity, a complementary t -set produces a rank- $(10 - d)$ shadow exactly when all its elements are coloops of the target matroid. A loopless rank- $(10 - d + t)$ matroid on eleven elements has at most $9 - d + t$ coloops, or all remaining elements would be loops. Hence the target contains at most $Q_{d,t}$ source shadows. The two double counts give (MS). $\square$

Corollary 51.24 (All-step kernel-shadow capacity cut). *The individual ambient/record caps, full containment, and the 28 inequalities (MS) exclude the dominant kernel lane for*

$$10 \leq K' \leq 18158.$$

At $K' = 18158$, integer demand exceeds floored capacity by

$$289110608820324799941118306538399899258195112067661304310498.$$

At $K' = 18159$, capacity exceeds demand by

$$20286290696334777989469267474876769475675508046109372076445,$$

so no larger interval is claimed by this method.

Proof. On all 58 replayed rows $18102 \leq K' \leq 18159$, the exact certificate uses the hierarchy tree

$$(t, d) = (2, 3), (2, 4), (2, 6), (2, 8), (3, 5), (2, 7), (2, 9).$$

The corank-one cap and full-containment resource bind, all nine coranks are positive, and every other cap and the rank-preserving shadow resource are slack. Seventeen hierarchy rows bind; the ten tight rows outside the tree are checked as exact cycle identities. The primary replay solves the nine dual equations by rational Gaussian elimination. An independently written replay reconstructs both primal factors and dual multipliers by forward and backward tree recurrences. Both check all 28 inequalities, feasibility, and strong duality on every row and obtain the displayed opposite boundary signs. $\square$

Proposition 51.25 (Corank-one projective-normal pair cap). *After canonical rank-nine basis cancellation, the complete shortened corank-one chart has*

$$(n, K, m, s) = (1048577, 1, 67473, 1).$$

Every selected record owns at least $2(m - 1) = 134944$ ordered coordinate pairs whose agreement normals are linearly independent. Consequently the number of records in the chart is at most

$$(PP) \quad M_1 \leq \left\lfloor \frac{n(n-1)}{2(m-1)} \right\rfloor = 8147918.$$

This improves the support-local transversality cap 16295594 by 8147676.

Proof. Represent the affine explanation family by parameter points in $\mathbb{F}^2$. Agreement at a coordinate is an affine line with a normal in $\mathbb{F}^2$. The common-zero bound gives $z \leq K - s = 0$, so all m normals incident with one selected point are nonzero. Same-support pair noncontainment forces them to span $\mathbb{F}^2$.

Let their nonempty projective-class sizes be $a_1, \dots, a_c$. Then $c \geq 2$, $\sum_i a_i = m$, and the number of ordered independent pairs is

$$m^2 - \sum_i a_i^2.$$

Convexity gives $\sum_i a_i^2 \leq (m-1)^2 + 1$, with equality for class sizes $m-1, 1$. Thus every selected point owns at least $2(m-1)$ ordered independent coordinate pairs. Two independent affine lines meet in at most one parameter point, so a coordinate pair is owned by at most one record. Double counting the $n(n-1)$ ordered coordinate pairs proves (PP); the integer division has remainder 29760. $\square$

Corollary 51.26 (Projective-pair kernel-capacity cut). *The projective cap (PP), the other individual ambient/record caps, and the 28 inequalities (MS) exclude the dominant kernel lane for*

$$10 \leq K' \leq 377673.$$

At $K' = 377673$, integer demand exceeds floored capacity by

$$608290099077401798561583762592584078050381528604243813748500153228.$$

At $K' = 377674$, capacity exceeds demand by

$$1089804128361045148874283346879615159892995682385275039289561845323,$$

so no larger interval is claimed by this resource system.

Proof. The predecessor handles $K' \leq 18158$. On the remaining exact replay the corank-one and corank-two individual caps bind, both shared nine-shadow resources are strict, and every corank is positive. The seven hierarchy tree edges are

$$(t, d) = (2, 3), (3, 4), (2, 5), (2, 6), (2, 7), (2, 8), (2, 9).$$

They determine all nonroot counts. Twenty-two of the 28 hierarchy rows are equalities. Backward tree recurrence gives positive edge prices and positive prices for both root caps, with dual objective equal to the primal sum.

The source-pinned exact replay checks all 359516 rows $18159 \leq K' \leq 377674$ in 64 bounded workers. The packet verifiers independently reconstruct the projective partition bound and the replay start, endpoint, and adjacent wall from exact tree-path products. These give the displayed opposite integer signs. $\square$

Proposition 51.27 (Corank-two projective-normal basis cap). *After canonical rank-eight basis cancellation, the complete shortened corank-two chart has*

$$(n, K, m, s) = (1048578, 2, 67474, 2).$$

Every selected record owns at least $3(m-1)(m-2) = 13657614768$ ordered coordinate triples whose agreement normals are linearly independent. Consequently the number of records in the chart is at most

$$(PB) \quad M_2 \leq \left\lfloor \frac{n(n-1)(n-2)}{3(m-1)(m-2)} \right\rfloor = 84416263.$$

This improves the support-local transversality cap 253241283 by 168825020.

Proof. Represent the affine explanation family by parameter points in $\mathbb{F}^3$. The common-zero bound gives $z \leq K - s = 0$. For any one-dimensional normal span, the support-local MDS step leaves at least

$$w + s - 1 = 67472 + 2 - 1 = 67473 = m - 1$$

incident normals outside it. Thus no two of the m incident nonzero normals are projectively equal. Same-support pair noncontainment gives full incident rank, so the resulting m points in $\text{PG}(2)$ are not all collinear.

Choose a projective line containing $q \geq 2$ of the points and put $r = m - q \geq 1$. A collinear triple not contained in that line has zero or one point on it. There are at most $\binom{r}{3}$ triples of the

first type, while each pair off the line determines at most one triple of the second type. Hence the number of collinear triples is at most

$$\binom{q}{3} + \binom{r}{3} + \binom{r}{2} = \binom{q}{3} + \binom{r+1}{3} \leq \binom{m-1}{3}.$$

The last inequality follows from

$$\binom{m-1}{3} - \binom{q}{3} - \binom{r+1}{3} = (r-1) \left(\binom{q}{2} - 1 \right) + (q-2) \binom{r-1}{2} \geq 0.$$

Therefore the number of ordered independent triples is at least

$$m(m-1)(m-2) - 6 \binom{m-1}{3} = 3(m-1)(m-2).$$

Three independent affine hyperplanes in $\mathbb{F}^3$ meet in at most one parameter point. Double counting the $n(n-1)(n-2)$ ordered coordinate triples proves (PB); the integer division has remainder 2935655472. $\square$

Proposition 51.28 (Projective-paving caps and the shortening gap). *In the completely shortened corank- d chart, $1 \leq d \leq 9$, the incident normal matroid has rank $d+1$ and every set of at most d normals is independent. It therefore has at least $\binom{m-1}{d}$ unordered bases, and projective-basis double counting gives the record cap*

$$P_d = \left\lfloor \frac{(1048576 + d)_{d+1}}{(d+1)(67472 + d - 1)_d} \right\rfloor.$$

For $d = 1, \dots, 9$ these caps are

$$\begin{aligned} &8147918, 84416263, 983902549, 12232092309, \\ \text{(PC)} \quad &158406193634, 2109949210211, 28689347099870, \\ &396280526311830, 5542092977392141. \end{aligned}$$

They are complete-chart caps, not uniform shortened-chart caps.

More precisely, put $R = 1048576$, $w = 67472$, and, after rank- $(10-d)$ shortening, let

$$t = K' - 10 - z, \quad 0 \leq t \leq K' - 10,$$

where z is the number of global zero normals. For $d \geq 2$ the support-local endpoint envelope is (IG)

$$M_d(K') = \lfloor \max\{P_d, F_d(1), F_d(K' - 10)\} \rfloor,$$

where, before flooring,

$$P_d = \frac{(R+d)_{d+1}}{(d+1)(w+d-1)_d}, \quad F_d(t) = \frac{(R+d+t)_{d+1}}{(w+d+t)(w+1)^{\overline{d-1}}}.$$

At $K' = 377674$ this gives $M_2 = 253238254$ and $M_3 = 3935391907$. Corank one is exceptional: its projective-pair argument proves the uniform cap $M_1 = 8147918$. This endpoint envelope is a valid generic bound, not an obstruction to a stronger multiplicity-aware count; the corank-two theorem below supplies such a count.

Proof. The matroid assertion follows by deletion-contraction. If an element is a coloop, adjoining it to every $(d-1)$ -subset of the deletion already gives $\binom{m-1}{d}$ bases. Otherwise deletion and contraction preserve the required independence hypotheses, and Pascal's identity closes the induction. Ordering the bases and double counting their affine-hyperplane intersections proves (PC).

For $t \geq 1$, the pointwise support-local transversality count before endpoint maximization gives $F_d(t)$. The sign of $F_d(t+1)/F_d(t) - 1$ is the sign of

$$dt + (d+1)(w+d) - R.$$

This expression increases with t , so F_d has one turn and its maximum on the integer interval is at 1 or $K' - 10$. The complete chart $t = 0$ is controlled separately by P_d , proving (IG). The

same one-turn argument for the rank-two projective-pair count is maximized at $t = 0$, proving the uniform corank-one assertion. Exact integer evaluation at $K' = 377674$ gives the displayed audit values. $\square$

Proposition 51.29 (Rank-three bounded-parallel basis floor). *Let M be a loopless rank-three matroid on $m \geq 3$ elements, and suppose every parallel class has size at most a . If $b(M)$ is its number of unordered bases, then*

$$(BP) \quad 2b(M) \geq (m-1)(m-1-a).$$

If $a \mid (m-1)$ and $m-1 \geq 2a$, equality is attained by one coloop over a rank-two matroid whose parallel classes all have size a .

Proof. Induct on m . If e is a coloop, the deletion is a loopless rank-two matroid on $q = m-1$ elements. Writing its parallel-class sizes as q_i ,

$$2b(M) = q^2 - \sum_i q_i^2 \geq q^2 - aq = (m-1)(m-1-a).$$

Otherwise choose e in a smallest parallel class P , of size c . Deletion preserves rank three, so induction gives

$$2b(M \setminus e) \geq (m-2)(m-2-a).$$

In M/e , the other elements of P are loops. Every nonloop parallel class is a union of original parallel classes and hence has size at least c . There are at least two such classes and their total size is $m-c$. Merging classes decreases the number of independent pairs, so

$$b(M/e) \geq c(m-2c) \geq m-2.$$

Indeed the simplification of M has at least three points, hence $m \geq 3c$, and the residual at $m = 3c$ is $(c-1)(c-2) \geq 0$. Deletion-contraction now gives

$$2b(M) \geq (m-2)(m-2-a) + 2(m-2) = (m-1)(m-1-a) + (a-1),$$

which proves (BP). Equal rank-two classes in the stated coloop construction give equality. $\square$

Proposition 51.30 (Uniform corank-two projective-normal basis cap). *After canonical rank-eight basis cancellation and deletion of all global zero normals, put $t = K' - 10 - z$. Every official corank-two chart has*

$$(n, K, m, s) = (1048578 + t, 2 + t, 67474 + t, 2), \quad 0 \leq t \leq 1048566,$$

and at most

$$(UPB) \quad M_2 \leq 84416263$$

records. The complete chart $t = 0$ attains the largest value of the proved envelope.

Proof. Write $R = 1048576$ and $w = 67472$. Support-local transversality leaves at least $w+1$ incident normals outside every one-dimensional normal span. Thus the loopless rank-three incident normal matroid has maximum parallel class size

$$a \leq m - (w+1) = t+1.$$

By Proposition 51.29, every record owns at least

$$6b(M) \geq 3w(w+t+1)$$

ordered independent coordinate triples. Double counting gives

$$H(t) = \frac{(R+t)(R+t+1)(R+t+2)}{3w(w+t+1)}$$

as a record envelope. The sign of $H(t+1)/H(t) - 1$ is the sign of

$$2t + 3w + 3 - R.$$

Hence H has one turn and its maximum on the official integer interval is at an endpoint. Exact division gives

$$\lfloor H(0) \rfloor = 84416263, \quad \lfloor H(1048566) \rfloor = 40828171,$$

with positive next-integer gaps. This proves (UPB) uniformly. $\square$

Corollary 51.31 (Projective-basis kernel-capacity cut). *The projective caps for coranks one and two, the other individual ambient/record caps, and the 28 inequalities (MS) exclude the dominant kernel lane for*

$$10 \leq K' \leq 568338.$$

At $K' = 568338$, integer demand exceeds floored capacity by

$$38432453444617070485037263551626410396462586389410416394578520596038.$$

At $K' = 568339$, capacity exceeds demand by

$$36180877960369511460476382880286784896208001102094988739728829832800,$$

so no larger interval is claimed by this resource system.

Proof. The predecessor handles $K' \leq 377673$. On the remaining exact replay the corank-one and corank-two individual caps bind, both shared nine-shadow resources are strict, every nonroot individual cap is strict, and every corank is positive. The seven hierarchy tree edges are

$$(t, d) = (2, 3), (2, 4), (2, 6), (2, 8), (3, 5), (2, 7), (2, 9).$$

They give the two rooted components

$$1 \longrightarrow 3, qquad 2 \longrightarrow 4 \longrightarrow 6 \longrightarrow 8, qquad 2 \longrightarrow 5 \longrightarrow 7 \longrightarrow 9.$$

Seventeen of the 28 hierarchy rows are equalities. Backward tree recurrence gives positive edge prices and positive prices for both root caps, with dual objective equal to the primal sum.

The source-pinned exact replay checks all 190666 rows $377674 \leq K' \leq 568339$ in 64 bounded workers. The independent verifier enumerates every projective line split in the basis bound and reconstructs the replay start, endpoint, and adjacent wall from exact tree-path products. These give the displayed opposite integer signs. $\square$

Proposition 51.32 (Corank-three projective-frame basis cap). *After canonical rank-seven basis cancellation, the complete shortened corank-three chart has*

$$(n, K, m, s) = (1048579, 3, 67475, 3).$$

Every selected record owns at least $4(m-1)(m-2)(m-3) = 1228711865141376$ ordered coordinate quadruples whose agreement normals are linearly independent. Consequently the number of records in the chart is at most

$$(PF) \quad M_3 \leq \left\lfloor \frac{n(n-1)(n-2)(n-3)}{4(m-1)(m-2)(m-3)} \right\rfloor = 983902549.$$

This improves the support-local transversality cap 3935435218 by 2951532669.

Proof. Represent the affine explanation family by parameter points in $\mathbb{F}^4$. The common-zero bound gives $z \leq K - s = 0$. For every one-dimensional normal span, the support-local MDS step leaves at least

$$w + s - 1 = 67472 + 3 - 1 = 67474 = m - 1$$

incident normals outside it, so the incident projective points are distinct. For every two-dimensional normal span, the same step leaves at least

$$w + s - 2 = 67473 = m - 2$$

normals outside it. Thus no three incident points are collinear. Same-support pair noncontainment gives full incident rank, so these points span $\text{PG}(3)$.

Choose three incident points and let H be their projective plane. Put $q = |S \cap H| \geq 3$ and $r = m - q \geq 1$. Partition coplanar unordered quadruples by the number of points in H . Four points in H contribute $\binom{q}{4}$, while three in H and one outside contribute none. For a fixed outside pair, its line meets H at a point $x \notin S$, since $x \in S$ would give three collinear selected points. Secants of $S \cap H$ through x use disjoint point pairs, giving at most $\lfloor q/2 \rfloor \binom{r}{2}$ quadruples with two points in H . Every outside triple spans a plane meeting H in a line with at most two selected

points, giving at most $2\binom{r}{3}$ quadruples with one point in H . All-outside quadruples contribute at most $\binom{r}{4}$. Hence the dependent quadruple count is at most

$$B(q, r) = \binom{q}{4} + \left\lfloor \frac{q}{2} \right\rfloor \binom{r}{2} + 2\binom{r}{3} + \binom{r}{4}.$$

Vandermonde expansion gives

$$\binom{m-1}{4} - B(q, r) = (q-3)\binom{r}{3} + \left(\binom{q-1}{2} - \left\lfloor \frac{q}{2} \right\rfloor \right) \binom{r}{2} + (r-1)\binom{q-1}{3} \geq 0.$$

Therefore the number of ordered independent quadruples is at least

$$(m)_4 - 24\binom{m-1}{4} = 4(m-1)(m-2)(m-3).$$

Four independent affine hyperplanes in $\mathbb{F}^4$ meet in at most one parameter point. Double counting the $(n)_4$ ordered coordinate quadruples proves (PF); the integer division has remainder 1056607358217600. $\square$

Proposition 51.33 (Rank-four bounded-point/line basis floor). *Let M be a loopless rank-four matroid on $m = a + r$ elements, where $a \geq 1$ and $r \geq 3$. Suppose every parallel class has size at most a and every rank-two flat has size at most $a + 1$. Put*

$$h_a(r) = \min \left\{ \left\lfloor \frac{a+1}{2} \right\rfloor, \left\lfloor \frac{a+r}{4} \right\rfloor \right\},$$

$$C_a(r) = (a+r-1)(r-1)(r-2), \quad L_a(r) = 3(a+r-h_a(r)-1)(r-2),$$

and define $Q_a(3) = 6$ and, for $r \geq 4$,

$$(BPL) \quad Q_a(r) = \min\{C_a(r), Q_a(r-1) + L_a(r)\}.$$

Then the number $b(M)$ of unordered bases satisfies

$$6b(M) \geq Q_a(r).$$

Proof. Induct on r . At $r = 3$, rank four supplies one basis. If M has a coloop e , then $M \setminus e$ has rank three on $q = m - 1$ elements. Sequentially choosing an ordered independent triple gives at least q , $q - a = r - 1$, and $q - (a + 1) = r - 2$ choices, respectively. Hence $6b(M) \geq C_a(r)$.

Otherwise choose e in a smallest parallel class P , of size c . The simplification has at least four points, so $4c \leq m$. A rank-two flat through P and any other parallel class contains at least $2c$ elements, so $2c \leq a + 1$. Thus $c \leq h_a(r)$. Deletion preserves rank four and the two ceilings. After contracting e and deleting the loops $P \setminus e$, the resulting rank-three matroid has $a + r - c$ elements and parallel-class ceiling $a + 1 - c$. Proposition 51.29 therefore gives

$$6b(M/e) \geq 3(a+r-c-1)(r-2) \geq L_a(r).$$

Deletion-contraction gives the second branch of (BPL).

For exact evaluation, unfold the recurrence. Besides the no-reset path, a last coloop reset at $j \geq 4$ contributes

$$C_a(j) + \sum_{x=j+1}^r L_a(x).$$

The difference between the candidates at $j + 1$ and j is $(j-1)(3h_a(j+1) - a - 2)$. Since h_a is nondecreasing, the reset candidates have one turn; the no-reset path is no larger than the $j = 4$ candidate by $6(h_a(4) - 1) \geq 0$. Hence only the no-reset path and the reset adjacent to that sign change can minimize. Splitting $h_a(x) = \min\{\lfloor (a+1)/2 \rfloor, \lfloor (a+x)/4 \rfloor\}$ at its constant branch and by $a + x$ modulo four gives an exact integer evaluator. $\square$

Corollary 51.34 (Uniform corank-three projective-frame cap). *Every official corank-three shortened chart satisfies*

$$(UPF) \quad M_3 \leq 983902549.$$

Proof. After canonical rank-seven basis cancellation, delete all global zero normals and put $t = K' - 10 - z$, $a = t + 1$. The remaining chart has

$$(n, K, m, s) = (1048579 + t, 3 + t, 67475 + t, 3), \quad r = m - a = 67474.$$

Support-local transversality leaves at least 67474 incident normals outside every one-dimensional span and at least 67473 outside every two-dimensional span. Thus the incident loopless rank-four matroid has parallel-class ceiling a and rank-two-flat ceiling $a + 1$. Proposition 51.33 gives at least $4Q_{t+1}(67474)$ ordered independent coordinate quadruples per record, so

$$M_3(t) \leq \left\lfloor \frac{(1048576 + t)(1048577 + t)(1048578 + t)(1048579 + t)}{4Q_{t+1}(67474)} \right\rfloor.$$

The one-turn exact evaluator checks every $0 \leq t \leq 1048566$. Its maximum is 983902549 at $t = 0$; the values at $t = 1$ and $t = 1048566$ are 983891721 and 951742008. The primary verifier exhausts all 1048567 rows with exact integers, while the independent verifier unfolds the recurrence directly at five separated rows. A source-pinned 512 MB, 60-second-limited worker independently reports the same all-row maximum. $\square$

Corollary 51.35 (Projective-frame kernel-capacity cut). *The projective caps for coranks one, two, and three, the other individual ambient/record caps, and the 28 inequalities (MS) exclude the dominant kernel lane for*

$$10 \leq K' \leq 796598.$$

At $K' = 796598$, integer demand exceeds floored capacity by

$$1063274038253455766288412818872693782800681544679740581002823089126086.$$

At $K' = 796599$, capacity exceeds demand by

$$670721678337441589385303494237372283642375643589068751593971045368244,$$

so no larger interval is claimed by this resource system.

Proof. The predecessor handles $K' \leq 568338$. On the remaining exact replay the corank-one, corank-two, and corank-three individual caps bind, both shared nine-shadow resources are strict, every nonroot individual cap is strict, and every corank is positive. The six hierarchy forest edges are

$$(t, d) = (2, 4), (2, 5), (3, 6), (4, 7), (5, 8), (6, 9).$$

They give the components

$$1, \quad 2 \longrightarrow 4, \quad 3 \longrightarrow 5, 6, 7, 8, 9.$$

Twelve of the 28 hierarchy rows are equalities. Backward forest recurrence gives positive edge prices and positive prices for all three root caps, with dual objective equal to the primal sum.

The source-pinned exact replay checks all 228261 rows $568339 \leq K' \leq 796599$ in 64 bounded workers. The independent verifier enumerates every projective plane split in the frame bound and reconstructs the replay start, endpoint, and adjacent wall from direct forest products. These give the displayed opposite integer signs. $\square$

Proposition 51.36 (Shortening-weighted record/extension cap). *Put $R = 1048576$, $w = 67472$, $S = K' - 10$, and let P_d be the printed complete-chart cap in corank d . For one fixed basis of a decorated corank- d kernel incidence, define*

$$F_d(t) = \frac{(R + d + t)^{\underline{d+1}}}{(w + d + t)(w + 1)^{\underline{d-1}}}.$$

On $796599 \leq K' \leq 1048576$, its record/extension contribution is at most

$$(SWE) \quad U_d(K') = \begin{cases} P_d\left(\begin{smallmatrix} S \\ d+1 \end{smallmatrix}\right), & 1 \leq d \leq 3, \\ F_d(1)\left(\begin{smallmatrix} S-1 \\ d+1 \end{smallmatrix}\right), & 4 \leq d \leq 9. \end{cases}$$

After summing over bases, every undecorated incidence is counted at least $d + 2$ times.

Proof. Let z be the number of common zero normals outside the fixed basis and put $t = S - z$. Every extension of the basis uses $d + 1$ of those zero normals, so one record has at most $\binom{S-t}{d+1}$ extensions. The same z is deleted in the support-local chart. Thus the complete chart contributes at most $P_d\binom{S}{d+1}$, while a noncomplete chart contributes at most $F_d(t)\binom{S-t}{d+1}$.

For $q = d + 1$ and every nonzero term with $t \geq 1$, the successive ratio of the unfloored noncomplete bounds is

$$\frac{(R + d + t + 1)(w + d + t)(S - t - q)}{(R + t)(w + d + t + 1)(S - t)} < \left(1 + \frac{q}{R + t}\right) \left(1 - \frac{q}{S - t}\right) < 1.$$

The final inequality uses $R + t > S - t$. Hence the noncomplete maximum is at $t = 1$. For $d \leq 3$, the already proved uniform caps P_d and the larger complete extension count give the first branch of (SWE). For $4 \leq d \leq 9$, exact cross multiplication at $K' = 796599$ shows that the $t = 1$ branch exceeds the complete branch; their ratio increases with S . The six comparisons are replayed with exact rationals. The all-bases decoration theorem supplies the final divisor $d + 2$. $\square$

Corollary 51.37 (Terminal shortening-weighted kernel cut). *The fixed-kernel branch is impossible on every official residual dimension*

$$10 \leq K' \leq 1048576.$$

Proof. The projective-frame capacity cut handles $K' \leq 796598$. On the remaining interval, Proposition 51.36 gives the rational capacity

$$C(K') = \sum_{d=1}^9 \frac{\binom{R+K'}{10-d} U_d(K')}{d+2}.$$

Compare it with the exact unrounded demand

$$D(K') = \frac{495405467}{10^9} 274980728111260126 \binom{w + K'}{11},$$

and put $G = D - C$. This is a polynomial of degree at most eleven. Exact forward differencing at $K_0 = 796599$ gives

$$\Delta^j G(K_0) > 0 \quad (0 \leq j \leq 11).$$

For every integer $s \geq 0$, Newton interpolation therefore gives

$$G(K_0 + s) = \sum_{j=0}^{11} \binom{s}{j} \Delta^j G(K_0) > 0.$$

This covers the complete upper interval. Floors only decrease capacity. The primary verifier reconstructs and hashes all twelve exact Newton coefficients. The independent verifier instead expands $G(K_0 + s)$ in the ordinary power basis; all twelve coefficients are positive and have SHA-256 digest 85f67eade03100e7fc6f3ef0443a3a787f1133427a0f9d4030d638eca9080082. $\square$

Proposition 51.38 (Rank-eight minimal-shortening exclusion). *At residual dimension $K' = 10$, the rank-eight affine-owner alternative of Proposition 51.8 is empty. Every nine-coordinate restriction of the correction space has rank nine. This conclusion is sharp for the dimension-equality argument: the adjacent row $K' = 11$ admits the rank-eight chart in Proposition 51.40.*

Proof. The dense-root saturation theorem places the ten-dimensional correction space V' inside the residual degree-below- K' Reed–Solomon space. At $K' = 10$ both spaces have dimension ten, so

$$V' = \{f \in \mathbb{F}[X] : \deg f < 10\}.$$

For any nine distinct coordinates $b_1, \dots, b_9$, the evaluation minor formed by $1, X, \dots, X^8$ has Vandermonde determinant

$$\prod_{1 \leq i < j \leq 9} (b_j - b_i) \neq 0.$$

Thus evaluation on the nine-set has rank nine, contradicting the rank-eight target. At $K' = 11$, the space $\langle 1, X, \dots, X^7, L_B, XL_B \rangle$ from the next proposition has rank eight on B , so no adjacent-row propagation is claimed. $\square$

Proposition 51.39 (Rank-eight codimension-one circuit-shadow census). *At $K' = 11$, fix a nine-set B in the rank-eight affine-owner target. There is one circuit $C_B \subset B$, independent of every extension pair, with $2 \leq c := |C_B| \leq 9$. Every eleven-set $T = B \cup \{x, y\}$ has exactly*

$$\binom{11-c}{2}$$

rank-eight nine-shadows, $55 - \binom{11-c}{2}$ rank-nine nine-shadows, and c rank-ten bases. Moreover

$$L_{C_B}\{f : \deg f < 11 - c\} \subset V'.$$

Every $c \in \{2, \dots, 9\}$ occurs at the correction-space and chart-rank level. Thus no circuit-size branch is removed by local linear algebra.

Proof. Put $P = \{f : \deg f < 11\}$. It has dimension eleven, while V' is a ten-dimensional hyperplane. Since evaluation of V' on B has rank eight,

$$\ker(\text{ev}_B|_{V'}) = L_B\{a + bX : a, b \in \mathbb{F}\};$$

both sides have dimension two. A functional λ cutting out V' therefore annihilates the complete kernel of $\text{ev}_B : P \rightarrow \mathbb{F}^B$ and factors uniquely as

$$\lambda(f) = \sum_{b \in B} \mu_b f(b).$$

Let C_B be the support of $(\mu_b)_{b \in B}$. If $|C_B| = 1$, that coordinate is an evaluation loop of V' . For every one of the at least 2578110 distinct fixed-chart slopes, the normalized rich equation there reduces to $E_b(\gamma) = 0$. Since $E_b(Z)$ is affine, two slopes make it identically zero, producing a coordinate in the complete residual common support, contrary to maximal shortening. Hence $2 \leq |C_B| \leq 9$.

Evaluation of P on the eleven distinct points of T is an isomorphism. The unique relation on the rank-ten image of V' is consequently $(\mu_b)_{b \in B}$, extended by zero at x, y . Its support is the fixed circuit C_B . A nine-subset of T has rank eight exactly when its omitted pair misses this circuit; otherwise its nine columns are independent. This gives the two shadow counts. A ten-subset is a basis exactly when its omitted point lies in C_B , giving c bases. The factorization of λ also places every polynomial vanishing on C_B in V' , proving the locator-ideal inclusion.

For sharpness, choose any c points of B , nonzero weights on them, and take V' to be the kernel of the corresponding weighted evaluation functional. Its restriction to B has rank eight, and no coordinate is a loop because evaluation functionals at at most eleven distinct points are Vandermonde-independent. This realizes every $c = 2, \dots, 9$ at the stated linear-algebra level. The eight-petal space in Proposition 51.40 is the $c = 9$ endpoint, with one rank-eight and 54 rank-nine shadows. $\square$

Proposition 51.40 (Rank-eight fixed-chart local-payment fence). *At residual dimension $K' = 11$, the local rank-eight output of Proposition 51.8, even with the marked-extension unit retained, admits*

$$4070352 > 2578110$$

distinct exact-support records on one fixed chart. Each record has $\binom{67473}{2}$ rank-ten affine-owner component extensions, so the marked weight is

$$(R8F) \quad 9265216597693056 > 5869376383979174,$$

where the right side is the exact weighted-selector demand at $K' = 11$. The errors have affine rank at most two, every support is pair-noncontained, and the construction lifts reversibly to the official row.

Proof. On the residual domain, choose a nine-set B , put

$$u_0(X) = \prod_{b \in B} (X - b), \quad u_1(X) = Xu_0(X),$$

and take the ten-space

$$V' = \langle 1, X, \dots, X^7, u_0, u_1 \rangle.$$

Its evaluation rank on B is eight, with kernel $U = \langle u_0, u_1 \rangle$. For distinct $x, y \notin B$, the determinant of kernel evaluation on $\{x, y\}$ is

$$u_0(x)u_0(y)(y - x) \neq 0,$$

so $B \cup \{x, y\}$ has evaluation rank ten.

Here $(n', m') = (1048587, 67483)$. Partition the complement of B into eight petals P_e of size 67473 and a remainder R of size

$$1048587 - 9 - 8 \cdot 67473 = 508794.$$

Choose distinct t_e and owner pairs $(a_e, b_e) = (t_e u_0, 1)$. Define the received pair by

$$(r_0, r_1) = \begin{cases} (0, 1), & B, \\ (a_e, 1), & P_e, \\ (s_x, 0), & x \in R. \end{cases}$$

Choose the s_x successively so that all

$$\gamma_{e,x} = s_x - a_e(x)$$

are distinct and avoid the eighteen dense slopes. At the j -th remainder coordinate, at most $64j + 8 \cdot 18$ values are forbidden. The maximum

$$64(508794 - 1) + 144 = 32562896$$

is below the deployed base prime 2130706433, so the greedy choice exists.

At slope $\gamma_{e,x}$ use explanation $h_{e,x} = a_e + \gamma_{e,x}$. Its exact support is

$$B \dot{\cup} P_e \dot{\cup} \{x\},$$

of size 67483. Equality at another remainder coordinate would repeat a slope, while equality on another petal would force $(t_e - t_f)u_0 = 0$ off B . For every two-set in P_e , adjoining it to B gives a rank-ten component owned by $(a_e, 1)$. If any polynomial pair contained the complete support, its second component would equal one on $|B| + |P_e| = 67482 > 10$ coordinates, hence would be the constant one, but the received second column is zero at x . Thus the support is pair-noncontained. Finally

$$e_{e,x} = r_0 - a_e + \gamma_{e,x}(r_1 - 1),$$

so all error differences lie in $\langle u_0, r_1 - 1 \rangle$.

There are $8 \cdot 508794 = 4070352$ records and their marked weight is the left side of (R8F). Direct evaluation of the weighted selector at $K' = 11$ gives the right side. To lift, choose a common core J of size $K - 11$, multiply the residual received values, owners, explanations, and correction space by its locator, and set both received columns to zero on J . Evaluation ranks are unchanged off J , every support gains J , and the same root-bound argument preserves pair noncontainment.

This witness has only four million records and its normalized deviations do not span the complete ten-space. It therefore fences a payment from the printed chart-local output; it is not an unsafe line and does not realize the dense-anchor ancestors. $\square$

Proposition 51.41 (Rank-eight owner-pair capacity). *In the setup of Proposition 51.2, the rank-eight affine-owner fixed target is impossible for every*

$$37996 \leq K' \leq 1048576.$$

For a fixed rank-eight nine-subset, its marked component load is at most

$$(R8W) \quad W_B \leq 981105 \binom{n' - 9}{2}.$$

At $K' = 37996$, the weighted selector demand is 579191514708840299, while the right side of (R8W) is 579155144020629315, a gap of 36370688210984. At $K' = 37995$, the cap still exceeds demand by 18174297527234.

Proof. Fix a rank-eight nine-subset B and put $U = \ker(\text{ev}_B)$, so $\dim U = 2$. Owners agreeing with the received pair on B form one affine translate of U^2 .

Every marked extension has the form $T = B \cup \{x, y\}$ with $\text{rank}(\text{ev}_T) = 10$. Therefore evaluation on $\{x, y\}$ is an isomorphism from U to $\mathbb{F}^2$. The two equations for the first received column determine the first U -coordinate of the owner uniquely, and the two equations for the second column determine the second U -coordinate uniquely. Thus one unordered coordinate pair $\{x, y\}$ determines at most one owner point p .

Let q_p count the full-rank coordinate pairs outside B that determine p , and let t_p count the selected records owned by p . Coordinate-pair uniqueness and fixed-owner exception disjointness give

$$\sum_p q_p \leq \binom{n' - 9}{2}, \quad t_p \leq n' - m' + 1 = 981105.$$

Consequently

$$W_B \leq \sum_p t_p q_p \leq 981105 \binom{n' - 9}{2},$$

in the selector's marked (record, T) unit.

For uniformity, if $L(K')$ denotes the unrounded selector demand and $U(K')$ the displayed cap, then

$$\frac{L(K')}{U(K')} = \text{constant} \cdot \frac{\binom{m'}{11}}{\binom{n'}{11}},$$

using $\binom{n'}{9} \binom{n' - 9}{2} = 55 \binom{n'}{11}$. Each factor $(m' - i)/(n' - i)$, $0 \leq i \leq 10$, strictly increases with K' because $n' - m' = 981104$. The exact positive gap at $K' = 37996$ therefore persists through the deployed endpoint. The failed preceding row records the honest method wall. $\square$

Corollary 51.42 (Rank-eight bridge to the dense-owner terminal). *For every $22526 \leq K' \leq 37995$, a surviving rank-eight fixed chart contains one owner pair that owns at least 200632 selected records and whose complete pair core has deficiency at most four.*

Proof. With the notation in the preceding proof, let q_p count coordinate pairs determining owner p and t_p its selected records. Then

$$W_B \leq \sum_p t_p q_p, \quad \sum_p q_p \leq \binom{n' - 9}{2}.$$

At $K' = 22526$, exact replay gives

$$W_B - 200631 \binom{n' - 9}{2} = 11714977255865.$$

Thus some $t_p \geq 200632$. If that owner's pair-core deficiency were at least five, fixed-owner exception disjointness would give

$$t_p \leq 1 + \left\lfloor \frac{981104}{5} \right\rfloor = 196221,$$

a contradiction. Hence its deficiency is at most four.

The weighted-average ratio is again a positive constant times $\binom{m'}{11}/\binom{n'}{11}$, so the conclusion persists through $K' = 37995$. At $K' = 22525$, the comparison misses by 1170919108090, and no smaller interval is claimed. $\square$

Remark 51.43 (Typed scope of the rank-eleven packet). The density in Proposition 51.2 is a record/coordinate-subset incidence density, not a record or component count. The corollary is recordwise, and Proposition 51.5 is local to one fixed cell, and Proposition 51.7 remains local to one lifted owner plane. Proposition 51.8 selects one fixed populated chart, not the complete component lane. The deduplicated local conclusions alone do not pay that chart by Proposition 51.9. That fence does not carry the marked weight in Proposition 51.10, which eliminates the rank-nine alternative from $K' = 20618$ onward. Proposition 51.11 extends that elimination

down to $K' = 15635$ without transporting the original-row core floor. Proposition 51.12 first eliminates kernel-lane domination through $K' = 4598$. Retaining every basis decoration in Corollary 51.13 extends that interval through $K' = 11641$. The record-support/ambient hybrid of Corollary 51.14 extends it through $K' = 11772$. Proposition 51.15 couples the corank strata and Corollary 51.16 extends the exact interval through $K' = 15445$. Counting all contained nine-subsets in Proposition 51.17 supplies a second resource; Corollary 51.18 extends the interval through $K' = 15670$, but not above the printed 15671 method wall. Proposition 51.19 prices the rank-eight shadow extension deficit, and Corollary 51.20 extends the exact interval through $K' = 17608$, but not above the printed 17609 wall. Proposition 51.21 couples all remaining coranks two ranks at a time, and Corollary 51.22 extends the exact interval through $K' = 18101$, but not above the printed 18102 wall. Proposition 51.23 supplies all 28 step sizes, and Corollary 51.24 extends the exact interval through $K' = 18158$, but not above the printed 18159 wall. Proposition 51.25 sharpens the active corank-one cap, and Corollary 51.26 extends the exact interval through $K' = 377673$, but not above the printed 377674 wall. Proposition 51.27 sharpens the complete-chart corank-two cap. Proposition 51.28 records the weaker support-local shortening envelope, while Propositions 51.29 and 51.30 prove the stronger cap uniformly. Therefore Corollary 51.31 extends the exact interval through $K' = 568338$, but not above the printed 568339 wall. Proposition 51.32 sharpens the complete-chart corank-three cap. Proposition 51.33 and Corollary 51.34 prove that cap uniformly. Therefore Corollary 51.35 extends the exact interval through $K' = 796598$, but not above the printed 796599 wall. Proposition 51.36 couples each chart's record cap to its extension count instead of maximizing them separately. Corollary 51.37 then removes the fixed-kernel branch through the deployed endpoint. Proposition 51.41 eliminates the rank-eight target from $K' = 37996$ onward, but not below its printed wall. Corollary 51.42 routes the upper surviving rank-eight interval to the guarded dense-owner terminal, but does not assign chronology or coalesce the multiple owners allowed by the explicit fence. Proposition 51.38 removes rank eight from the shortest row. Thus the remaining intervals are rank nine only at $K' = 10$, rank nine or rank eight on $11 \leq K' \leq 15634$, rank eight only on $15635 \leq K' \leq 22525$, dense-owner chronology only on $22526 \leq K' \leq 37995$, and no rank-eleven component target on $37996 \leq K' \leq 1048576$. Proposition 51.40 shows that neither the distinct nor the weighted fixed-chart output alone can pay the lower rank-eight interval. Any continuation must retain the normalized full-span/anchor ancestry or couple charts and owners through a chronology-correct global incidence argument. Proposition 51.39 gives the exact first-row split: the full-circuit case has 54 rank-nine shadows, while every proper circuit gives additional rank-eight shadows and a locator ideal of dimension at least three. Summing fixed-chart caps without retaining owner labels is not asserted. The fence does not instantiate the fixed dense-anchor ancestors or an unsafe line. The packet does not recursively cover the remaining records, construct the chronology owner required by Remark 23.30, pay the remaining rank-eight or dense-owner chronology branch, move an active-v4 atom, or close KoalaBear.

52. THE NEW SPREAD RESIDUAL

Theorem 52.1 (spread residual after core-relative incidence). *Assume a deployed support-wise MCA-bad line contains a first-match primitive spread core and has more bad slopes than its row budget. Then all of the following must hold.*

- (R1) *At most 58 bad slopes use the core interpolant itself.*
- (R2) *The remaining corrections are not contained in one projective correction ray.*
- (R3) *If their linear correction span W has dimension at most 50 on KoalaBear or 15 on Mersenne-31, then some $s + 1$ coordinate equations have a positive-dimensional common component.*
- (R4) *If W does not absorb the high core, then its dimension and support excess lie beyond every transition in Corollary 50.5.*

- (R5) *Every low-dimensional surviving component is therefore an evaluation rank-flat or an exact polynomial clone component in the sense of [Lemma 50.3](#); it is not an unstructured failure of Bezout counting.*

Proof. Items (R1)–(R4) are the contrapositives of [Theorems 48.1](#), [49.3](#), [50.2](#) and [50.4](#). Item (R5) is [Lemma 50.3](#). $\square$

The theorem changes the spread-abundance target. It is no longer necessary to control arbitrary extensions of a core. The only remaining spread mechanism is

many projectively independent correction directions, together with
a high-dimensional correction span or a positive-dimensional
rank-flat/clone atlas surviving all first-match routes.

Problem 52.2 (next spread theorem). Intermediate status. *This is the local correction-component form of the final [Problem 54.1](#). It is retained to state the geometric alternatives exposed by [Theorem 52.1](#), but is not a separate terminal input.*

Prove that every positive-dimensional correction component from [Theorem 52.1](#) either

- (a) *yields a rational-owner, quotient, extension, field-drop, or quantified near-sunflower atom declared in [Theorems 39.1](#) and [42.2](#) and [Corollary 38.2](#);*
- (b) *has correction directions contained in a bounded union of projective rays, so [Theorem 49.3](#) pays it; or*
- (c) *admits a proper quotient correction space to which [Theorem 50.2](#) or [Theorem 50.4](#) applies.*

For Mersenne-31 it is enough that the total surviving spread contribution be at most 16,777,215; for KoalaBear the available budget is much larger.

Part 7. Completion certificates and remaining row-sharp inputs

53. INTEGRATED THEOREM STATUS

Theorem 53.1 (Audited theorem status). *All statements explicitly presented as theorems, propositions, lemmas, or corollaries in Parts I–VI are unconditional under their displayed hypotheses. In particular, the paper proves the complete support and profile compiler, the minimum-distance flag sharpening, the primitive- Q /Sidon implication and its shallow unconditional range, saturation, one-pencil moving-root BC , $Q \Rightarrow SP$, the order-32 rank dictionary, dyadic descent, rational-atom extraction and coherence, owner localization, the seven-point component-free owner-pencil bound, the 58-slope core-compatible bound, the complete correction-ray bound, and the displayed proper and clone-tolerant correction-space payments.*

No theorem in this paper proves either deployed adjacent MCA-safe row or either deployed adjacent list-safe row. The terminal open inputs used by the completion certificates are exactly [Problems 45.1](#), [54.1](#), [54.2](#) and [55.1](#). The earlier [Problems 41.2](#) and [52.2](#) are intermediate formulations subsumed by [Problem 54.1](#), not additional terminal cells.

Remark 53.2 (Status of the v3 saturated-BC problem). The v3 problem labelled `prob:saturated-bc` is superseded in v4, not silently discharged. Its primitive one-parameter MCA clause is proved by [Corollary 23.5](#). Its higher-dimensional MCA clause remains open and is represented jointly by [Problems 45.1](#), [54.1](#) and [54.2](#); in particular, v4 still requires exhaustive same-owner routing from every surviving balanced-core component into that endgame. Its separate arbitrary-word list-interior clause is contained in [Problem 55.1](#). Thus no archived problem is being banked as a live theorem, and no higher-dimensional balanced-core payment is claimed merely from the one-pencil moving-root bound.

54. THE FINAL MCA CERTIFICATE

The order-32 method replaces the old sum over many independently budgeted MCA cells by a maximum-type endgame.

Problem 54.1 (Complete spread-component routing). *This is the final v_4 successor to the higher-dimensional MCA clause of the v_3 `prob:saturated-bc`. Together with Problems 45.1 and 54.2, it must provide the exhaustive same-owner routing that the old balanced-core cell left open. Its row-sharp conclusion subsumes the intermediate Problems 41.2 and 52.2 for purposes of the final certificate.*

Prove that every positive-dimensional correction component surviving Theorem 52.1 is routed, with the same received-line owner retained, into one of the following disjoint outcomes:

- (a) *a rational-owner, quotient, extension, field-drop, or quantified near-sunflower atom;*
- (b) *a bounded union of complete projective correction rays paid by Theorem 49.3;*
- (c) *a proper quotient correction space paid by Theorem 50.2 or Theorem 50.4;*
- (d) *an explicitly named residual component with a literal row-sharp slope bound.*

The total spread contribution must fit the row budget.

Problem 54.2 (Complete MCA exception routing). *Verify the exception-routing input of Definition 46.1: denominator-root, extension, coordinate-clone, imperfect-quotient, and quantified near-sunflower failures must either be absorbed into the spread or large-owner bounds with a common owner retained, or occupy a set of at most 31 slopes.*

Theorem 54.3 (Exact MCA completion certificate). *Fix either deployed MCA row. If Problems 45.1, 54.1 and 54.2 are proved with the literal row constants in Parts V–VI, then*

$$B_{\text{MCA}}(m) \leq B^*.$$

Together with the unsafe witness at agreement $m - 1$ from CAP25 v13.2, the exact adjacent MCA threshold is m .

Proof. The hypotheses supply the spread-abundance and exception-routing inputs of Theorem 46.2; Problem 45.1 supplies its atom input. Apply that theorem and then the imported unsafe witness. $\square$

55. THE FINAL LIST CERTIFICATE

For a list row retain the disjoint summed ledger

$$(55.1) \quad U_{\text{list}} = U_{\text{paid}} + U_Q + U_{\text{list-int}} + U_{\text{ext}} + U_{\text{new}}.$$

Here U_{paid} is the exact first-match sum of all tangent, common-support, quotient, planted, field-drop, fixed-union, affine-span, common-zero, rank-flat, codimension-one, support-flag, minimum-distance, puncturing, and other certified cells; U_Q is the row-sharp pruned boundary-prefix maximum; $U_{\text{list-int}}$ is the high-nullity arbitrary-word interior codeword-ray term; U_{ext} is the exact extension projection payment; and U_{new} is zero unless an explicitly named residual cell survives.

Problem 55.1 (Row-sharp list completion). *This is the final v_4 successor to the arbitrary-word list-interior clause of the v_3 `prob:saturated-bc`; it also includes the remaining prefix and extension terms required by the v_4 list ledger.*

For each deployed list row, prove exact integer bounds for the remaining high-nullity prefix fiber, arbitrary-word interior, and extension projection terms so that

$$U_{\text{list}} \leq B^*.$$

The prefix estimate must use the actual residual mass and the exact budget left after the certified compiler cells are charged.

Theorem 55.2 (Exact summed completion certificate). *Fix one deployed list row. If an exhaustive first-match compiler proves, with no unresolved cell,*

$$U_{\text{list}} \leq B^*,$$

then the first safe integer agreement is the adjacent row $a_0 + 1$.

Proof. CAP25 v13.2 proves the unsafe inequality at a_0 . Exhaustive first-match coverage and the displayed upper bound give safety at $a_0 + 1$. Monotonicity identifies the first safe integer agreement. $\square$

56. CONCLUSION

This paper gives one unified Grande Finale manuscript. It contains the exact threshold refinements, profile and Fourier compiler, primitive $\mathbb{Q}$, support and minimum-distance recursions, saturation and moving-root geometry, the order-32 rank pivot, rational-atom extraction and coherence, owner localization, and the current spread-core incidence theory.

The remaining MCA problem has been compressed to three row-sharp inputs: route the positive-dimensional correction components isolated by [Theorem 52.1](#), prove the large-owner image bound of [Problem 45.1](#), and complete the finite exception routing of [Problem 54.2](#). The binding numerical atom target is the Mersenne-31 full-owner factor 9.57219783037. The remaining list problem is the exact high-nullity prefix/interior/extension sum of [Problem 55.1](#). Until those inputs are proved, no adjacent safe row is claimed.

REFERENCES

- [Cho26] P. Chojecki, Universal field-size caps and a two-sided sandwich for mutual correlated agreement on smooth Reed–Solomon domains, Version 13.2, manuscript, July 2026.
- [BS94] A. Balog and E. Szemerédi, A statistical theorem of set addition, *Combinatorica* 14 (1994), 263–268.
- [BGM23] J. Brakensiek, S. Gopi, and V. Makam, Generic Reed–Solomon codes achieve list-decoding capacity, *Proceedings of STOC 2023*; arXiv:2206.05256.
- [BCHKS25] E. Ben-Sasson, D. Carmon, U. Haböck, S. Kopparty, and S. Saraf, On proximity gaps for Reed–Solomon codes, ECCC TR25-169, 2025.
- [GG25] R. Goyal and V. Guruswami, Optimal proximity gaps for subspace-design codes and (random) Reed–Solomon codes, ECCC TR25-166, 2025.
- [GGSW26] R. Goyal, V. Guruswami, Y. Sun, and M. Wootters, Locality of curve-decoding and improved proximity gaps, arXiv:2607.08516, 2026.
- [GMR+22] B. Green, D. Matolcsi, I. Z. Ruzsa, G. Shakan, and D. Zhelezov, A weighted Prékopa–Leindler inequality and sumsets with quasicubes, in *Analysis at Large*, Springer, 2022, 125–129.
- [LW10] J. Li and D. Wan, A new sieve for distinct coordinate counting, *Sci. China Math.* 53 (2010), 2351–2362.
- [Tao10] T. Tao, Sumset and inverse sumset theorems for Shannon entropy, *Combin. Probab. Comput.* 19 (2010), 603–639.
- [TV06] T. Tao and V. Vu, *Additive Combinatorics*, Cambridge University Press, 2006.
- [Gow98] W. T. Gowers, A new proof of Szemerédi’s theorem for arithmetic progressions of length four, *Geom. Funct. Anal.* 8 (1998), 529–551.
- [GMT23] B. Green, F. Manners, and T. Tao, Sumsets and entropy revisited, arXiv:2306.13403, 2023.
- [Goh24] M. K. Goh, On an entropic analogue of additive energy, arXiv:2406.18798, 2024.
- [MRSZ22] D. Matolcsi, I. Z. Ruzsa, G. Shakan, and D. Zhelezov, An analytic approach to cardinalities of sumsets, *Combinatorica* 42 (2022), 203–236.
- [Wei48] A. Weil, On some exponential sums, *Proc. Natl. Acad. Sci. USA* 34 (1948), 204–207.